

The Ibis GMI (Group Method Invocation) Programmers Manual

The Ibis Group

November 16, 2009

1 Introduction

For many parallel and distributed applications, the simple synchronous unicast communication model (one-to-one communication with a reply) offered by RMI is inadequate. Applications often require more different, more complex forms of communication, such as asynchronous unicast (one-to-one communication without a reply), broadcast (one-to-all communication), multicast (one-to-many communication), or data reduction operations (where data must be collected from multiple machines). Although these alternative forms of communication can be implemented using RMI, this is often complex and inefficient. For example, a simple way to implement multicast is to perform multiple RMI calls, one after the other. Unfortunately, this is very inefficient, since each RMI must wait until the previous RMI has finished completely, including waiting for the (unused) result to be returned. More efficient implementations use threads to perform multiple RMIs simultaneously or create distributed multicast trees which use multiple machines to forward calls. These implementations are complex, however, and often suffer from performance problems caused by the synchronous nature of RMI. In Ibis, we offer a new programming model called GMI (Group Method Invocation). GMI is an extension of RMI designed to be flexible enough to express the complex forms of communication needed by many parallel applications. Nevertheless, GMI is as easy to use as RMI, hiding the complex details of the communication in its implementation. The flexible model of GMI also allows its implementation to use efficient communication algorithms for multicast, data reduction, etc. The GMI model will be explained in detail below.

2 The GMI model

The GMI model generalizes the RMI model in three ways. First, it introduces the notion of a 'group', a set of objects which all implement the same interface. The objects in a group may be distributed over a number of JVMs. A group can be addressed using a single 'group reference'. An example is shown in Figure 1, where an application thread uses a single group reference to address a group of two objects (on different JVMs).

Like RMI, GMI uses compiler generated stub and skeleton objects which implement the necessary communication code. The application programmer does not have to write any communication code. Unlike RMI, however, GMI does not only generate




Figure 1: A group invocation with GMI.

code for synchronous unicast communication. Instead, many different types of communication code are generated, both for sending of method invocations and for returning of results. This is the second generalization introduced by GMI. Finally, through a simple API, the programmer can configure at run time which type of communication should be used to handle each individual method of a group reference. This API will be explained in more detail in the next section. Different ways of forwarding the invocation and handling the results can be combined, giving a rich variety of communication mechanisms. By configuring the methods at run time, the communication behavior of the application can easily be adapted to changing requirements.

The following types of method invocation forwarding are currently supported by GMI:

- single invocation: The method invocation is forwarded to a single object of the group, identified via a rank.
- group invocation: The invocation is forwarded to every object in the group.
- personalized group invocation: The invocation is forwarded to every object in the group, while the parameters are personalized for each destination using a user-defined method.
- combined invocation: Multiple application threads (possibly on multiple JVMs) invoke the same method on the same group. These invocations are combined into a single invocation using a user-defined method. This single invocation is forwarded to the group using one of the three other forwarding schemes.

The following types of method result handling are currently supported by GMI:

- discard results: No results are returned at all (including exceptions).
- return one result: A single result is returned, preselected via a rank if necessary.
- forward results: All results are returned, but they are forwarded to a user-defined object rather than being returned to the invoking thread.
- combine results: Combine all results into a single one using a user-defined method. The combined result is returned to the invoker.

```

import ibis.gmi.GroupMember;

class Implementation extends GroupMember implements Example {

    private String message = null;

    public synchronized void put(String message) {
        this.message = message;
        notify();
    }

    public synchronized String get() {
        while (message == null) {
            wait();
        }
        return message;
    }
}

```

Figure 2: Implementing a group object with GMI.

- **personalize result:** A result produced by one of the other result handling schemes is personalized using a user-defined method before being returned to each of the invokers (this is useful when a combined invocation is used). The four different forwarding schemes and five different result handling schemes can be combined orthogonally, resulting in a wide variety of useful communication patterns.

3 Hello world in GMI

We will now show a step by step example of how a GMI application can be written. The first step is to create an interface which will define the methods which can be invoked on the group of objects. Like in RMI, we use a special 'marker interface' to 'mark' group interfaces. Any interface extending `ibis.gmi.GroupInterface` will be recognized by the Ibis compiler as being a group interface. The Ibis compiler will then generate a stub object which contains the necessary communication code.

```

interface Example extends ibis.gmi.GroupInterface {
    public void put(String message);
    public String get();
}

```

In the example above, the 'Example' interface is turned into a group interface by extending `ibis.gmi.GroupInterface`. It defines just two methods, `put`, which can be used to store a string, and `get` which can be used to retrieve a stored string. After creating the group interface, an implementation of this interface is be created. This implementation object must implement the Example interface, and extend the `ibis.gmi.GroupMember` object, which contains some basic functionality needed to be part of a group (this is similar to the `UnicastRemoteObject` used in RMI).

In this implementation, shown in Figure 2, the `put` method stores the string it receives in the object, from where it can be retrieved using the `get` method. The standard synchronization primitives `synchronized`, `wait`, and `notify` are used to prevent the `get`

```

import ibis.gmi.*;

class BroadcastExample {

    public static void main(String[] args) {
        int size = Group.size();
        int rank = Group.rank();

        if (rank == 0) {
            // JVM 0 creates a new group.
            Group.create("ExampleGroup", Example.class, size);
        }

        // All JVMs create an implementation object.
        Implementation impl = new Implementation();

        // And join the group
        Group.join("ExampleGroup", impl);

        if (rank == size-1) {
            // The last JVM retrieves a group reference
            Example group = (Example) Group.lookup("ExampleGroup");

            // Then configures 'put' to be forwarded to the whole group
            GroupMethod m = Group.findMethod(group, "void put(java.lang.String)");
            m.configure(new GroupInvocation(), new DiscardReply());

            // Now invoke a method on the group
            group.put("Hello world!");
        }

        // All JVMs can now retrieve the data using a local(!) call
        String message = impl.get();
        System.out.println(message);

        // Done
        Group.exit();
    }
}

```

Figure 3: An example GMI application that uses a group object.

from returning before the string is available. Next, we will create an simple example application, `BroadcastExample`, which will use the group object. The source code is shown in Figure 3. This application is parallel; it is designed to be started simultaneously on multiple JVMs.

The main method of the application starts by invoking `Group.size` and `Group.rank`, two utility methods of GMI which can be used to find out how many JVMs are available (`size`), and what number is assigned to the current JVM (`rank`). The JVM with rank 0 then creates a new group using `Group.create`. This group will have the name `ExampleGroup`, use a group interface of the type `Example` and will contain 'size' objects. Each JVM then creates it's own `Implementation` object, and adds it to the group

using the `Group.join` method. `Group.join` will block until all 'size' objects have been added to the group. The last JVM (with rank 'size-1') then retrieves a group reference using `Group.lookup`. This method will return a stub generated by the Ibis compiler. This stub contains the communication code necessary to communicate with the objects in the group. Because the stub implements the `Example` interface, normal invocations of the `put` and `get` methods can be used and no communication code needs to be written by the programmer. But before the methods can be invoked, the group stub must first be configured. All methods in a group stub can be configured separately. In this example we will only use the `put` method. To configure `put`, we first perform a lookup of the method using `Group.findMethod`. A `GroupMethod` object will be returned, which represents the `put` method of the stub. Using the `configure` method in this `GroupMethod` object, it can be specified how the invocations of the `put` method should be handled. For this purpose the `configure` method takes two parameters, one describing how the invocation must be forwarded to the group, and one describing how the replies should be returned. In the example application we use `GroupInvocation` and `DiscardReply`, which indicates that invocations of `put` will be forwarded to all objects in the group, and that no replies will be returned. After the configuration is completed, the last JVM invokes the `put` method. The method invocation is then forwarded to all object in the group. All JVMs then retrieve their local copy of the string by directly invoking the `get` method on their implementation objects.

4 Further Reading

The Javadoc included in the `javadoc` directory has detailed information on all classes and their methods.

The Ibis web page <http://www.cs.vu.nl/ibis> lists all the documentation and software available for Ibis, including papers, and slides of presentations.

For detailed information on running a GMI application see the User's Manual, available in the `docs` directory of the Ibis GMI distribution.