



# ***Maestro: a Self-Organizing Dataflow Framework***

***Kees van Reeuwijk***

*reeuwijk@few.vu.nl*

***13 June 2009  
HPDC 2009***



# ***Motivation***

- Parallel systems are often inhomogeneous and unreliable
- Communication links are often inhomogeneous or imperfect too
- Parallelism is increasingly mainstream (multi-core, GPUs, specialized processors). Even a single consumer PC can be a heterogeneous system.
- Call it what you want: **distributed system**, grid, cloud, cluster...



# ***Distributed Systems Problems***



- Keeping an application running (efficiently) is hard!
  - Resources come and go
  - Resources crash
  - Heterogeneous: load balance??
- Any fixed use of resources is bound to fail



**Resource allocation must be  
dynamic and adaptive**



# Dataflow framework



- Jobs with one input, one output

```
interface Job {  
    Object run(Object in); }
```


- Jobs connected in series (pipeline) or in parallel
- Nested
- Predictable performance *per job*



# ***Maestro: self-organizing***


- Nodes with special tasks are failure points/bottlenecks
- In particular central nodes (scheduler!)

Solution: peer to peer  
⇒ self organizing



## ***Exception: work insertion***

- Currently there is one exception: only one node inserts work in the system, and handles final results
- Application specific







# Maestro Nodes

Any number, may join and leave any time

Each node contains:

- **Worker**: execute jobs from queue
- **Master**: distribute jobs over workers
- **Gossiper**: exchange performance info





## ***Scheduling policy***

- Each master tries to optimize for **total completion** time of all remaining steps
- Measured and gossiped:
  - Worker queue & compute stats
  - Master queue stats
  - Transmission time (not gossiped)
- Regulars are informed ASAP
- Efficient nodes are favored





# ***Learning strategy***

Emergent behavior: the system **learns**  
an efficient schedule:  
**reinforcement learning**

Consequences:

- In a homogeneous system the local node is favored
- New nodes should start with optimistic estimates




# ***Limited commitment***

Every worker should have one job waiting in its queue: no more, no less

- Limits commitment to one node, but reduces idle time
- Gives opportunities to less attractive nodes




# *Implemented on Ibis*

- A framework for distributed computing
- Based on Java (portable!) 
- Provides message passing, serialization (IPL layer)
- Join-Elect-Leave support (malleability)
- Robustness is central
  - Detect failed nodes
  - Circumvent NATs, firewalls, etc.
  - Handle multiple NICs (multi-homing)



# Benchmark


- Operations on video frames
  1. Generate 720x576 frame
  2. Scale to 1440x1152
  3. Sharpen (3x3 convolution)
  4. Compress (JPEG)
  5. Discard



# *Testbed*

VU cluster of the DAS3:

- 85 nodes:
  - 2x dual-core 2.4 GHz AMD Opteron
  - 4 GB memory
- Myrinet 10G interconnect
- In total there are 5 clusters with similar specs throughout the Netherlands




# ***Node configurations***

- Homogeneous
  - We expect:
    - Work is evenly divided over the nodes
    - All five stages of the video processing on the same node
- All steps in one job
  - We expect:
    - Work is evenly divided
  - Maestro is just used as master/worker



# *Homogeneous results*



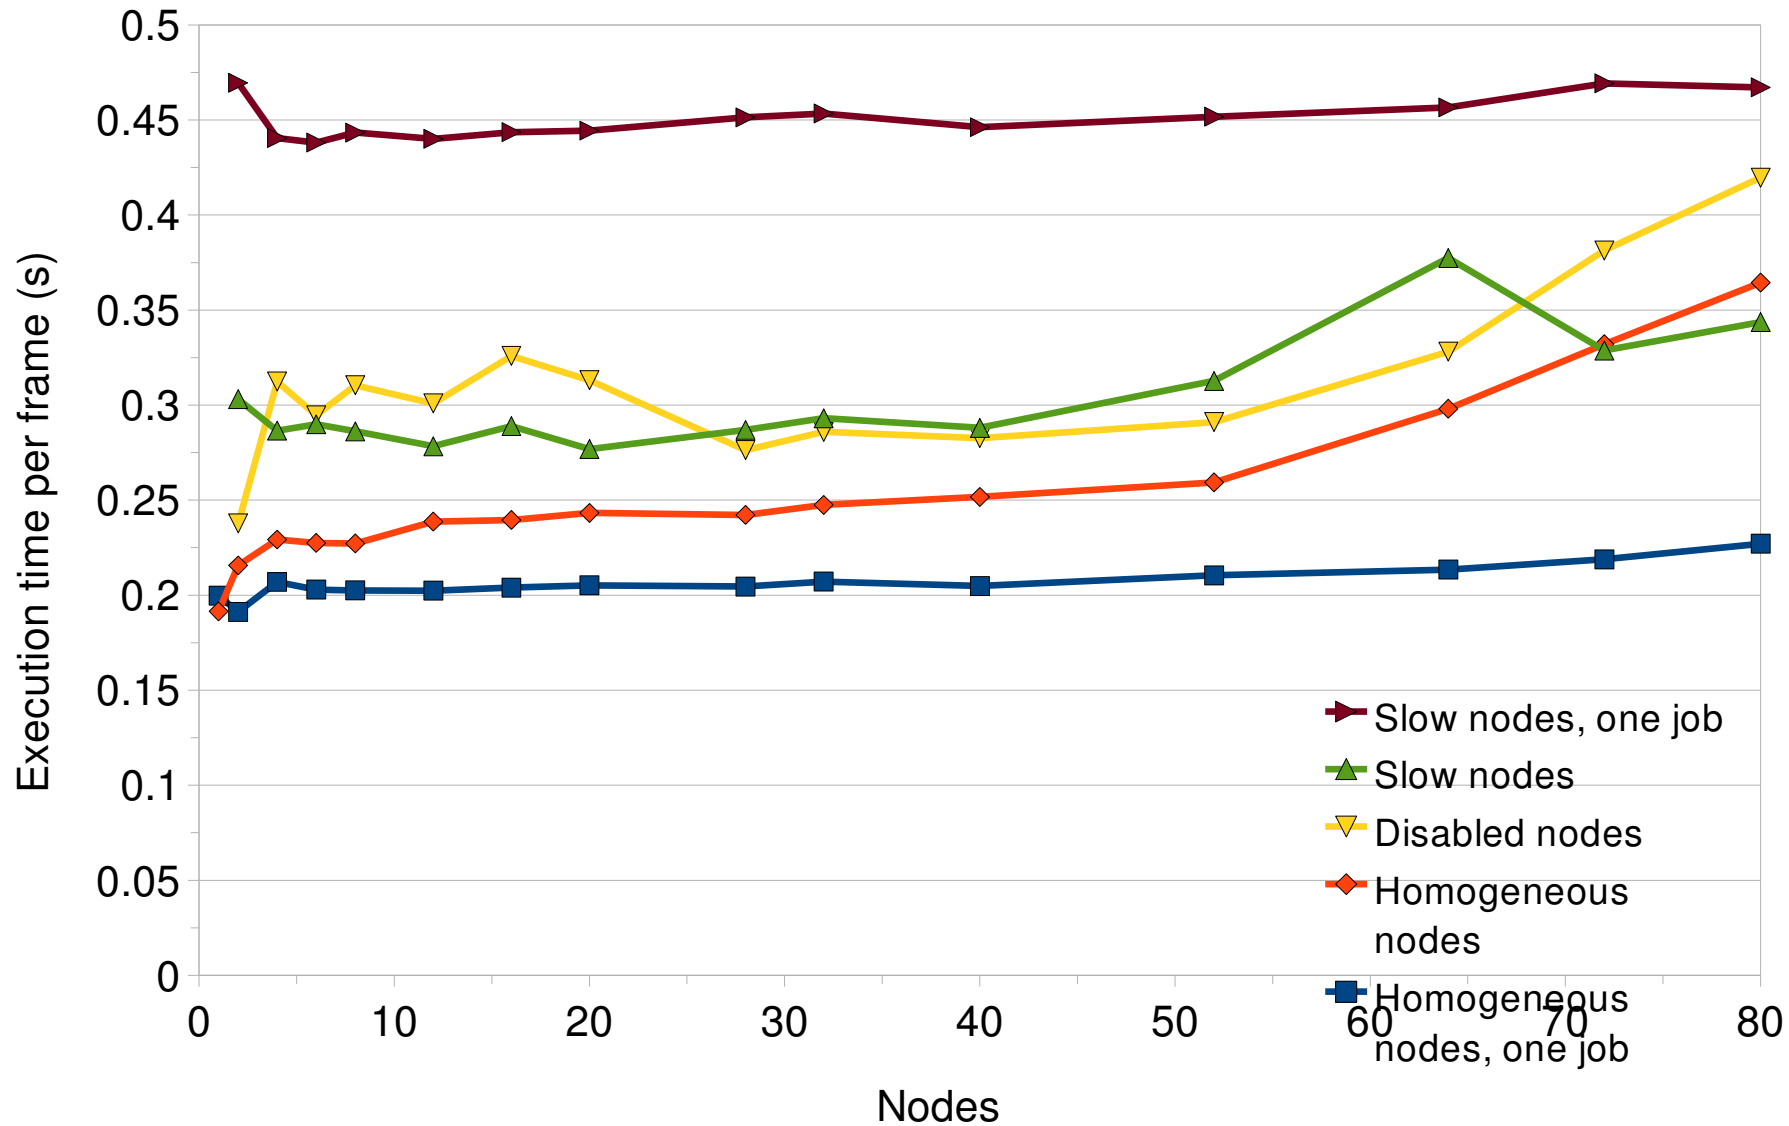
# ***Heterogeneous configurations***

- Half no scaling, half no sharpening
  - Now forced to `zigzag`
- Slow scaling, slow sharpening
  - At least the `zigzag`
- One job, slow scaling, sharpening
  - Slow computation unavoidable






# All results




# Work distribution

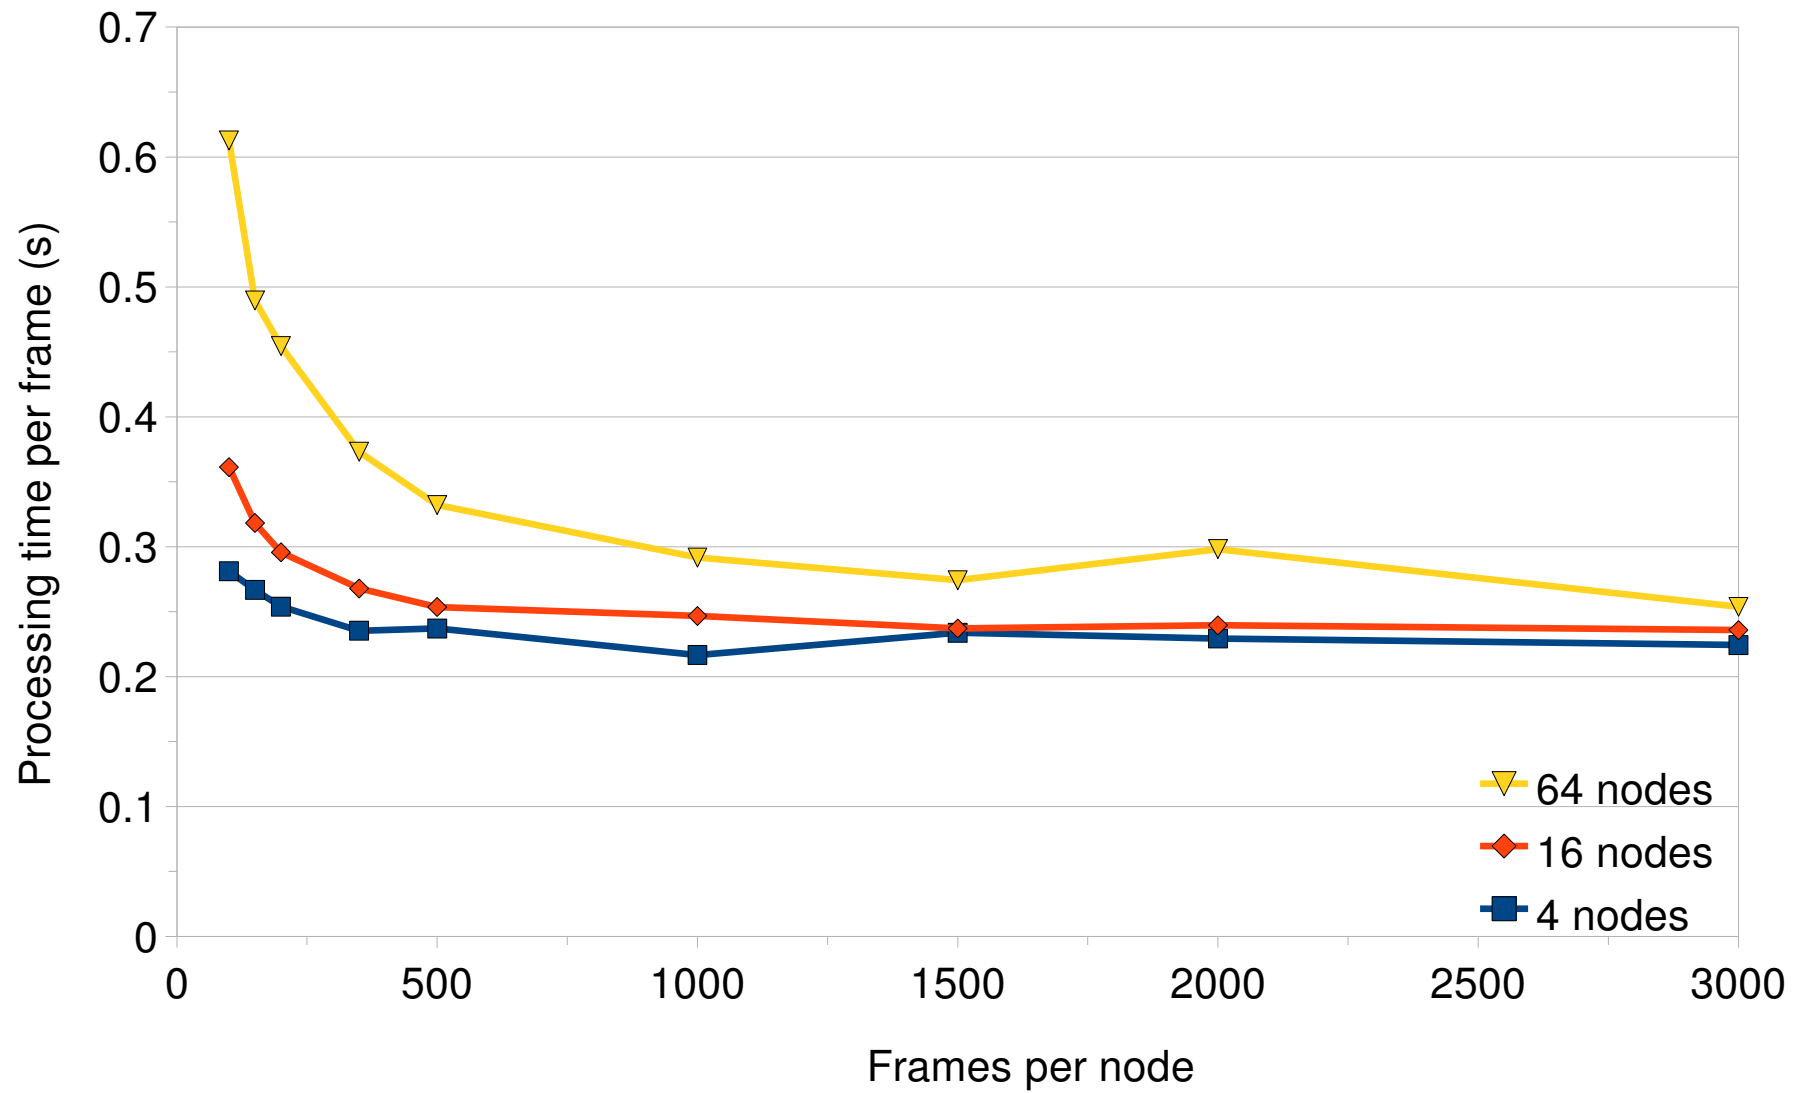
20 homogeneous nodes




10 nodes with slow scaling, 10 with slow sharpening



# ***Learning: homogeneous***

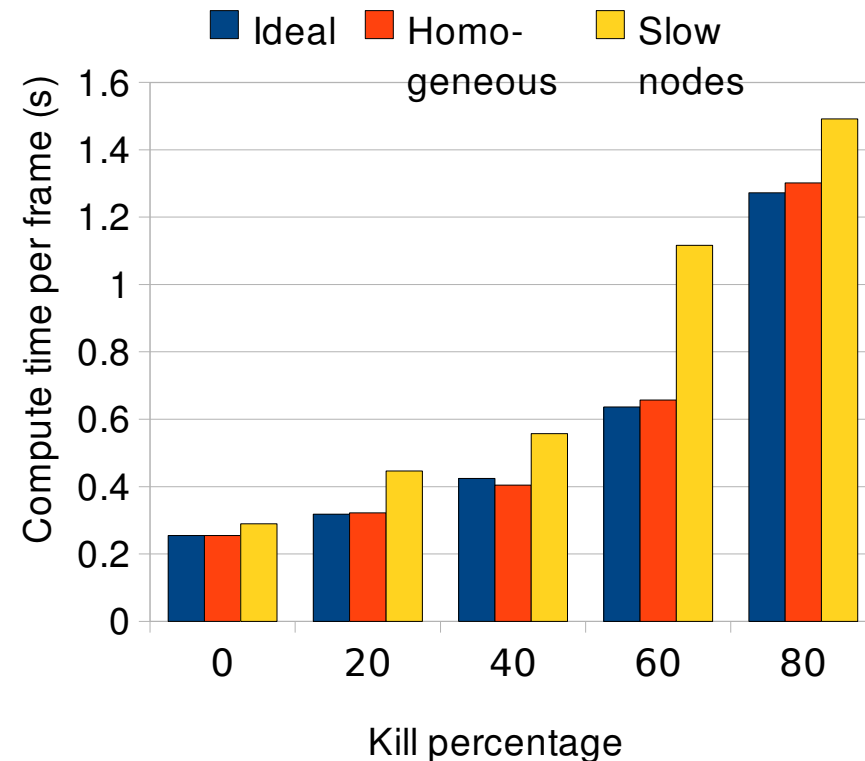


# *Learning: slow nodes*



# ***Fault tolerance***

- We start a run on 30 nodes
- After a few seconds kill some nodes
- Ideally, the rest of the nodes should take over the work
- All masters restart any work that was lost on the dead nodes
- Retry outstanding frames



# ***Conclusions & future work***

## Conclusions

- Self-organization of a data-flow computation works
- Can exploit strong points of non-homogeneous systems
- Extremely robust

## Future work

- Integrate with divide & conquer
- Scalability



# *Questions?*

?



`www.cs.vu.nl/ibis`

