



Ibis

Communication Library and Programming Models

Jason Maassen

jason@cs.vu.nl

Rob van Nieuwpoort

rob@cs.vu.nl



***Friday 27 April 2007
Vrije Universiteit Amsterdam***



vl·e

Overview

- Philosophy / design / implementation
 - Why Ibis ?
 - Design
 - Communication Library Model
 - Cool features
 - Programming Models
 - Satin
 - MPJ
 - GMI



We are interested in...

- Parallel applications on “the Grid”
 - Single site runs
 - Grid == big collection of clusters
 - Only communicate within cluster
 - Use fast local network (Myrinet/Infiniband/...)
 - Multi site runs
 - Grid == big processor pool
 - Communicate between clusters
 - Use regular network & internet



So why Ibis?

- Ideally, grid computing should be “fire and forget”
 - Develop application locally
 - Submit to some grid (using GAT)
 - Finds some suitable site(s)
 - Transfers your application and data to the sites, and runs it.
 - Returns the result



Problems

- Lots of problems
 - Resource selection
 - Data transfer
 - Security and authentication
 - ...
 - Heterogeneity
- GAT and SAGA solve part of this



Problems

- Grids are heterogeneous:
 - Intel / PowerPC / Mips / Arm / ...
 - Windows / Linux / Unix / OSX / ...
 - Different OS/library/tool versions
 - Different networks
- Compiled (C/MPI) apps. huge pain:
 - Need executable for every combination of CPU/network/OS/libraries etc.
 - Hard to connect sites together.
 - Makes 'fire & forget' runs really hard...



Solution (partly)

- So, we use Java instead C or Fortran
 - No recompilation required
 - Runs (almost) anywhere
 - Doesn't work on supercomputers such as Hitachi SR8000, IBM BlueGene ...
 - ... but most sites have clusters anyway!
 - Acceptable performance
- But ... only part of the solution!



How about 'portable' communication?

- Class libraries are portable, but ...
 - Sockets are too low-level
 - RMI model/performance is limited
- Most parallel libraries not suited ...
 - Example: MPIJava requires native code
 - needs recompilation
 - only supports static (fixed size) runs
 - multi-cluster MPI is a bit hard
 - not all applications are SPMD



Ibis

- Solution: Ibis Communication Library!
 - A “run-anywhere” communication library
 - Just send it along with your application!
- Plus: flexible communication models
 - Malleability & Fault-Tolerance
 - Change number of machines during the run
 - More than just unicast communication
 - More about this later



Ibis 2.0

- In this tutorial we describe Ibis 2.0
 - Ideas are the same, but interface is cleaned-up
 - Ibis 1.x interface stable for 4 years
 - Ibis 2.0 interface changed according to the lessons we learned from the previous versions



Ibis

- Portability vs. performance
 - On a single site run you often want to use the fast local network
- Ibis allows specialized implementation
 - Designed for Myrinet, Infiniband, etc.
 - Usually use native code
 - Installed in advance
 - not portable
 - cannot be shipped with application



Ibis

- As a result, there may be multiple Ibis' available on a site
 - Automatically choose 'best' at startup
 - Based on requirements specified by
 - Application & user (using properties)
- Not every appl. needs all features
 - Pick one at startup that suits your needs....




Ibis Design




Ibis Portability Layer (IPL)



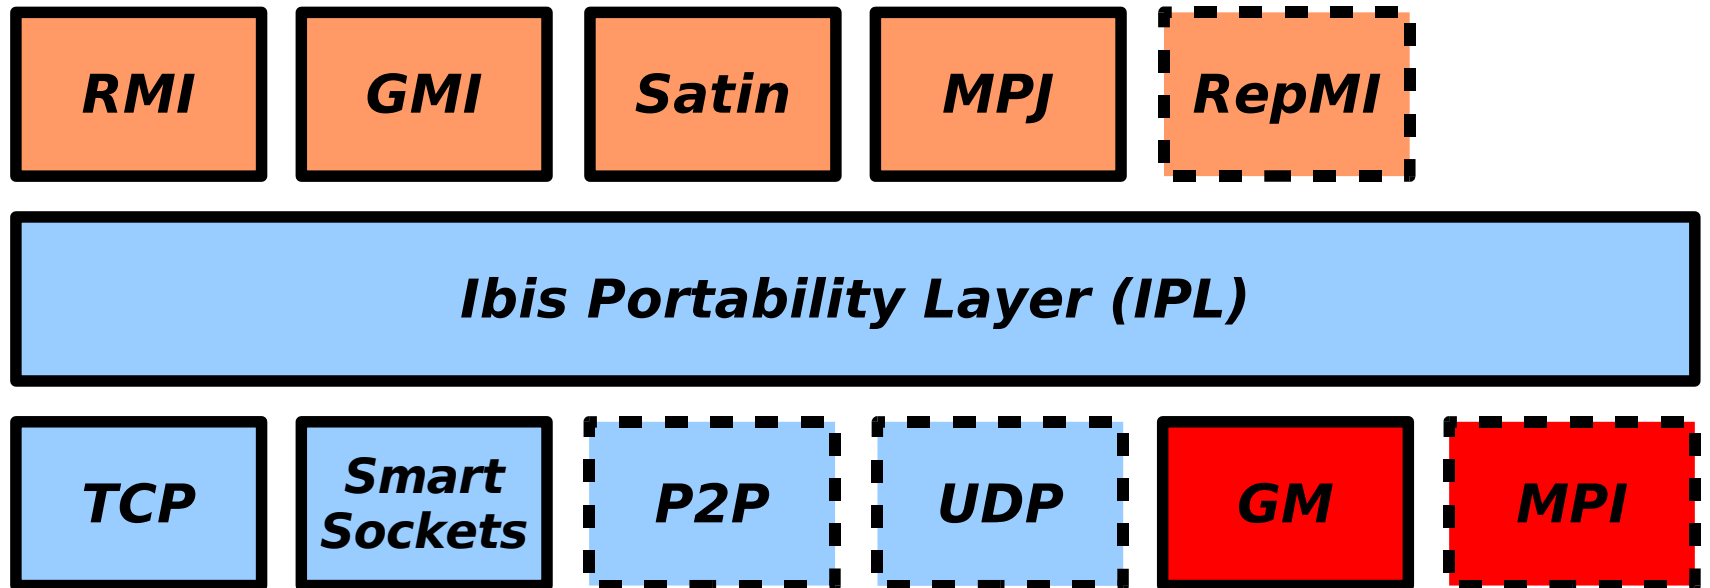
Ibis Design



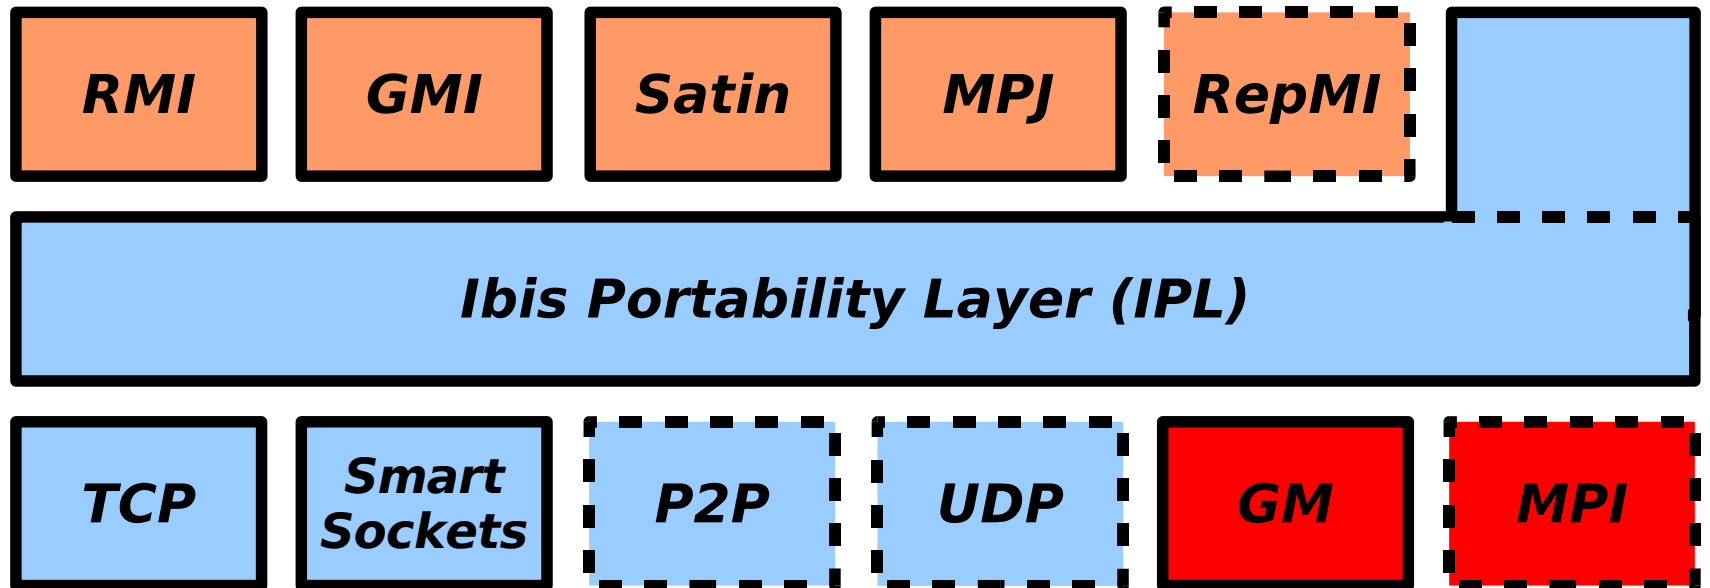
Ibis Design



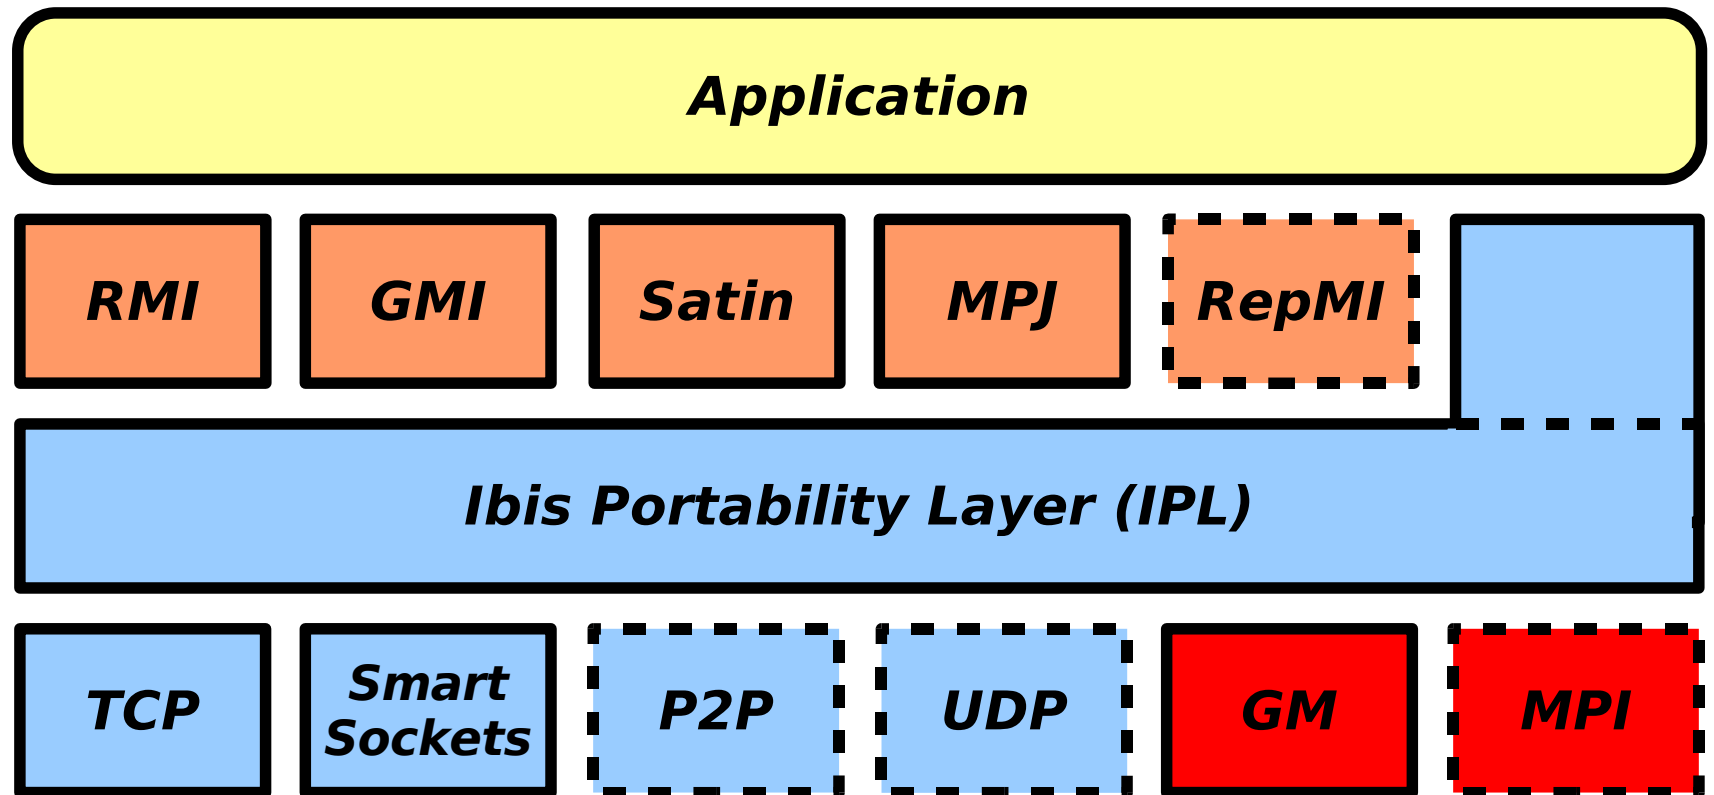
Ibis Design



Ibis Design



Ibis Design



Ibis Design

Application

ProActive

RMI

GMI

Satin

MPJ

RepMI

Ibis Portability Layer (IPL)

TCP

***Smart
Sockets***

P2P

UDP

GM

MPI



Ibis Portability Layer

- Basic Ibis programming interface
 - Reasonably simple (5 classes, 16 interfaces, 10 exceptions)
- Contains methods for
 - Loading an Ibis
 - Malleability/Membership
 - adding & removing machines
 - Connection handling
 - Communication primitives (low-level)



Basic Idea

- Abstract away from implementation
 - use abstract addressing scheme
 - hides real network addressing
 - use abstract communication primitives
 - hides real communication primitives
- Results in more portable applications
 - same application runs on sockets and MPI without changing any code



IbisIdentifiers

- In a parallel/distributed application
 - Each process has an Ibis instance
 - Each instance has an IbisIdentifier
- IbisIdentifier:
 - Uniquely identifies an Ibis instance
 - Abstracts away from the implementation
 - e.g. hostnames, IP addresses, MPI-ranks, etc.
 - Makes your application a bit more portable



Communication

- 'Low-level' communication model
- Unidirectional pipes
- Two end points
- Connection oriented (allows streaming)



Send & receive ports

- Can be connected in arbitrary ways




Send & receive ports

- Can be connected in arbitrary ways
- One to one (unicast) ...




Send & receive ports

- Can be connected in arbitrary ways
- ... one to many (multicast) ...




Send & receive ports

- Can be connected in arbitrary ways
- ... many to one ...



Send & receive ports

- Can be connected in arbitrary ways
- ... or some combination!




Send & receive ports

- Advantages:
 - Very simple & abstract model
 - Easy to implement using TCP/UDP/MPI/etc.
 - Allows multicast, many-to-one, etc.
 - Useful for parallel programs
 - Allows efficient implementation
 - Can be implemented using efficient low-level primitives (i.e., mpi-broadcast)
 - Other models do prevent this (e.g., RMI)



Send & receive ports

- Disadvantage:
 - Simplicity may cause some overhead...
 - Example: need two pairs for RPC / RMI




Port Types

- All ports have a type
 - Consists of a set of required capabilities:
 - Connection patterns
 - Unicast, many-to-one, one-to-many, many-to-many.
 - Communication properties:
 - Fifo ordering, numbering, reliability.
 - Serialization properties:
 - bytes, data, object
 - Message delivery:
 - Explicit receipt, automatic upcalls, polling




Port Types

- Defined at runtime
 - Specify set of capabilities
- Types must match when connecting!



Port Types

- Defined at runtime
 - Specify set of capabilities
- Types must match when connecting!



Port Types

- Forces programmer to specify how each communication channel is used
 - Prevents bugs
 - Exception when contract is breached
- Allows efficient impl. to be selected
 - Unicast only ?
 - Bytes only ?
 - Can save a lot complexity!




Connection setup

- Two options:
 - 1) Using a IbisIdentifier and a name
 - Name specifies the receiveport
 - Unique per Ibis instance
 - Human-readable (usually)
 - 2) Using a ReceivePortIdentifier
 - Uniquely identifies a receiveport
 - Created when ReceivePort is created
 - Can be passed around between Ibis instances.




Connection setup (1)

ibis-az33zx7



Create Ibis



ibis-34fdw21



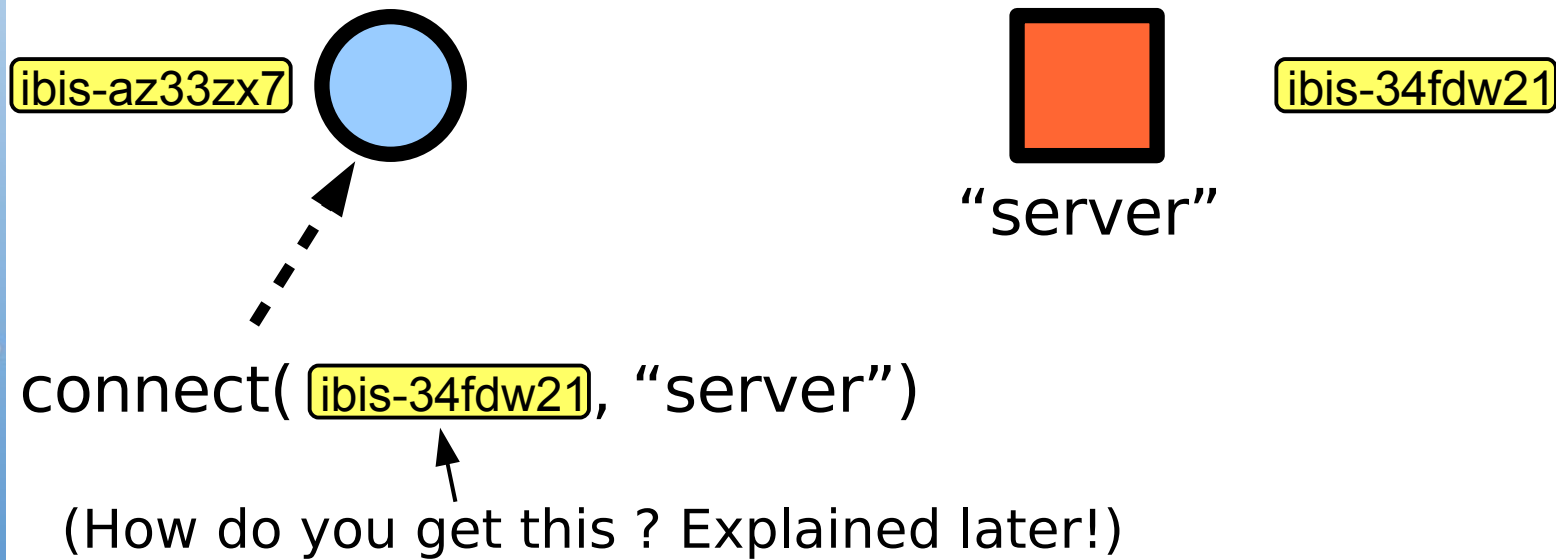
Create Ibis



Connection setup (1)



Connection setup (1)




Connection setup (1)




Connection setup (2)

Create anonymous
ReceivePort


Create SendPort




Connection setup (2)




Connection setup (2)



Connection setup (2)



Connection setup (2)




Messages

- Ports communicate using 'messages'
- Contain read or write methods for
 - Primitive types (byte, int, ...)
 - Object
 - Arrays slices (partial write / read in place)
- Unlimited message size




Messages

- Get WriteMessage from SendPort




Messages

- Write data into WriteMessage




Messages

- Finish the WriteMessage




Messages

- Data is send to ReceivePort




Messages

- ReceivePort produces ReadMessage
 - Explicit receive or callback (upcall)




Messages

- Read data from ReadMessage



Messages

- Finish the ReadMessage




Messages

- Done!



Messages or streams ?

- Message size is unlimited
 - Data may be forwarded at any time
 - Both S. & R. messages alive at same time
 - There's streaming!



Restrictions

- Must write and read data in same order
- A port can have only one message 'alive' at a time
- Messages are not thread safe (but ports are)
- The sender may block if the receiver is too slow (there may be flow-control)



Serialization

- Ibis supports 3/4 types of serialization
 - Bytes (no serialization at all)
 - Data (only primitive types/arrays)
 - Object (graphs of object + previous)
 - Sun (standard Sun serialization)
 - Ibis (efficient Ibis serialization)



Ibis Serialization

- Based on bytecode-rewriting
 - Adds serialization and deserialization code to serializable types
 - Prevents reflection overhead during (de-)serialization
 - Has fallback mechanism for non-rewritten classes
- Experimented with runtime rewriting



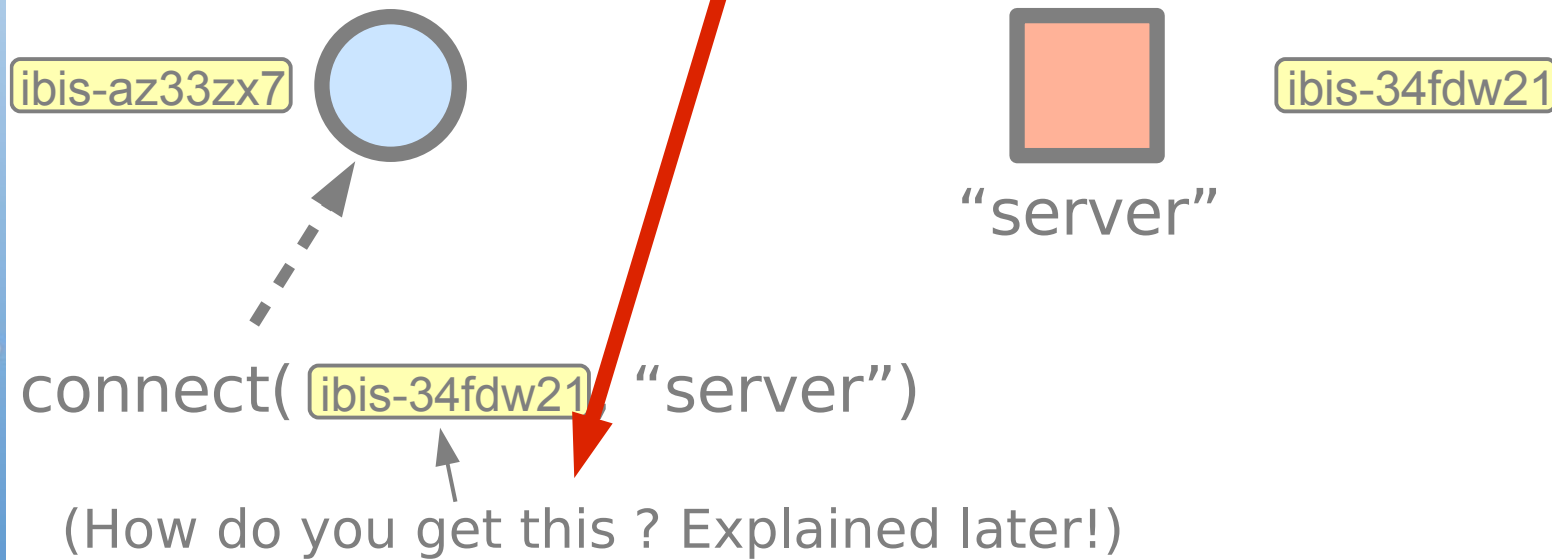
Short Recap

- First create PortType
- PortType creates Send & ReceivePort
 - Type is checked when connecting
- Several ways to connect
 - Abstract addressing
- Use Messages to communicate
 - Allows streaming
 - 3/4 types of serialization



Connection setup (1)

Remember this question ?



Pools & Malleability

- Ibis instances are part of a pool
 - Either variable size or fixed (create-once)
 - Fixed used by 'legacy' MPI-type applications
- Membership information
 - Can subscribe to pool information
 - Updates when Ibis instances join or leave pool
 - Useful for determining who's participating
 - Also used for fault-tolerance

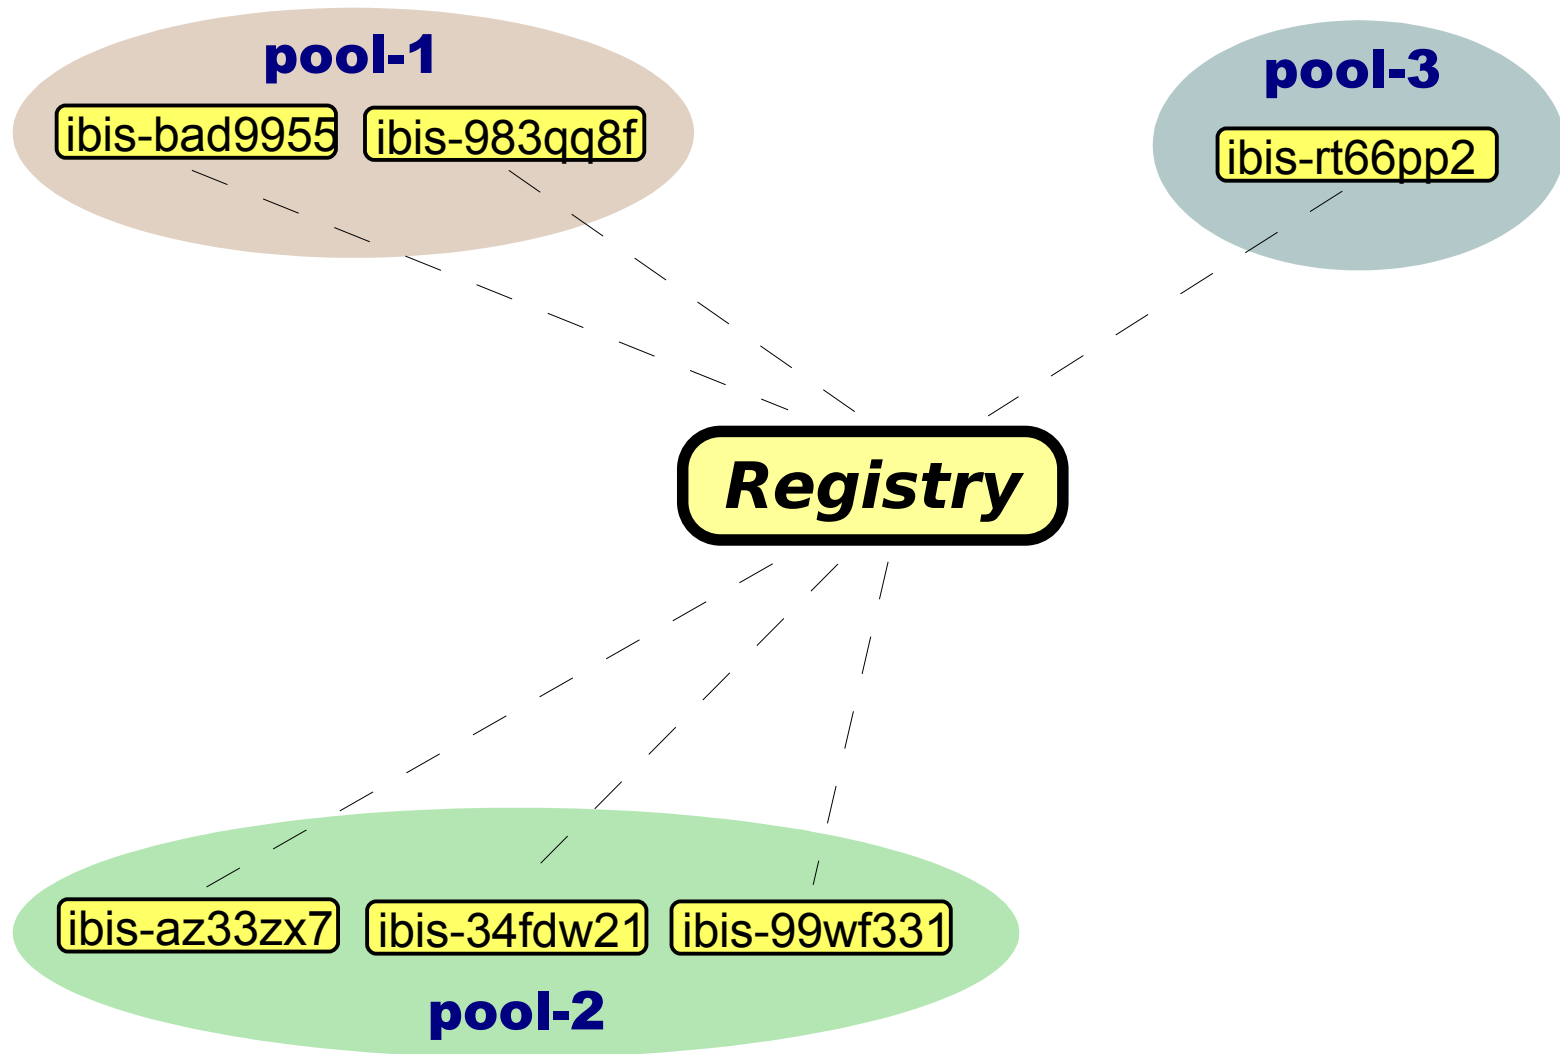


Membership updates

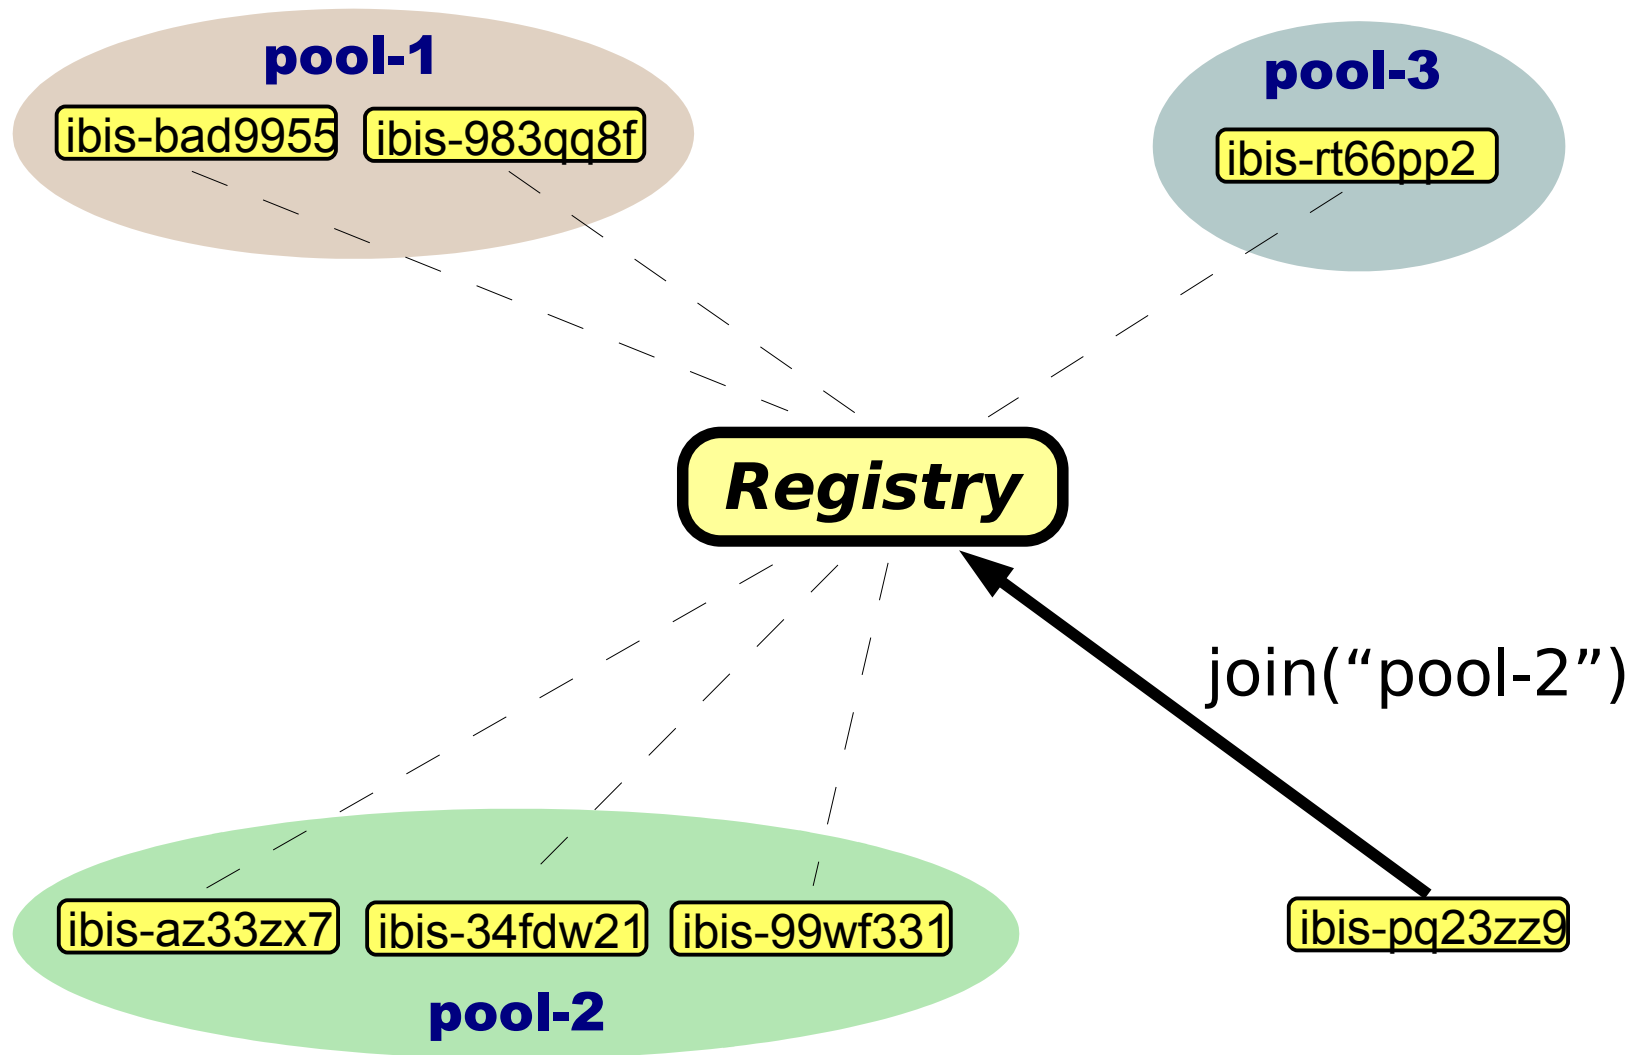
- Callbacks or explicit calls
 - If required, they can be delivered in the same order on all Ibis instances
- Uses external server (Registry)
 - Registry can handle multiple pools
 - No communication between pools



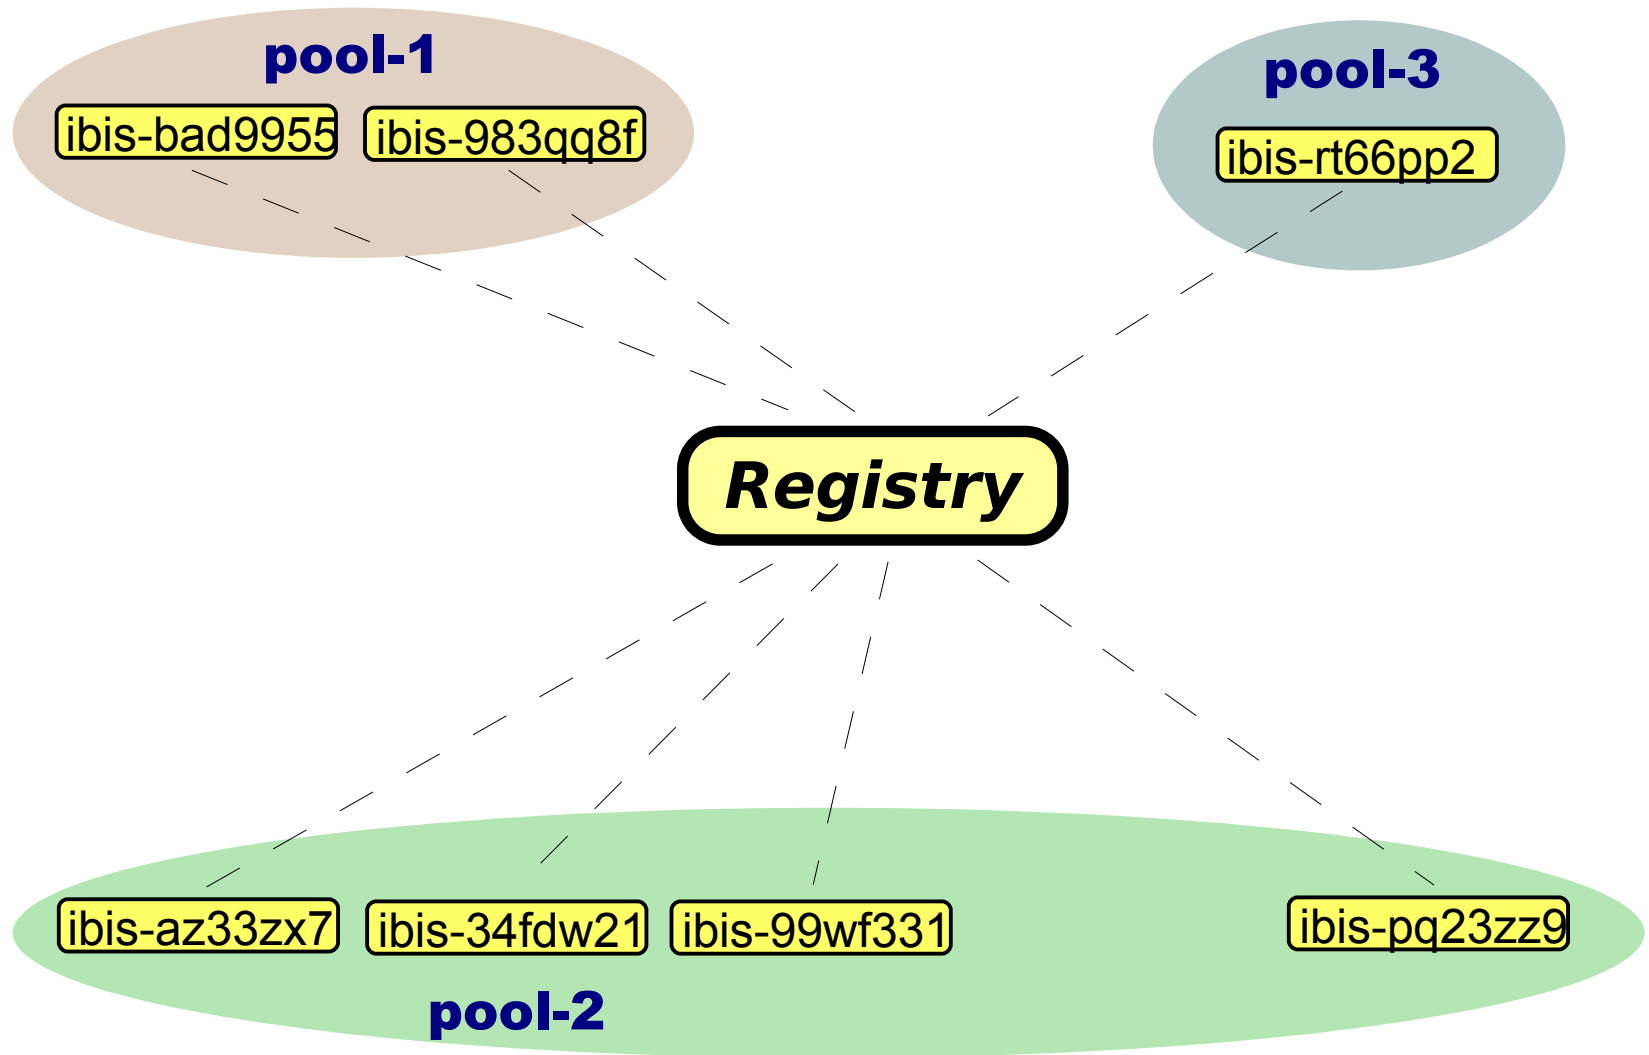
Pools & Malleability



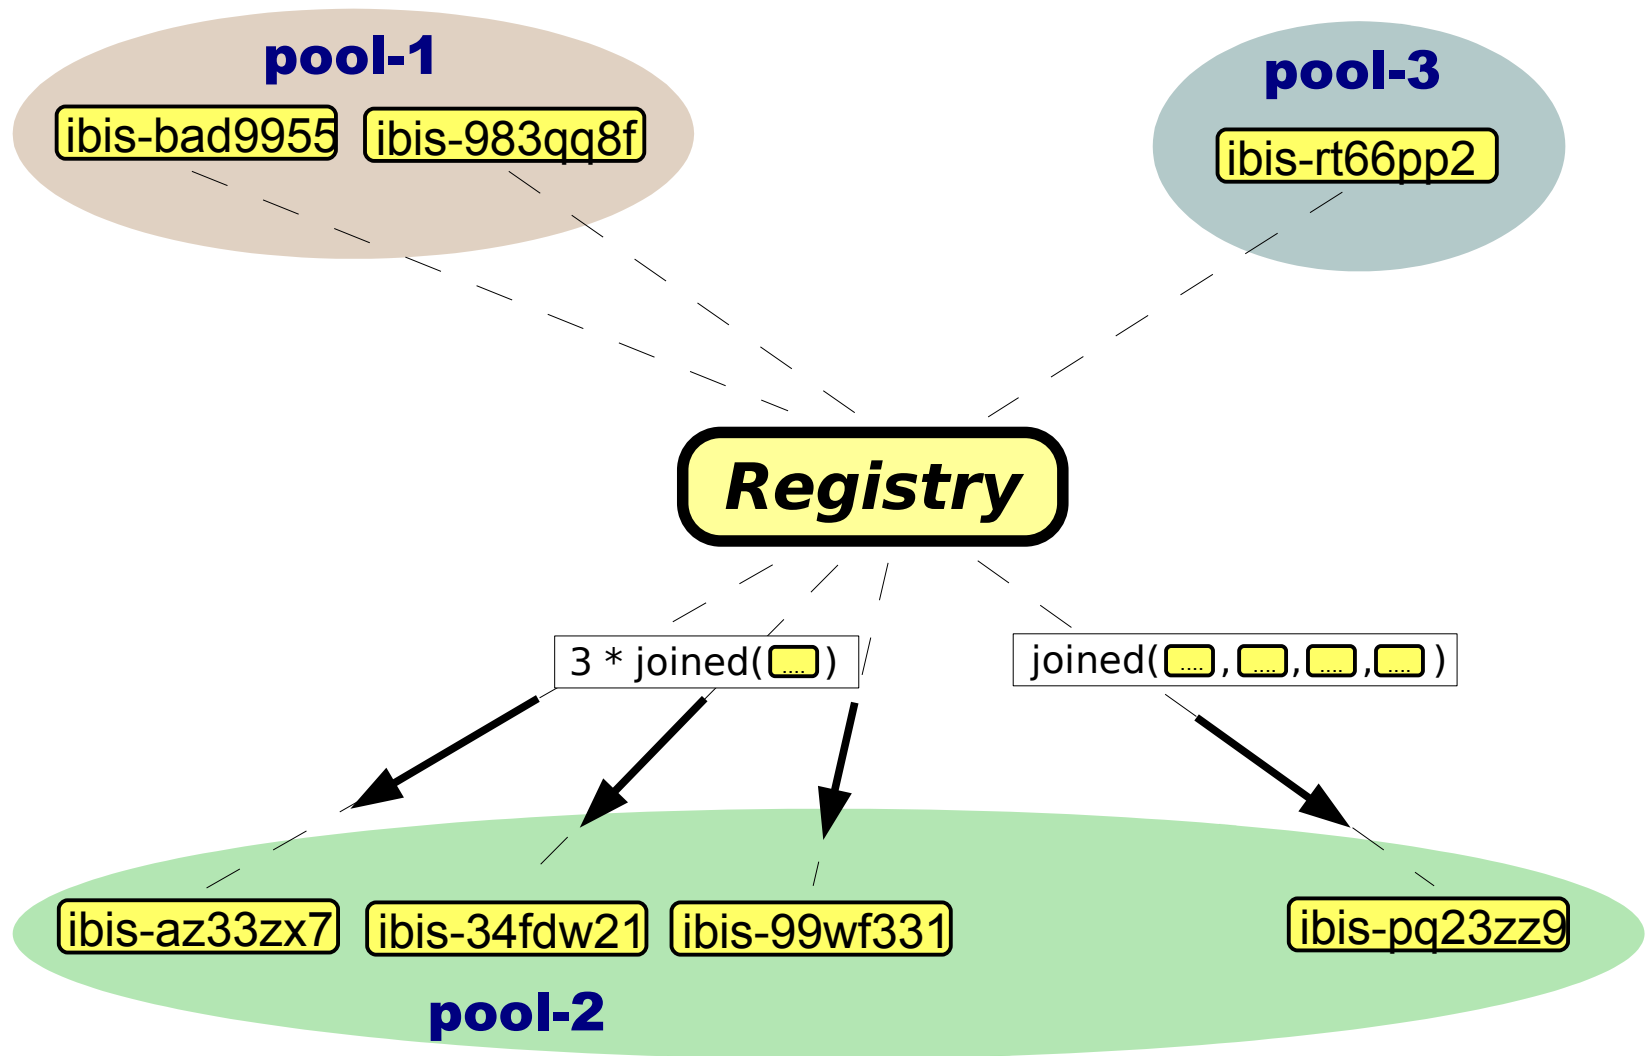
Pools & Malleability



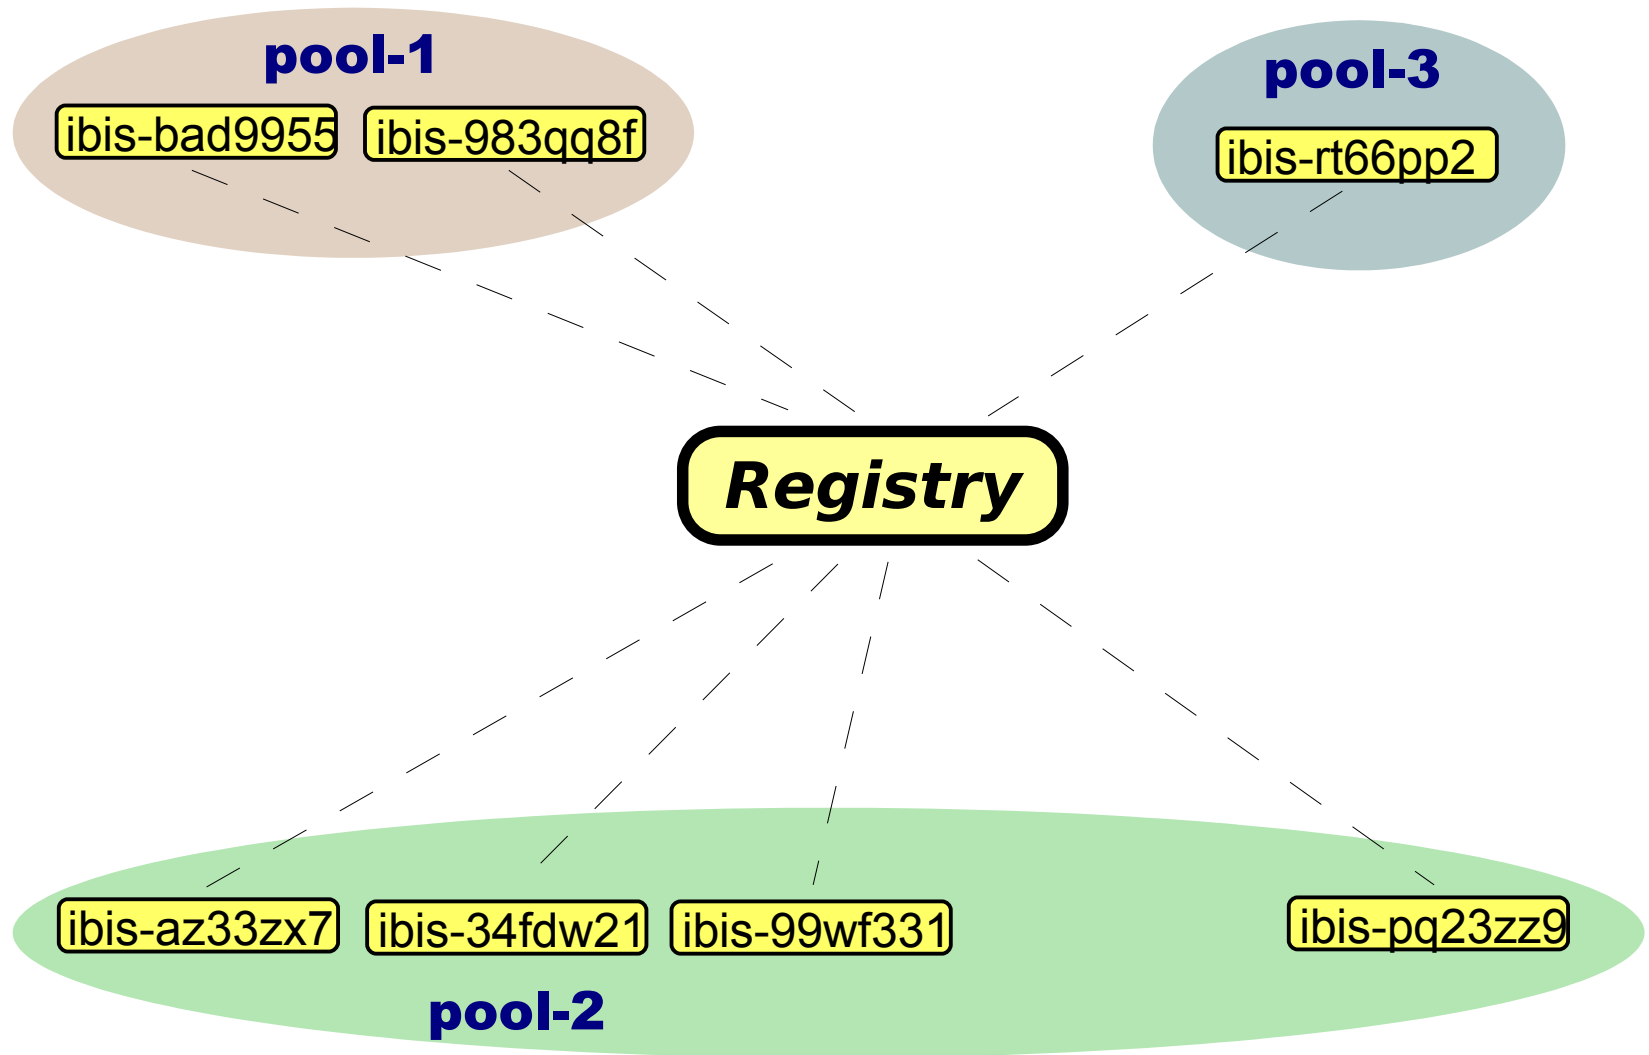
Pools & Malleability



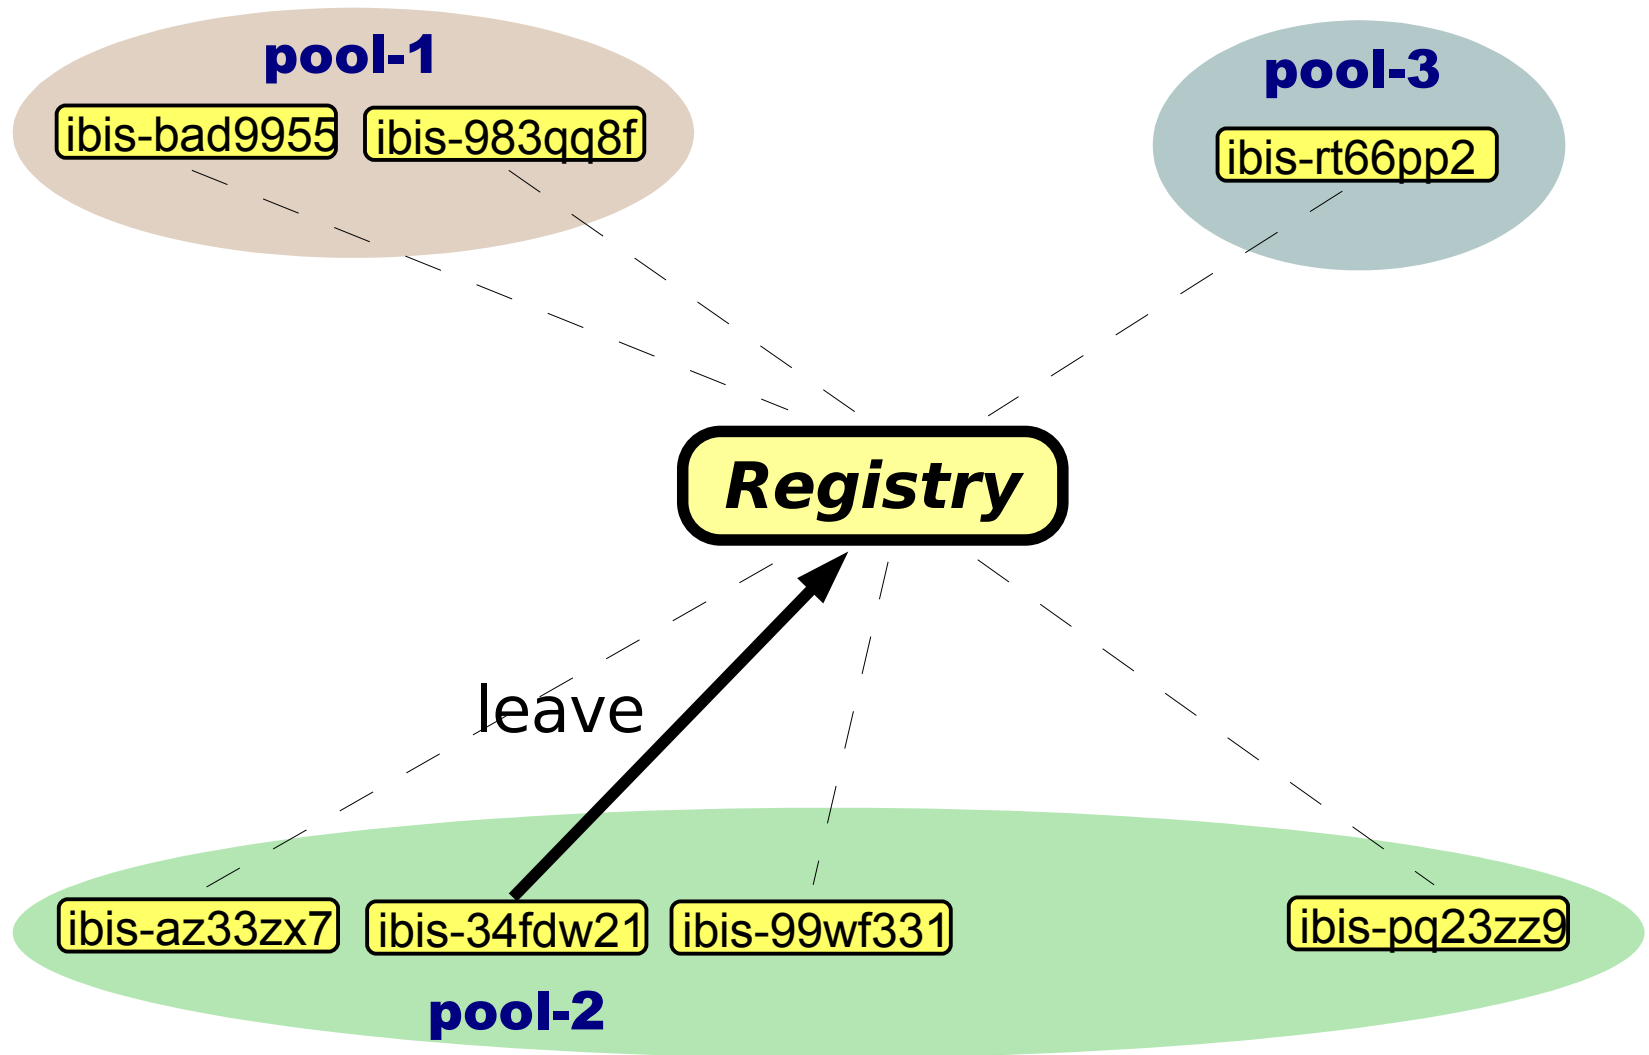
Pools & Malleability



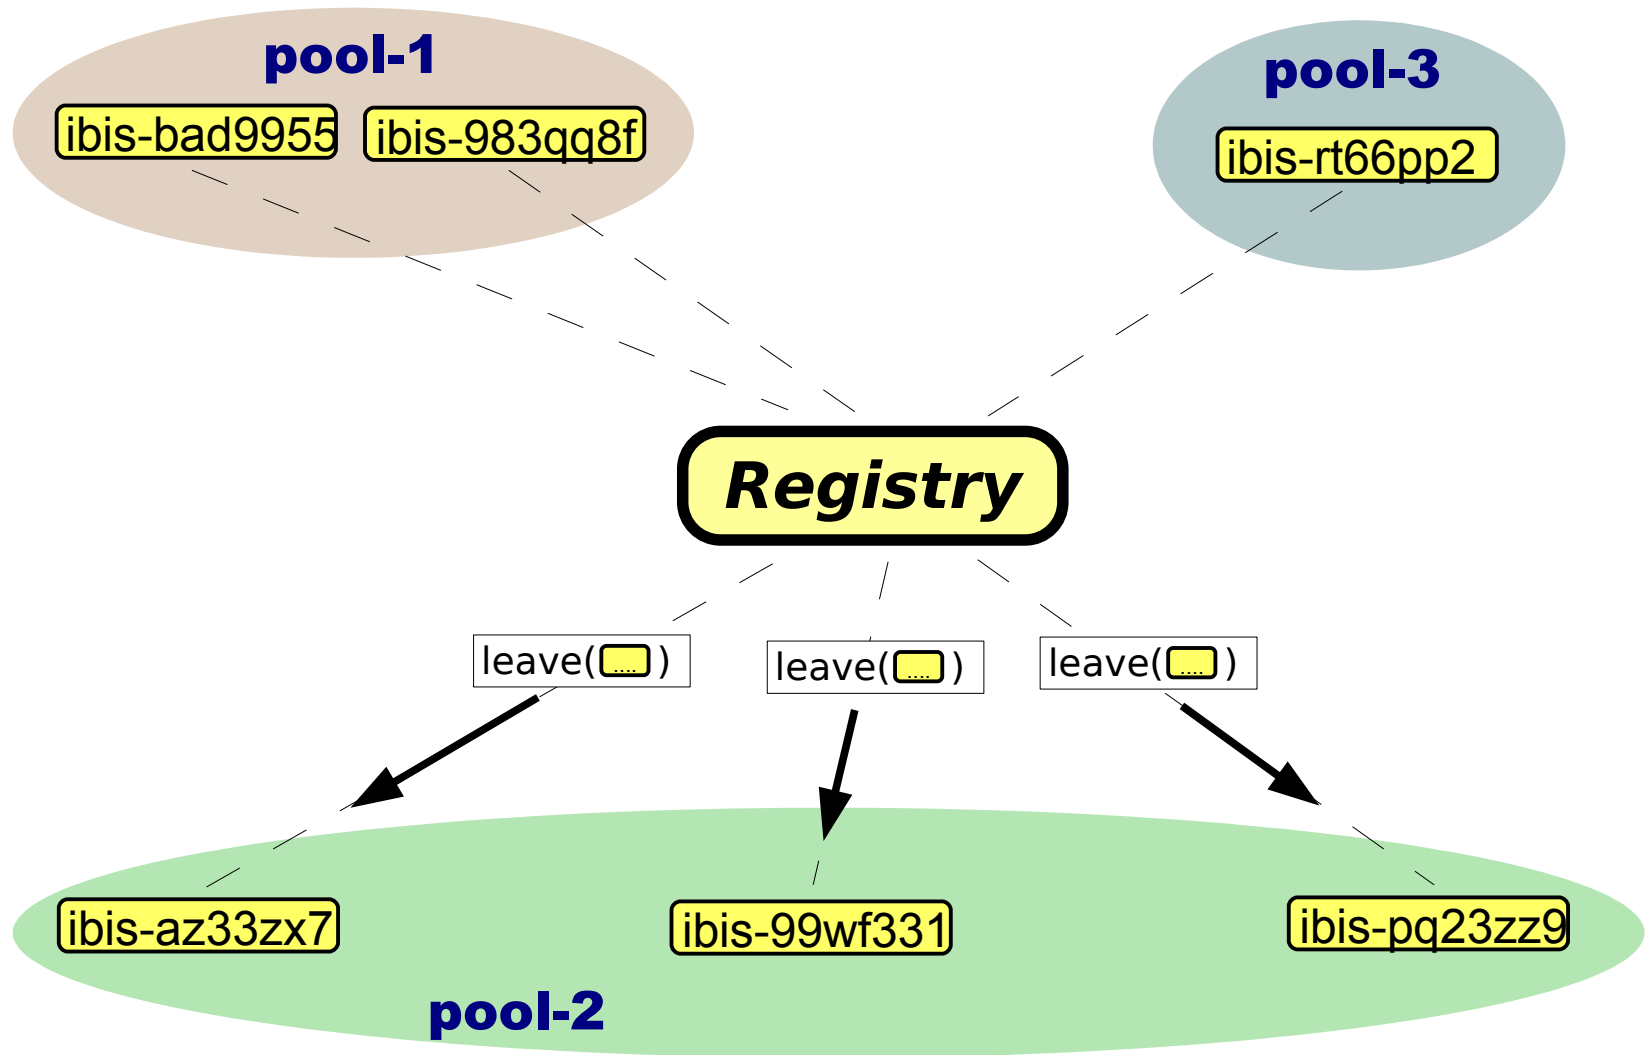
Pools & Malleability



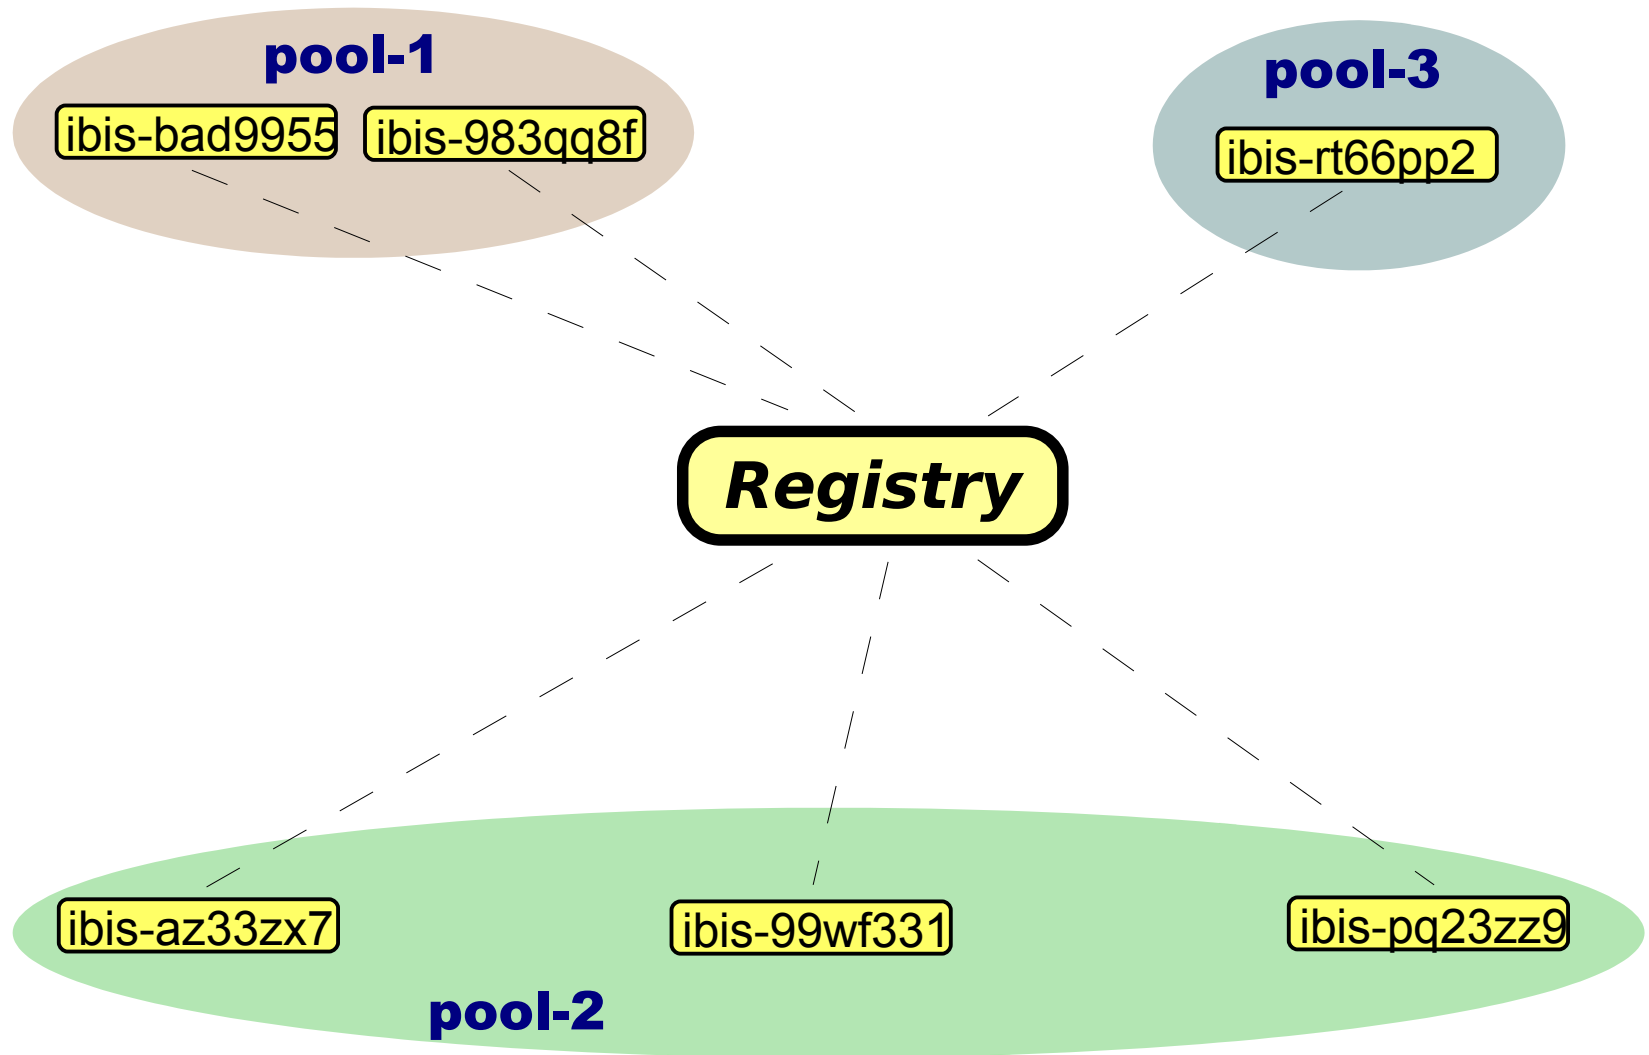
Pools & Malleability



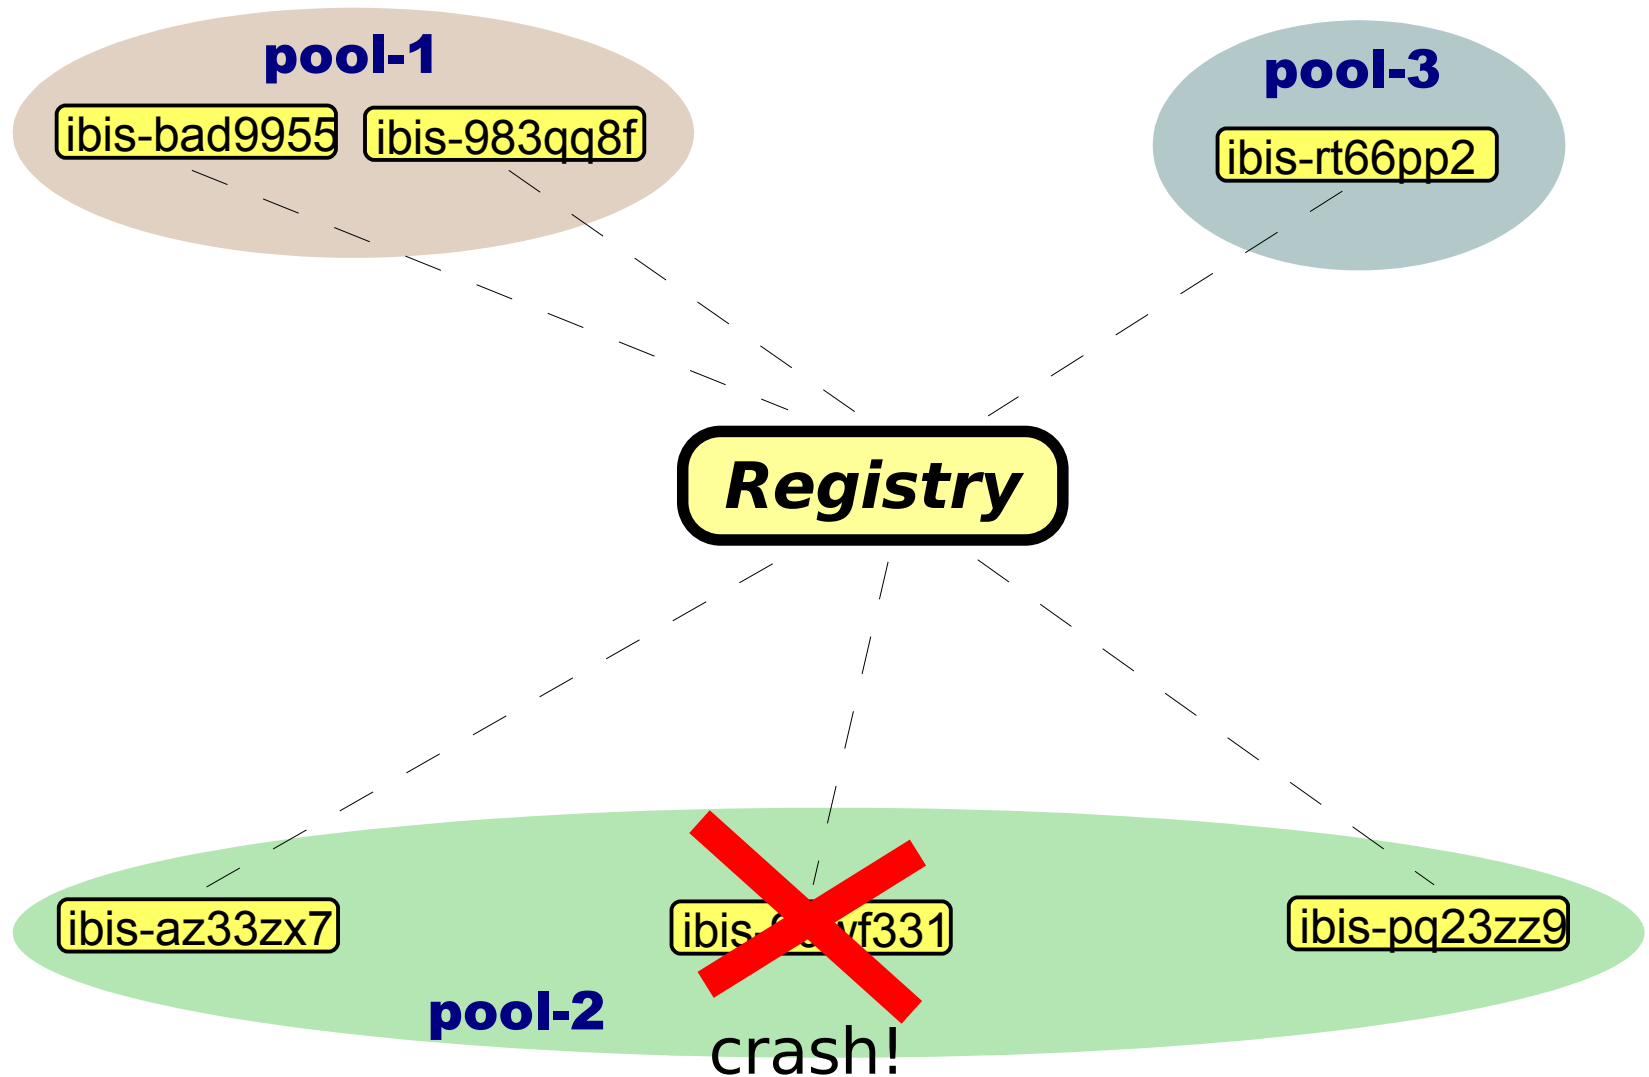
Pools & Malleability



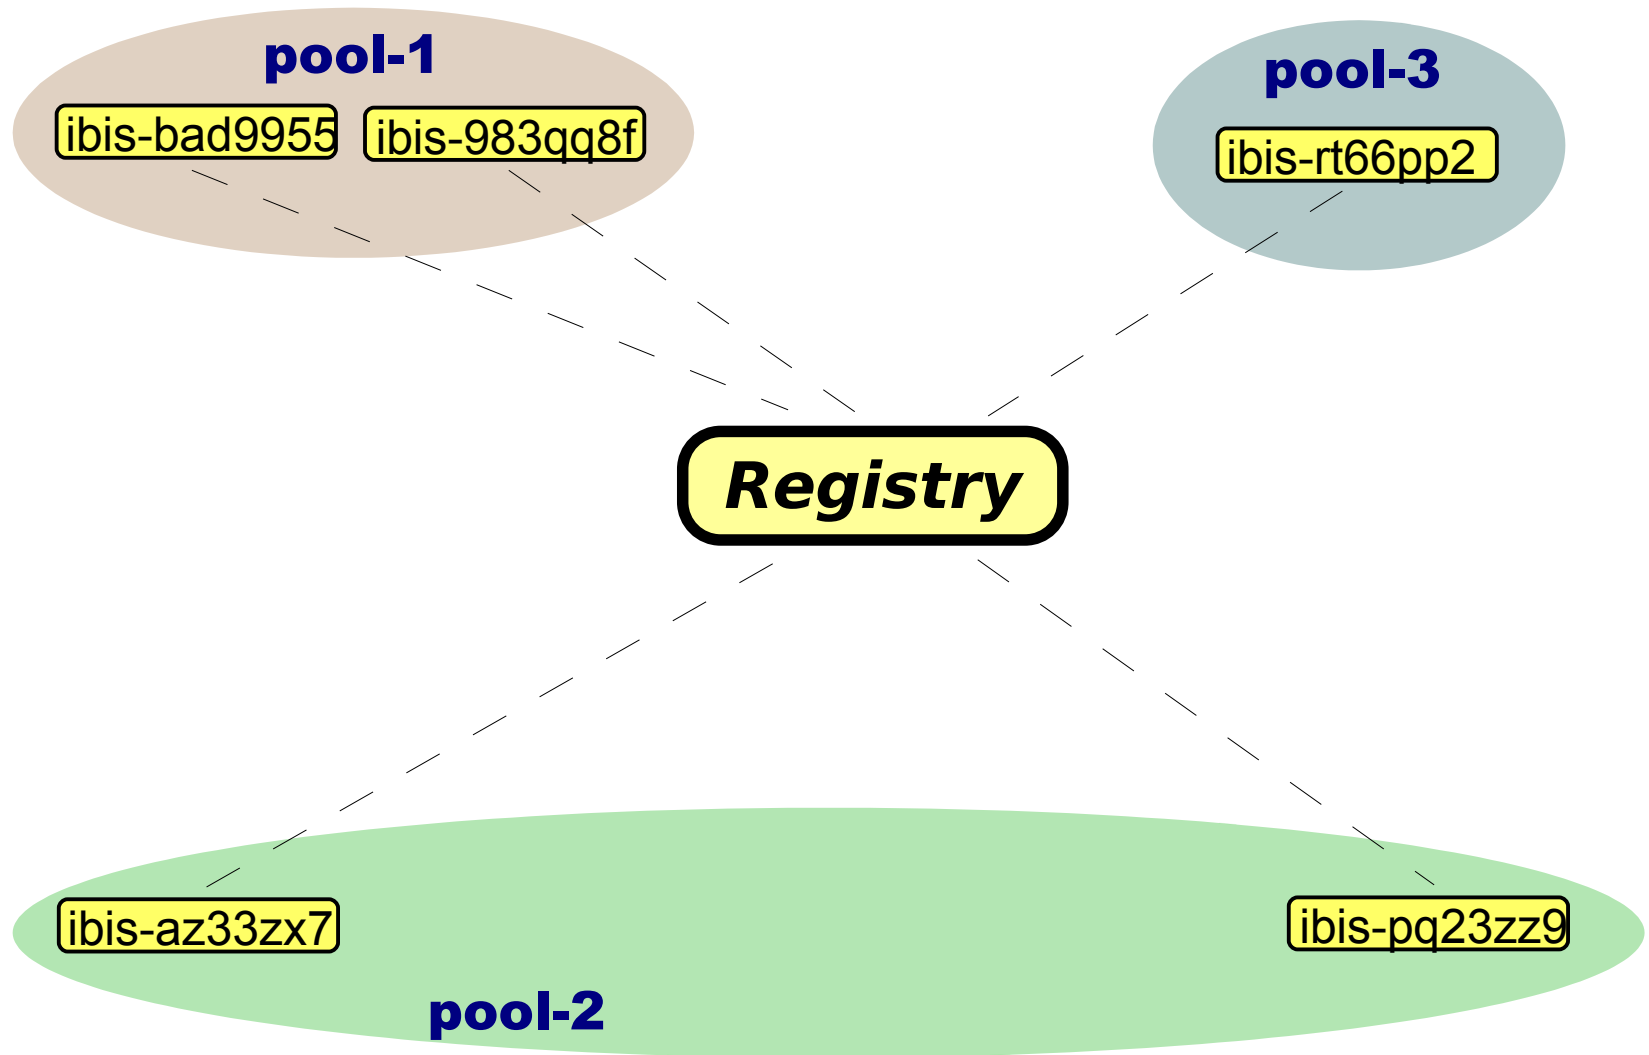
Pools & Malleability



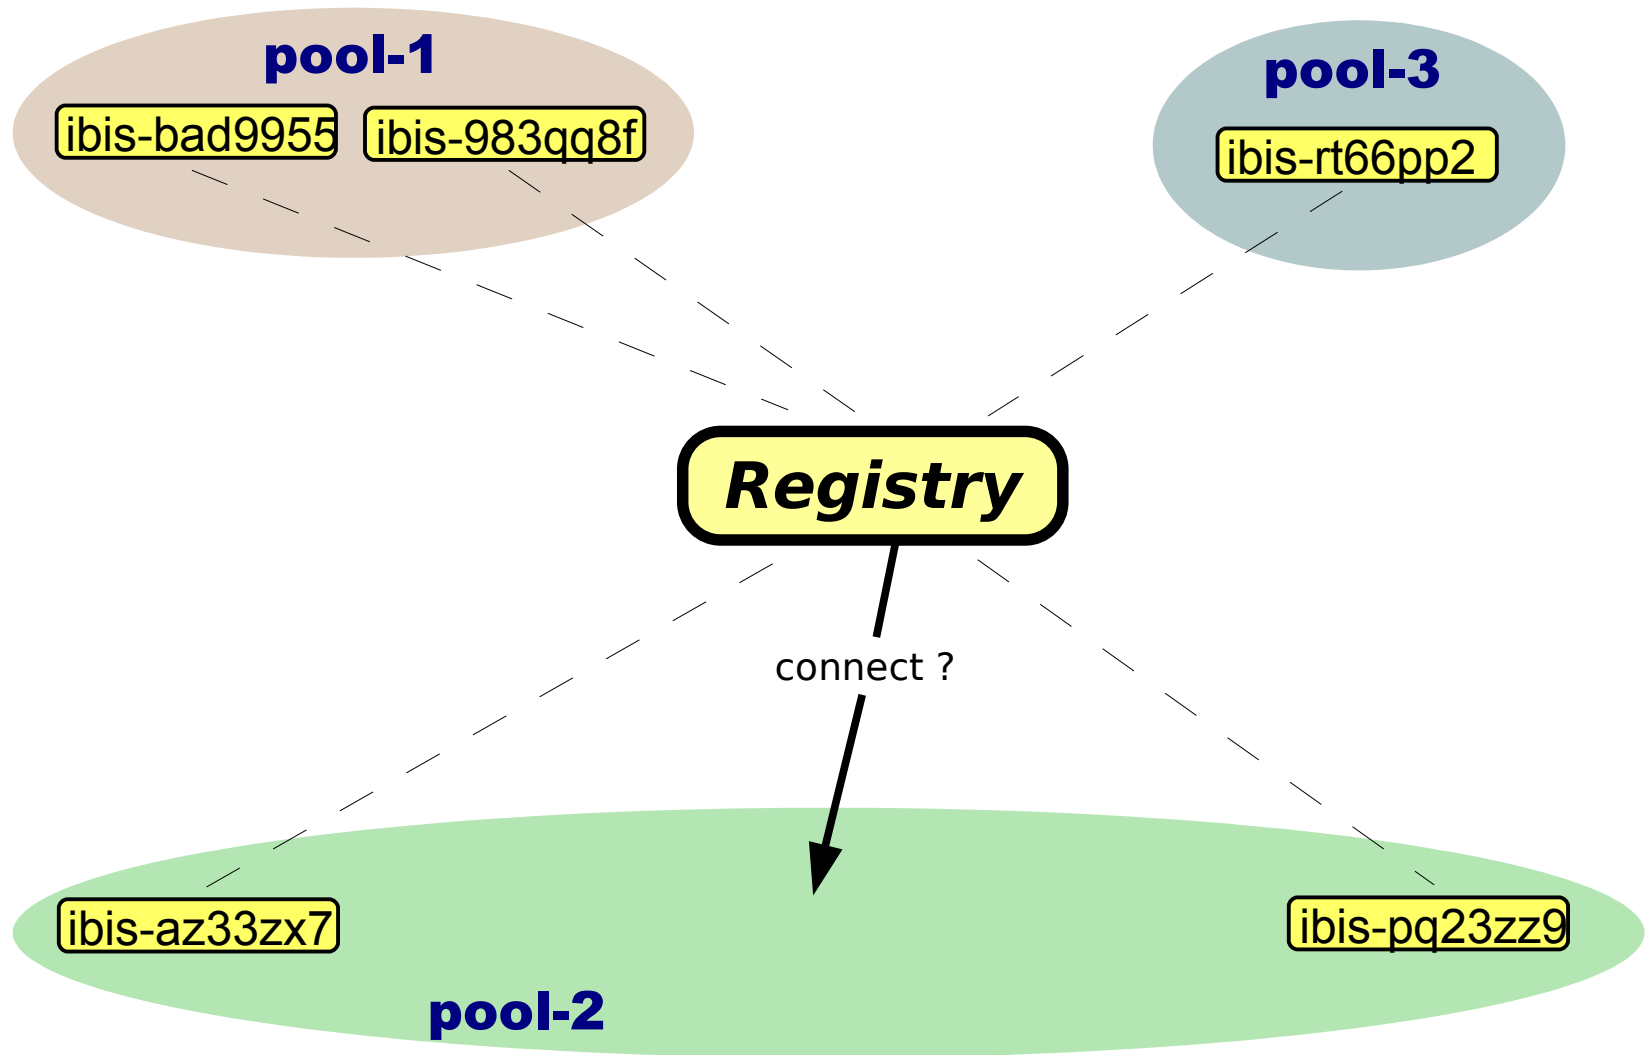
Pools & Malleability



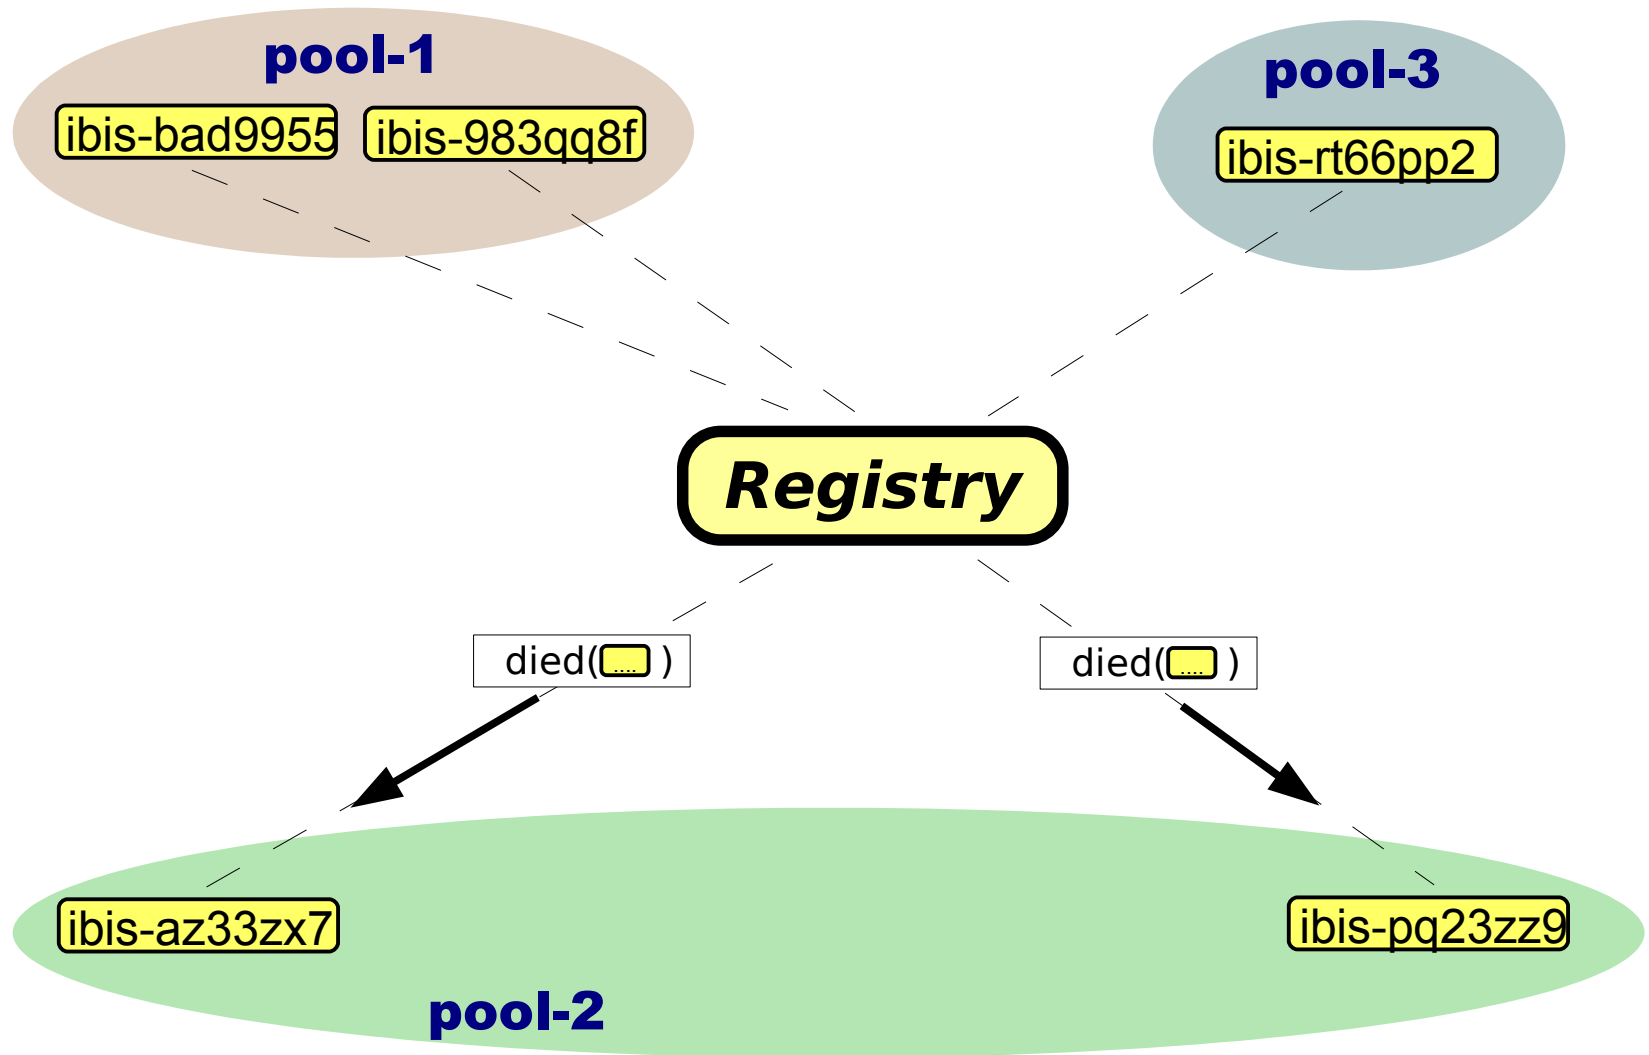
Pools & Malleability




Pools & Malleability



Pools & Malleability



Pools & Malleability



Elections

- Registry offers an 'election' mechanism
 - Allows a group of Ibisses to determine who's in charge
- Each election
 - Has a name (String)
 - Produces IbisIdentifier of the winner
 - Is not democratic
 - You can also be 'an observer'



Pools & Malleability

- This is just one example of a registry
 - Centralized implementation
 - Other implementations exist
 - none of them are interesting at the moment
- Ongoing research
 - scalability issues
 - peer-to-peer techniques
 - distributed election mechanisms

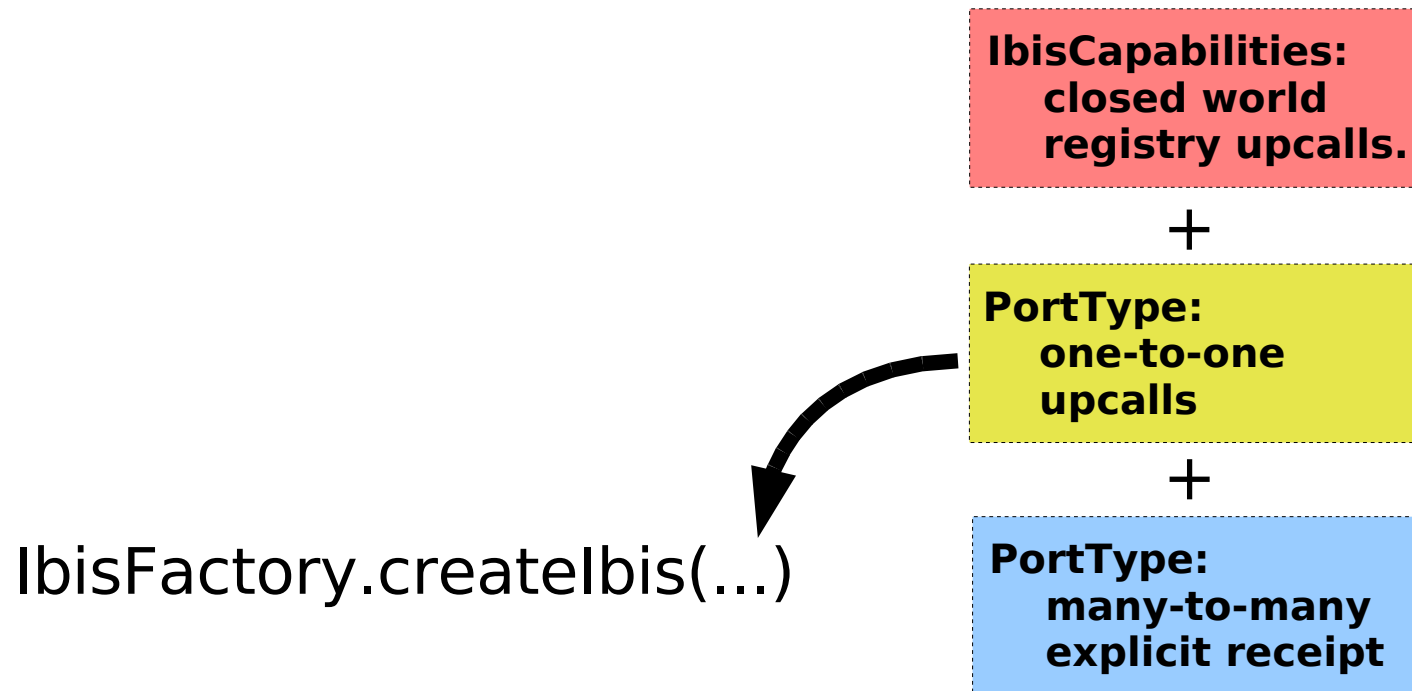


Creating an Ibis

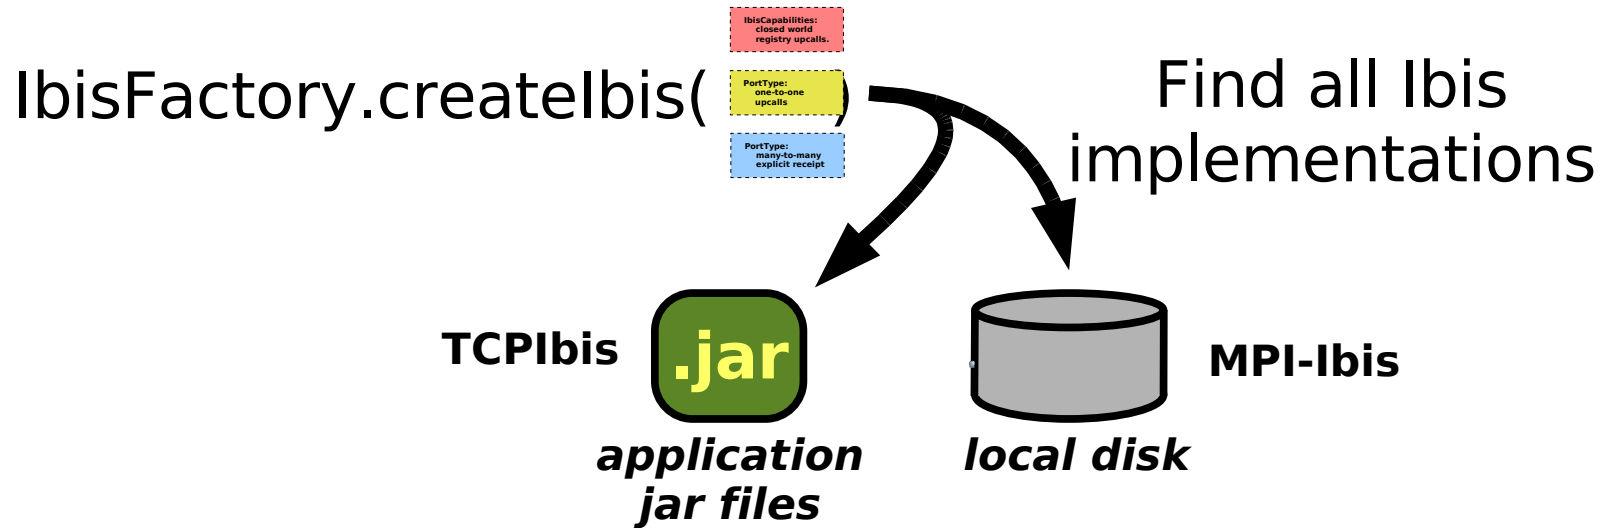
- First step in application
 - IPL is only abstract classes & interfaces
- Ibis selects implementation for you
 - Multiple may be available
 - Selected on the basis of required capabilities and port types
 - Specify the needs of the application



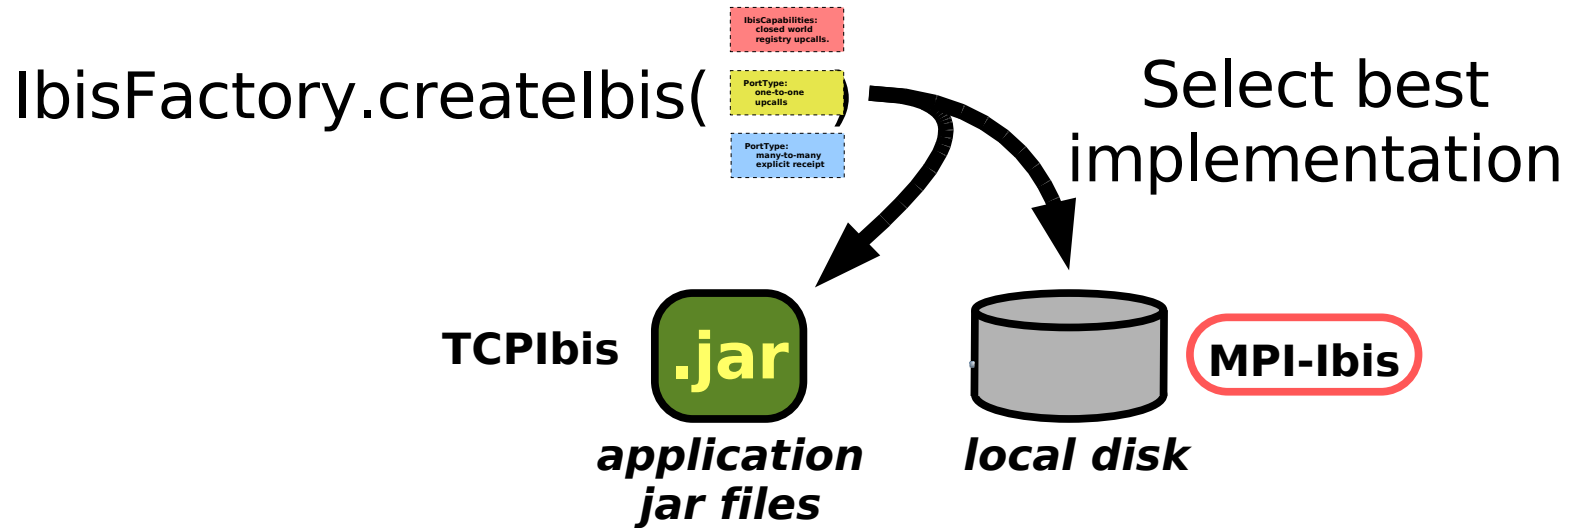
Selecting an Ibis



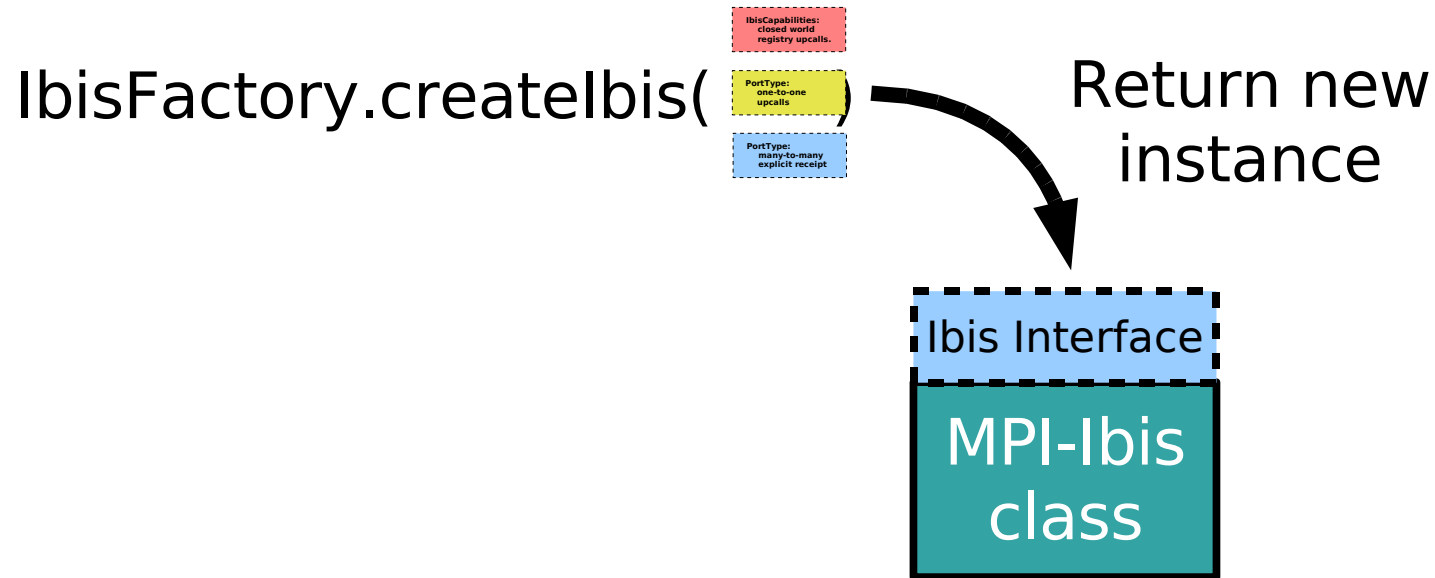
Selecting an Ibis



Selecting an Ibis



Selecting an Ibis



Capabilities

- Similar to Java properties
 - Set of boolean properties
 - “serialization.object”
 - “communication.reliable”
 - “connection.onetoone”
- Act as switches
 - Select which features in the API are required by the application



Capabilities

- Extensible
 - Introduce features without IPL changes
 - Just add more capabilities
 - Allows impl. specific capabilities
- Pitfalls
 - No compile time checks (only runtime)
 - Just strings
 - Sensitive to typos



PortType Capabilities

Capability	Description
connection.onetoone	Unicast
connection.onetomany	Multicast
connection.manytoone	Many to one
connection.manytomany	Multicast + many to one
communication.reliable	Reliable messages
communication.fifo	Fifo ordered messages
communication.numbered	Numbered messages
receive.explicit	Explicit receipt
receive.autoupcalls	Callback on receipt
receive.pollupcalls	Callback on receipt (polling required)
serialization.byte	Only send (arrays of) bytes
serialization.data	Only send (arrays of) primitive types
serialization.object	Send objects, don't care how
serialization.object.ibis	Send objects using Ibis serialization
serialization.object.sun	Send objects using standard ser.



Ibis Capabilities

Capability	Description
registry.elections	Support elections
registry.worldmodel.closed	Fixed set of machines (fixed pool)
registry.membership	Support membership updates
...	
ibis.malleable	Malleability support



Code example

```
package demo.ipl;

import ibis.ipl.*;

public class Example {

    public static void main(String args[]) throws Exception {

        // Step 1: create ibis
        PortType type = new PortType(
            PortType.COMMUNICATION_RELIABLE,
            PortType.SERIALIZATION_OBJECT,
            PortType.RECEIVE_EXPLICIT,
            PortType.CONNECTION_ONE_TO_ONE);

        IbisCapabilities cap = new IbisCapabilities(
            IbisCapabilities.ELECTIONS,
            IbisCapabilities.MALLEABLE);

        Ibis ibis = IbisFactory.createIbis(cap, null, null, type);

        // Step 2: elect server
        Registry reg = ibis.registry();
        IbisIdentifier server = reg.elect("Server");
        boolean amServer = server.equals(ibis.identifier());

        if (amServer) {
            // Step 3, create port
            ReceivePort rp = ibis.createReceivePort(type, "server");
            rp.enableConnections();

            // Step 4, receive message and read data
            ReadMessage rm = rp.receive();
            String tmp = (String) rm.readObject();
            rm.finish();

            System.out.println("Client says: " + tmp);

            // Step 5, close port
            rp.close();
        } else {
            // Step 3, create send port and connect
            SendPort sp = ibis.createSendPort(type);
            sp.connect(server, "server");

            // Step 4, get message and write data
            WriteMessage wm = sp.newMessage();
            wm.writeObject("Hello World");
            wm.finish();

            // Step 5, close port
            sp.close();
        }

        // Step 6, clean up
        ibis.end();
    }
}
```



Code example

```
package demo.ipl;  
  
import ibis.ipl.*;  
  
public class Example {  
  
    public static void main(String args[]) throws Exception {
```

```
        // Step 1: create ibis  
        PortType type = new PortType(  
            PortType.COMMUNICATION_RELIABLE,  
            PortType.SERIALIZATION_OBJECT,  
            PortType.RECEIVE_EXPLICIT,  
            PortType.CONNECTION_ONE_TO_ONE);  
  
        IbisCapabilities cap = new IbisCapabilities(  
            IbisCapabilities.ELECTIONS,  
            IbisCapabilities.MALLEABLE);  
  
        Ibis ibis = IbisFactory.createIbis(cap, null, null, type);
```

```
        // Step 2: select server
```

// Step 1: create ibis

```
PortType type = new PortType(  
    PortType.COMMUNICATION_RELIABLE,  
    PortType.SERIALIZATION_OBJECT,  
    PortType.RECEIVE_EXPLICIT,  
    PortType.CONNECTION_ONE_TO_ONE);
```

```
IbisCapabilities cap = new IbisCapabilities(  
    IbisCapabilities.ELECTIONS,  
    IbisCapabilities.MALLEABLE);
```

```
Ibis ibis = IbisFactory.createIbis(cap, null, null, type);
```

```
        server");  
        .identifier());  
  
        eReceivePort(type, "server");  
  
        and read data  
        ();  
        adObject();  
  
        says: " + tmp);  
  
        and connect  
        ndPort(type);  
        );  
  
        write data  
        sage();  
        ");
```



Code example

```
package demo.ipl;

import ibis.ipl.*;

public class Example {

    public static void main(String args[]) throws Exception {

        // Step 1: create ibis
        PortType type = new PortType(
            PortType.COMMUNICATION_RELIABLE,
            PortType.SERIALIZATION_OBJECT,
            PortType.RECEIVE_EXPLICIT,
            PortType.CONNECTION_ONE_TO_ONE);

        IbisCapabilities cap = new IbisCapabilities(
            IbisCapabilities.ELECTIONS,
            IbisCapabilities.MALLEABLE);

        Ibis ibis = IbisFactory.createIbis(cap, null, null, type);

        // Step 2: elect server
        Registry reg = ibis.registry();
        IbisIdentifier server = reg.elect("Server");
        boolean amServer = server.equals(ibis.identifier());

        if (amServer) {
            // Step 3: create port
            SendPort sp = ibis.createSendPort(type);
            sp.connect(server, "server");

            // Step 4: get message and write data
            WriteMessage wm = sp.newMessage();
            wm.writeObject("Hello World");
            wm.finish();

            // Step 5: close port
            sp.close();
        }

        // Step 6: clean up
        ibis.end();
    }
}
```

// Step 2: elect server

Registry reg = ibis.registry();

IbisIdentifier server = reg.elect("Server");

boolean amServer = server.equals(ibis.identifier());



Code example

```
package demo.ipl;

import ibis.ipl.*;

public class Example {

    public static void main(String args[]) throws Exception {

        // Step 1: create ibis
        PortType type = new PortType(
            PortType.COMMUNICATION_RELIABLE,
            PortType.SERIALIZATION_OBJECT,
            PortType.RECEIVE_EXPLICIT,
            PortType.CONNECTION_ONE_TO_ONE);

        IbisCapabilities cap = new IbisCapabilities(
            IbisCapabilities.ELECTIONS,
            IbisCapabilities.MALLEABLE);

        Ibis ibis = IbisFactory.createIbis(cap, null, null, type);

        // Step 2: elect server
        Registry reg = ibis.registry();
        IbisIdentifier server = reg.elect("Server");
        boolean amServer = server.equals(ibis.identifier());

        if (amServer) {
            // Step 3, create port
            ReceivePort rp = ibis.createReceivePort(type, "server");
            rp.enableConnections();

            // Step 4, receive message and read data
            ReadMessage rm = rp.receive();

            WriteMessage wm = sp.newMessage();
            wm.writeObject("Hello World");
            wm.finish();

            // Step 5, close port
            sp.close();
        }

        // Step 6, clean up
        ibis.end();
    }
}
```

```
if (amServer) {
    // Step 3, create port
    ReceivePort rp = ibis.createReceivePort(type, "server");
    rp.enableConnections();
}
```



Code example

```
package demo.ipl;

import ibis.ipl.*;

public class Example {

    public static void main(String args[]) throws Exception {

        // Step 1: create ibis
        PortType type = new PortType(
            PortType.COMMUNICATION_RELIABLE,
            PortType.SERIALIZATION_OBJECT,
            PortType.RECEIVE_EXPLICIT,
            PortType.CONNECTION_ONE_TO_ONE);

        IbisCapabilities cap = new IbisCapabilities(
            IbisCapabilities.ELECTIONS,
            IbisCapabilities.MALLEABLE);

        Ibis ibis = IbisFactory.createIbis(cap, null, null, type);

        // Step 2: elect server
        Registry reg = ibis.registry();
        Server server = reg.elect("Server");
        server.equals(ibis.identifier());

        // Step 3, create send port and connect
        SendPort sp = ibis.createSendPort(type);
        sp.connect(server, "server");

        // Step 4, receive message and read data
        ReceiveMessage rm = rp.receive();
        String tmp = (String) rm.readObject();
        rm.finish();

        System.out.println("Client says: " + tmp);

        // Step 5, close port
        rp.close();

    } else {

        // Step 3, create send port and connect
        SendPort sp = ibis.createSendPort(type);
        sp.connect(server, "server");

        // Step 4, get message and write data
        WriteMessage wm = sp.newMessage();
        wm.writeObject("Hello World");
        wm.finish();

        // Step 5, close port
        sp.close();

    }

    // Step 6, clean up
    ibis.end();
}
```

// Step 3, create send port and connect
SendPort sp = ibis.createSendPort(type);
sp.connect(server, "server");



Code example

```
package demo.ipl;

import ibis.ipl.*;

public class Example {

    public static void main(String args[]) throws Exception {

        // Step 1: create ibis
        PortType type = new PortType(
            PortType.COMMUNICATION_RELIABLE,
            PortType.SERIALIZATION_OBJECT,
            PortType.RECEIVE_EXPLICIT,
            PortType.CONNECTION_ONE_TO_ONE);

        IbisCapabilities cap = new IbisCapabilities(
            IbisCapabilities.ELECTIONS,
            IbisCapabilities.MALLEABLE);

        Ibis ibis = IbisFactory.createIbis(cap, null, null, type);

        // Step 2: elect server
        Registry reg = ibis.registry();
        server = reg.elect("Server");
        r = server.equals(ibis.identifier());

        // Step 3, create port
        SendPort sp = ibis.createSendPort(type, "server");
        ReceivePort rp = ibis.createReceivePort(type, "server");
        rp.enableConnections();

        // Step 4, receive message and read data
        Message rm = rp.receive();
        String tmp = (String) rm.readObject();
        rm.finish();

        System.out.println("Client says: " + tmp);

        // Step 5, close port
        rp.close();
    } else {
        // Step 3, create send port and connect
        SendPort sp = ibis.createSendPort(type);
        sp.connect(server, "server");

        // Step 4, get message and write data
        WriteMessage wm = sp.newMessage();
        wm.writeObject("Hello World");
        wm.finish();

        // Step 5, close port
        sp.close();
    }

    // Step 6, clean up
    ibis.end();
}
```

// Step 4, get message and write data
WriteMessage wm = sp.newMessage();
wm.writeObject("Hello World");
wm.finish();



Code example

```
package demo.ipl;

import ibis.ipl.*;

public class Example {

    public static void main(String args[]) throws Exception {

        // Step 1: create ibis
        PortType type = new PortType(
            PortType.COMMUNICATION_RELIABLE,
            PortType.SERIALIZATION_OBJECT,
            PortType.RECEIVE_EXPLICIT,
            PortType.CONNECTION_ONE_TO_ONE);

        IbisCapabilities cap = new IbisCapabilities(
            IbisCapabilities.ELECTIONS,
            IbisCapabilities.MALLEABLE);

        Ibis ibis = IbisFactory.createIbis(cap, null, null, type);

        // Step 2: elect server
        Registry reg = ibis.registry();
        IbisIdentifier server = reg.elect("Server");
        boolean amServer = server.equals(ibis.identifier());

        if (amServer) {
            // Step 3, create port
            ReceivePort rp = ibis.createReceivePort(type, "server");
            rp.enableConnections();

            // Step 4, receive message and read data
            ReadMessage rm = rp.receive();
            String tmp = (String) rm.readObject();
            rm.finish();

            System.out.println("Server says: " + tmp);

            // Step 5, send message and write data
            SendPort sp = ibis.createSendPort(type, "server");
            sp.enableConnections();

            // Step 6, send message and write data
            WriteMessage wm = sp.writeObject("World");
            wm.finish();
        } else {
            // Step 7, receive message and read data
            ReadMessage rm = rp.receive();
            String tmp = (String) rm.readObject();
            rm.finish();

            System.out.println("Client says: " + tmp);
        }
    }
}
```



Code example

```
package demo.ipl;

import ibis.ipl.*;

public class Example {

    public static void main(String args[]) throws Exception {

        // Step 1: create ibis
        PortType type = new PortType(
            PortType.COMMUNICATION_RELIABLE,
            PortType.SERIALIZATION_OBJECT,
            PortType.RECEIVE_EXPLICIT,
            PortType.CONNECTION_ONE_TO_ONE);

        IbisCapabilities cap = new IbisCapabilities(
            IbisCapabilities.ELECTIONS,
            IbisCapabilities.MALLEABLE);

        Ibis ibis = IbisFactory.createIbis(cap, null, null, type);

        // Step 2: elect server
        Registry reg = ibis.registry();
        IbisIdentifier server = reg.elect("server");
        boolean amServer = server.equals(ibis);

        if (amServer) {
            // Step 3, create port
            ReceivePort rp = ibis.createReceivePort("rp");
            rp.enableConnections();

            // Step 4, receive message and read data
            ReadMessage rm = rp.receive();
            String tmp = (String) rm.readObject();
            rm.finish();

            System.out.println("Client says: " + tmp);

            // Step 5, close port
            rp.close();
        } else {
            // Step 3, create send port
            SendPort sp = ibis.createSendPort("sp");
            sp.connect(server, "server");

            // Step 4, get message and write
            WriteMessage wm = sp.newMessage();
            wm.writeObject("Hello World");
            wm.finish();

            // Step 5, close port
            sp.close();
        }

        // Step 6, clean up
        ibis.end();
    }
}
```

//Step 5, close port
rp.close();

//Step 5, close port
sp.close();



Code example

```
package demo.ipl;

import ibis.ipl.*;

public class Example {

    public static void main(String args[]) throws Exception {

        // Step 1: create ibis
        PortType type = new PortType(
            PortType.COMMUNICATION_RELIABLE,
            PortType.SERIALIZATION_OBJECT,
            PortType.RECEIVE_EXPLICIT,
            PortType.CONNECTION_ONE_TO_ONE);

        IbisCapabilities cap = new IbisCapabilities(
            IbisCapabilities.ELECTIONS,
            IbisCapabilities.MALLEABLE);

        Ibis ibis = IbisFactory.createIbis(cap, null, null, type);

        // Step 2: elect server
        Registry reg = ibis.registry();
        IbisIdentifier server = reg.elect("Server");
        boolean amServer = server.equals(ibis.identifier());

        if (amServer) {
            // Step 3, create port
            ReceivePort rp = ibis.createReceivePort(type, "server");
            rp.enableConnections();

            // Step 4, receive message and read data
            ReadMessage rm = rp.receive();
            String tmp = (String) rm.readObject();
            rm.finish();

            System.out.println("Client says: " + tmp);

            // Step 5, close port
            rp.close();
        } else {
            // Step 3, create send port and connect
            SendPort sp = ibis.createSendPort(type);
            sp.connect(server, "server");

            // Step 4, get message and write data
            WriteMessage wm = sp.newMessage();
            wm.writeObject("Hello World");
            wm.finish();

            // Step 5, close port
            sp.close();
        }

        // Step 6, clean up
        ibis.end();
    }
}
```

**// Step, clean up
ibis.end();**



Code example

- Live demo

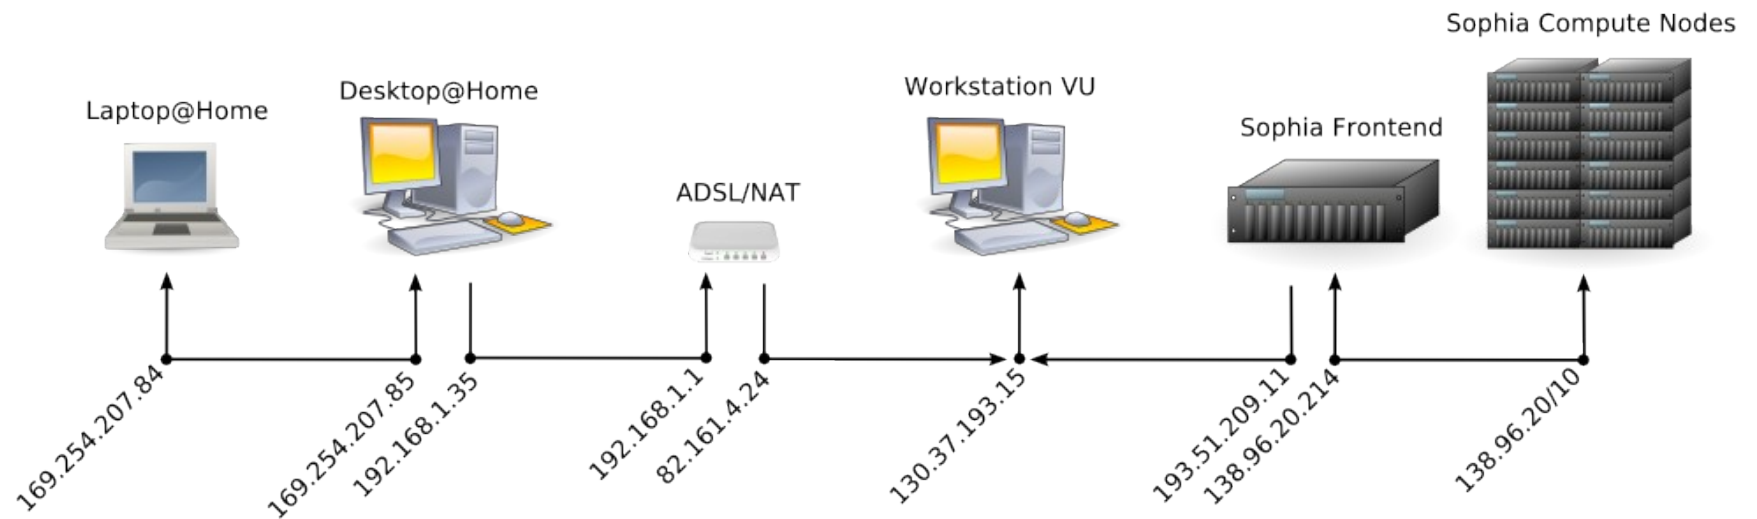



Other cool features

- Ibis can handle 'complicated network setups' using SmartSockets
- Real-world examples:
 - Multiple networks
 - Mix local and global IPs depending on target
 - NAT / firewalls
 - Using UPnP, STUN, TCP splicing, etc.
 - Routing messages through external points
 - Non-routed networks
 - Routing messages through external points



Example



Programming models

- Remote Method Invocation (RMI)
- Group Method Invocation (GMI)
- Satin (Divide & Conquer)
- MPJ (MPI Java 'standard')
- Others are being developed
 - Balutek (data parallel)
 - Replicated Method Invocation (RepMI)



Satin

- Parallel Divide-and-conquer
 - Divide work into independent parts
 - Spawn sub-jobs
 - Combine sub-results
 - Repeat recursively
- Master-Worker is a subset of this
 - Only one level of recursion
- Targeted at the grid (and clusters)



Sequential Fibonacci

```
public long fib(int n) {  
    if (n < 2) return n;  
  
    long x = fib(n - 1);  
    long y = fib(n - 2);  
  
    return x + y;  
}
```



Parallel Fibonacci

```
interface FibInterface extends ibis.satin.Spawnable {  
    public long fib(int n);  
}
```

```
public long fib(int n) {  
    if (n < 2) return n;  
  
    long x = fib(n - 1);  
    long y = fib(n - 2);  
    sync();  
    return x + y;  
}
```



Parallel Fibonacci

```
interface FibInterface extends ibis.satin.Spawnable {  
    public long fib(int n);  
}
```

```
public long fib(int n) {  
    if (n < 2) return n;  
  
    long x = fib(n - 1);  
    long y = fib(n - 2);  
    sync();  
    return x + y;  
}
```

Mark methods as
Spawnable.
They can run in parallel.



Parallel Fibonacci

```
interface FibInterface extends ibis.satin.Spawnable {  
    public long fib(int n);  
}
```

```
public long fib(int n) {  
    if (n < 2) return n;
```

```
    long x = fib(n - 1);
```

```
    long y = fib(n - 2);
```

```
    sync();
```

```
    return x + y;
```

```
}
```

Mark methods as

Spawnable.

They can run in parallel.

Wait until spawned
methods are done.



Satin features

- Satin distributes jobs across machines
- Load-balancing is done automatically
 - Uses random stealing
 - Algorithm has been proven to be optimal on homogeneous systems
 - Additional highly-efficient grid-aware algorithms



Satin features

- Malleability
 - Add/remove machines on the fly
- Fault-tolerance
 - When a machine leave suddenly (crashes) the others continue the computation and automatically recompute the lost work
- Shared Objects (added recently)
 - Allows machines to share 'global data'



Satin Applications

- More interesting applications
 - Satisfiability solver
 - Gene sequencing
 - N-body simulations
 - Grammar-based text analysis
 - Game-tree search
 - Raytracing
 - Numerical functions
 - ...




Satin Applications

- More interesting applications
 - Satisfiability solver
 - Gene sequencing
 - N-body simulations ← Demo!
 - Grammar-based text analysis
 - Game-tree search
 - Raytracing
 - Numerical functions
 - ...




Ibis RMI

- Replacement for Sun RMI
 - Has the same interface
 - Used different stub compiler (rmic)
 - Generates Ibis specific stubs/skeletons



Ibis RMI

- Replacement for Sun RMI
 - Has the same interface
 - Used different stub compiler (rmic)
 - Generates Ibis specific stubs/skeletons



Ibis RMI

- Not interoperable with Sun RMI
 - uses a different protocol
- No socket factories
 - Ibis doesn't have to use sockets!
- No activatable objects




GMI

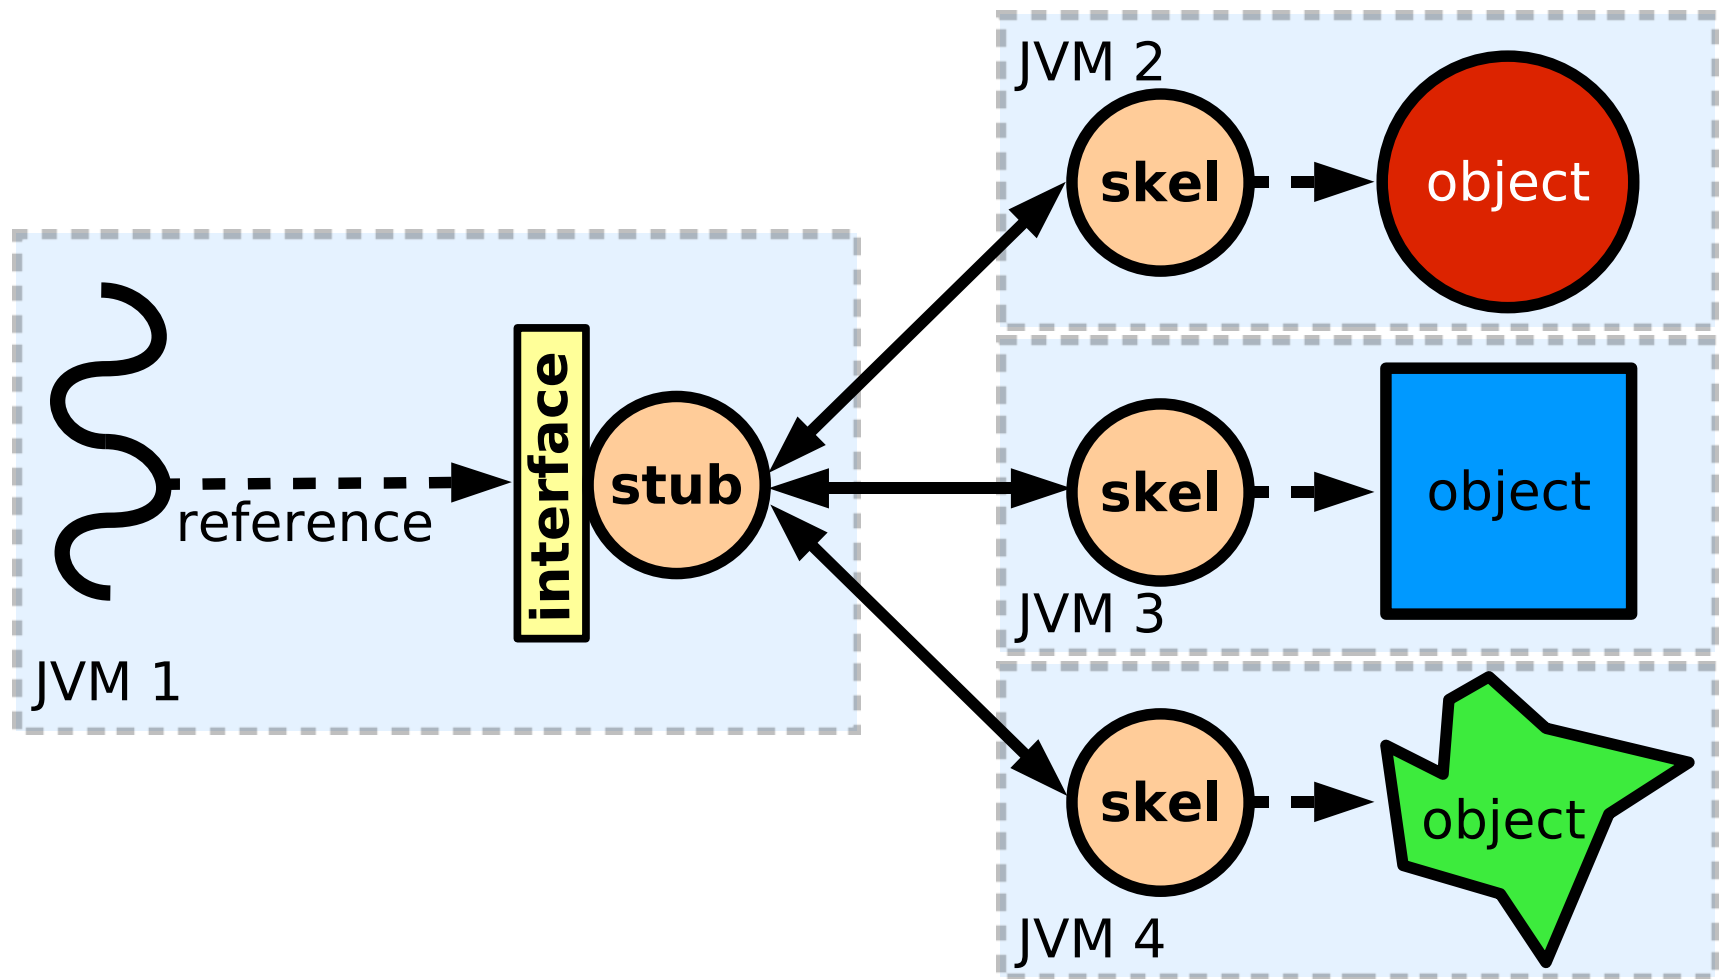
- Generalized RMI model
 - Allows communication with groups
 - A single stub refers to an entire group
 - Allows more 'advanced' communication
 - By offering different ways of forwarding a method invocation and handling the reply



GMI Example



GMI Implementation



Group operations

- The group reference can be configured
 - How is a method invocation handled
 - How is the method result handled
 - Configuration per method
- Implemented by selecting different communication code in the generated stubs and skeletons




Invocation Schemes


- **Single**
 - Forward to 1 object in group
- **Group**
 - Forward to all objects in group
- **Personalized**
 - Forward to all objects, but personalize parameters for each target
- **Combined**
 - Combine several invocation into one, then forward to the group using one of the above




Single




Single




Group




Group







Personalized




Personalized




Combined



Combined



Combined




Reply handling schemes


- **Discard**
- **Return**
- **Forward**
 - Reply is forwarded to a separate object
- **Combine**
 - Multiple replies are combined into one
- **Personalize**
 - A personalized result is returned to each participant of a combined invocation




Result Combining




Result Combining




Result Combining




Result Combining



Result Combining



Result Combining



GMI Communication

<i>Operation</i>	<i>Invocation</i>	<i>Reply</i>
RMI	Single	Return
Async. RMI	Single	Discard
Future	Single	Forward
Broadcast	Group	Discard
Scatter	Personalized	Discard
Reduce results	Group	Combine (binomial)
Gather results	Group	Combine (flat)
Reduce inv.	Combine + Single	Discard
Gather inv.	Combine + Single	Discard



Code example

```
public interface i_SimpleGroup extends GroupInterface {
    void ping();
}

public class SimpleGroup extends GroupMember implements i_SimpleGroup {

    public SimpleGroup() {
        super();
    }

    public void ping() {
        System.out.println("ping");
    }
}

public class MulticastNoReply {
    public static void main(String[] args) throws Exception {

        int rank = Group.rank();
        int size = Group.size();

        // Create the group
        if (rank == 0) {
            Group.create("GroupNoReply", i_SimpleGroup.class, size);
        }

        // Everyone adds an object
        SimpleGroup s = new SimpleGroup();
        Group.join("GroupNoReply", s);

        if (rank == 0) {
            // Perform lookup to get group reference
            i_SimpleGroup g = (i_SimpleGroup) Group.lookup("GroupNoReply");

            // Configure reference to perform group invocation
            GroupMethod m = Group.findMethod(g, "void ping()");
            m.configure(new GroupInvocation(), new DiscardReply());

            // Perform the invocation
            g.ping();
        }

        // Done
        Group.exit();
    }
}
```



Code example

```
public interface i_SimpleGroup extends GroupInterface {  
    void ping();  
}
```

```
public class SimpleGroup extends GroupMember implements i_SimpleGroup {  
  
    public SimpleGroup() {  
        super();  
    }  
}
```

```
public interface i_SimpleGroup extends GroupInterface {  
    void ping();  
}
```

```
public static void main(String[] args) throws Exception {  
  
    int rank = Group.rank();  
    int size = Group.size();  
  
    // Create the group  
    if (rank == 0) {  
        Group.create("GroupNoReply", i_SimpleGroup.class, size);  
    }  
  
    // Everyone adds an object  
    SimpleGroup s = new SimpleGroup();  
    Group.join("GroupNoReply", s);  
  
    if (rank == 0) {  
        // Perform lookup to get group reference  
        i_SimpleGroup g = (i_SimpleGroup) Group.lookup("GroupNoReply");  
  
        // Configure reference to perform group invocation  
        GroupMethod m = Group.findMethod(g, "void ping()");  
        m.configure(new GroupInvocation(), new DiscardReply());  
  
        // Perform the invocation  
        g.ping();  
    }  
  
    // Done  
    Group.exit();  
}
```



Code example

```
public interface i_SimpleGroup extends GroupInterface {  
    void ping();  
}
```

```
public class SimpleGroup extends GroupMember implements i_SimpleGroup {  
  
    public SimpleGroup() {  
        super();  
    }  
  
    public void ping() {  
        System.out.println("ping");  
    }  
}
```

```
public class MulticastNoReply {  
    public static void main(String[] args) throws Exception {  
  
        int rank = Group.rank();  
        int size = Group.size();
```

```
public class SimpleGroup extends GroupMember implements i_SimpleGroup {
```

```
    public SimpleGroup() {  
        super();  
    }
```

```
    public void ping() {  
        System.out.println("ping");  
    }
```

```
}
```

```
        Group.exit();  
    }  
}
```



Code example

```
public interface i_SimpleGroup extends GroupInterface {
    void ping();
}

public class SimpleGroup extends GroupMember implements i_SimpleGroup {

    public SimpleGroup() {
        super();
    }

    public void ping() {
        System.out.println("ping");
    }
}

public class MulticastNoReply {
    public static void main(String[] args) throws Exception {

        int rank = Group.rank();
        int size = Group.size();

        // Create the group
        if (rank == 0) {
            Group.create("GroupNoReply", i_SimpleGroup.class, size);
        }
    }
}
```

```
public class MulticastNoReply {  
    public static void main(String[] args) throws Exception {
```

```
        int rank = Group.rank();  
        int size = Group.size();
```

```
        // Perform the invocation  
        g.ping();  
    }  
  
    // Done  
    Group.exit();  
}
```



Code example

```
public interface i_SimpleGroup extends GroupInterface {
    void ping();
}

public class SimpleGroup extends GroupMember implements i_SimpleGroup {

    public SimpleGroup() {
        super();
    }

    public void ping() {
        System.out.println("ping");
    }
}

public class MulticastNoReply {
    public static void main(String[] args) throws Exception {

        int rank = Group.rank();
        int size = Group.size();

        // Create the group
        if (rank == 0) {
            Group.create("GroupNoReply", i_SimpleGroup.class, size);
        }

        // Everyone adds an object
        SimpleGroup s = new SimpleGroup();
        Group.join("GroupNoReply", s);

        if (rank == 0) {
            // Perform lookup to get group reference
        }

        // Done
        Group.exit();
    }
}
```

```
// Create the group
if (rank == 0) {
    Group.create("GroupNoReply", i_SimpleGroup.class, size);
}
```



Code example

```
public interface i_SimpleGroup extends GroupInterface {
    void ping();
}

public class SimpleGroup extends GroupMember implements i_SimpleGroup {

    public SimpleGroup() {
        super();
    }

    public void ping() {
        System.out.println("ping");
    }
}

public class MulticastNoReply {
    public static void main(String[] args) throws Exception {

        int rank = Group.rank();
        int size = Group.size();

        // Create the group
        if (rank == 0) {
            Group.create("GroupNoReply", i_SimpleGroup.class, size);
        }

        // Everyone adds an object
        SimpleGroup s = new SimpleGroup();
        Group.join("GroupNoReply", s);

        if (rank == 0) {
            // Perform lookup to get group reference
            i_SimpleGroup g = (i_SimpleGroup) Group.lookup("GroupNoReply");

            // Configur
            GroupMet
            m.configu

            // Perform
            g.ping();
        }

        // Done
        Group.exit();
    }
}
```

```
// Everyone adds an object
SimpleGroup s = new SimpleGroup();
Group.join("GroupNoReply", s);
```



Code example

```
public interface i_SimpleGroup extends GroupInterface {  
    void ping();  
}
```

```
public class SimpleGroup extends GroupMember implements i_SimpleGroup {
```

```
    public SimpleGroup() {
```

```
        if (rank == 0) {  
            // Perform lookup to get group reference  
            i_SimpleGroup g = (i_SimpleGroup) Group.lookup("GroupNoReply");  
  
            // Configure reference to perform group invocation  
            GroupMethod m = Group.findMethod(g, "void ping()");  
            m.configure(new GroupInvocation(), new DiscardReply());  
  
            // Perform the invocation  
            g.ping();  
        }  
    }
```

```
        if (rank == 0) {  
            // Perform lookup to get group reference  
            i_SimpleGroup g = (i_SimpleGroup) Group.lookup("GroupNoReply");  
  
            // Configure reference to perform group invocation  
            GroupMethod m = Group.findMethod(g, "void ping()");  
            m.configure(new GroupInvocation(), new DiscardReply());  
  
            // Perform the invocation  
            g.ping();  
        }  
    }
```

```
        // Done  
        Group.exit();  
    }
```

```
}
```



Code example

```
public interface i_SimpleGroup extends GroupInterface {
    void ping();
}

public class SimpleGroup extends GroupMember implements i_SimpleGroup {

    public SimpleGroup() {
        super();
    }

    public void ping() {
        System.out.println("ping");
    }
}

public class MulticastNoReply {
    public static void main(String[] args) throws Exception {

        int rank = Group.rank();
        int size = Group.size();

        // Create the group
        if (rank == 0) {
            Group.create("GroupNoReply", i_SimpleGroup.class, size);
        }

        // Everyone adds an object
        SimpleGroup s = new SimpleGroup();
        Group.join("GroupNoReply", s);

        if (rank == 0) {
            // Perform lookup to get group reference
            i_SimpleGroup g = (i_SimpleGroup) Group.lookup("GroupNoReply");

            // Configure reference to perform group invocation
            GroupMethod m = Group.findMethod(g, "void ping()");
            m.configure(new GroupInvocation(), new DiscardReply());

            // Perform the invocation
            g.ping();
        }
        // Done
        Group.exit();
    }
}
```

// Done
Group.exit();



Code Example

- Live demo



GMI

- MPI-style programming
- But not necessarily SPMD
 - Can also do client-server style applications
- Does not have Grid optimizations yet
 - Can be done
- Fault-tolerance support is harder
 - Model isn't as 'clean' as Satin



After the Break

- Hands-on session
 - Installing Ibis.
 - Running applications
 - Writing your own applications



END



Function Objects

- Some operations need user defined functions
 - Personalizing a method invocation
 - Combining a result or invocation
 - Forwarding of results
- GMI uses function objects
 - Extend a class from the GMI package



Result Combiners

- Use 'combiner' to merge the results of an invocation
- FlatCombiner
 - Combines all results in one go
 - Similar to 'gather' operation of MPI
- BinomialCombiner
 - Pairwise combines results
 - Similar to 'reduce' operation of MPI



FlatCombiner

```
public class FlatCombiner {  
  
    public boolean combine(boolean[] results, Exception[] ex)  
    public byte combine(byte[] results, Exception[] ex)  
    public char combine(char[] results, Exception[] ex)  
    public short combine(short[] results, Exception[] ex)  
    public int combine(int[] results, Exception[] ex)  
    public long combine(long[] results, Exception[] ex)  
    public float combine(float[] results, Exception[] ex)  
    public double combine(double[] results, Exception[] ex)  
  
    public Object combine(Object[] results, Exception[] ex)  
  
    public void combine(Exception[] exceptions)  
}
```



FlatCombiner

```
public class MyCombiner extends FlatCombiner {  
  
    public int combine(int[] results, Exception[] ex) {  
  
        int sum = 0;  
  
        for (int i=0;i<results.length;i++) {  
            sum += results[i];  
        }  
  
        return sum;  
    }  
}
```



FlatCombiner

// Get a group reference

```
X g = (X) Group.lookup("your group");
```

// Configure reference to perform group invocation,

// and combine the replies using 'MyCombiner'

```
GroupMethod m = Group.findMethod(g, "int get()");
```

```
m.configure(new GroupInvocation(),  
            new CombineReply(new MyCombiner()));
```

// Perform the invocation

```
int result = g.get();
```



FlatCombiner Demo

- Live demo



Overview

- Programming models
 - IPL (bare bones)
 - RMI (remote invocation)
 - GMI (group communication)
 - Satin (divide and conquer)
 - MPJ (MPI to Java binding)
- Hands-on session
 - How to roll your own Ibis applications



Connection setup

- Ibis- and ReceivePortIdentifiers
 - Hide implementation details
 - Independent of
 - IP-addresses
 - Host names
 - Port numbers
 - MPI-ranks, etc...
- Abstract way of addressing machines and connection endpoints



GMI

- Group
 - Contains 1 or more objects
 - Fixed size (set when it is created)
 - All objects must implement the same group interface
 - But objects may have different type!
 - Unique name
 - Used in lookup (produces group reference)
 - Group members have rank
 - Ranks are 'per-group'

