

Ibis: an Efficient Java-based Grid Programming Environment



**Rob V. van Nieuwpoort, Jason Maassen
Rutger Hofman, Thilo Kielmann, Henri E. Bal**

Introduction

- **Goal: distributed supercomputing**
- **Grid programming environments**
 - **Portable, support for heterogeneity**
 - **Flexible (multiple protocols, malleable)**
 - **Convenient programming model(s)**
 - **Efficient**
- **Problem: No single programming environment combines all features**

C/Fortran and MPI on the grid

- **Portable only in the traditional sense**
- **Limited to a SPMD-style programming model**
 - **No RPC (implicit receive)**
 - **No dynamic data structures**
 - **Does not fit in object oriented programming model**
- **But efficient !**

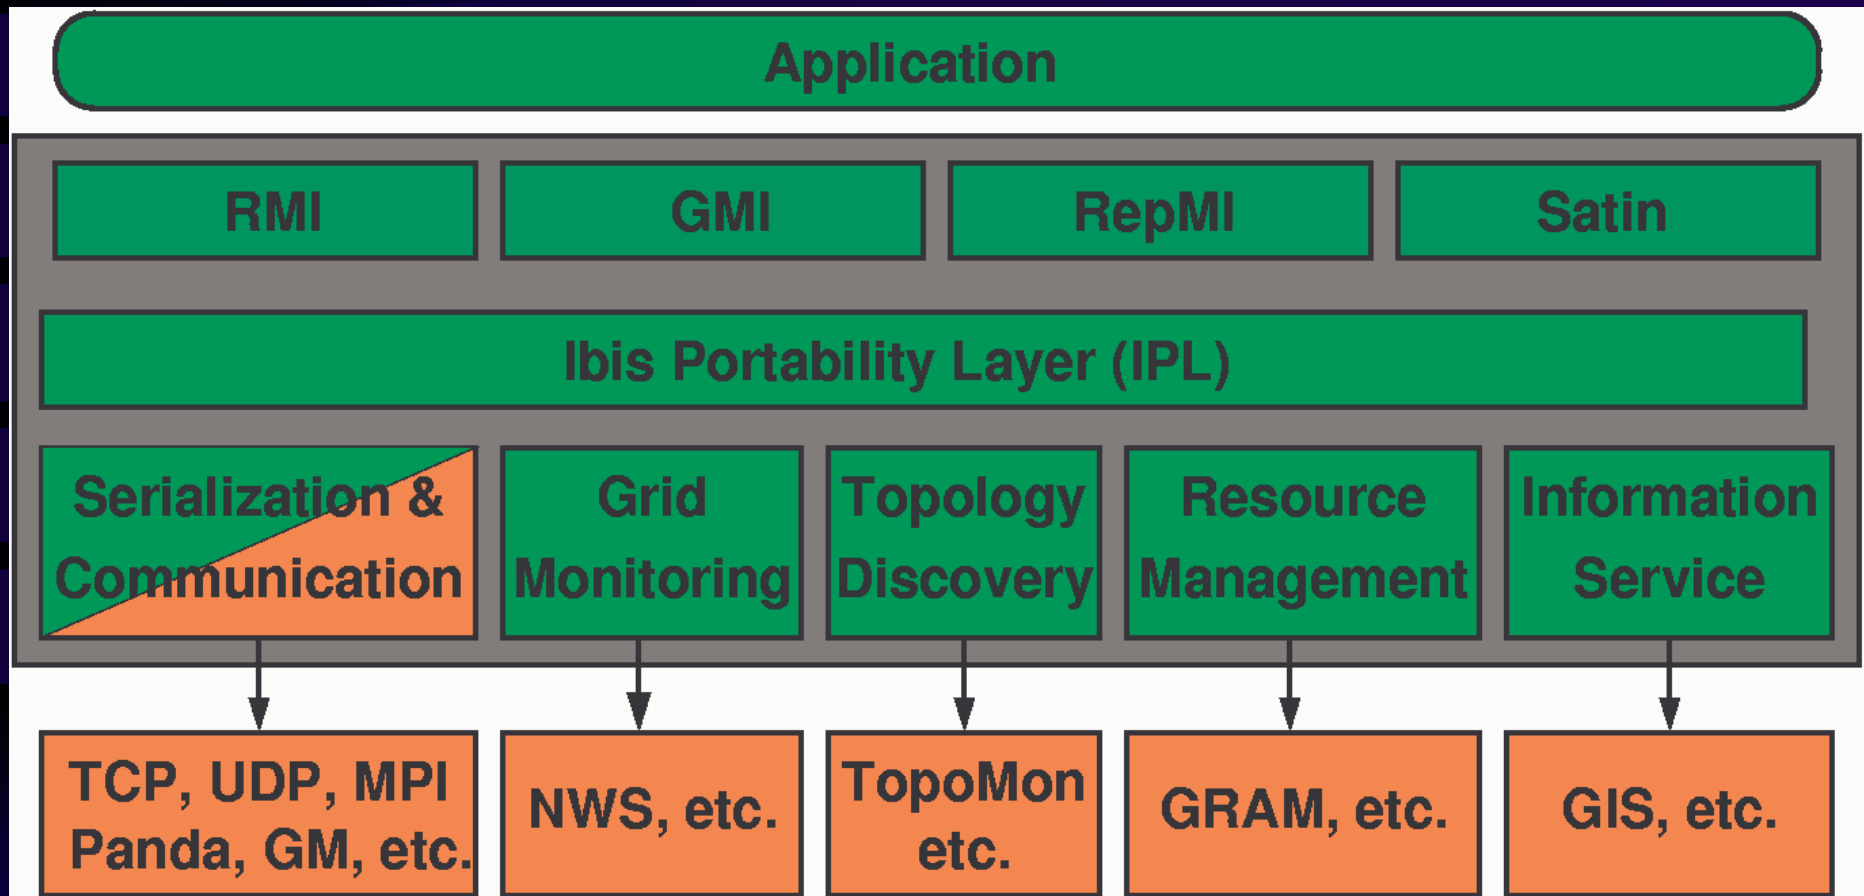
Java on the grid

- **Virtual machine approach / set of libraries**
 - **“Write once, run everywhere” portability**
- **Communication**
 - **RMI is flexible: serialization, polymorphism**
 - **Only synchronous and point-to-point**
 - **Sub-optimal performance**
- **Many Java-based grid projects, but they lack either flexibility or portability or performance**

The Ibis approach

- **Java-based middleware and programming models, combining in one system:**
 - **Reasonable efficient “run everywhere” solution**
 - **Very efficient solutions for special cases**
 - **Native Myrinet communication**
 - **Manta native compiler**
- **One common interface to the hardware (IPL)**
- **Negotiate properties at runtime and dynamically load the best components**


Ibis structure




Challenges

- **How to make the system flexible enough**
 - Run seamlessly on different hardware / protocols
 - When efficient hardware primitives are available, make it possible to use them
- **Make the `□□□□` Java solution efficient enough**
 - Serialization and communication
- **Study which additional optimizations can be done to improve performance in special cases**
 - JIT is black box

IPL: send ports and receive ports

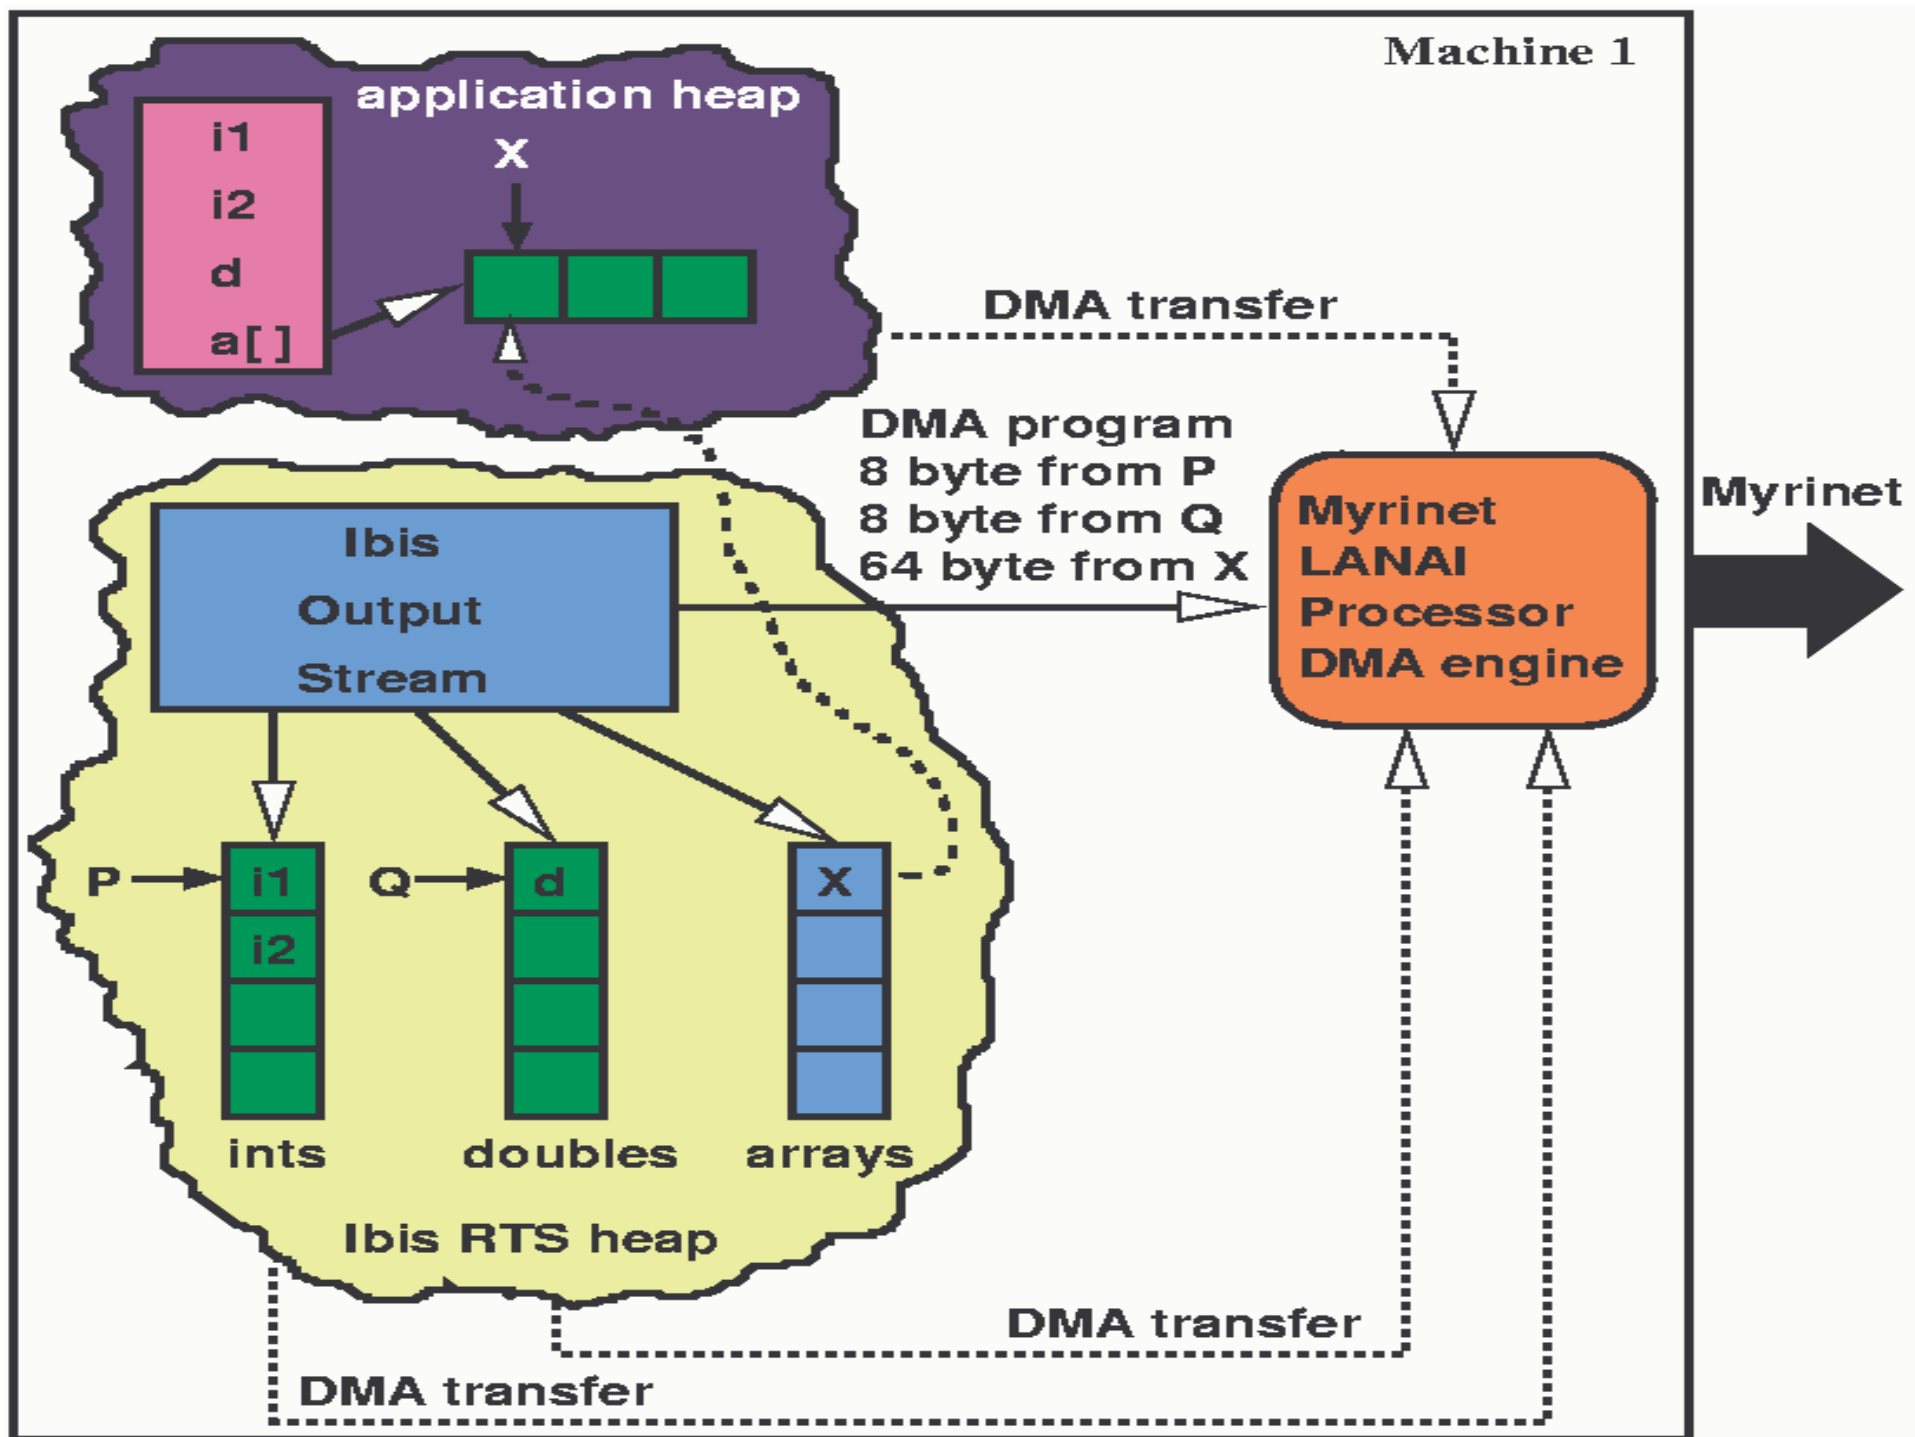


Other communication patterns



Serialization and communication

- **void runtime type inspection**
 - **Compiler generated serialization**
 - **uses Java bytecode for portability**
- **Reduce data copying and conversion with typed buffers for native implementations**
 - **zero copy implementation for primitive arrays**
 - **One copy for objects**
 - **Depending on JIT**



Communication performance

network		low level		RMI performance		
		Ibis	MPI/C	Java	KaRMI	Ibis
Fast	latency	126	161	218.3	127.9	131.3
Ethernet	throughput int []	9.6	11.1	9.5	9.6	9.6
	throughput tree	5.8	5.4	2.2	2.3	4.3
	Latency	33	22	n.a.	32.2	42.2
Myrinet	throughput int []	122	143	n.a.	45.6	76.0
	throughput tree	23.6	16.6	n.a.	2.5	22.9

Conclusions

- **Ibis: high performance computing for grids**
 - **Efficient object-based communication combined with “write one, run everywhere”**
 - Reasonable efficient “run everywhere” solution
 - Very efficient solutions for special cases
- **Interface that allows dynamic selection of exactly the right functionality**
- **Streaming, compiler generated serialization**
- **Zero copy for native implementations**

Future work

- **Write a \square DP implementation for even better “run everywhere performance”**
- **More research on “special cases”**
- **Investigate adaptivity, malleability and fault tolerance for the Ibis programming models**
- **Integrate with existing grid software**

www.cs.vu.nl/ibis