



Distributed Supercomputing in Java

Henri Bal


Vrije Universiteit Amsterdam

Keynote talk 11-th Euromicro Conference on Parallel,
Distributed and Network based Processing, Genoa, Italy,
(February, 5-7, 2003)

Distributed supercomputing

- Parallel processing on geographically distributed computing systems (grids)
- Examples:
 - SETI@home (, RSA-155, Entropia, Cactus 
- Currently limited to trivially parallel applications
- Our goals:
 - Generalize this to more HPC applications
 - Provide high-level programming support

Grids versus supercomputers



- Performance/scalability
 - Speedups on geographically distributed systems?
- Heterogeneity
 - Different types of processors, operating systems, etc.
 - Different networks (Ethernet, Myrinet, WANs)
- General grid issues
 - Resource management, co-allocation, firewalls, security, monitoring, authorization, accounting,

Our approach



- Performance/scalability
 - Exploit hierarchical structure of grids (Albatross project)
- Heterogeneity
 - Use Java + JVM (Java Virtual Machine) technology
- General grid issues
 - Import knowledge from elsewhere (GGF, GridLab)





Speedups on a grid?

- Grids usually are hierarchical
 - Collections of clusters, supercomputers
 - Fast local links, slow wide-area links
- Can optimize algorithms to exploit this hierarchy
 - Message combining + latency hiding on wide-area links
 - Collective operations for wide-area systems
 - Load balancing
- Successful for many applications
 - Did many experiments on a homogeneous wide-area test bed (DAS) [HPCA 1999, IEEE TPDS 2002]






The Ibis system

- High-level & efficient programming support for distributed supercomputing on heterogeneous grids
- Use Java-centric approach + JVM technology
 - Inherently more portable than native compilation
"Write once, run everywhere"
 - Requires entire system to be written in Java
- Use special-case (native) optimizations on demand

Outline

- Programming support
- Highly portable & efficient implementation
- Experiences on DAS-2 and EC GridLab testbeds

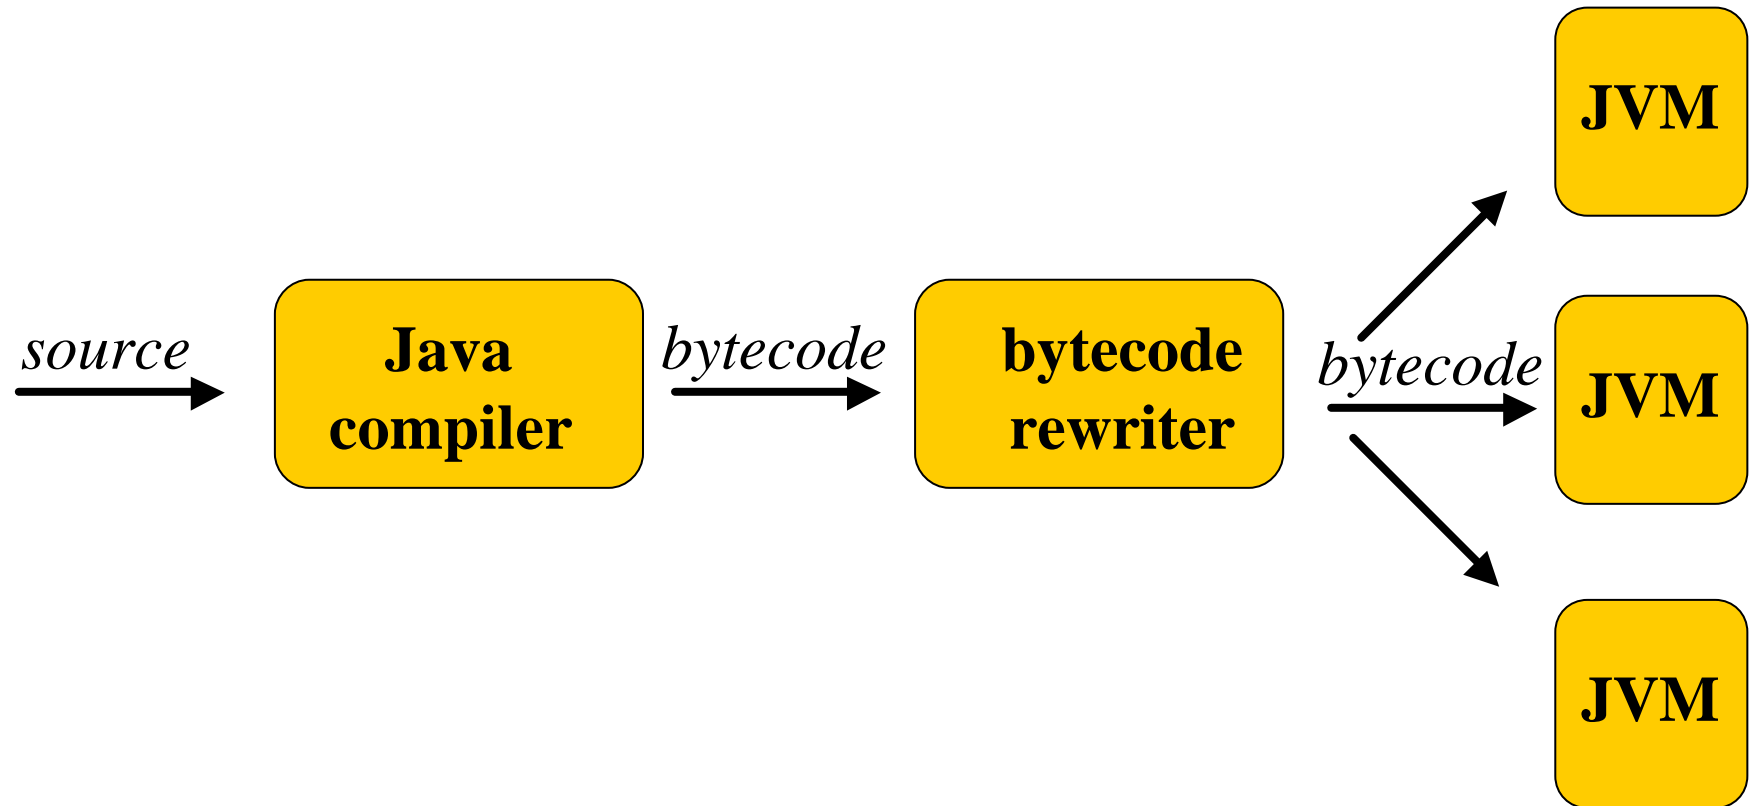


Ibis programming support



- Ibis provides
 - Remote Method Invocation (RMI)
 - Replicated objects (RepMI) - as in Orca
 - Group/collective communication (GMI) - as in MPI
 - Divide & conquer (Satin) - as in Cilk
- All integrated in a clean, object-oriented way into Java, using special “marker” interfaces
 - Invoking native library (e.g. MPI) would give up Java’s “run everywhere” portability

Compiling Ibis programs




GMI (group communication)

- Generalizes Remote Method Invocation
 - Modify how invocations & results are handled
 - Invoke multiple objects, combine/gather results, etc.
 - Expresses many forms of group communication

```
interface Example extends GroupInterface {  
    public int get() throws ...;    ...  
}  
Example e = ;                      // get group stub  
m = findMethod("int get()", e);    // method to configure  
m.useGroupInvocation();           // get() will be multicast  
m.useCombineResult(...);          //results are combined  
int result = e.get();              // example invocation
```

Divide-and-conquer parallelism

- Divide-and-conquer is inherently hierarchical



- Satin
 - Cilk-like primitives (spawn/sync)
- New load balancing algorithm
 - Cluster-aware random work stealing [PPoPP'01]

Example

```
interface FibInter {  
    public int fib(long n);  
}  
  
class Fib implements FibInter {  
    int fib (int n) {  
        if (n < 2) return n;  
        return fib(n-1) + fib(n-2);  
    }  
}
```


Single-threaded Java

```
interface FibInter
    extends ibis.satin.Spawnable {
        public int fib(long n);
    }
```

```
class Fib
    extends ibis.satin.SatinObject
    implements FibInter {
        public int fib (int n) {
            if (n < 2) return n;
            int x = fib (n - 1);
            int y = fib (n - 2);
            sync();
            return x + y;
        }
    }
```

Java + divide&conquer

Example

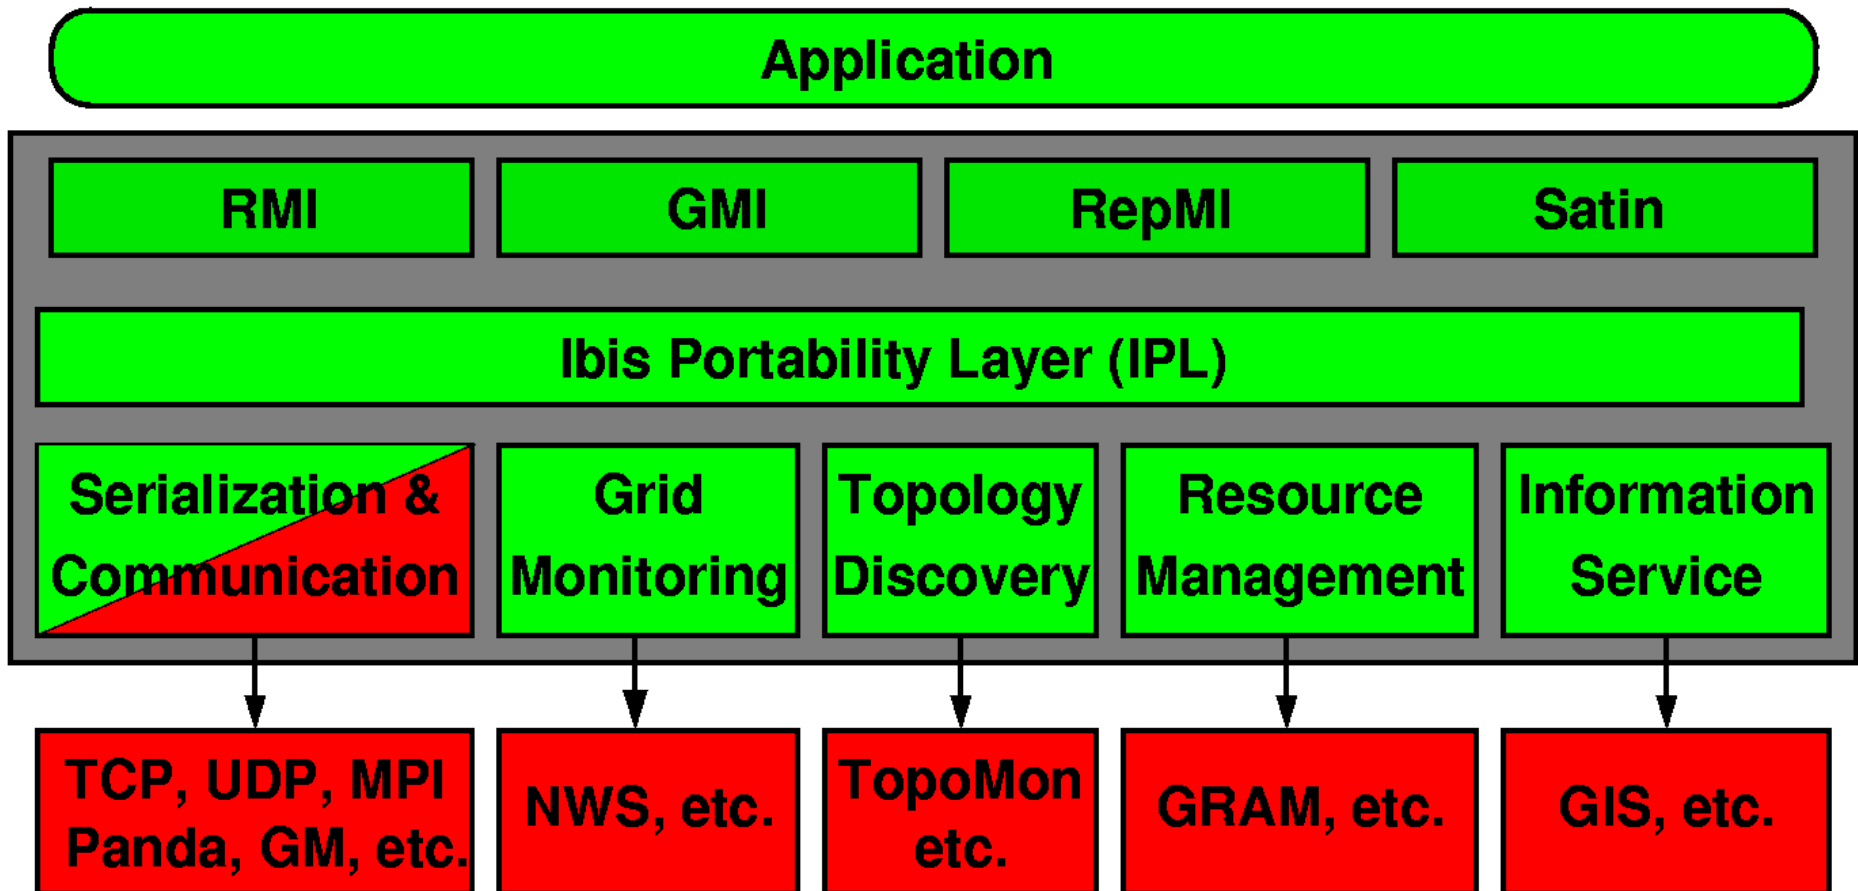


GridLab testbed


Ibis implementation

- Want to exploit Java's "run everywhere" property, **but**
 - That requires 100% pure Java implementation, **no** single line of native code
 - Hard to use native communication (e.g. Myrinet) or native compiler/runtime system
- Ibis approach:
 - Reasonably efficient pure Java solution (for any JVM)
 - Optimized solutions with native code for special cases

Ibis design



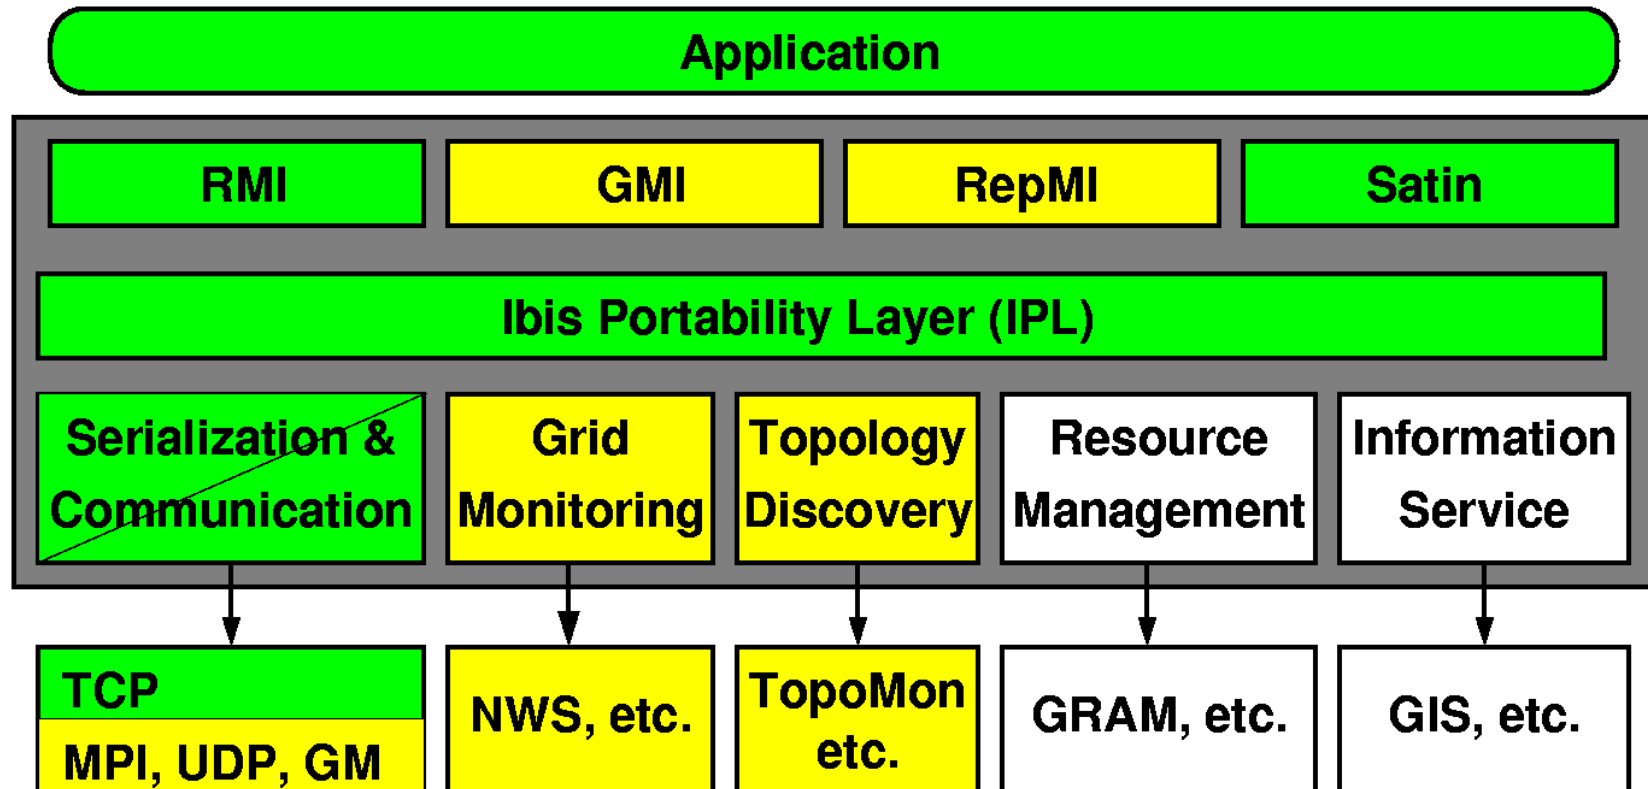
Native code



Pure Java code




Current status



done



in progress



to do

Challenges

- How to make the system flexible enough
 - Run seamlessly on different hardware / protocols
- Make the pure-Java solution efficient enough
 - Need fast local communication even for grid applications
- Special-case optimizations



Flexibility

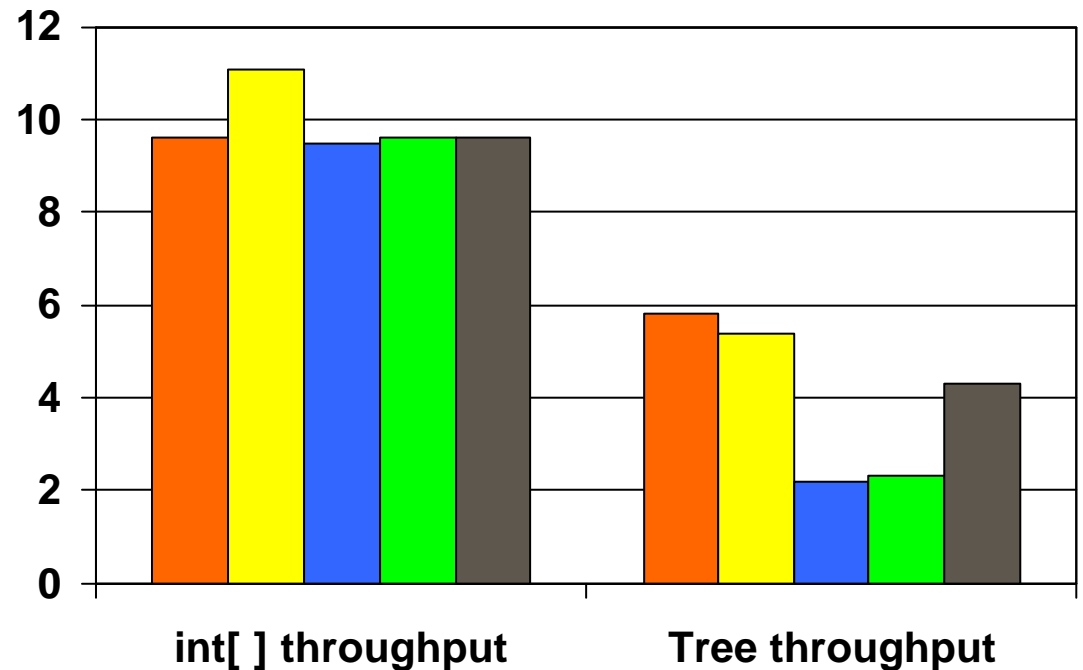


- Support different communication substrates
- IPL just defines an *interface* between high-level programming systems and underlying platforms
- Higher levels can ask IPL to load different *implementations* at runtime, using class loading
 - Eg FIFO ordering, reliable communication

Fast communication in pure Java




- Manta system [ACM TOPLAS Nov. 2001]
 - RMI at RPC speed, but using native compiler & RTS
- Ibis does similar optimizations, but in pure Java
 - Compiler-generated serialization at bytecode level
 - 5-9x faster than using runtime type inspection
 - Reduce copying overhead
 - Zero-copy native implementation for primitive arrays
 - Pure-Java requires type-conversion (=copy) to bytes

Communication performance on Fast Ethernet



Latency (μ s) & throughput (MB/s), measured on 1 GHz Pentium-IIIs (KaRMI = Karlsruhe RMI)

Communication performance on Myrinet

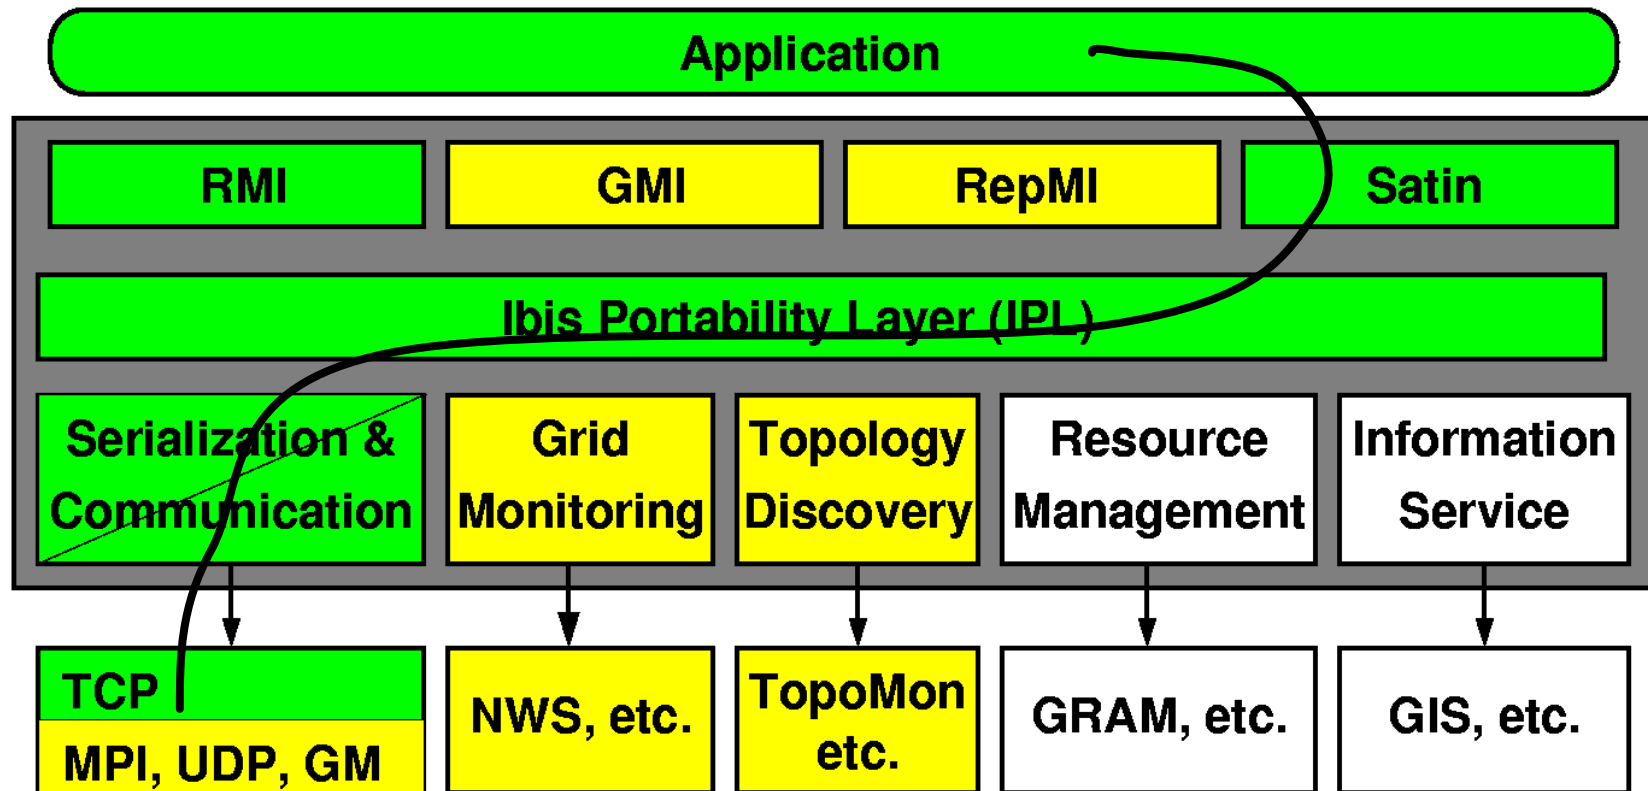


Early Grid experiences with Ibis

- Using Satin divide-and-conquer system
 - Implemented with Ibis in pure Java, using TCP/IP
- Application measurements on
 - DAS-2 (homogeneous)
 - Testbed from EC GridLab project (heterogeneous)




Layers involved



done



in progress




to do


Distributed ASCI Supercomputer (DAS) 2

Node configuration


Dual 1 GHz Pentium-III
≥ 1 GB memory
Myrinet
Linux



Advanced School for Computing and Imaging



Satin on wide-area DAS-2





Satin on GridLab: Theory versus practice

- No support for co-allocation yet (done manually)
- Firewall problems everywhere
 - Currently: use a range of open ports
 - Future: use multiplexing & ssh
- Java indeed runs everywhere
modulo bugs in (old) JVMs
 - IBM 1.3.1 JIT: bug in versioning mechanism
 - Origin2000 JDK: bug in thread synchronization


But it works ...

- Satin/Ibis program was run simultaneously on

Type	OS	CPU	Location
Cluster	Linux	Pentium-3	VU Amsterdam
Server	Solaris	Sparc	VU Amsterdam
Origin 2000	Irix64	MIPS	AEI Potsdam
Cluster	Linux	Pentium-3	PSNC Poznan

Performance on GridLab

performance relative to DAS2 cluster



16-node DAS-2 cluster + 16-node Origin

Summary

- Ibis: a programming environment for grids
 - RMI, group communication, divide&conquer
- Portable
 - Using Java's "write once, run everywhere" property
- Efficient
 - Reasonably efficient "run everywhere" solution
 - Optimized solutions for special cases
- Experience with prototype system on GridLab

Acknowledgements

Rob van Nieuwpoort

Jason Maassen

Thilo Kielmann

Rutger Hofman

Ceriel Jacobs

Gosia Wrzesinska

Olivier Aumage

Kees Verstoep



vrije Universiteit

web site: www.cs.vu.nl/ibis