

Fault tolerance, malleability and migration for divide-and-conquer applications on the Grid


Gosia Wrzesińska, Rob V. van Nieuwpoort,
Jason Maassen, Henri E. Bal



Ibis

Distributed supercomputing

- Parallel processing on geographically distributed computing systems (grids)
- Needed:
 - Fault-tolerance: survive node crashes (also entire clusters)
 - Malleability: add or remove machines at runtime
 - Migration: move a running application to another set of machines (comes for free with malleability)
- We focus on divide-and-conquer applications



Outline

- The **Ibis** grid programming environment
- **Satin**: a divide-and-conquer framework
- Fault-tolerance, malleability and migration in Satin
- Performance evaluation

The Ibis system


- Java-centric => portability
 - „write once, run anywhere”
- Efficient communication
 - Efficient pure Java implementation
 - Optimized solutions for special cases with native code
- High level programming models:
 - Divide & Conquer (Satin)
 - Remote Method Invocation (RMI)
 - Replicated Method Invocation (RepMI)
 - Group Method Invocation (GMI)

<http://www.cs.vu.nl/ibis/>



Satin: divide-and-conquer on the Grid

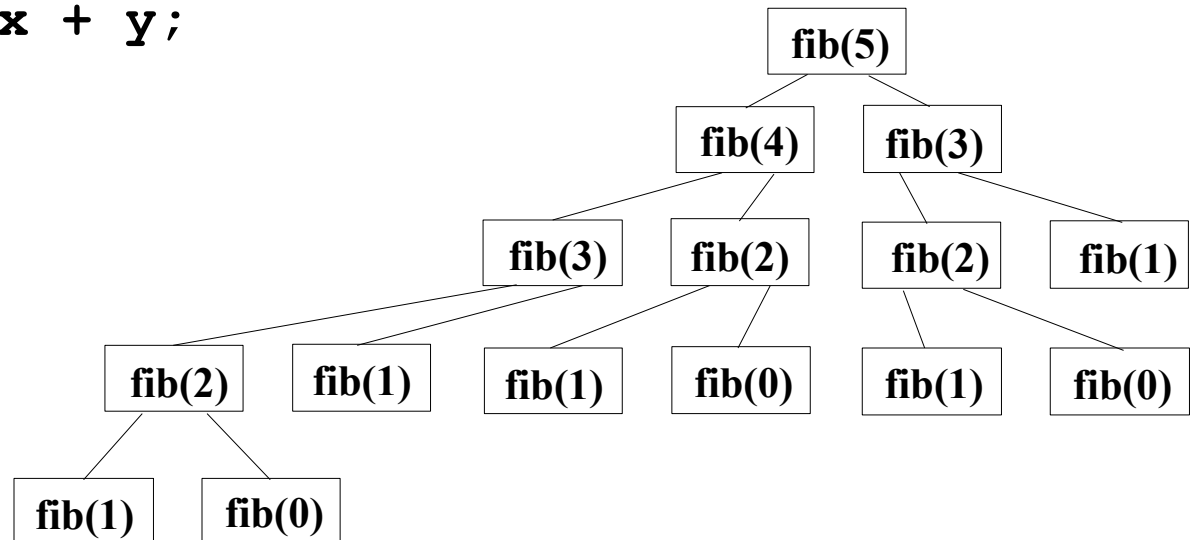
- Effective paradigm for Grid applications (hierarchical)
- Satin: Grid-friendly load balancing (aware of cluster hierarchy)
- Missing support for
 - Fault tolerance
 - Malleability
 - Migration



Example: Fibonacci

```
class Fib {  
    int fib (int n) {  
        if (n < 2) return n;  
        int x = fib(n-1);  
        int y = fib(n-2);  
        return x + y;  
    }  
}
```



**Single-threaded
Java**



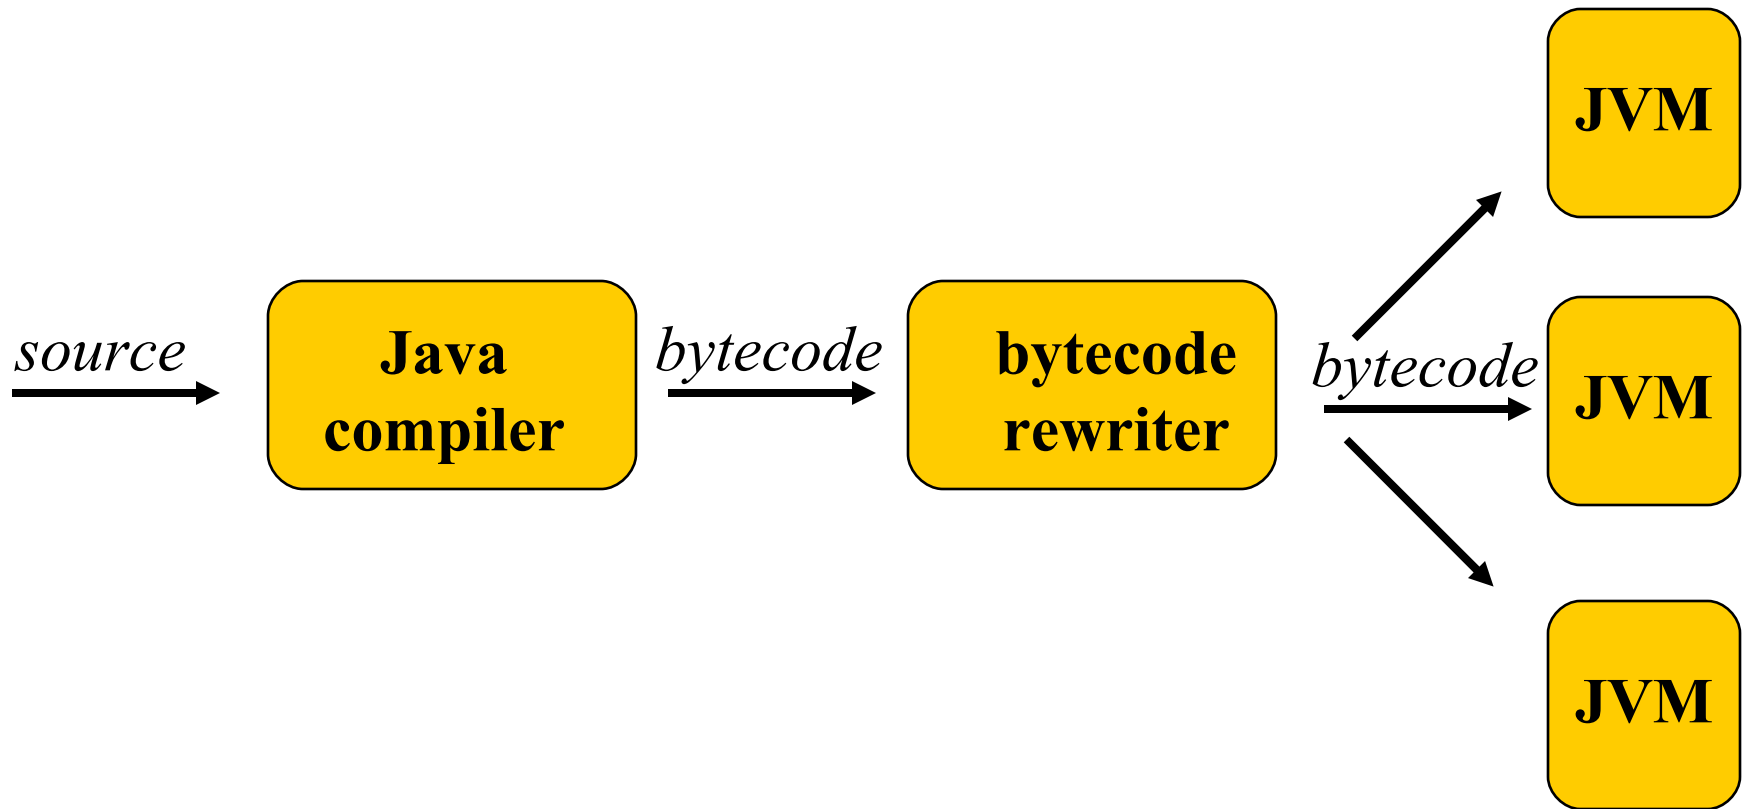
Example: Fibonacci

```
public interface FibInter extends  
ibis.satin.Spawnable {  
    public int fib (int n);  
}
```

```
class Fib extends ibis.satin.SatinObject  
implements FibInter {  
    public int fib (int n) {  
        if (n < 2) return n;  
        int x = fib(n-1); /*spawned*/  
        int y = fib(n-2); /*spawned*/  
        sync();  
        return x + y;  
    }  
}
```



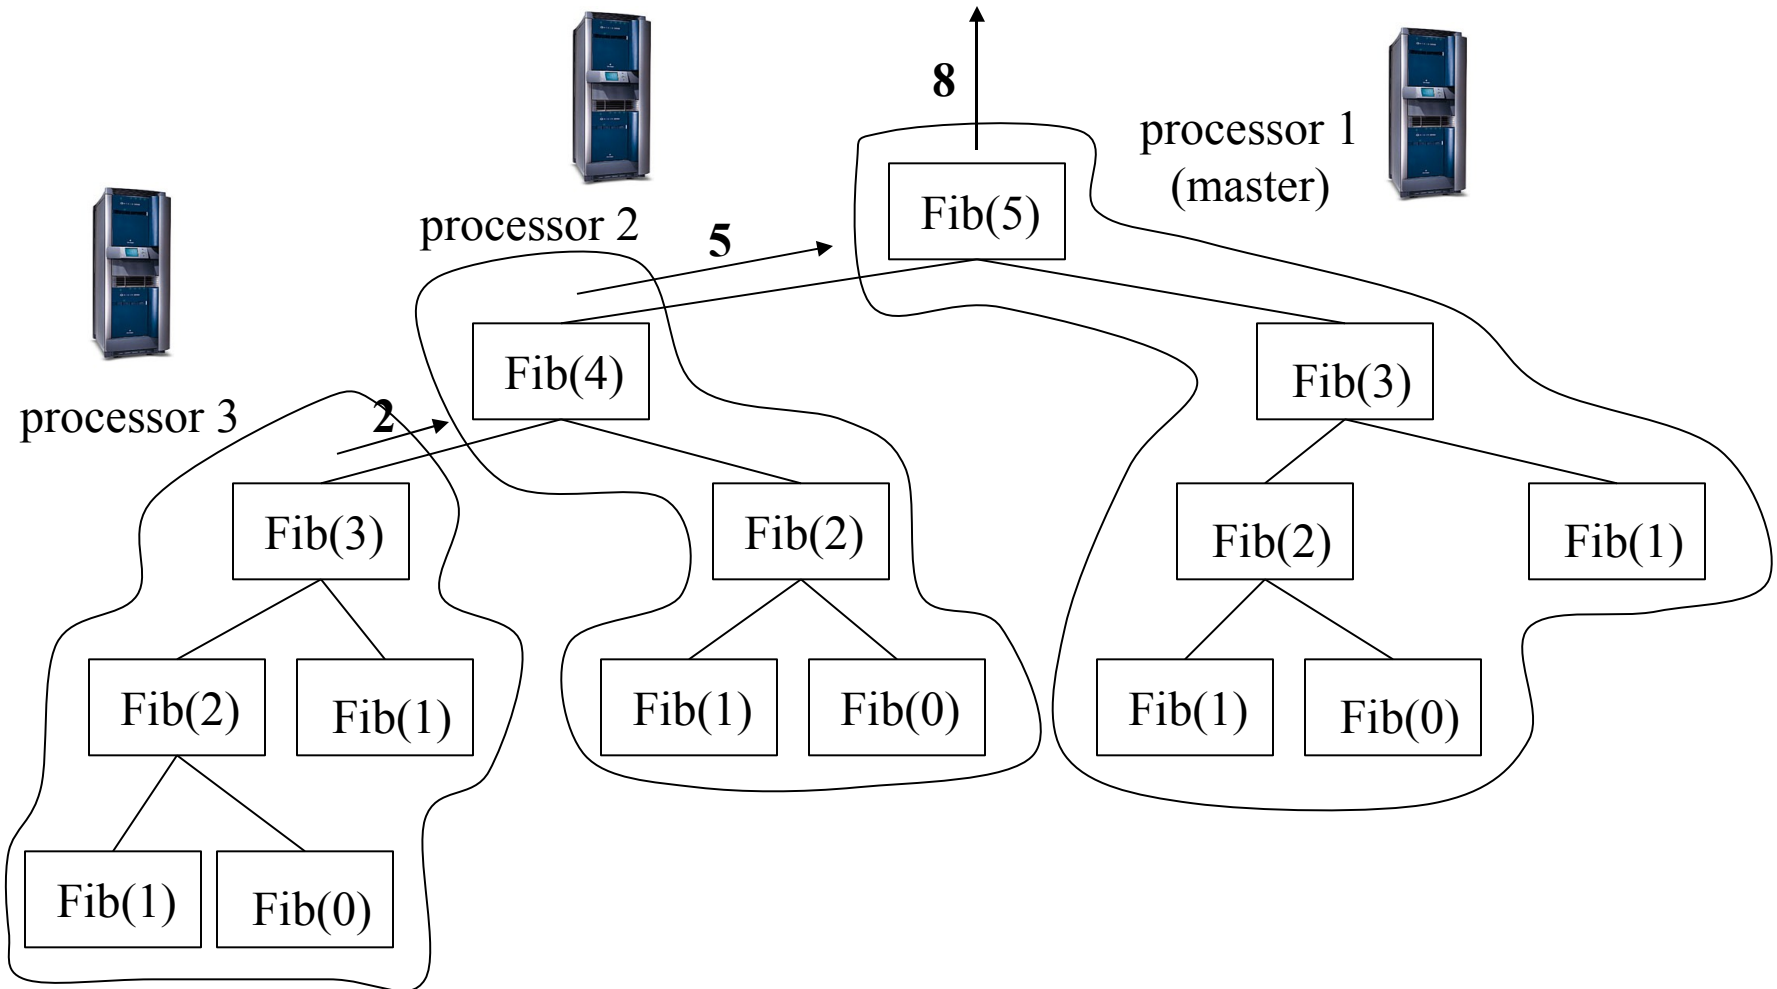
Compiling Satin programs



Executing Satin programs

- Spawn: put work in work queue
- Sync:
 - Run work from queue
 - If empty: steal (load balancing)

Example application: Fibonacci



Satin: load balancing for Grids

- Random Stealing (RS)
 - Pick a victim at random
 - Provably optimal on a single cluster (Cilk)
 - Problems on multiple clusters:
 - $(C-1)/C$ % stealing over WAN
 - Synchronous protocol


Grid-aware load balancing

- Cluster-aware Random Stealing (CRS)
[van Nieuwpoort et al., PPoPP 2001]
 - When idle:
 - Send **asynchronous** steal request to random node in **different cluster**
 - In the meantime steal locally (synchronously)
 - Only one wide-area steal request at a time

Satin raytracer on a grid



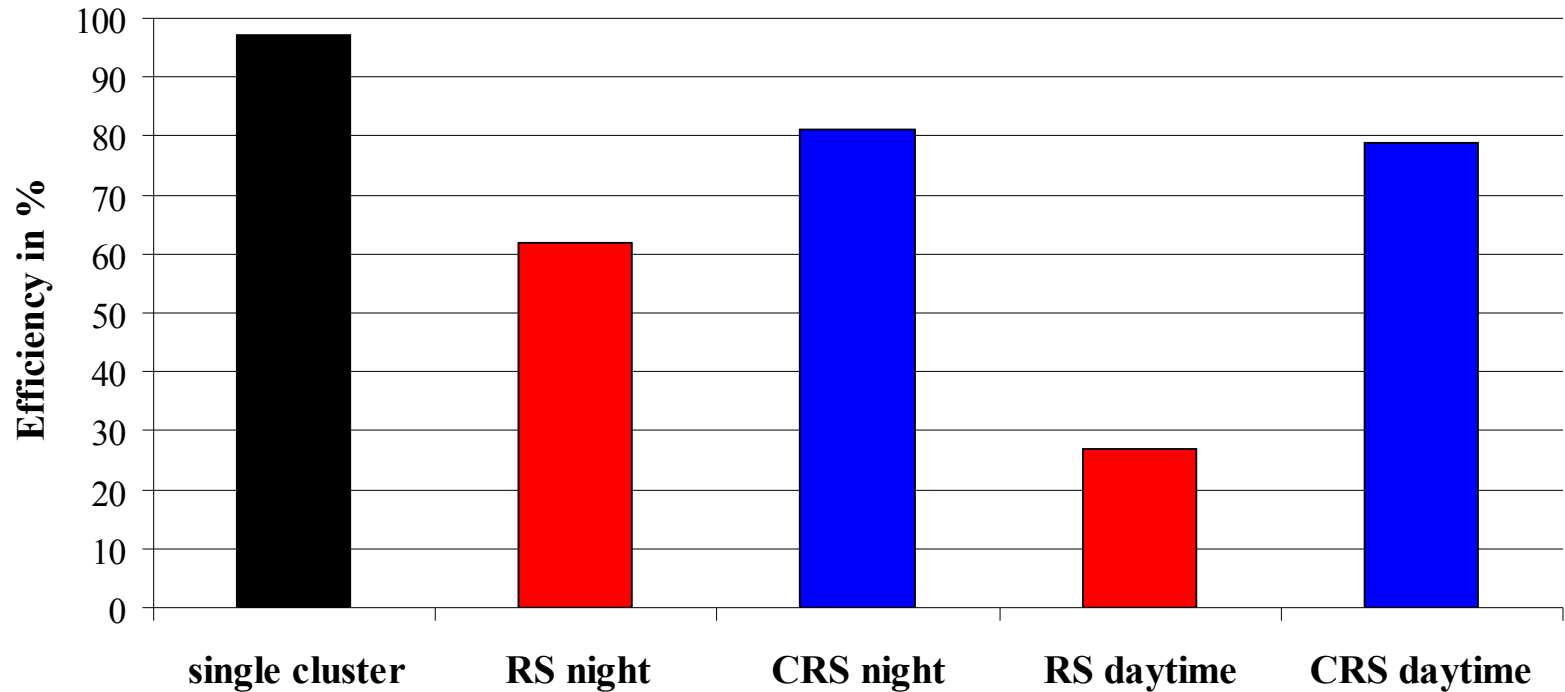
- GridLab testbed: 5 cities in Europe
- 40 cpus
- Distance up to 2000km
- Factor of 10 difference in CPU speeds
- Latencies:
 - 0.2 – 210 ms daytime
 - 0.2 – 66 ms night
- Bandwidth:
 - 9KB/s – 11MB/s
- Three orders of magnitude difference in communication speeds



Configuration

Location	Type	OS	CPU	CPUs
Amsterdam, The Netherlands	Cluster	Linux	Pentium-3	8 x 1
Amsterdam, The Netherlands	SMP	Solaris	Sparc	1 x 2
Brno, Czech Republic	Cluster	Linux	Xeon	4 x 2
Cardiff, Wales, UK	SMP	Linux	Pentium-3	1 x 2
ZIB Berlin, Germany	SMP	Irix	MIPS	1 x 16
Lecce, Italy	SMP	Tru64	Alpha	1 x 4

CRS performance on GridLab testbed



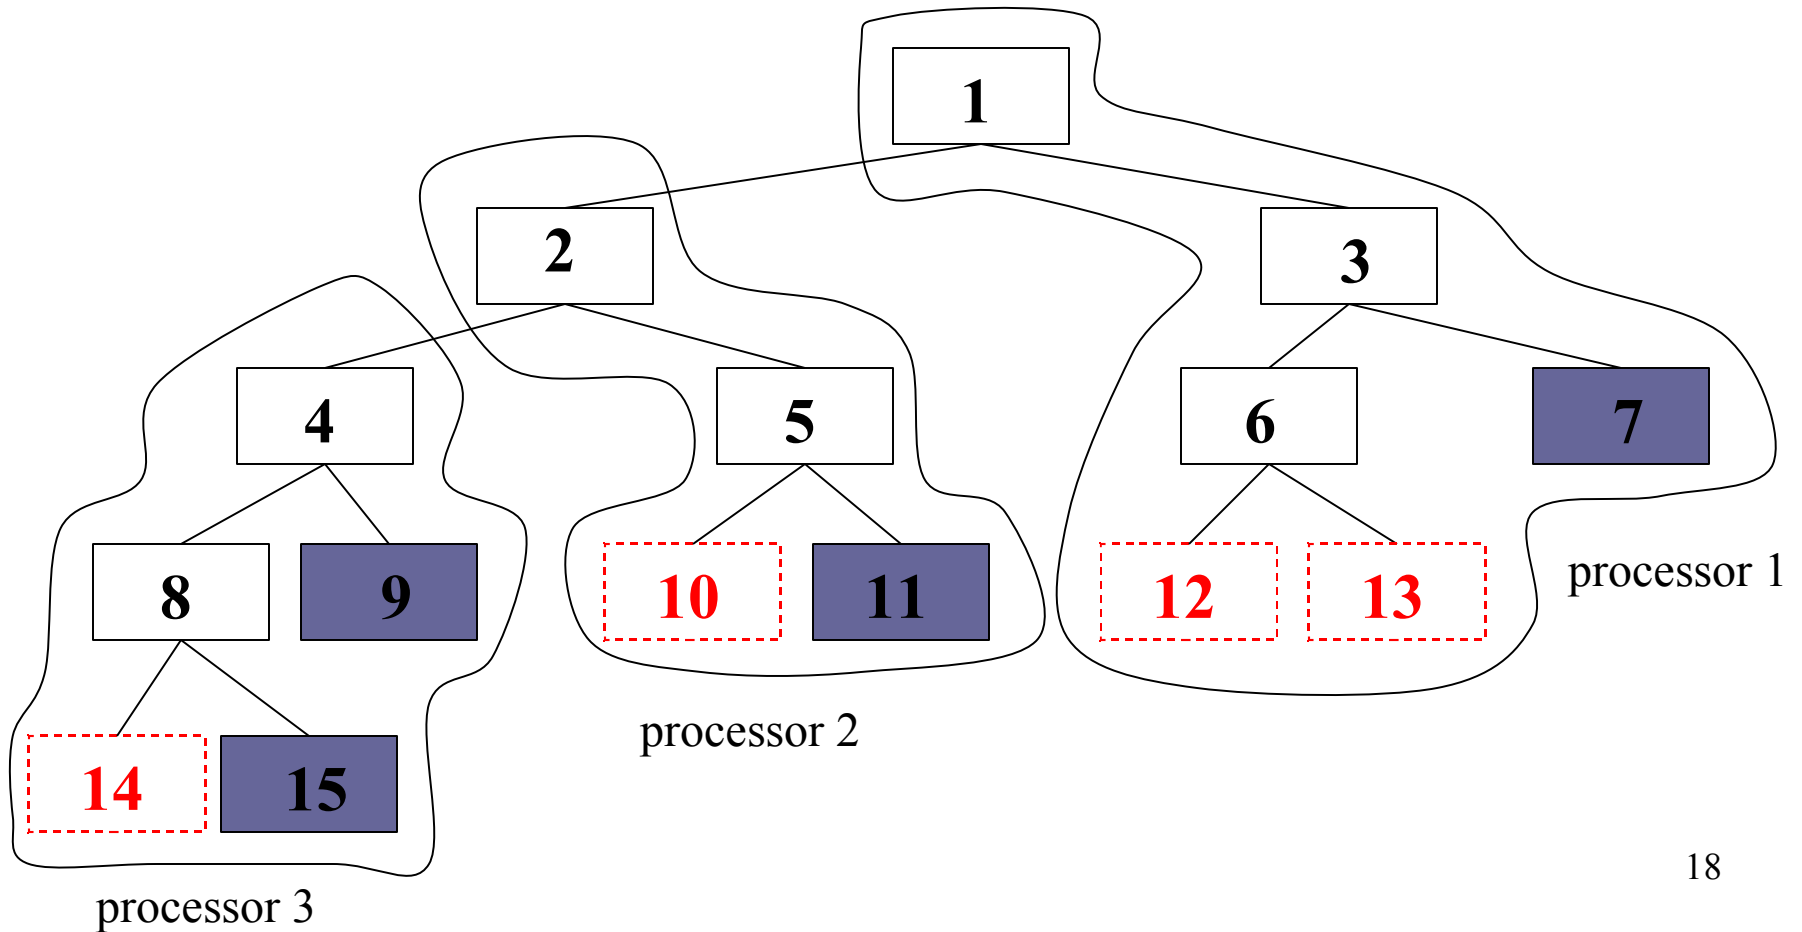
Satin summary

- Satin allows rapid development of parallel applications which are able to run with high efficiency in geographically distributed and highly heterogeneous environments
- Applications:
 - Barnes-Hut, Raytracer, SAT solver, TSP, Knapsack
 - All master-worker algorithms
- Still missing: FT, malleability, migration



Fault-tolerance, malleability, migration

- Can be implemented by handling processors joining or leaving the ongoing computation
- Processors may leave either unexpectedly (crash) or gracefully
- Handling joining processors is trivial:
 - Let them start stealing jobs
- Handling leaving processors is harder:
 - **Recompute** missing jobs
 - **Problems**: orphan jobs, partial results from gracefully leaving processors


Crashing processors



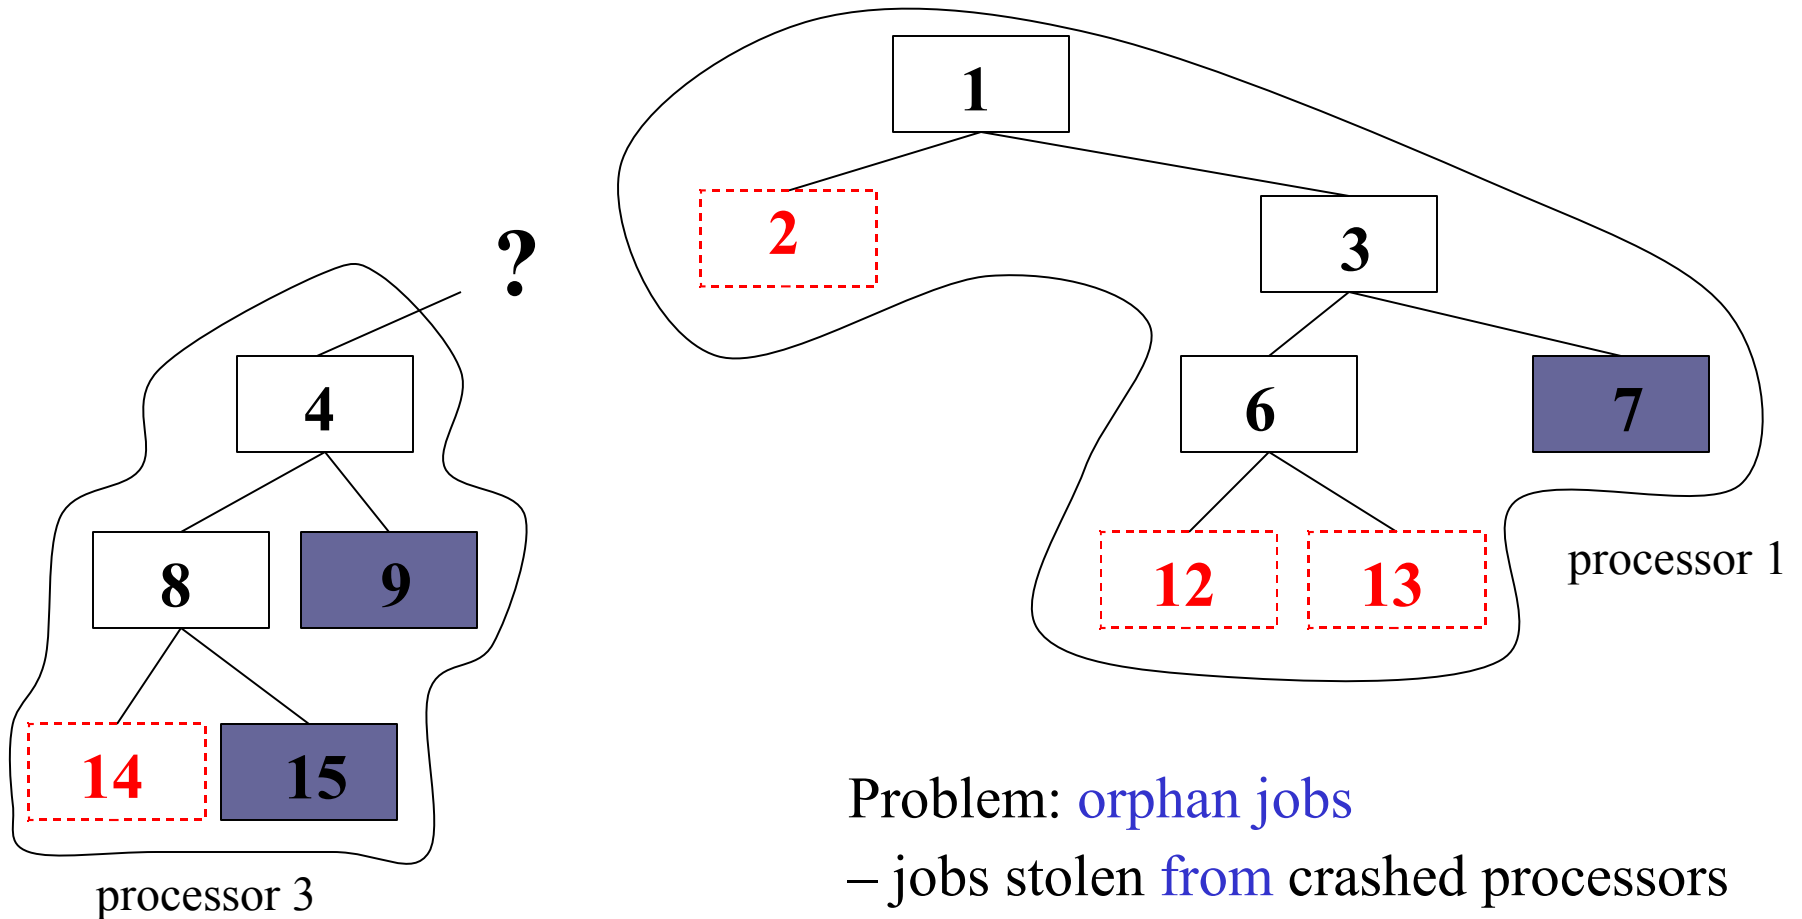
Crashing processors




Crashing processors



Crashing processors




Crashing processors

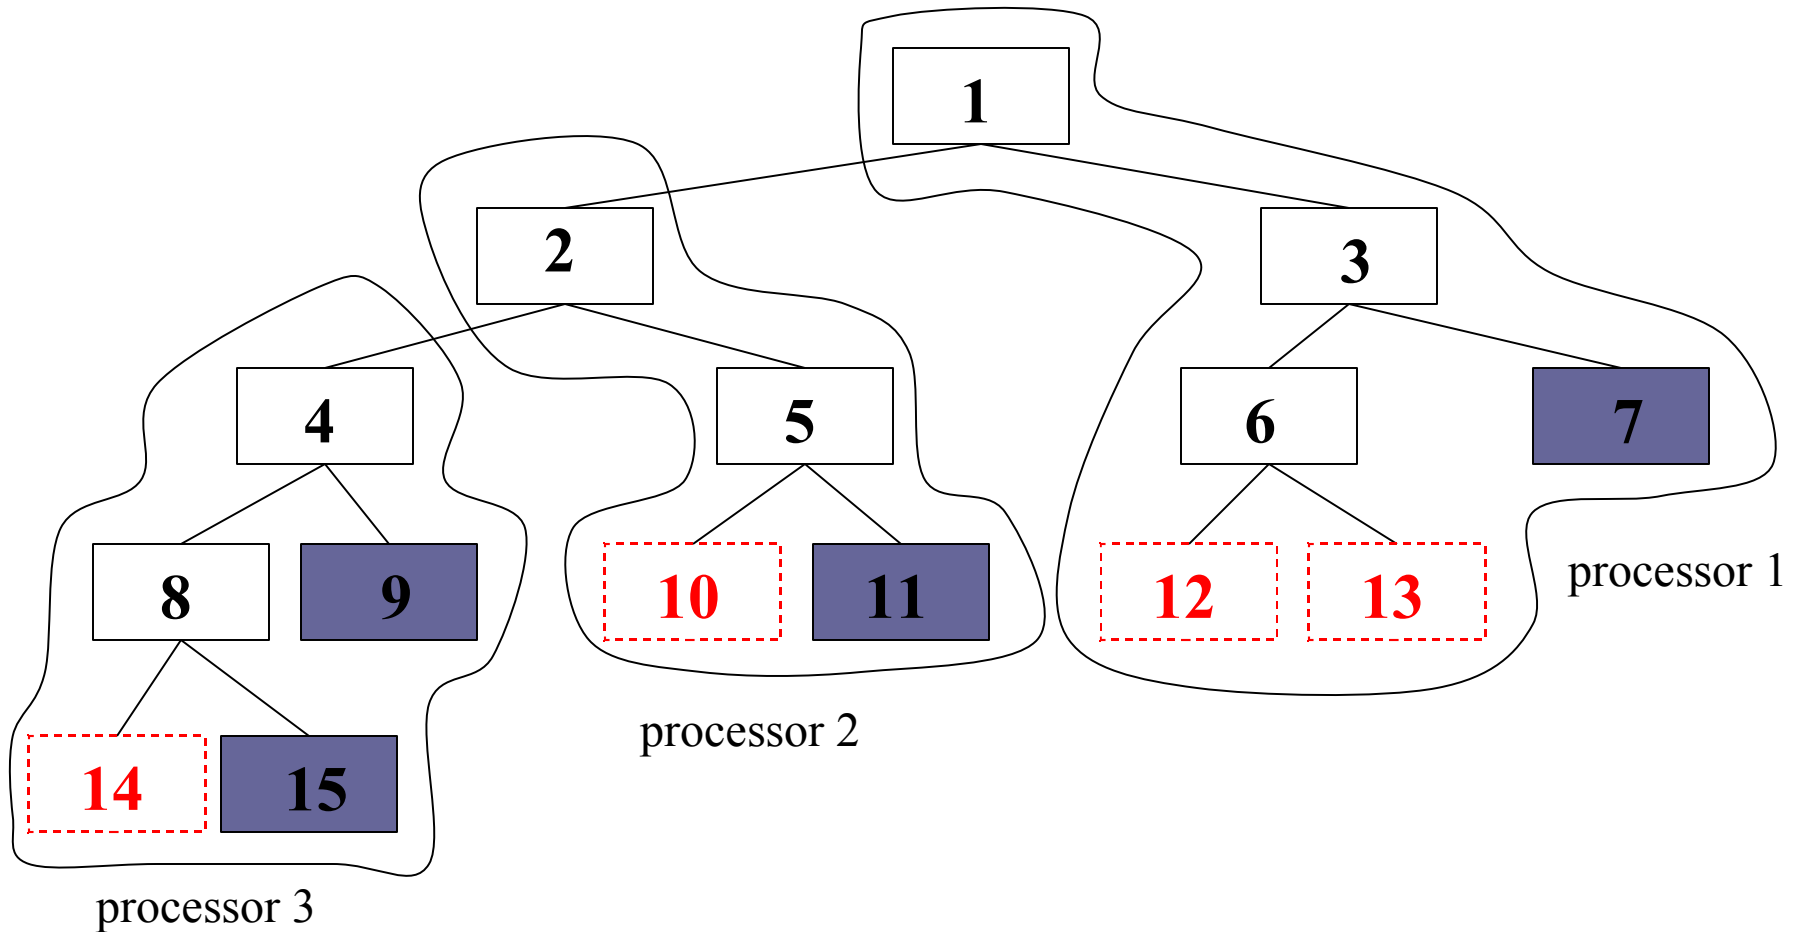


Handling orphan jobs



- For each **finished** orphan, broadcast (jobID,processorID) tuple; abort the rest
- All processors store tuples in **orphan tables**
- Processors perform lookups in orphan tables for each **recomputed** job
- If successful: send a **result request** to the owner (async), put the job on a stolen jobs list




Handling orphan jobs - example



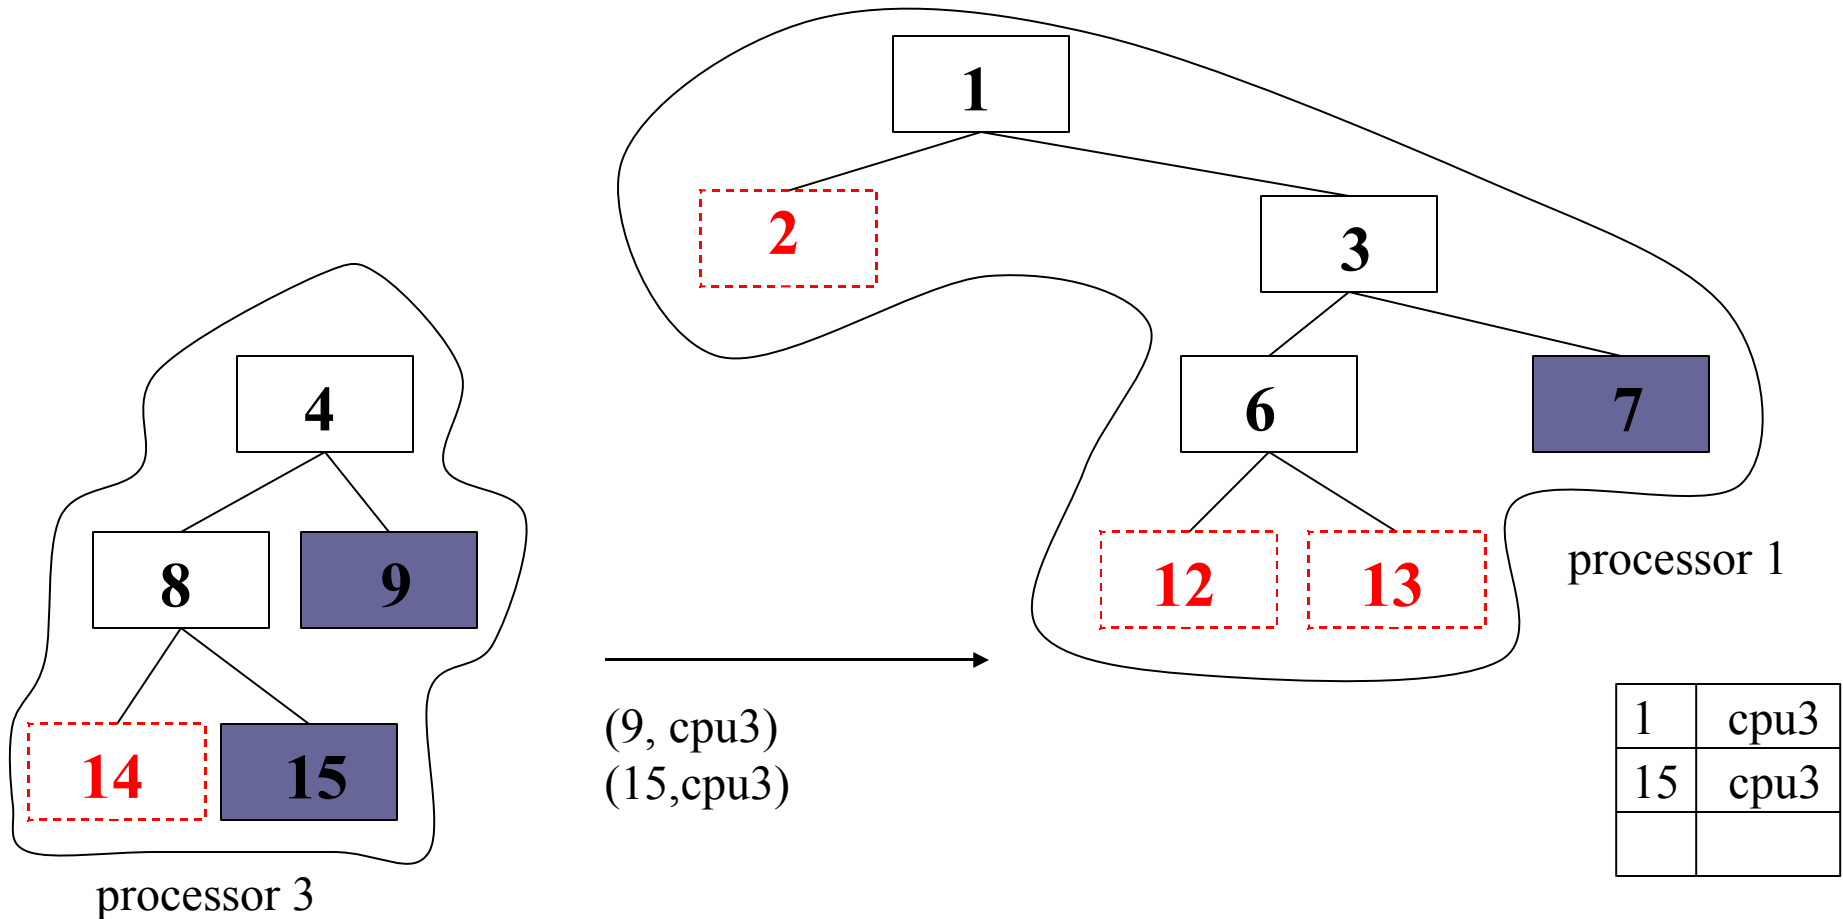
Handling orphan jobs - example




Handling orphan jobs - example




Handling orphan jobs - example




Handling orphan jobs - example



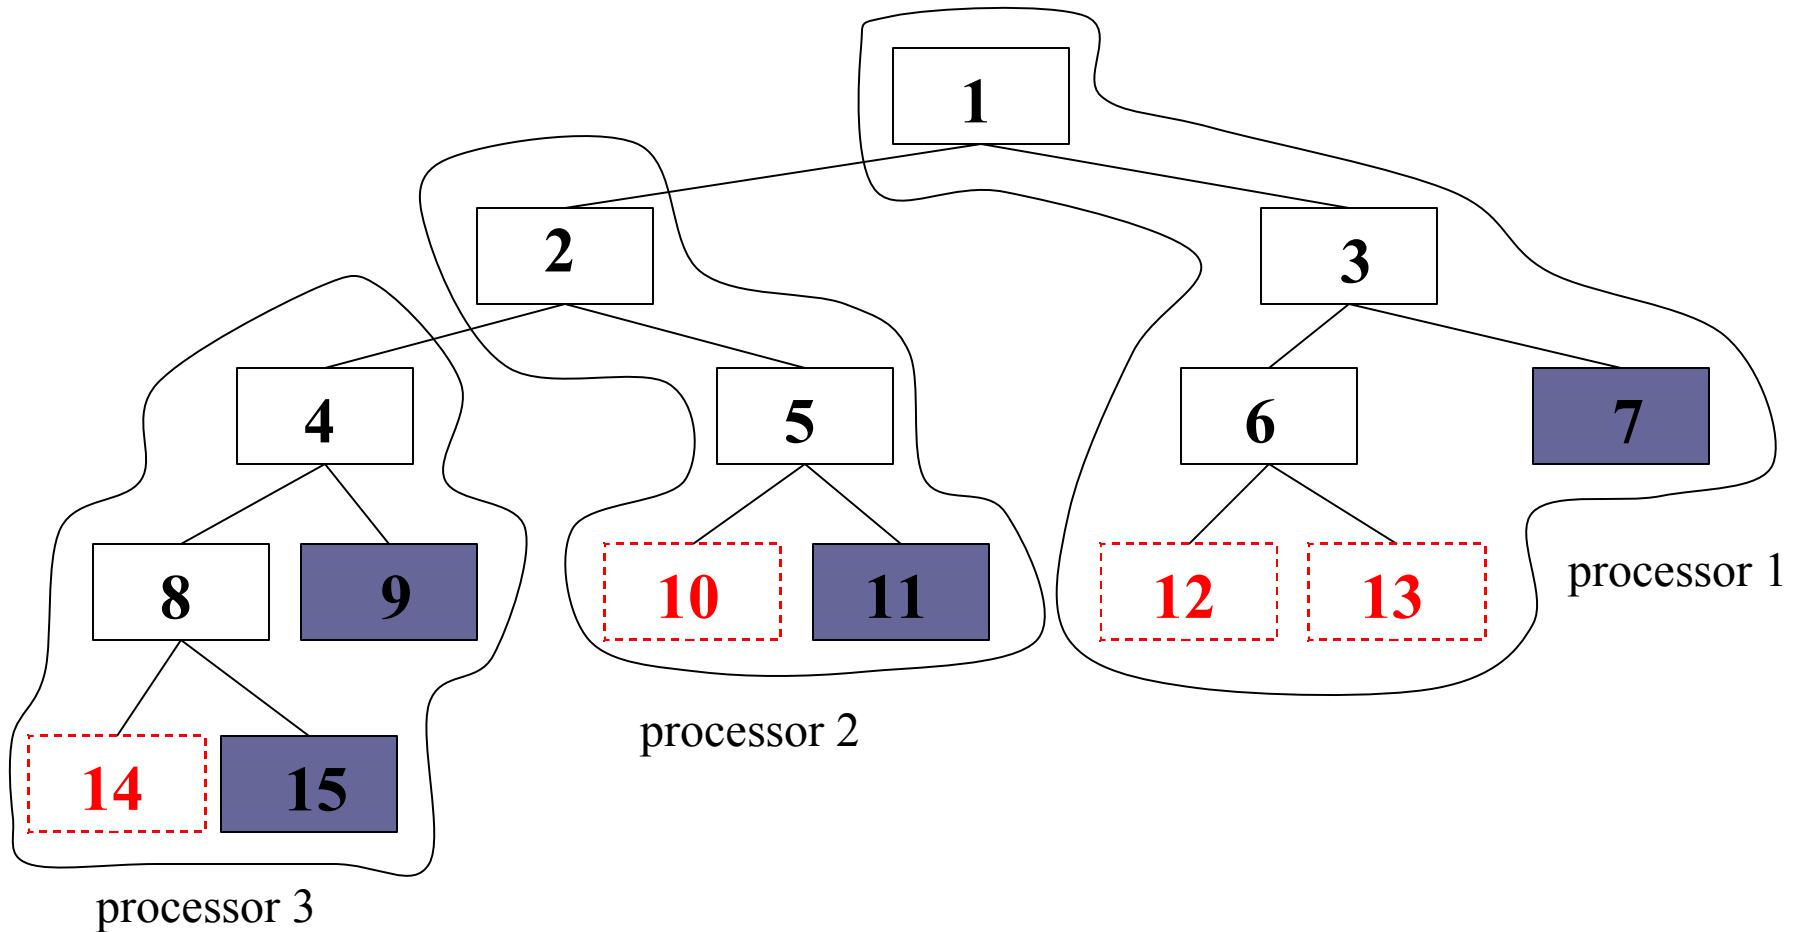
Handling orphan jobs - example



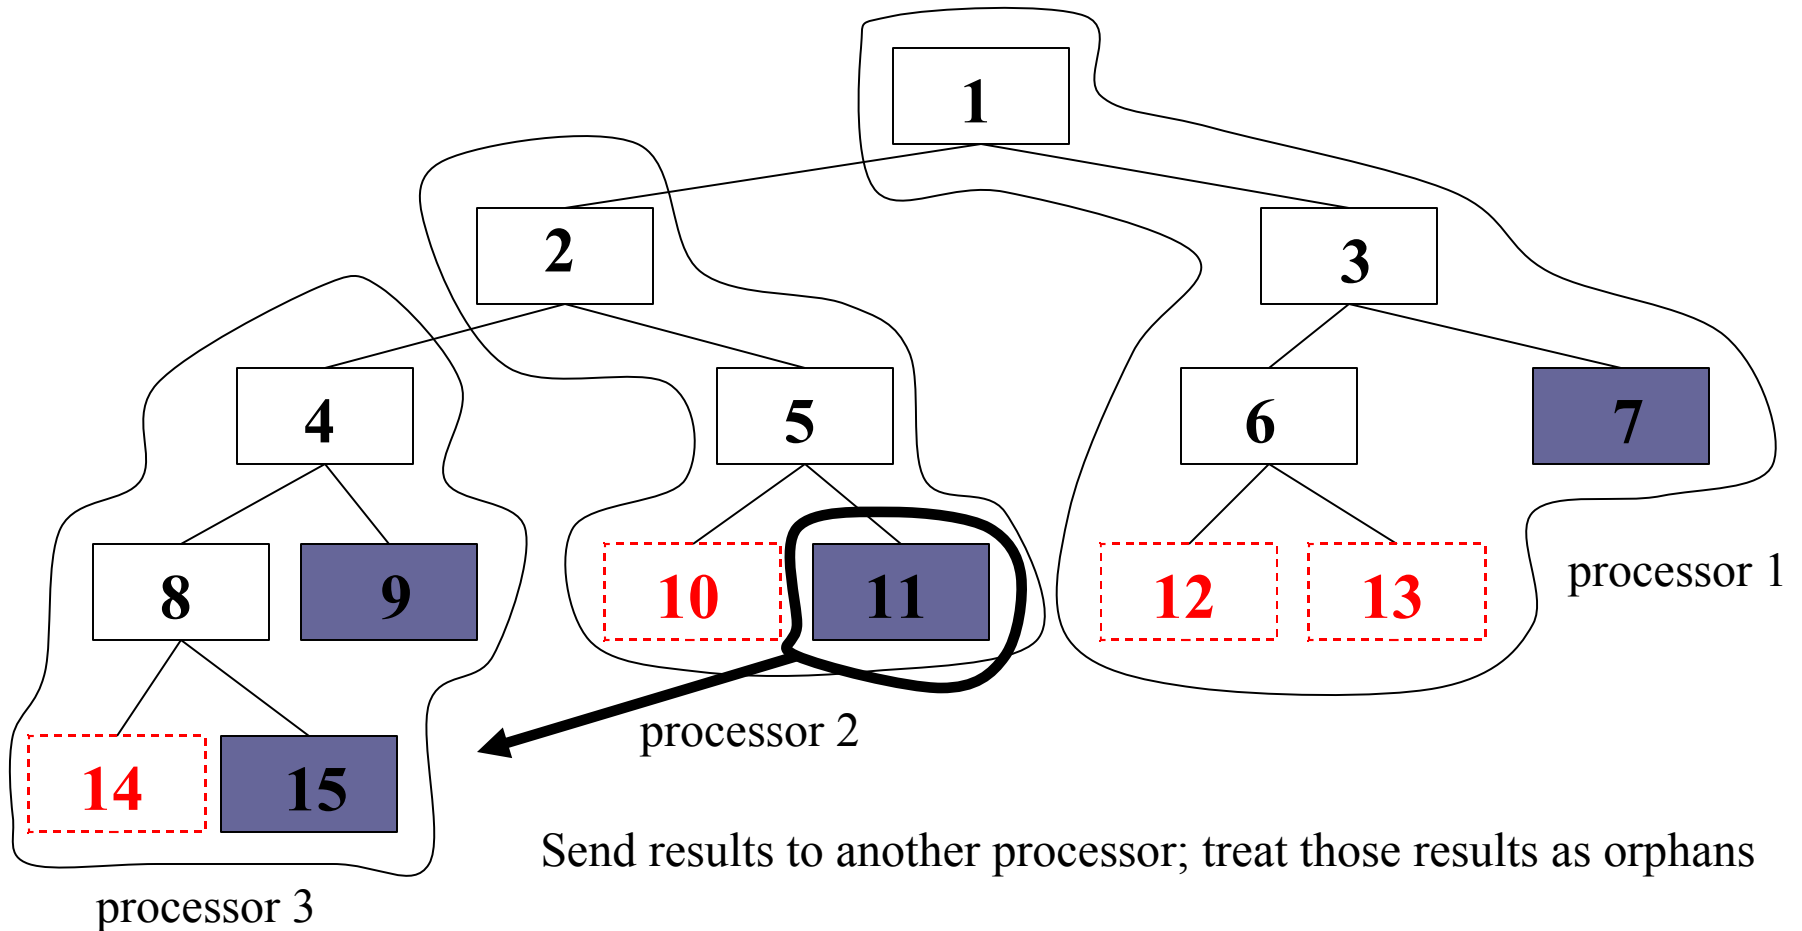
Handling orphan jobs - example



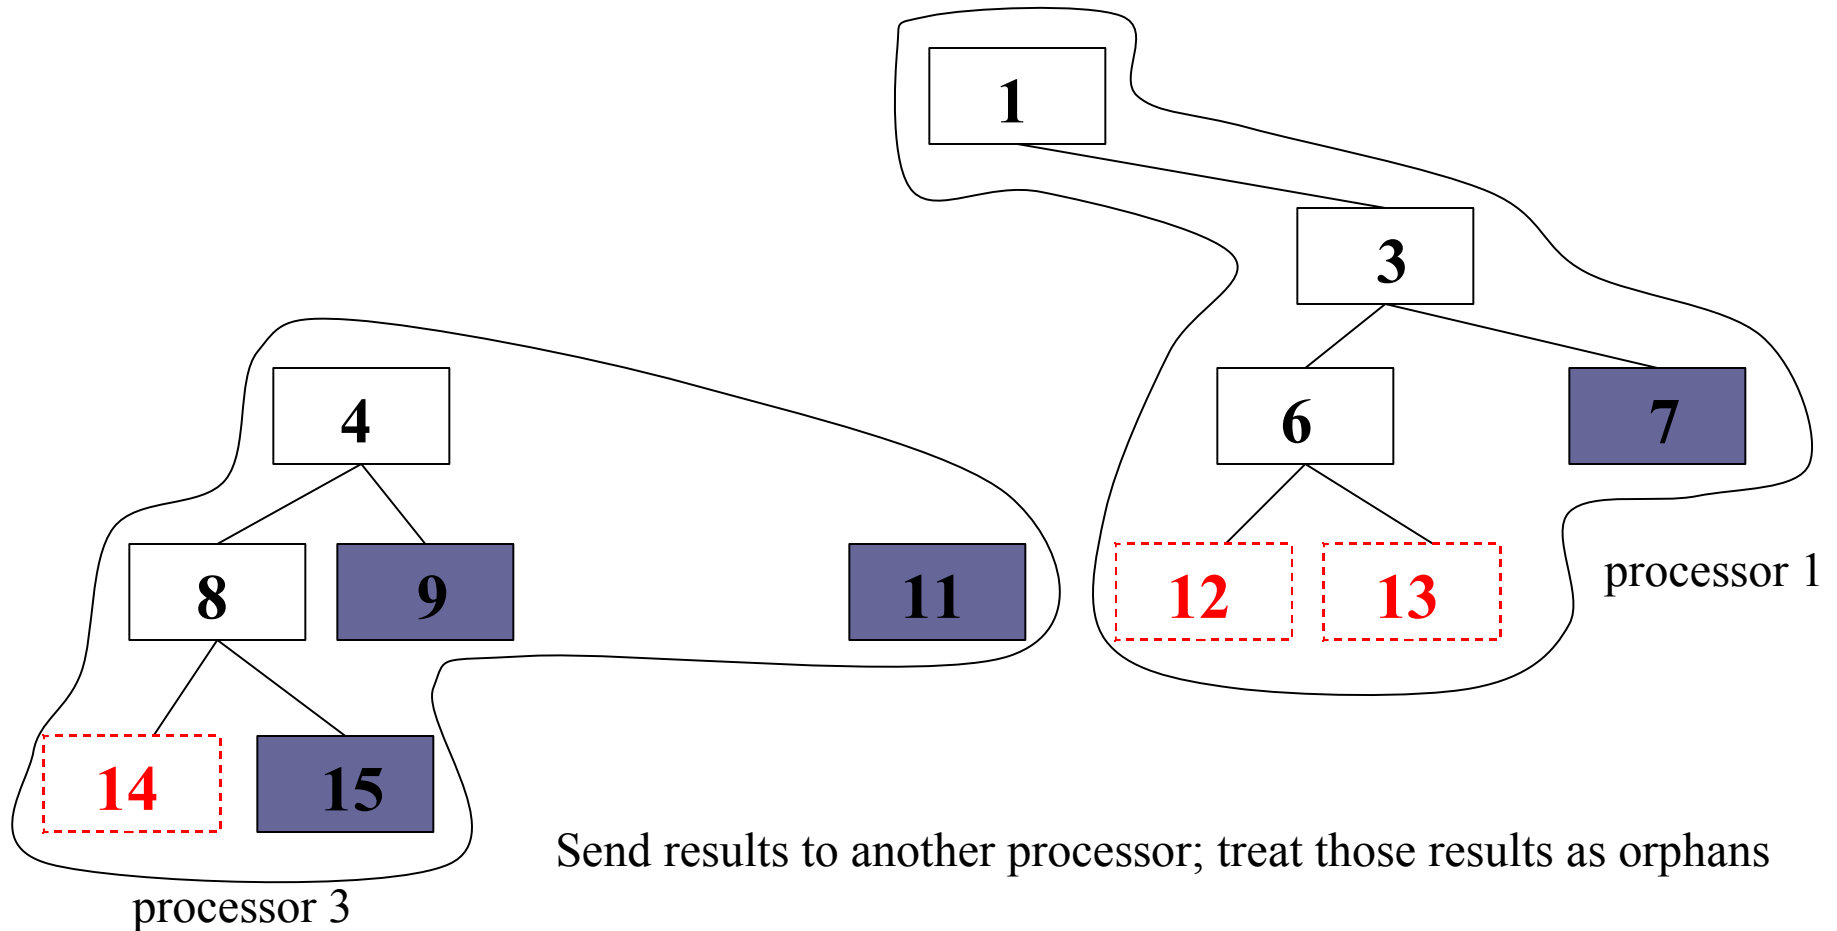
Processors leaving gracefully




Processors leaving gracefully




Processors leaving gracefully




Processors leaving gracefully



Processors leaving gracefully



Processors leaving gracefully



A crash of the master


- Master: the processor that started the computation by spawning the root job
- If master crashes:
 - Elect a new master
 - Execute normal crash recovery
 - New master restarts the applications
 - In the new run, all results from the previous run are reused

Some remarks about scalability

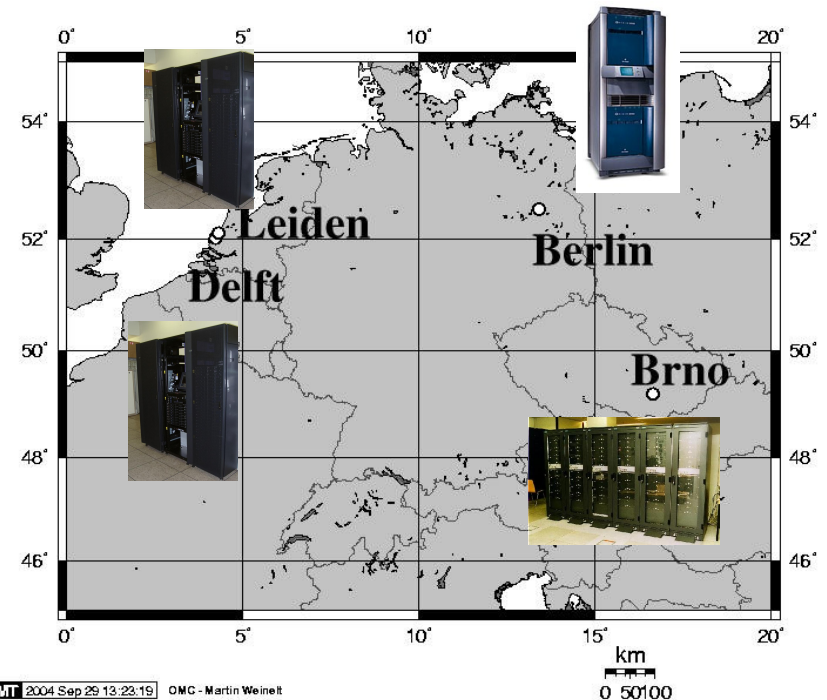
- Little data is broadcast ($< 1\%$ jobs, pointers)
- Message combining
- Lightweight broadcast: no need for reliability, synchronization, etc.

Performance evaluation


- Leiden, Delft (DAS-2) + Berlin, Brno (GridLab)
- Bandwidth:
62 – 654 Mbit/s
- Latency:
2 – 21 ms



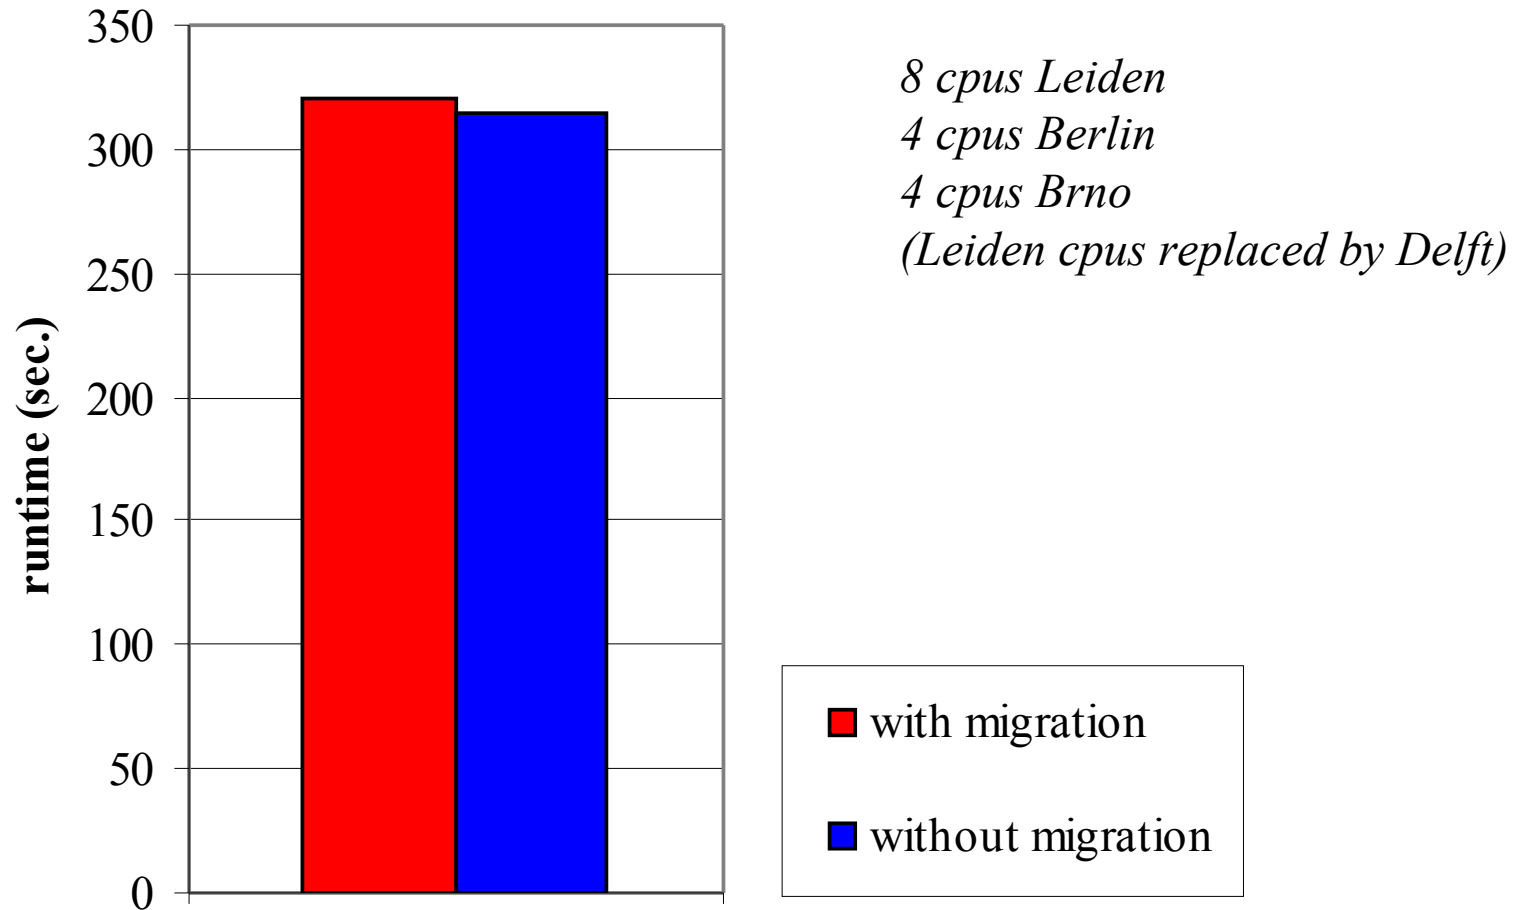
Advanced School for Computing and Imaging



Impact of saving partial results




Migration overhead



Crash-free execution overhead

Used: 32 cpus in Delft



Summary

- Satin implements fault-tolerance, malleability and migration for divide-and-conquer applications
- Save partial results by repairing the execution tree
- Applications can adapt to changing numbers of cpus and migrate without loss of work (overhead $< 10\%$)
- Outperform traditional approach by 25%
- No overhead during crash-free execution

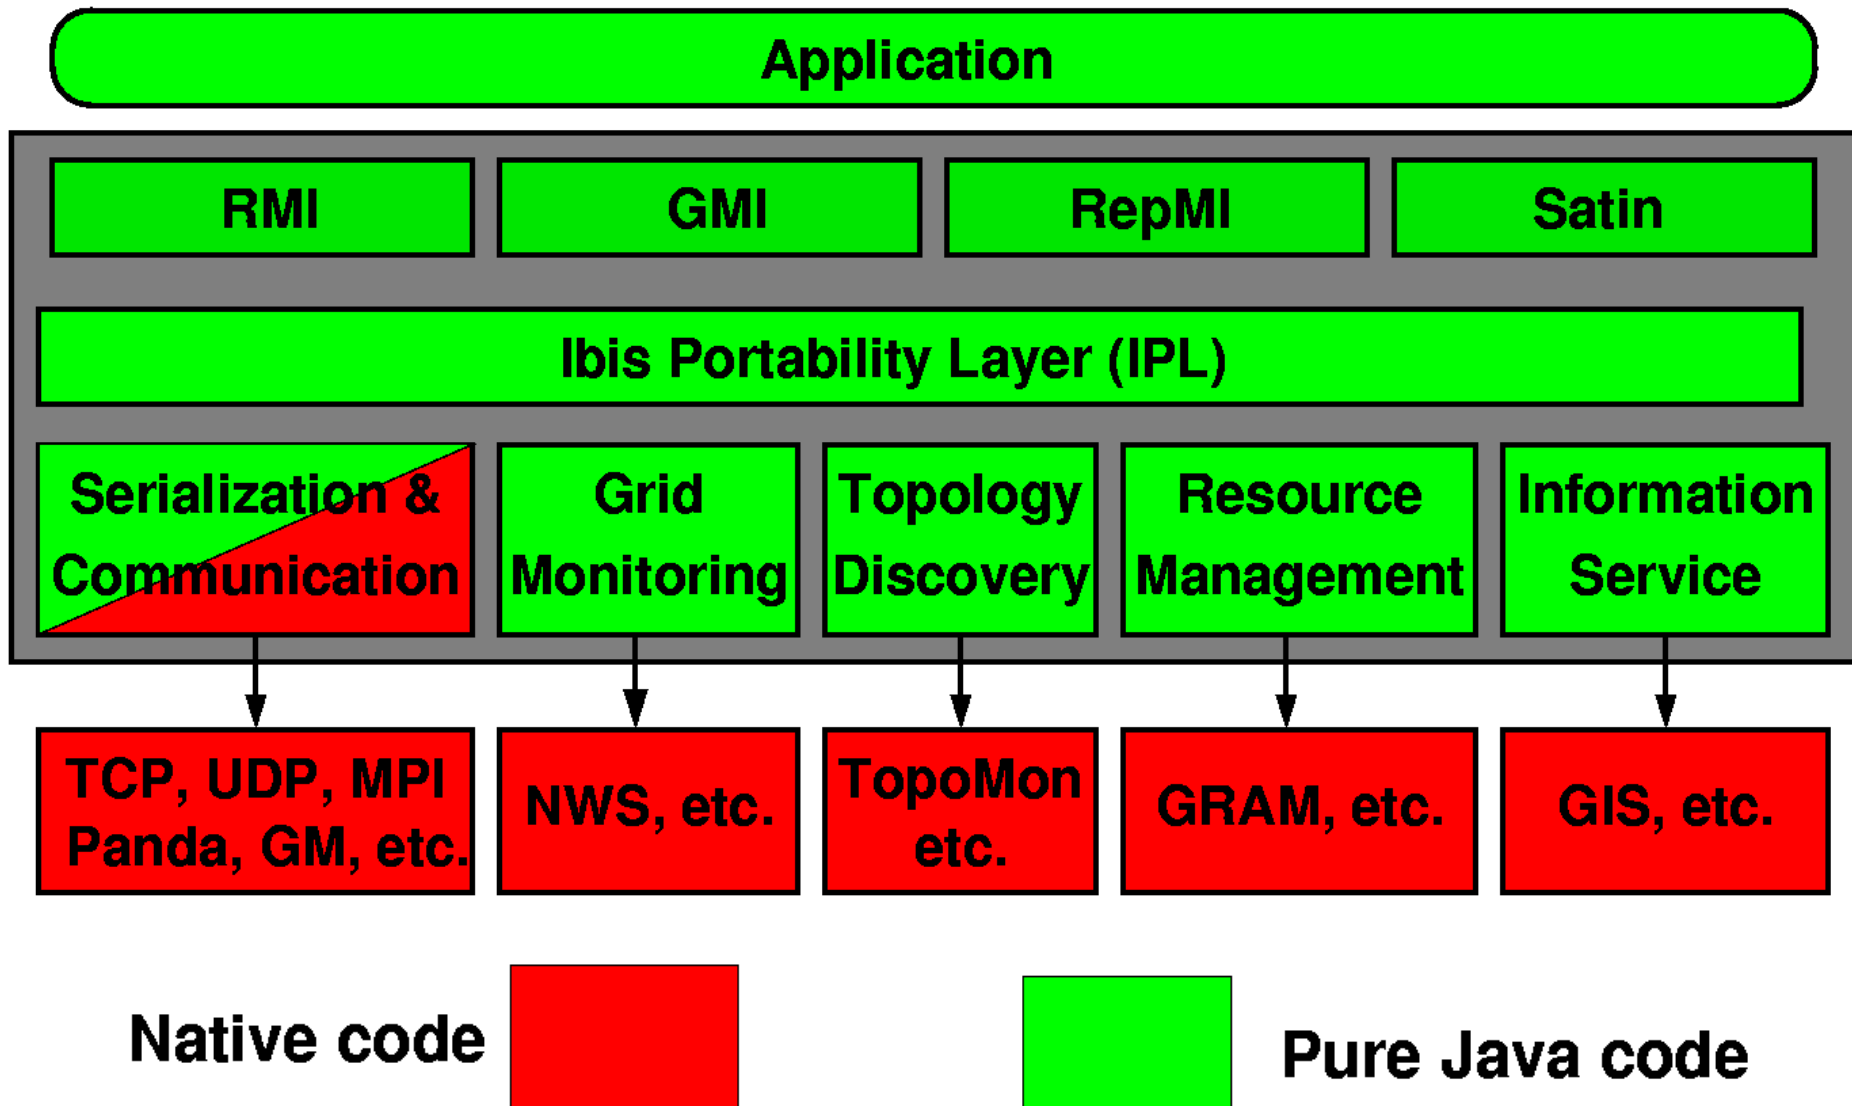
Further information

Publications and a software distribution available at:

<http://www.cs.vu.nl/ibis/>

Additional slides

Ibis design



Partial results on leaving cpus

If processors leave **gracefully**:

- Send all **finished** jobs to another processor
- Treat those jobs as orphans = broadcast (jobID, processorID) tuples
- Execute the normal crash recovery procedure


Job identifiers

- $\text{rootId} = 1$
- $\text{childId} = \text{parentId} * \text{branching_factor} + \text{child_no}$
- Problem: need to know maximal branching factor of the tree
- Solution: strings of bytes, one byte per tree level

Distributed ASCI Supercomputer (DAS) – 2

Node configuration

Dual 1 GHz Pentium-III
>= 1 GB memory
100 Mbit Ethernet +
(Myrinet)
Linux




```
interface FibInter
    extends ibis.satin.Spawnable {
        public int fib(long n);
    }
```

```
class Fib
    extends ibis.satin.SatinObject
    implements FibInter {
        public int fib (int n) {
            if (n < 2) return n;
            int x = fib (n - 1);
            int y = fib (n - 2);
            sync();
            return x + y;
        }
    }
```

Java + divide&conquer

Example



GridLab testbed

Grid results

Program	sites	CPUs	Efficiency
Raytracer	5	40	81 %
SAT-solver	5	28	88 %
Compression	3	22	67 %

- Efficiency based on normalization to single CPU type (1GHz P3)