



Grid Proof Programming Models

Jason Maassen
jason@cs.vu.nl



vrije Universiteit



vl·e

Before the break ...

- We looked at how Ibis makes Grids user friendly:
 - **JavaGAT** provides an easy-to-use API for the various flavours of Grid middleware
 - **Zorilla** provides a configuration free alternative to existing Grid middleware
 - **SmartSockets** provides an easy-to-use library that solves connectivity problems



Remaining problems

- However, grids are still:
 - **Heterogeneous**: many differences in hardware, performance, OS, libraries, etc.
 - **Faulty**: machines may be claimed by others, lose contact, or simply crash
 - **Malleable**: the set of available machines varies constantly
- Therefore we need ways to make applications **Grid Proof**



Grid Proof

- **Heterogeneity:**
 - Solved by using modern languages that run in a managed execution environment:
 - Safer and more portable
 - Easy to deploy and less dependencies
 - no compilation on grid sites
 - no libraries, headers, scripts, compilers, build tools...
 - **Java**, Python, Fortress, C#, ...
 - Language level virtualization
 - Alternative: system level virtualization

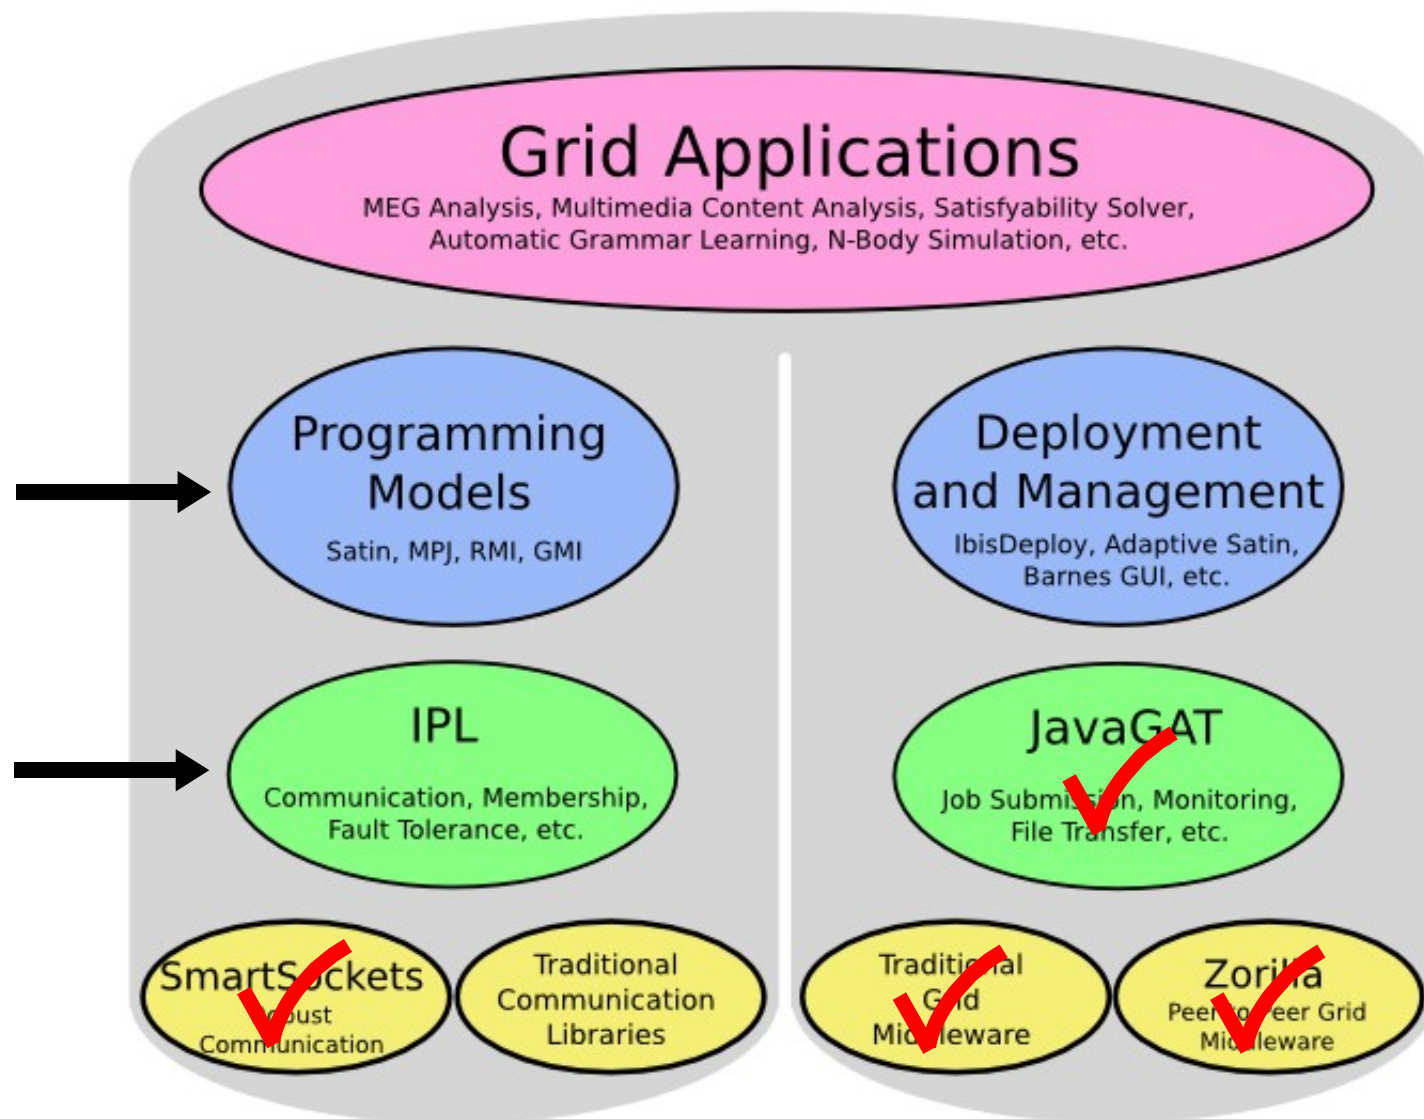


Grid Proof -- cont'd

- **Fault tolerance:**
 - We cannot prevent machines crashing
 - However, we can provide mechanisms to detect crashes ...
 - .. and use those to implement fault-tolerant programming models!
- **Malleability:**
 - Machines joining and leaving cleanly
 - Handled by same mechanism



Overview



Ibis Portability Layer (IPL)

- Simple API for Grid Communication
 - Flexible communication model
 - connection oriented messaging
 - abstract addressing scheme
 - Malleability/Fault Tolerance
 - notifications when machines join/leave
 - open & closed world (not just SMPD)
 - Serialization
 - send bytes, doubles, objects, etc.

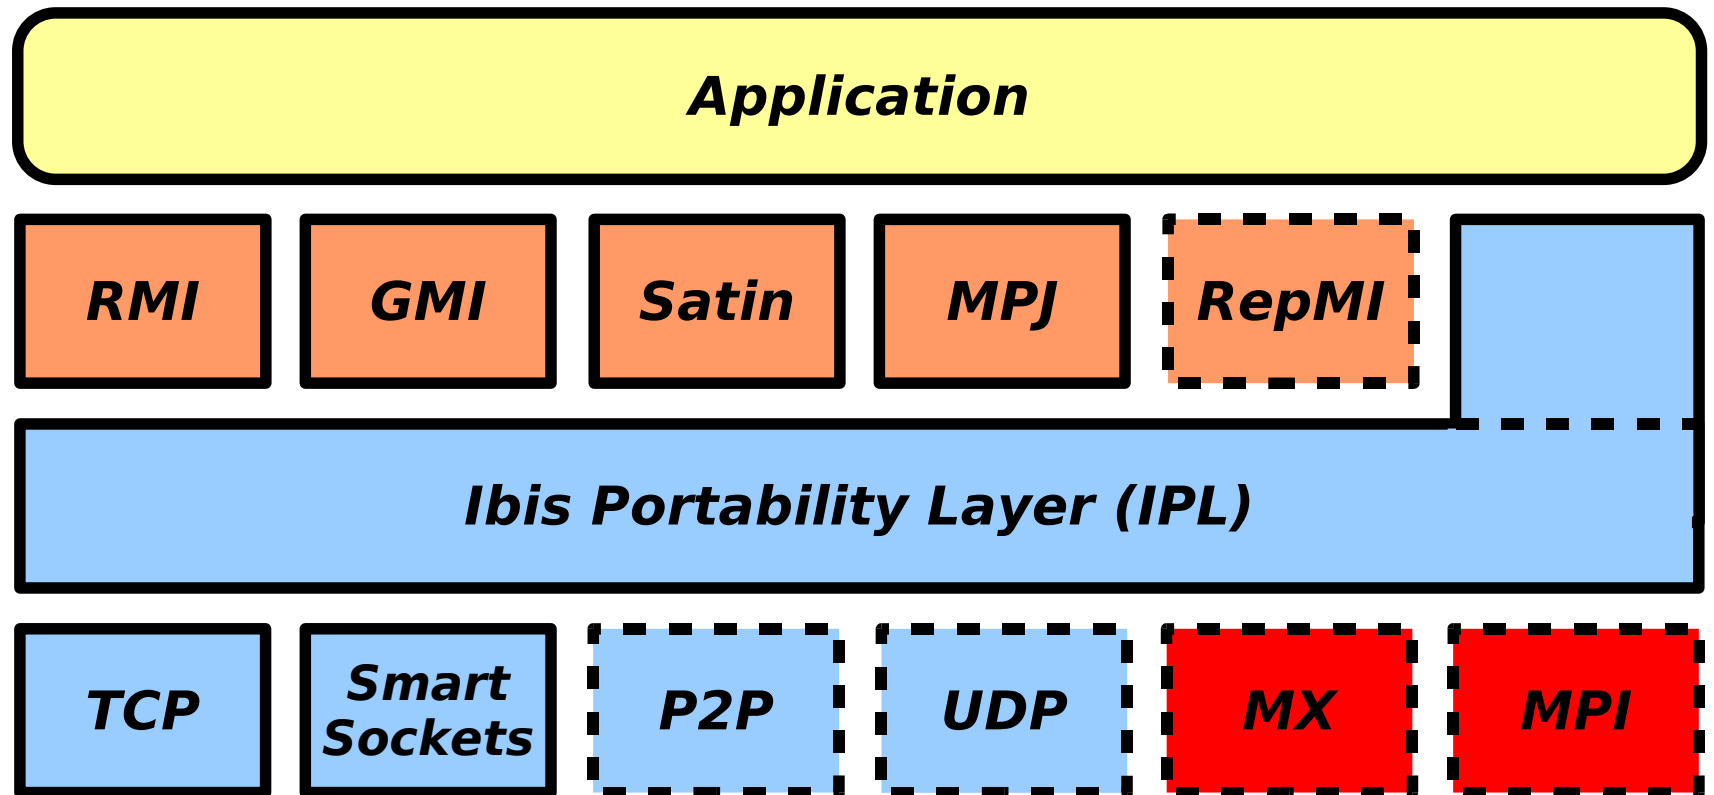


IPL

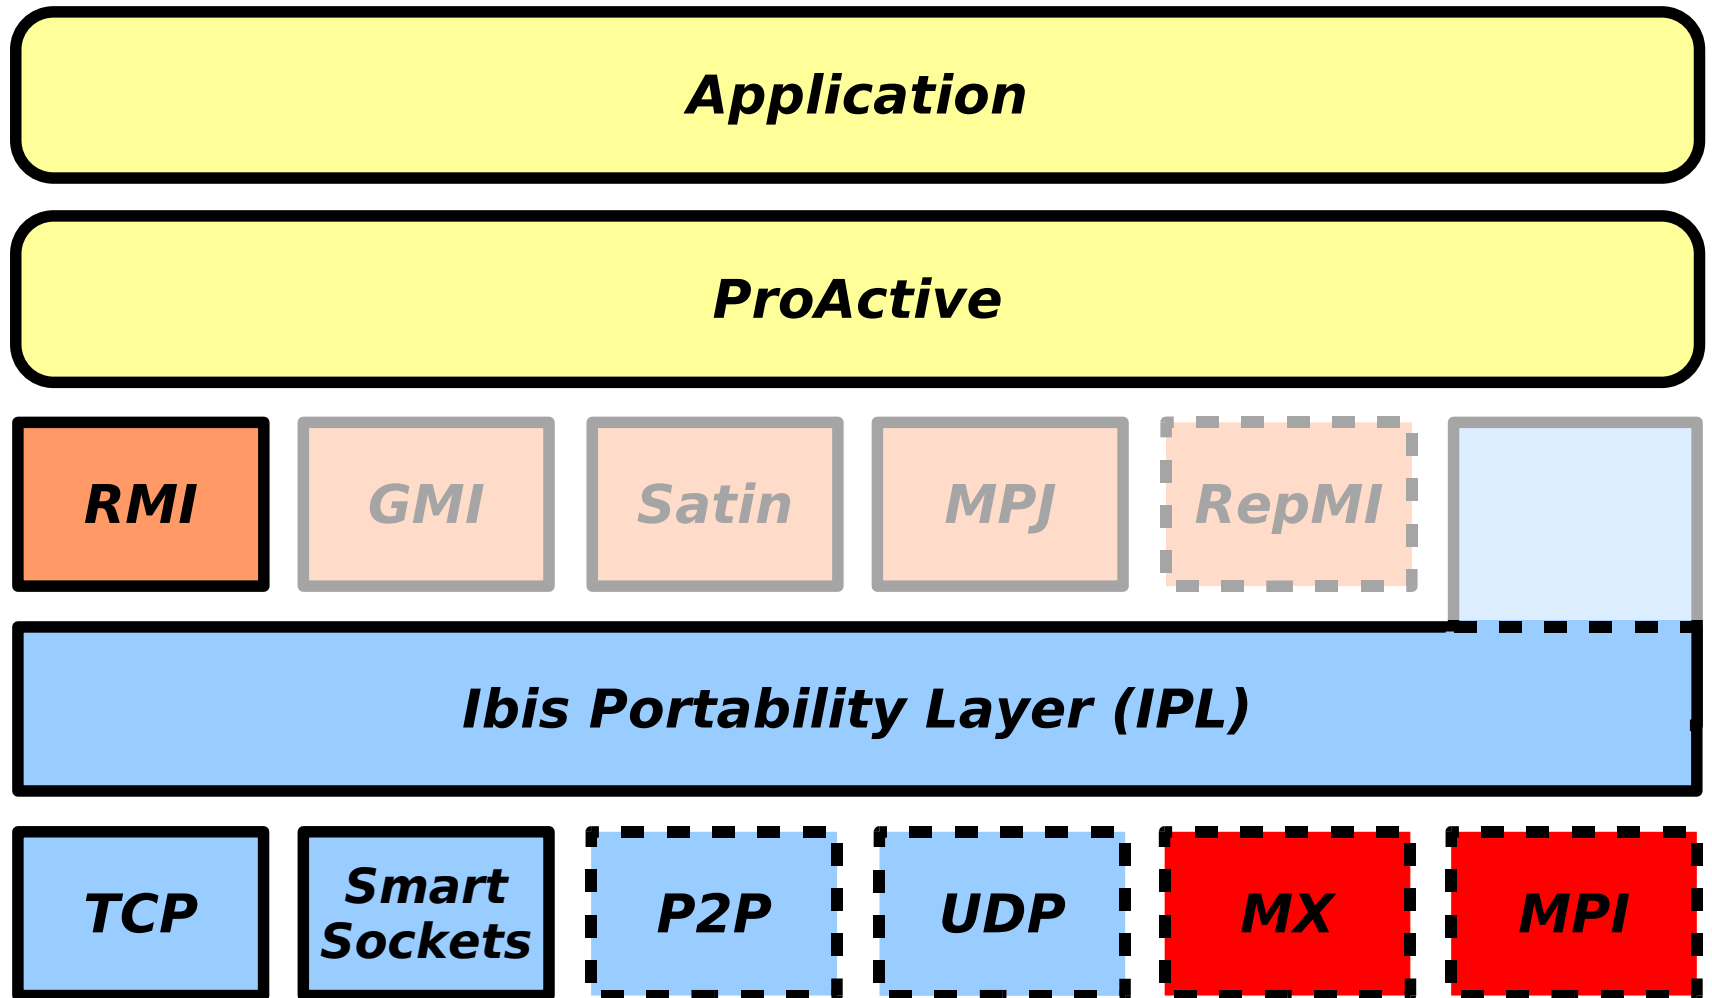
- Clean & abstract API
 - hides network specific details
 - hostnames, IP addresses, MPI ranks, etc.
 - easy to implement on TCP, UDP, MPI, MX...
- Hides network peculiarities
 - Results in more portable applications
 - Suitable for Grids



IPL and Friends



IPL and Friends






Communication

- 'Low-level' communication model
- Unidirectional pipes
- Two end points
- Connection oriented (allows streaming)




Send & receive ports

- Can be connected in arbitrary ways



Send & receive ports

- Simplicity may cause some overhead...
 - Example: need two pairs for RPC / RMI




Port Types

- All ports have a **type**
 - Consists of a set of required capabilities:
 - Connection patterns
 - Unicast, many-to-one, one-to-many, many-to-many.
 - Communication properties:
 - Fifo ordering, numbering, reliability.
 - Serialization properties:
 - bytes, data, object
 - Message delivery:
 - Explicit receipt, automatic upcalls, polling



Port Types

- Defined at runtime
 - Specify set of capabilities
- Types must match when connecting!



Port Types

- Forces programmer to specify how each communication channel is used
 - Prevents bugs
 - Exception when contract is breached
- Allows efficient impl. to be selected
 - Unicast only ?
 - Bytes only ?
 - Can save a lot complexity!



IbisIdentifiers

- In a parallel/distributed application
 - Each process has an Ibis instance
 - Each instance has an **IbisIdentifier**
- IbisIdentifier:
 - Uniquely identifies an Ibis instance
 - Abstracts away from the implementation
 - e.g. hostnames, IP addresses, MPI-ranks, etc.
 - Makes your application a bit more portable



Connection setup WEG ?


- Two options:
 - 1) Using a IbisIdentifier and a name
 - Name specifies the receiveport
 - Unique per Ibis instance
 - Human-readable (usually)
 - 2) Using a ReceivePortIdentifier
 - Uniquely identifies a receiveport
 - Created when ReceivePort is created
 - Can be passed around between Ibis instances.



Connection setup (1)




ibis-az33zx7




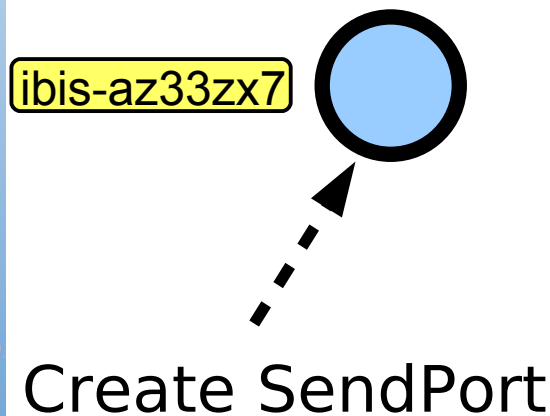
Create Ibis

ibis-34fdw21

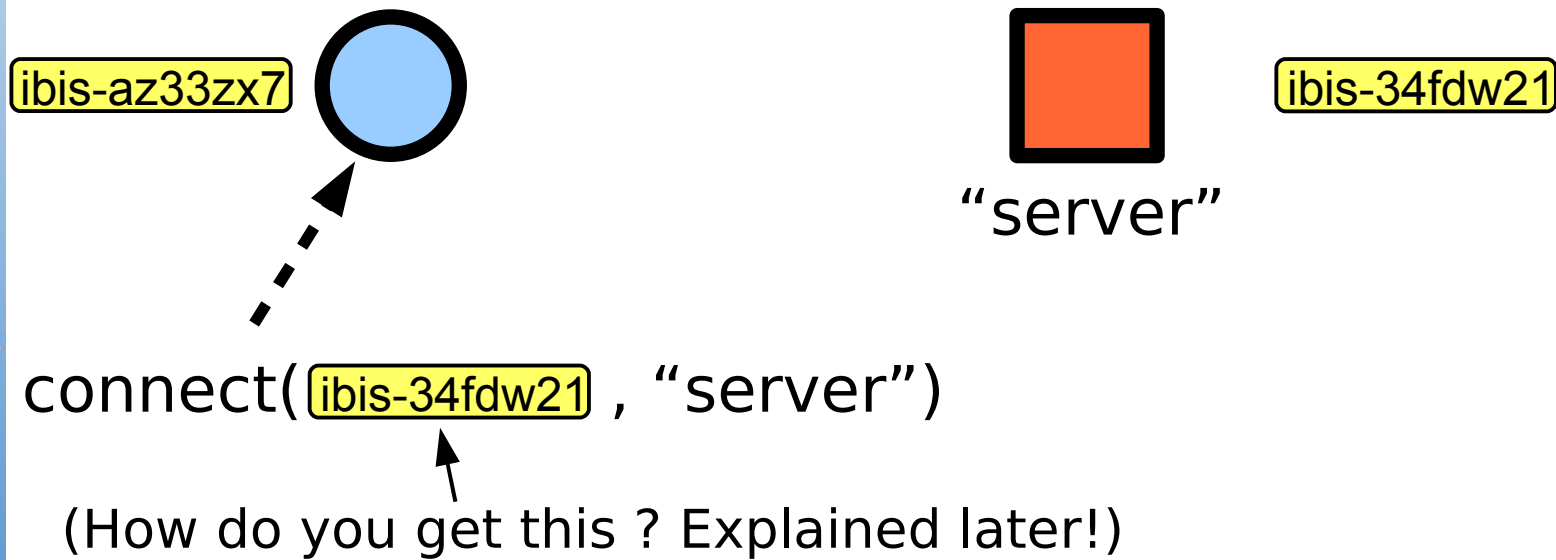


Create Ibis

Connection setup (1)



Connection setup (1)




Connection setup (1)




Connection setup (2)

Create anonymous
ReceivePort


Create SendPort




Connection setup (2)




Connection setup (2)



Connection setup (2)



Connection setup (2)



Messages

- Ports communicate using 'messages'
- Contain read or write methods for
 - Primitive types (byte, int, ...)
 - Object
 - Arrays slices (partial write / read in place)
- Unlimited message size




Ibis Serialization

- Based on bytecode-rewriting
 - Adds serialization and deserialization code to serializable types
 - Prevents reflection overhead during (de-)serialization
 - Has fallback mechanism for non-rewritten classes
- Experimented with runtime rewriting




Messages

- Get WriteMessage from SendPort




Messages

- Write data into WriteMessage




Messages

- Finish the WriteMessage




Messages

- Data is send to ReceivePort




Messages

- ReceivePort produces ReadMessage
 - Explicit receive or callback (upcall)




Messages

- Read data from ReadMessage



Messages

- Finish the ReadMessage




Messages

- Done!



Messages or streams ?

- Message size is unlimited
 - Data may be forwarded at any time
 - Both S. & R. messages alive at same time
 - There's streaming!



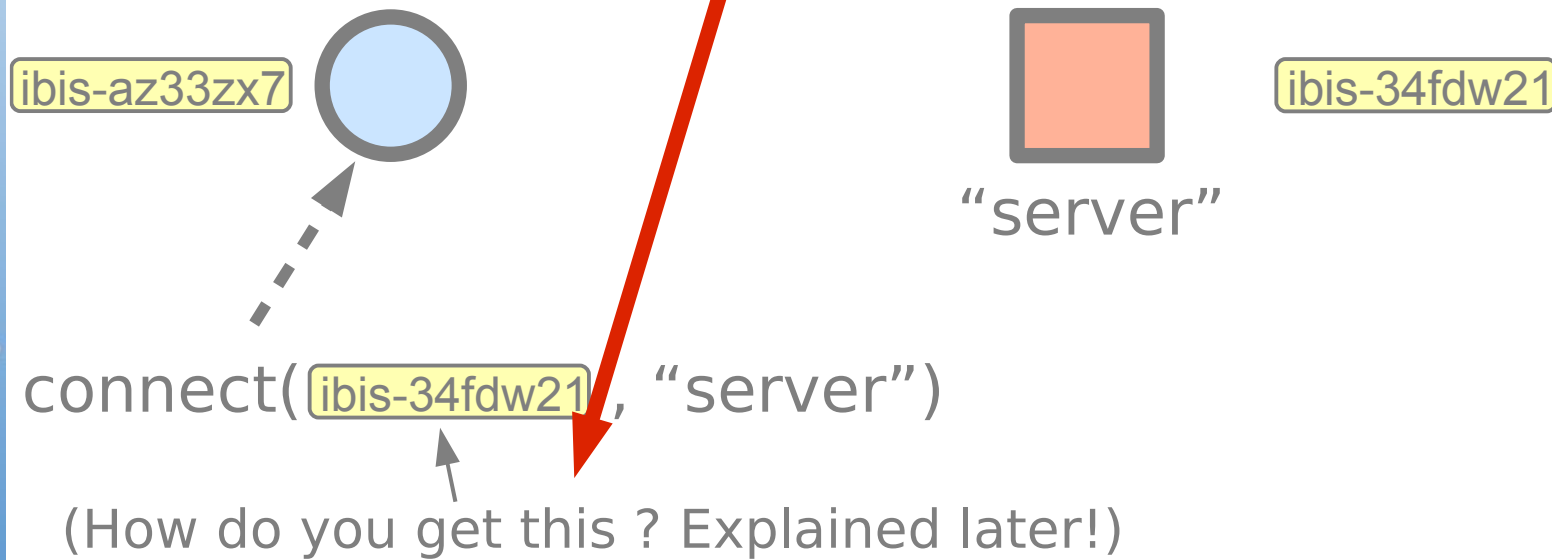
Short Recap

- First create PortType
- PortType creates Send & ReceivePort
 - Type is checked when connecting
- Several ways to connect
 - Abstract addressing
- Use Messages to communicate
 - Allows streaming
 - 3/4 types of serialization



Connection setup (1)

Remember this question ?

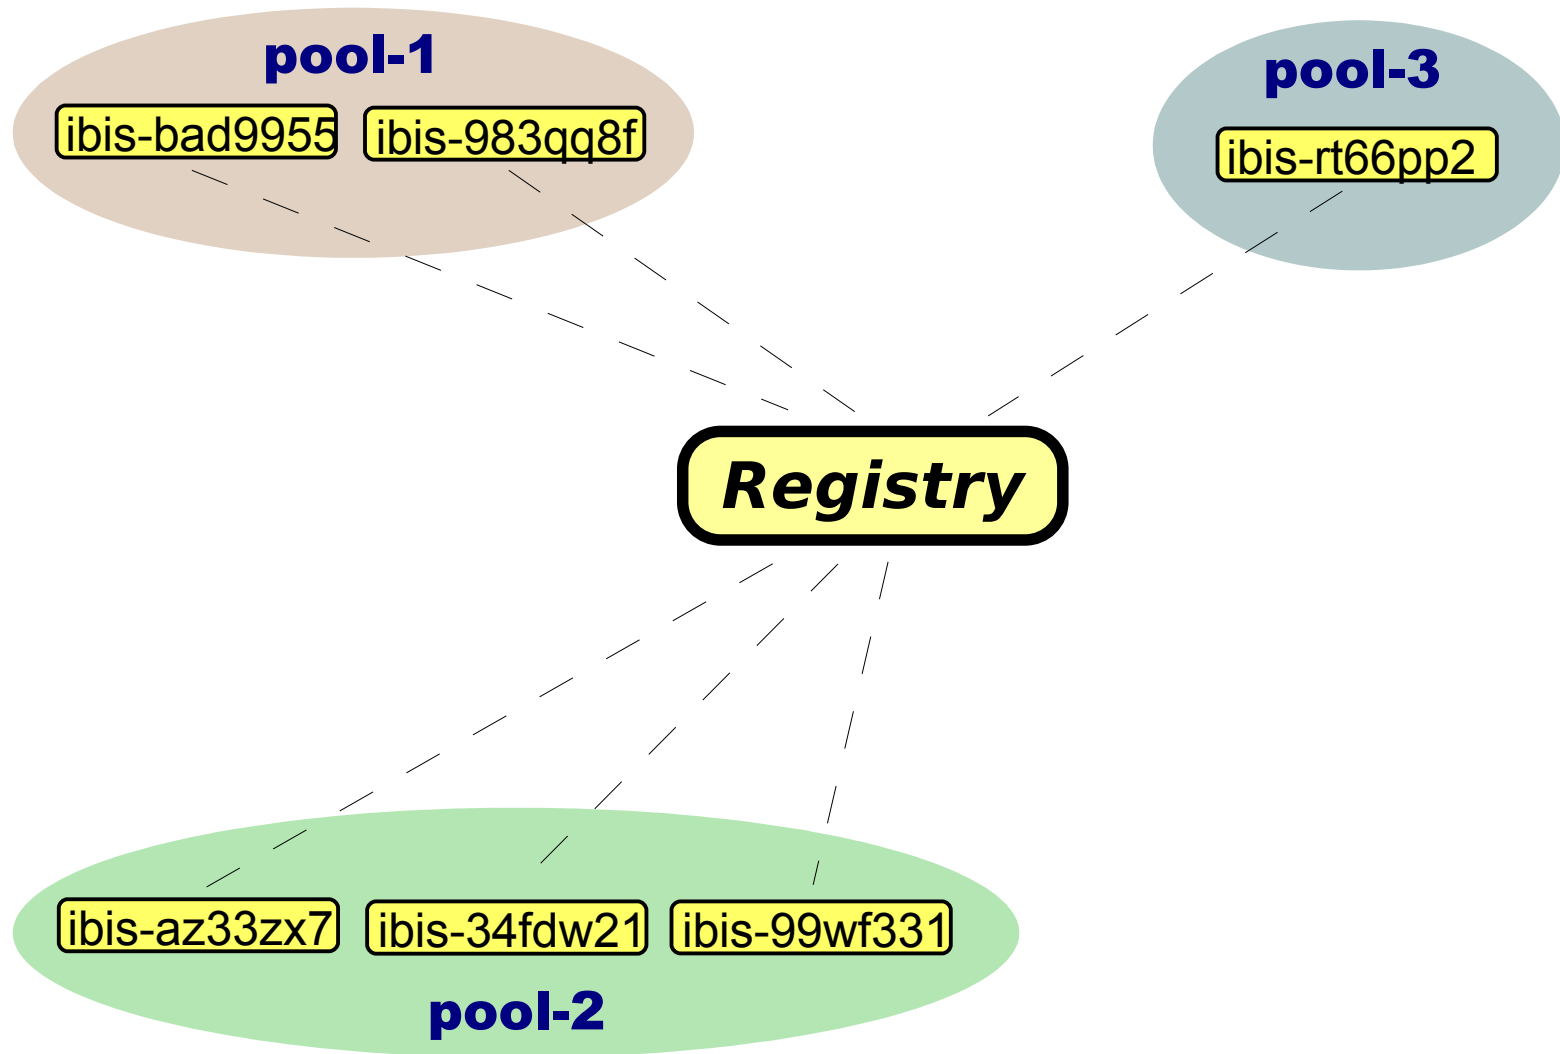


Join Elect Leave (JEL)

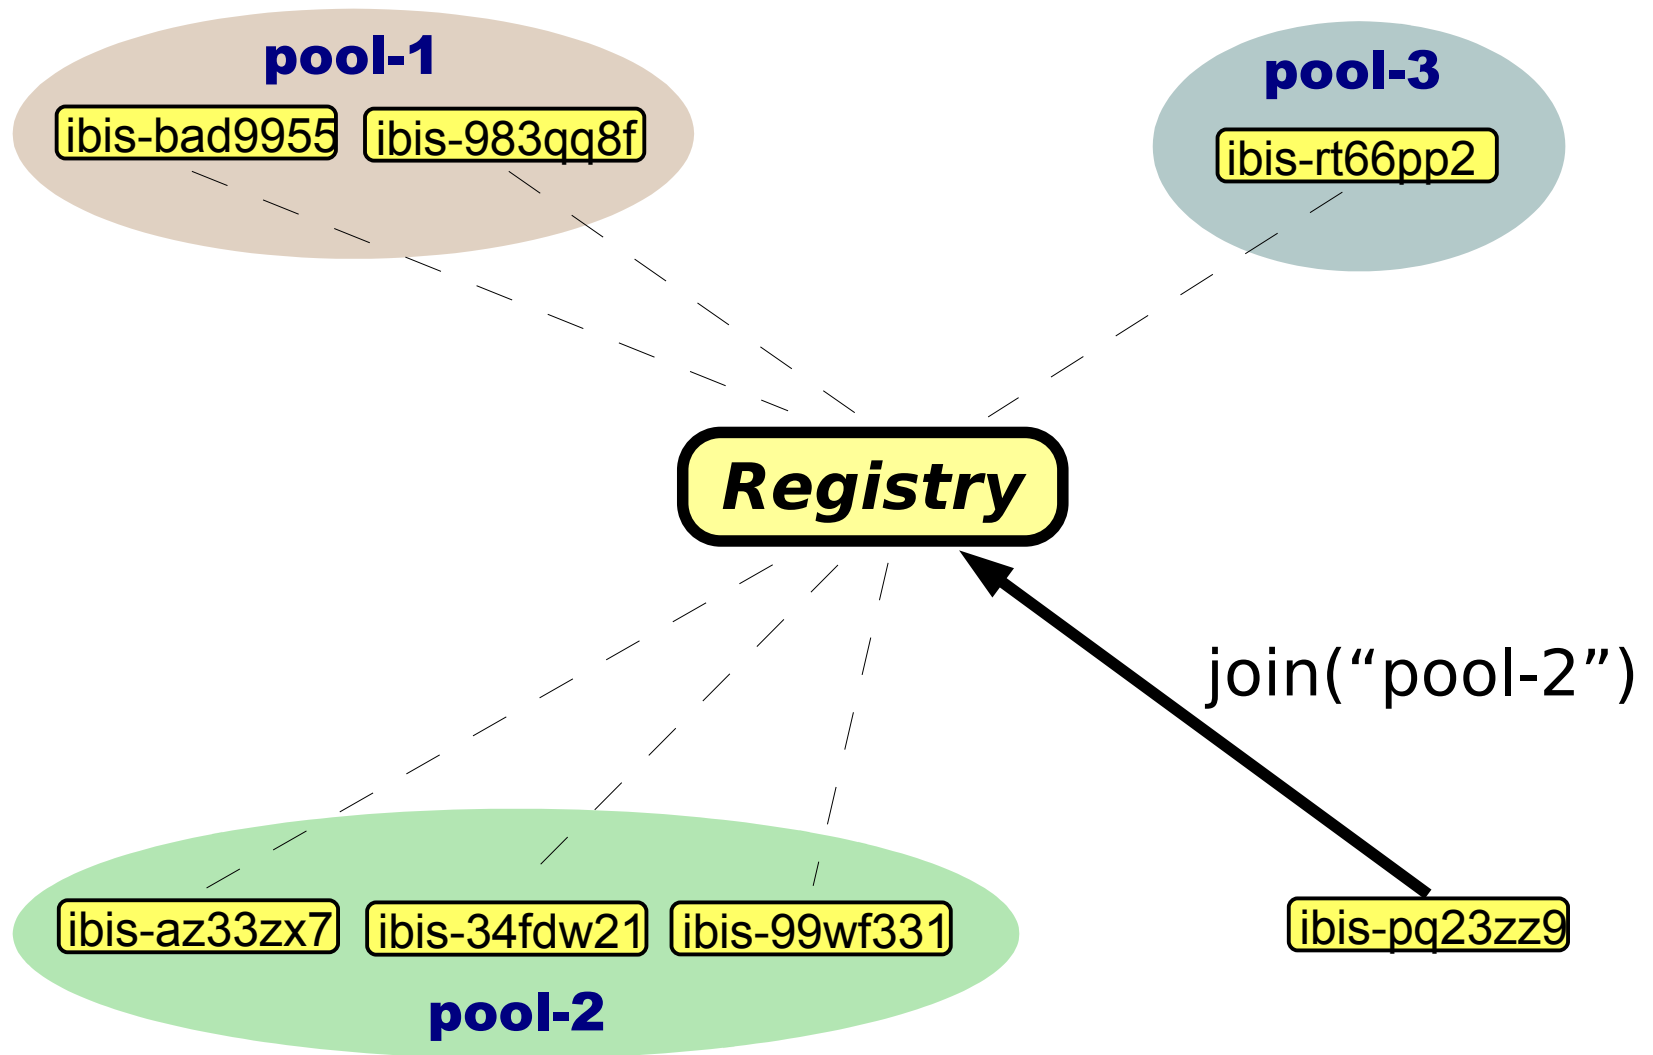
- Membership information
 - Can subscribe to information
 - Updates when Ibis instances join or leave
 - Useful for determining who's participating
 - Also used for fault-tolerance
- Ibis instances are part of a **pool**
 - Either variable size or fixed (create-once)
 - Fixed used by 'legacy' MPI-type applications



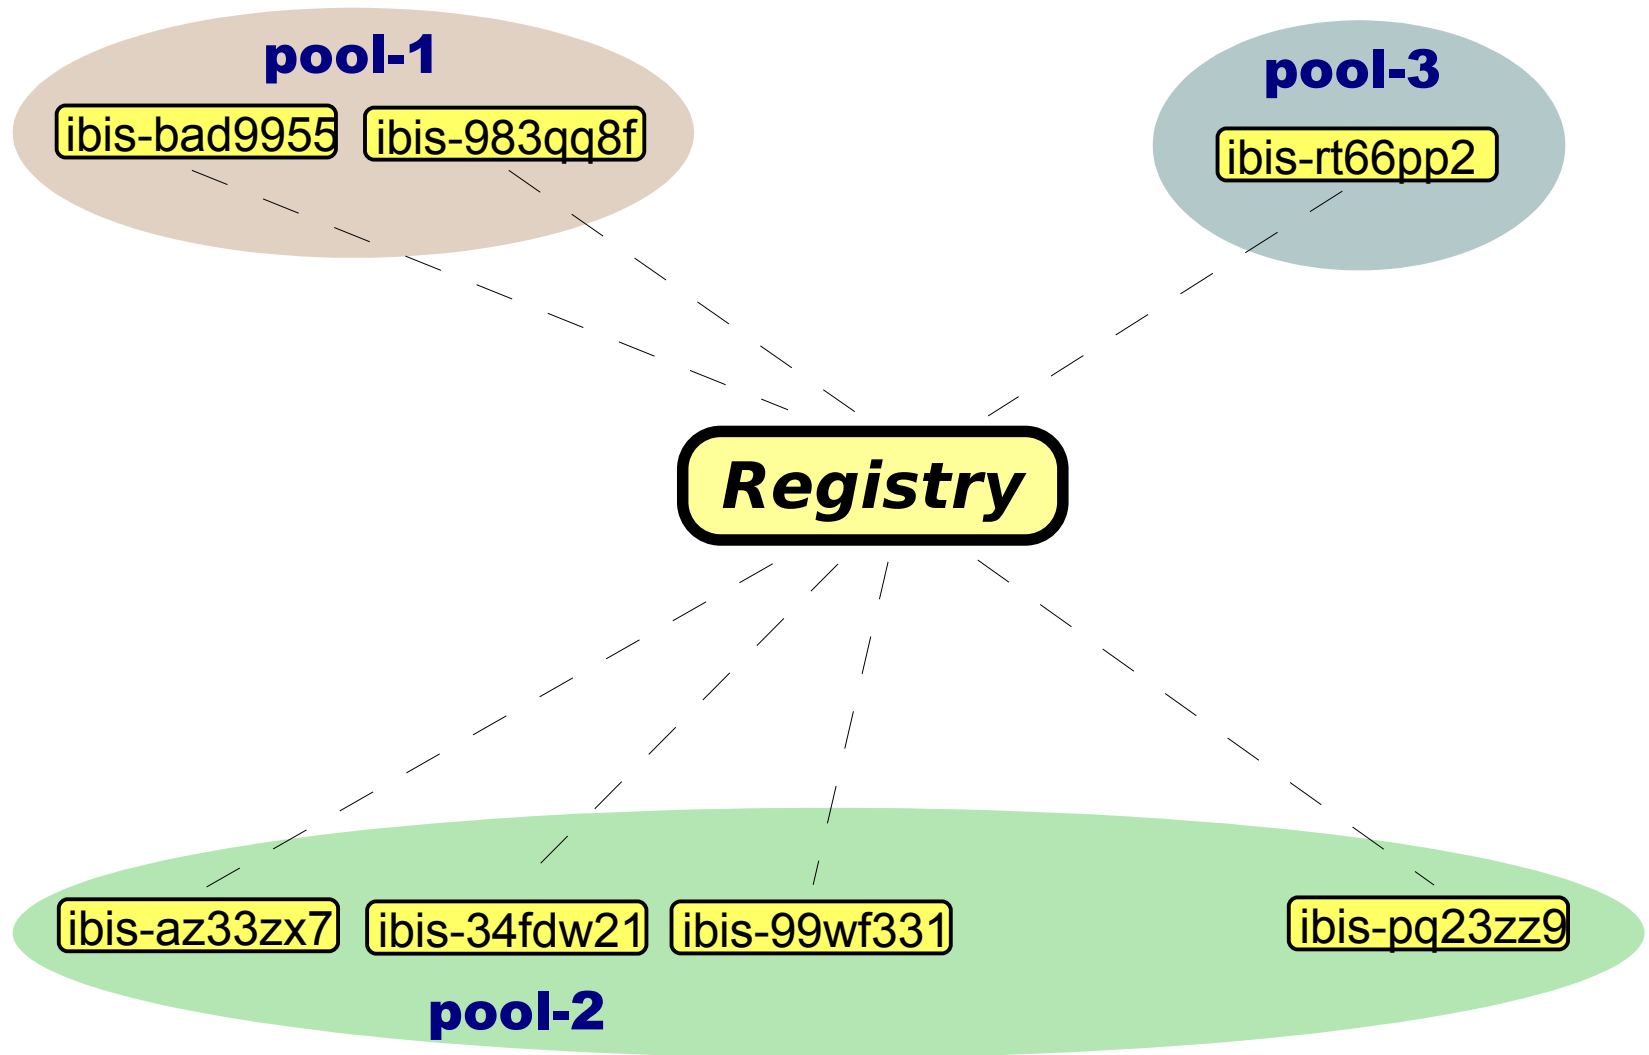
Pools & Malleability



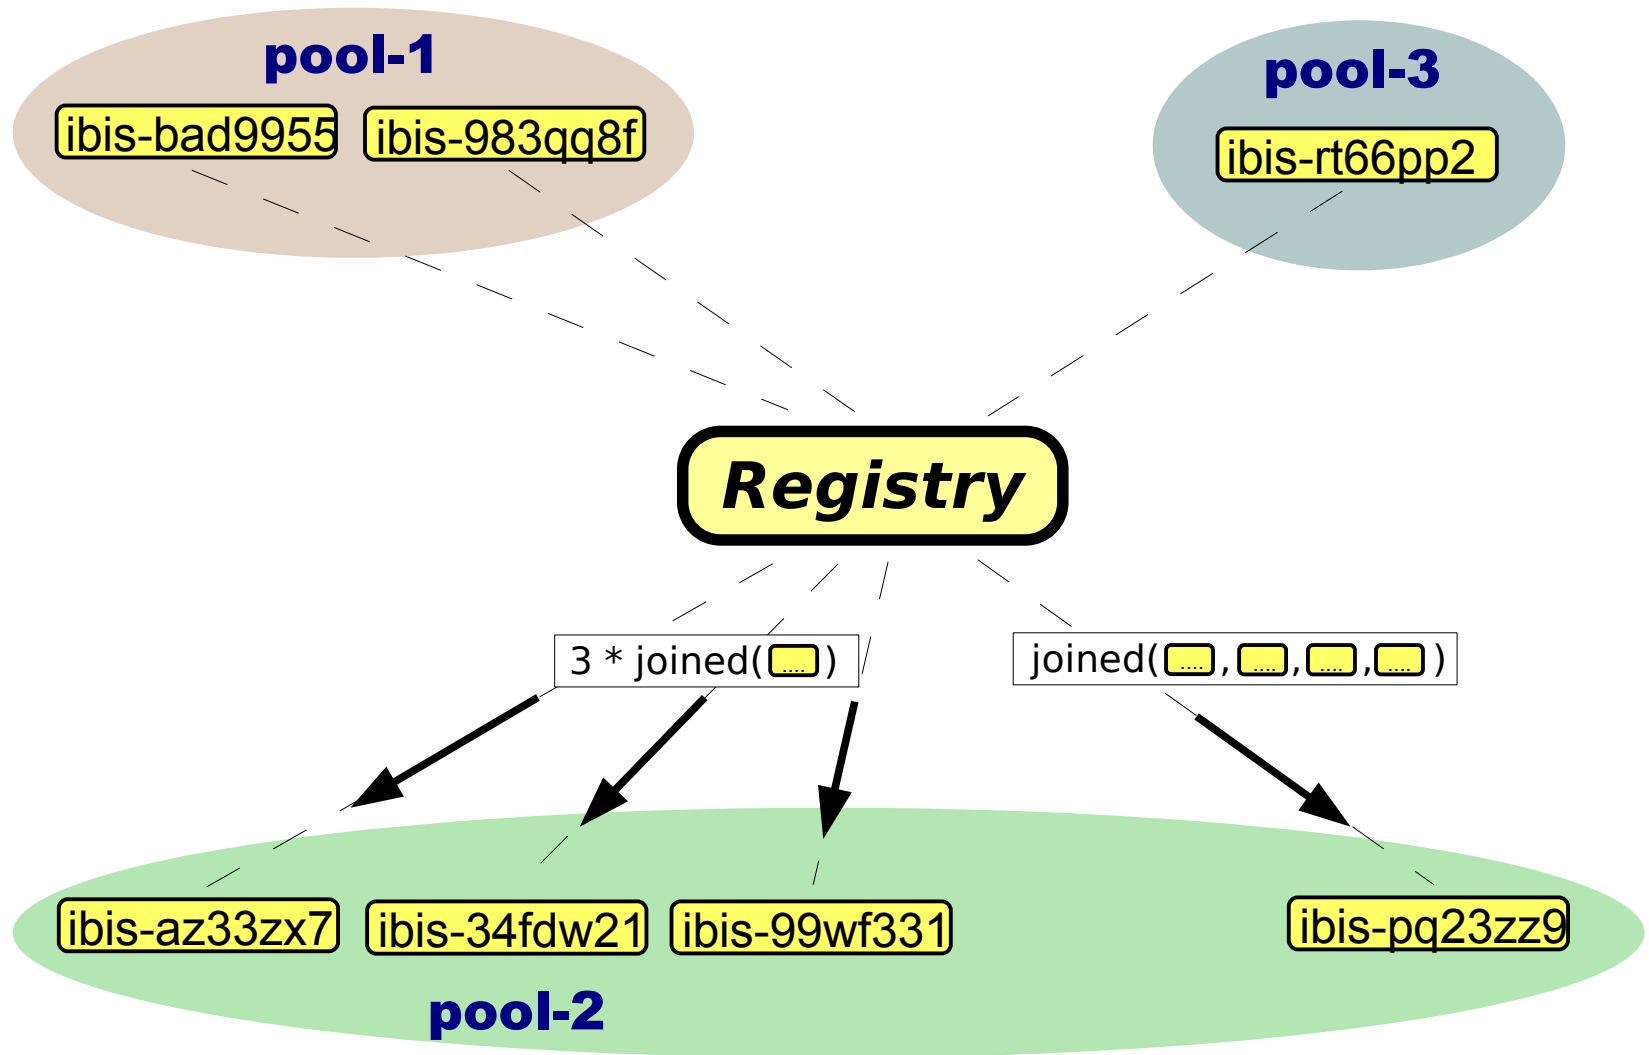
Pools & Malleability



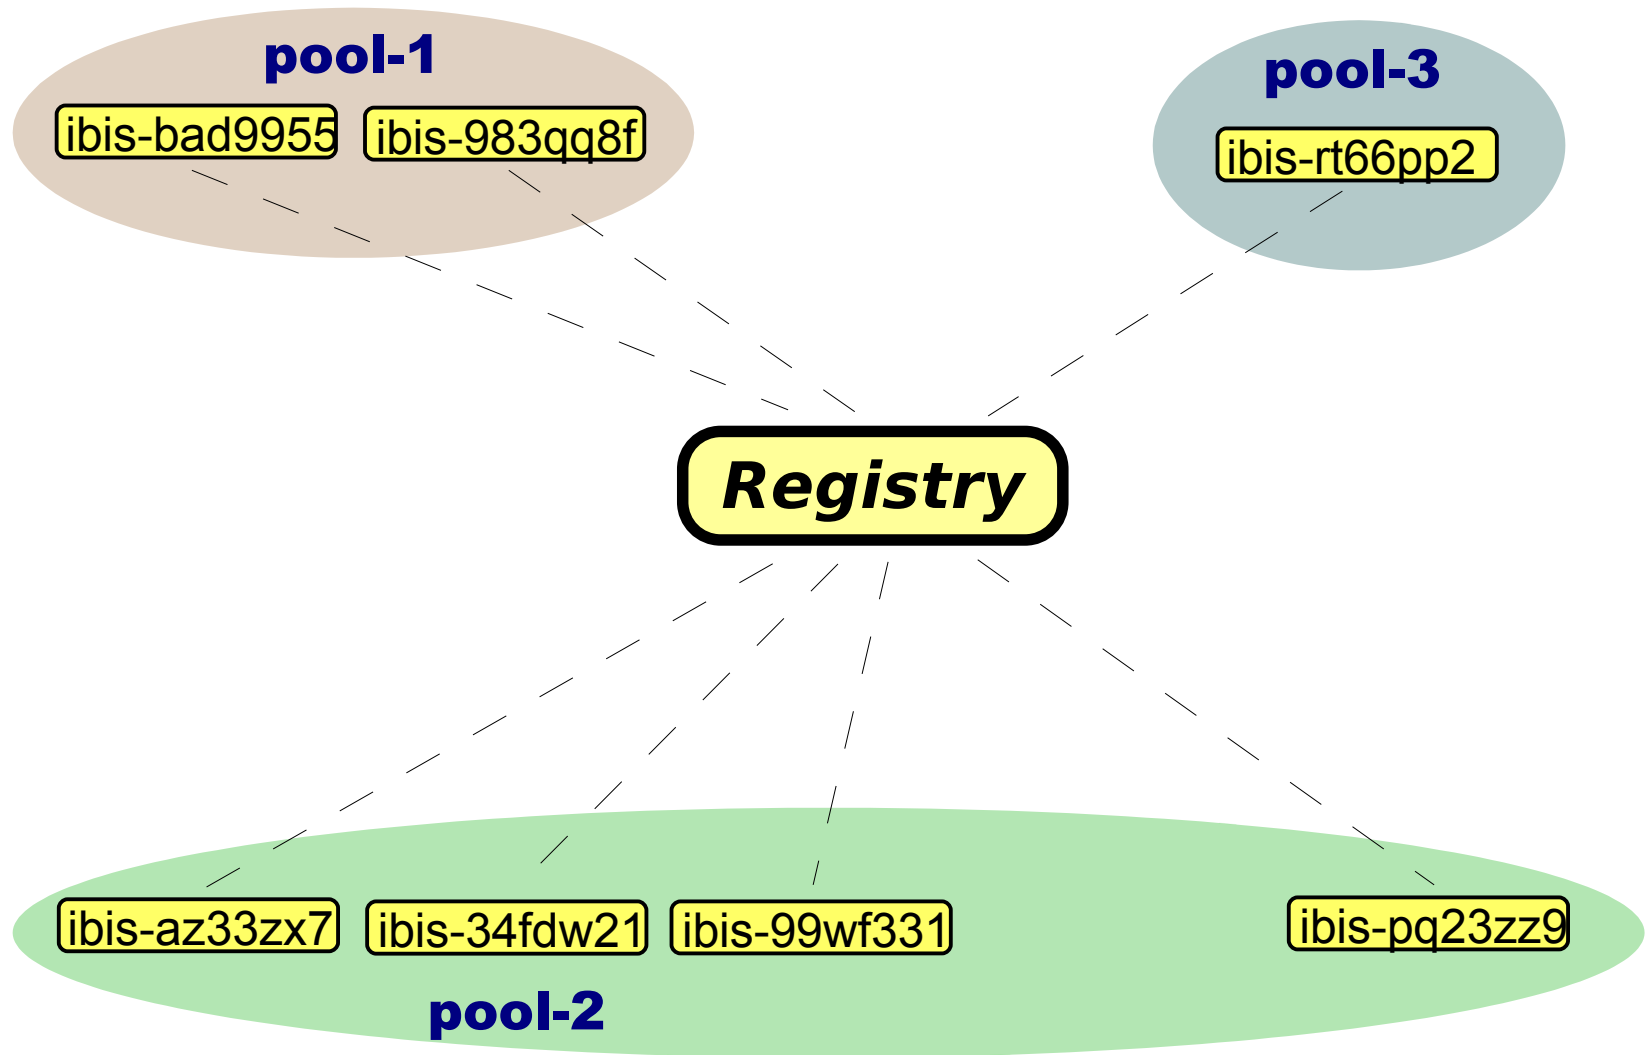
Pools & Malleability



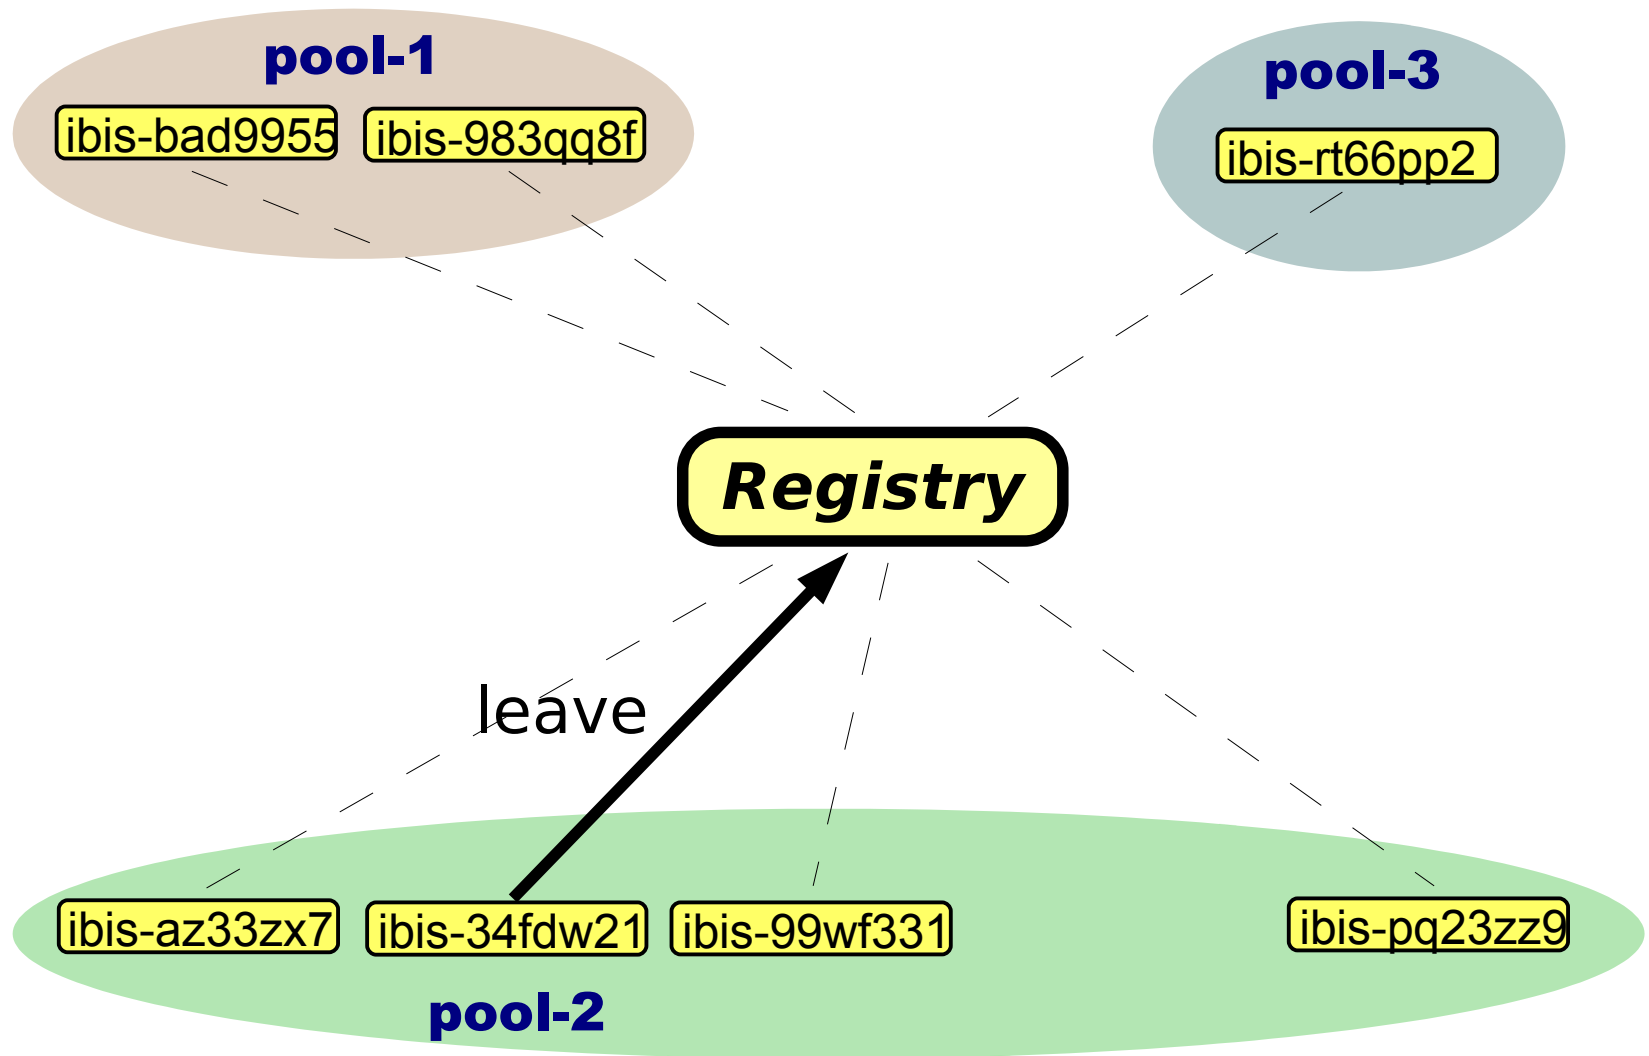
Pools & Malleability



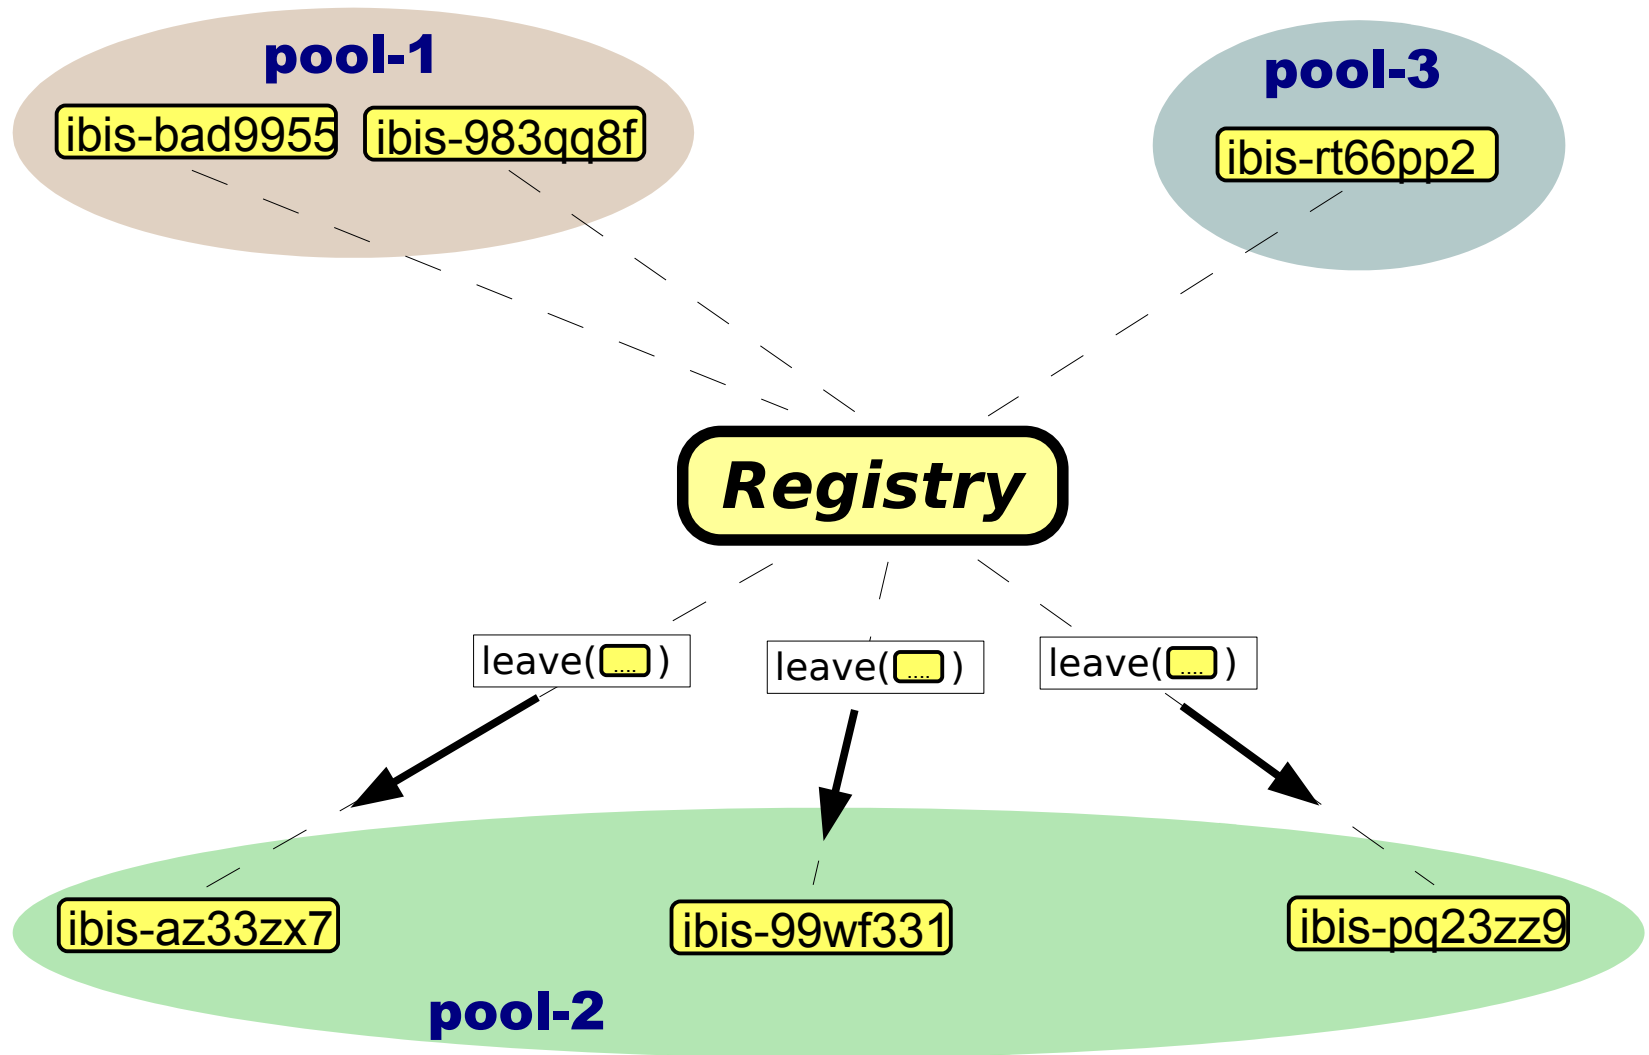
Pools & Malleability



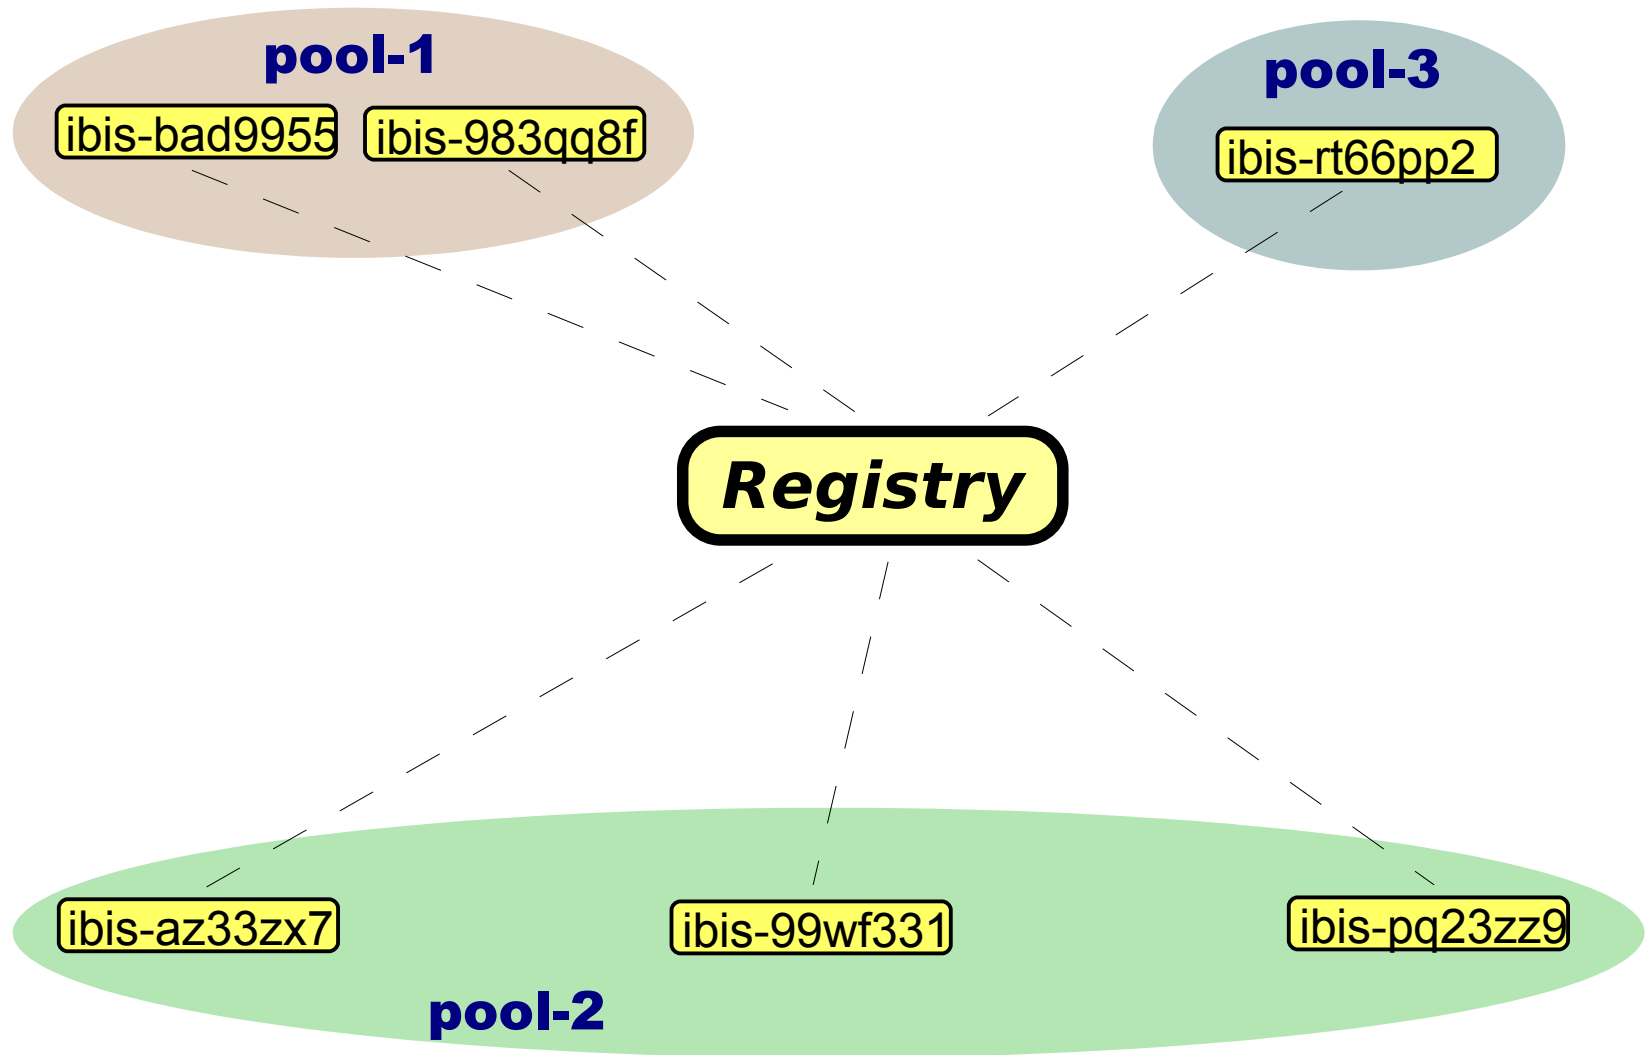
Pools & Malleability



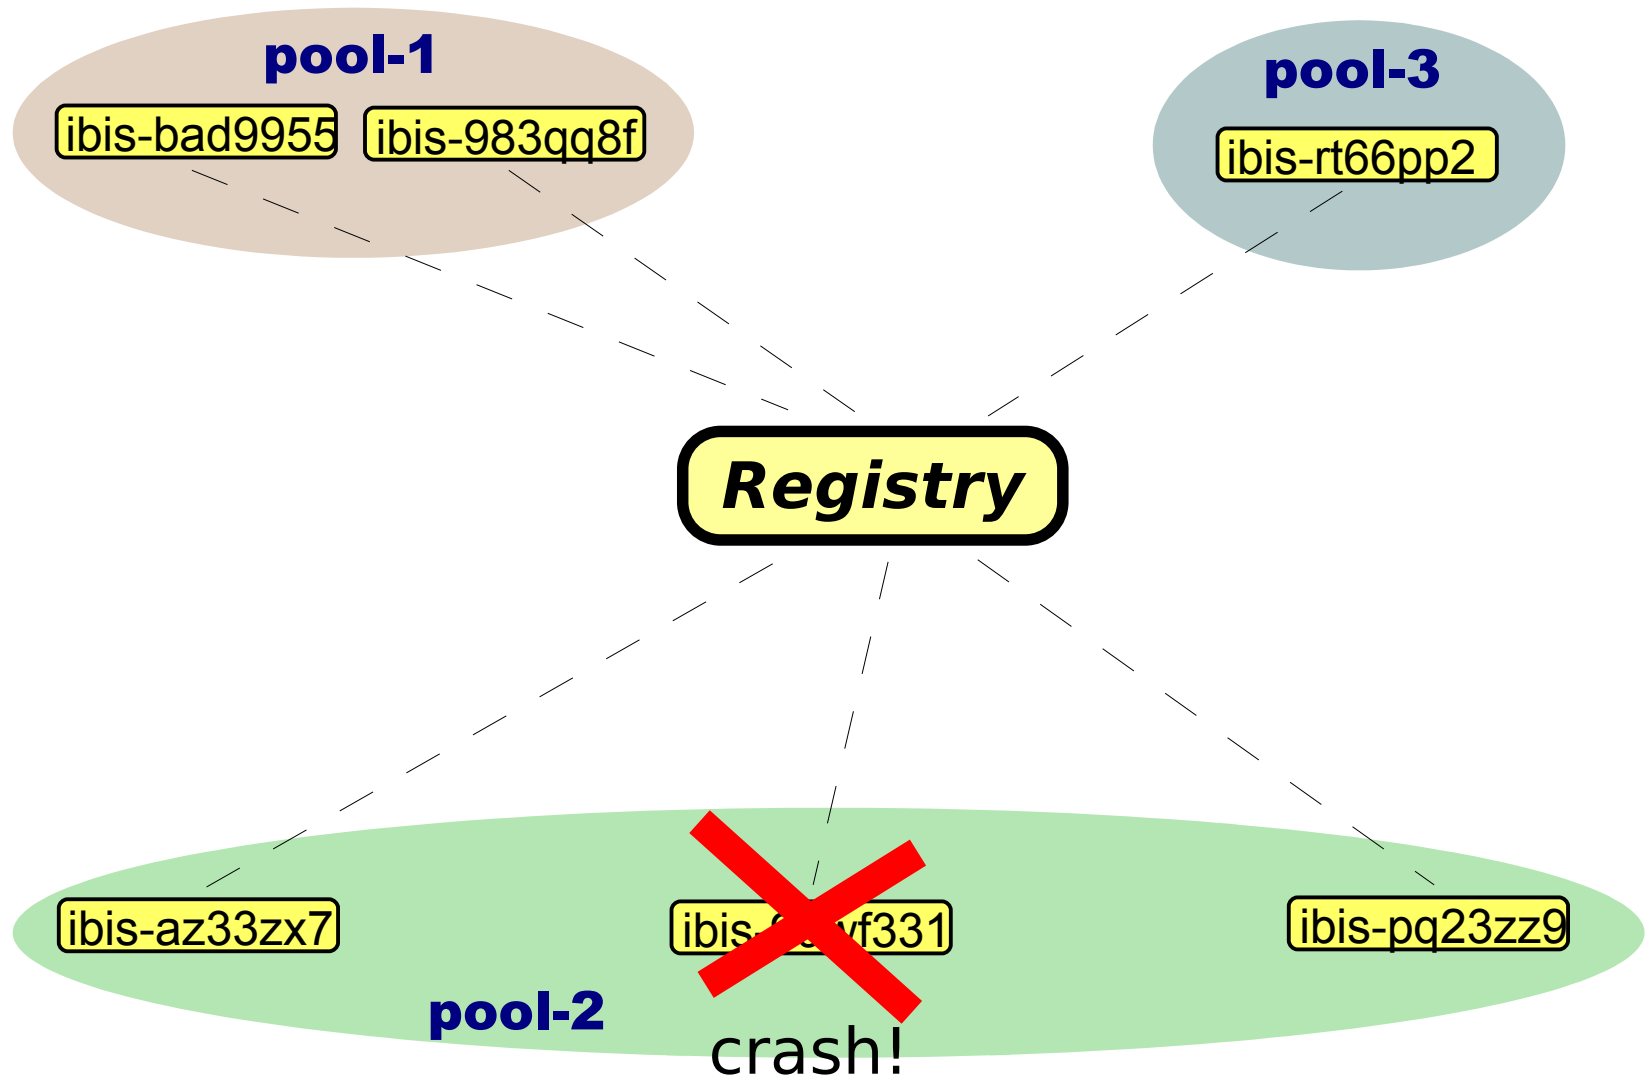
Pools & Malleability



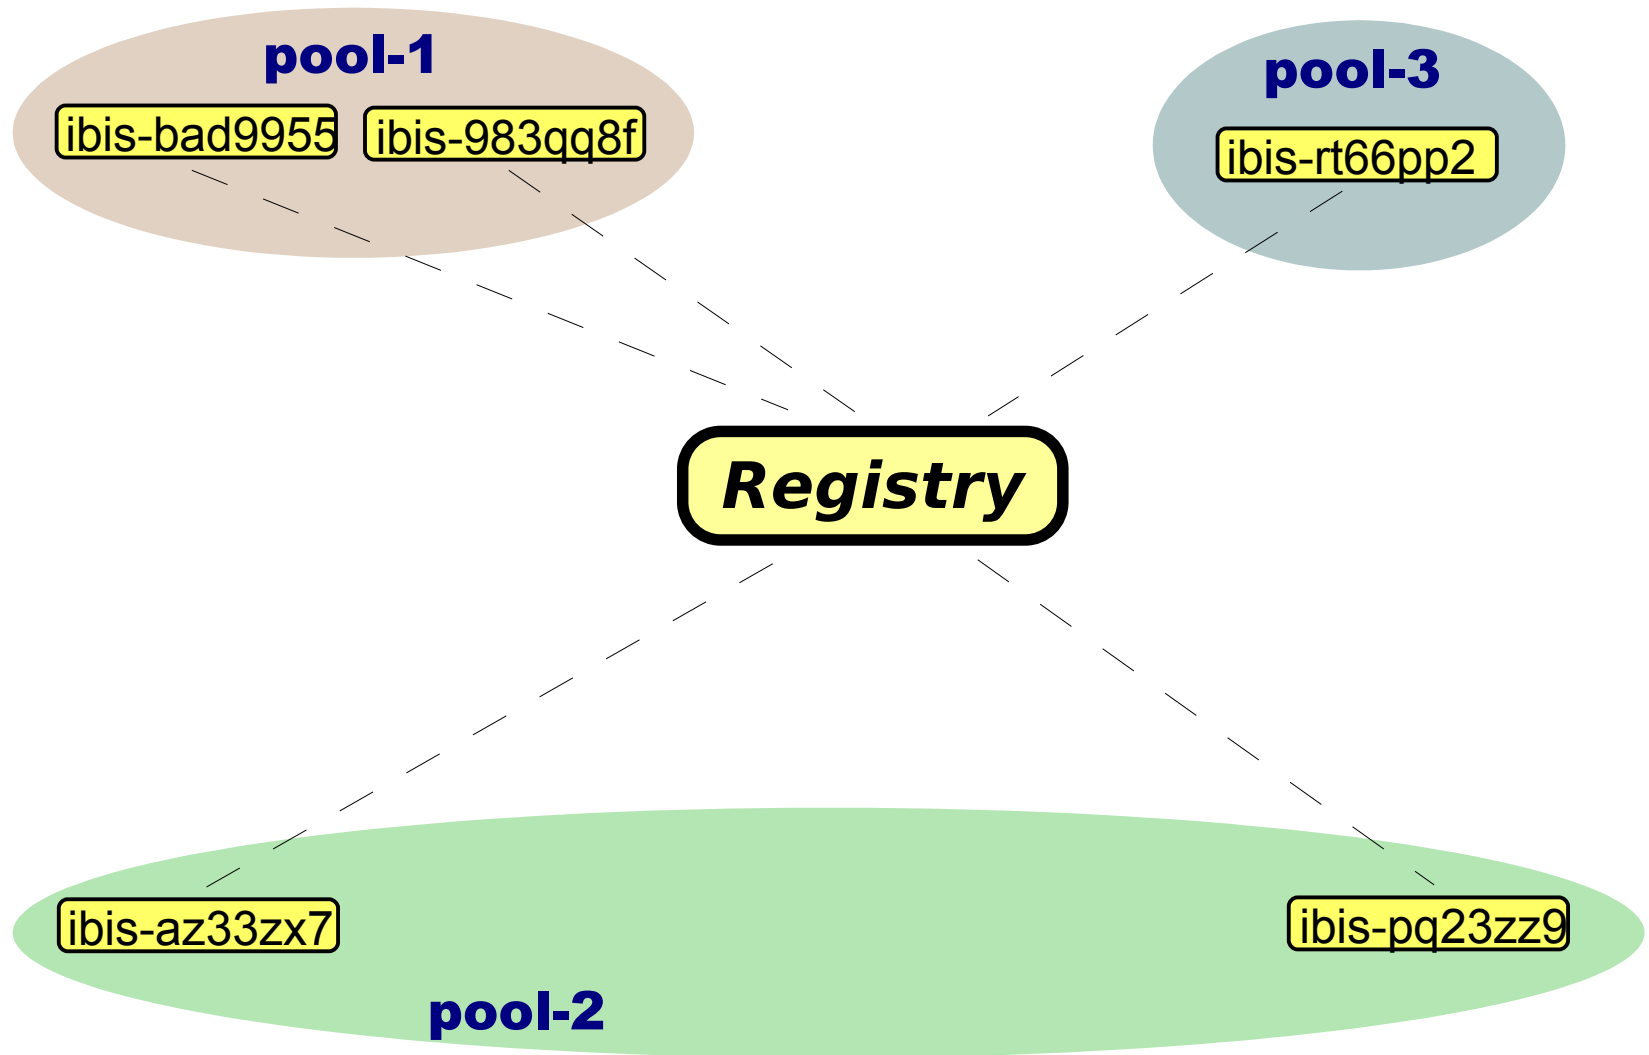
Pools & Malleability



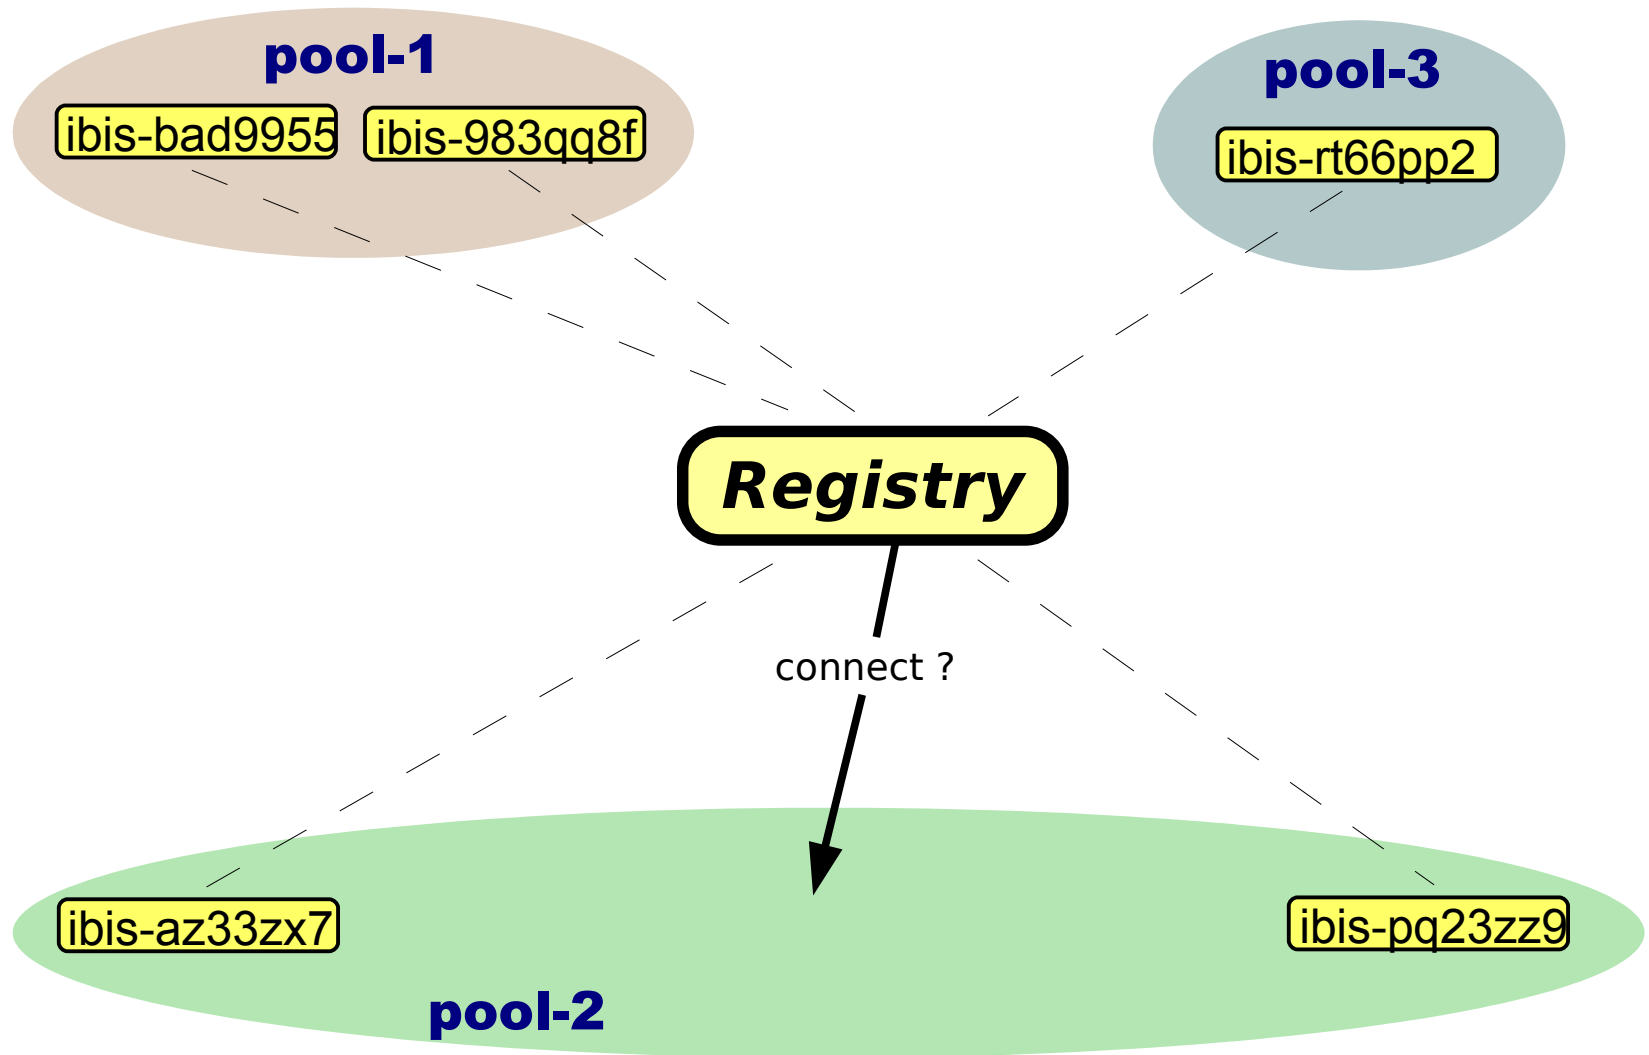
Pools & Malleability



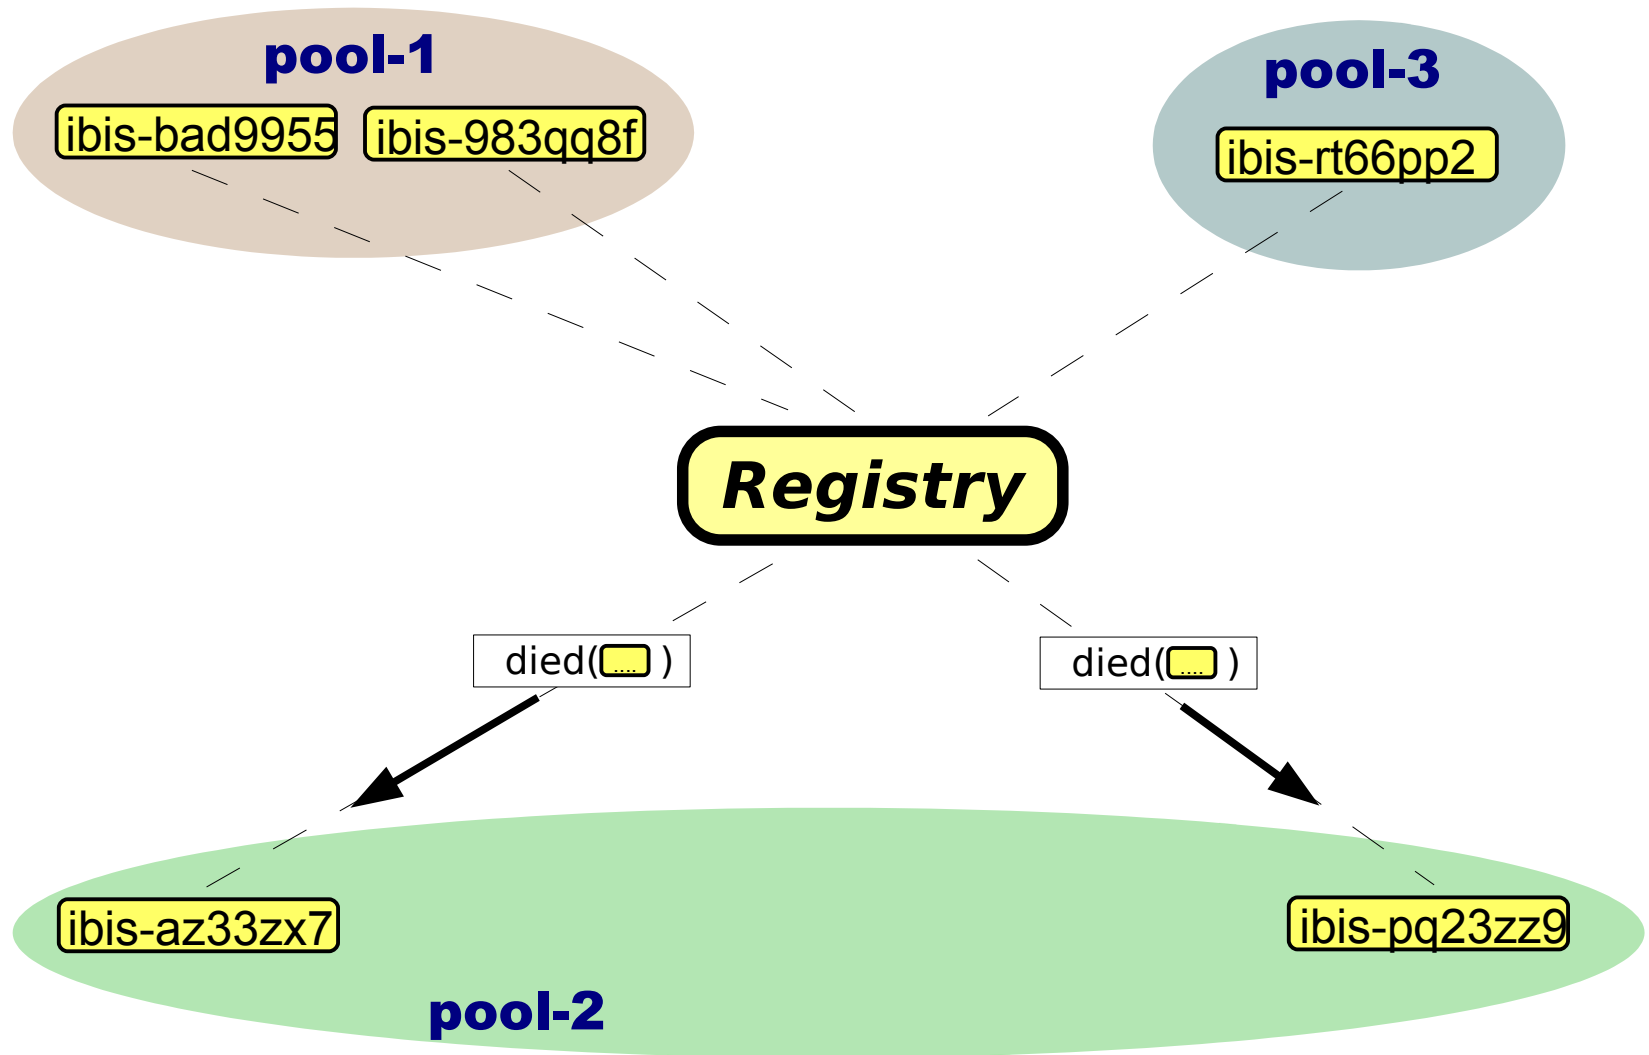
Pools & Malleability




Pools & Malleability



Pools & Malleability



Pools & Malleability



Elections

- Registry offers an 'election' mechanism
 - Allows a group to determine who's **special**
- Each election
 - Has a name (String)
 - Produces IbisIdentifier of the winner
 - Is not democratic
 - You can also be 'an observer'



Registry

- Example shows centralized version
 - also have broadcast tree and gossiping implementations (improve scalability)
- You can select the functionality and consistency that is needed
 - reducing functionality or consistency further improves scalability



Summary

- IPL offers an abstract model
 - Connection oriented message passing
 - Hides network details (for portability)
- Supports fault tolerance / malleability
 - **No application-level fault tolerance!**
 - Only offers the means to implement this!
- Higher level models (Satin) do offer this at application level

