



Ibis as Glue

Jason Maassen

Computer Systems Group
Department of Computer Science
VU University, Amsterdam, The Netherlands



Short recap...

- JavaGAT acts as **master key** (or passepartout):
 - Able to access many different types of resources.
 - List of supported middleware is still increasing!
 - Abstracts away from (most) details of the middleware.
 - Improves portability of applications!
 - Simple and easy to use API.
- Easy way to deploy applications.




However...

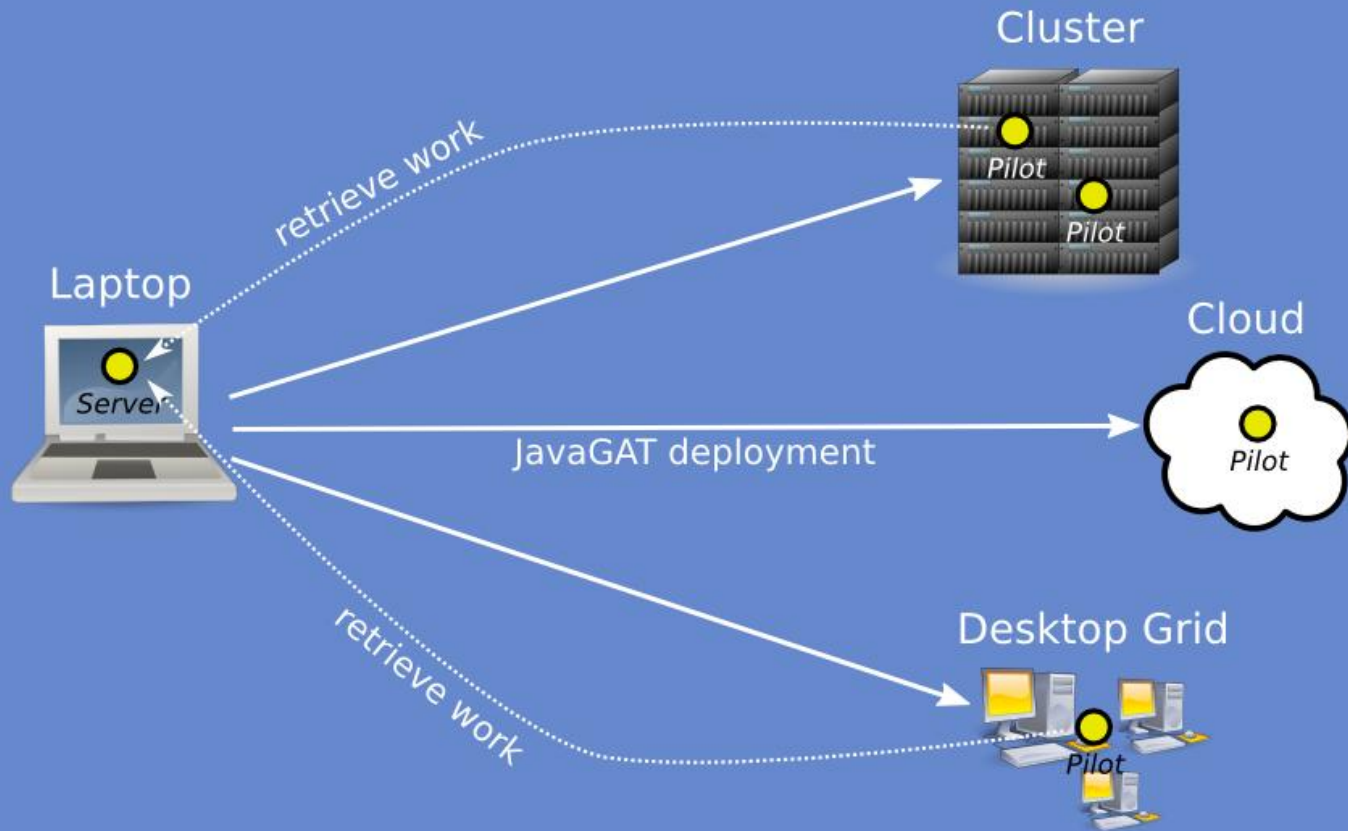
- With many jobs performance may be an issue:
 - Each job is submitted individually
 - Expensive with small jobs and long queueing times
- Solution: **pilot jobs** (many systems exists)
 - First **acquire** resources by submitting a generic job
 - Then use own software infrastructure to run jobs on these resources, **bypassing** the queues altogether
 - **Decouples** workload submission from resource selection and job execution



Example (pilot jobs)



Example (pilot jobs)



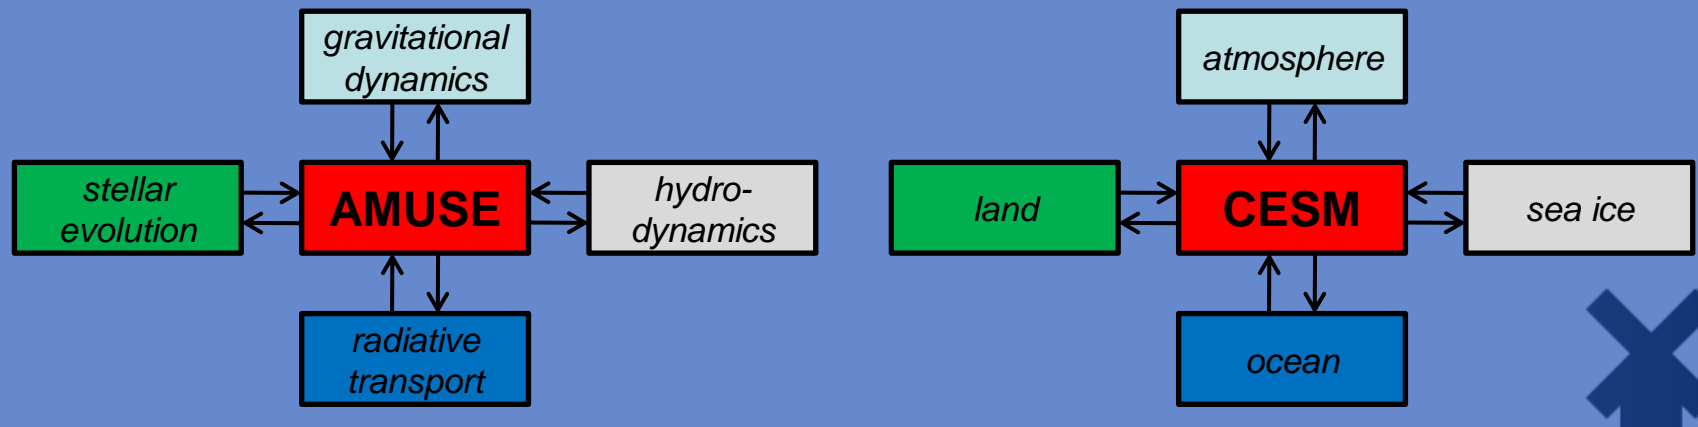
Pilot Job

- Need communication between resources!
- Nothing new about pilot jobs
 - Many frameworks exists
- However: it is a good example to illustrate the “Ibis as glue” part of our tutorial!
 - Glue a jungle of resources into a single resource pool

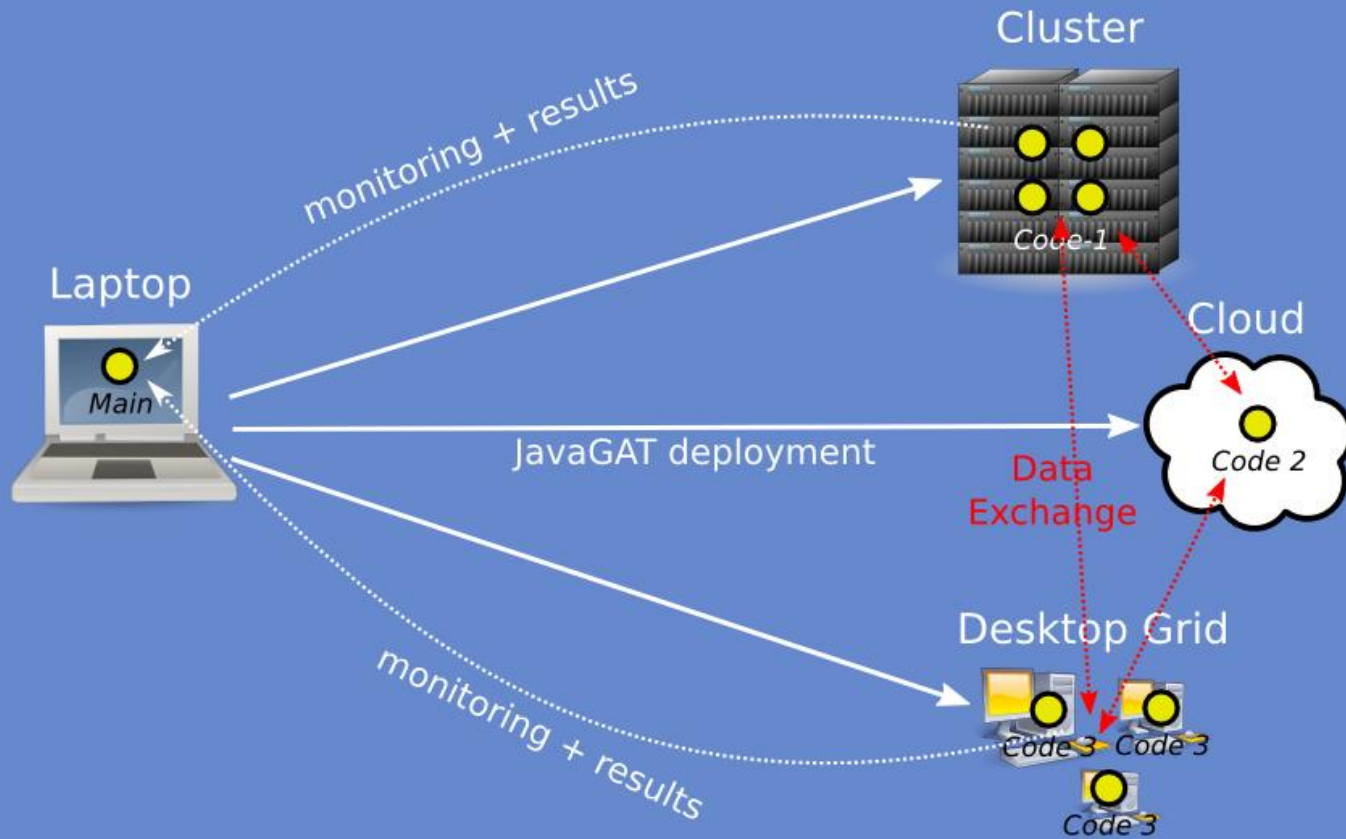


In addition ...

- Not all applications use task farming!
- Many recent applications require **multiple codes** to **run concurrently** on a collection of resources
 - Coupled codes
 - AMUSE (astrophysics), CESM (climate modelling), ...



Example (coupled codes)



We need communication!

- Both pilot jobs and coupled codes need communication between resources
 - Connecting clusters, grids, clouds, ... is hard!
- Many sites have **connectivity issues**
 - Firewalls
 - Network Address Translation (NAT)
 - Non-routed networks
 - Multi homing
 - Mis-configured machines
 - ...



Existing communication libraries

- **Sockets** is too low-level for daily use
 - Only point-to-point
 - No resource management
- **MPI** is too inflexible
 - Focus on SPMD model
 - Little support for malleability or fault tolerance
 - Hard to use in heterogeneous environments
- Neither can handle firewalls/NAT/etc.



What do we need

- Something that solves the connectivity issues
 - Less help from the user is better!
- Better resource tracking
 - **Malleability**: resources come and go
 - **Fault Tolerance**: resources may crash at any time
 - **Robust** and **globally unique naming**
- Flexible communication primitives
 - **Multicast** or **many-to-one** communication
 - Efficient serialization of complex data structures

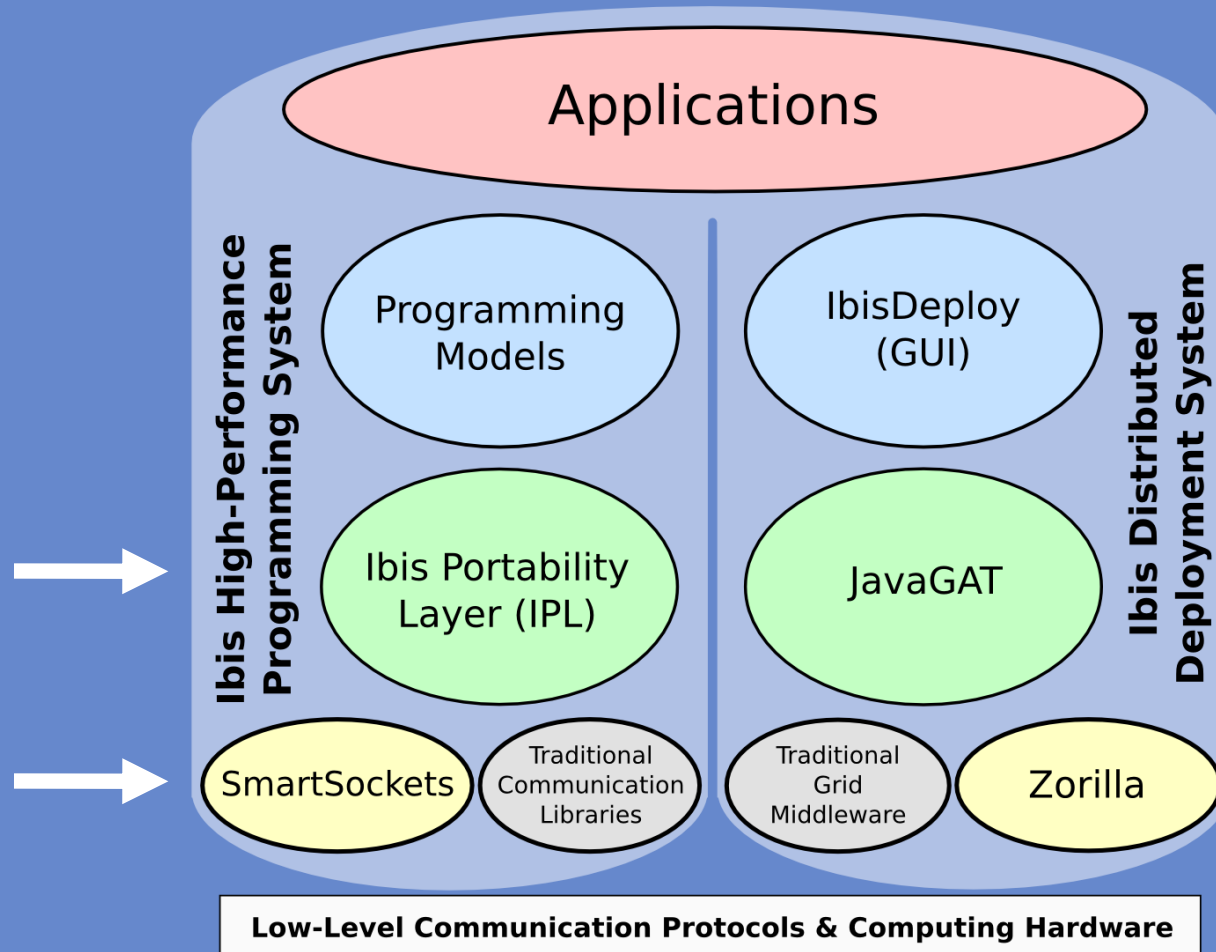


Ibis

- Ibis offers “Jungle proof” communication:
 - SmartSockets
 - Sockets library (on top of regular TCP/IP)
 - Solves low-level connectivity problems
 - Ibis Portability Layer (IPL)
 - Offers high-level communication primitives and resource tracking




Where are we ?



SmartSockets

What problems does it solve ?




- Unreachable machines:
 - Behind firewall / NAT or on private network
- Machine identification:
 - Machines have multiple IPs
 - Multiple machines have the same (private) IP



Problem 1: Firewalls


- Blocks 'inappropriate' connections
 - Usually only blocks incoming connections
 - Some also block outgoing connection



Problem 2:


Network Address Translation

- Allows multiple machines to share an IP address



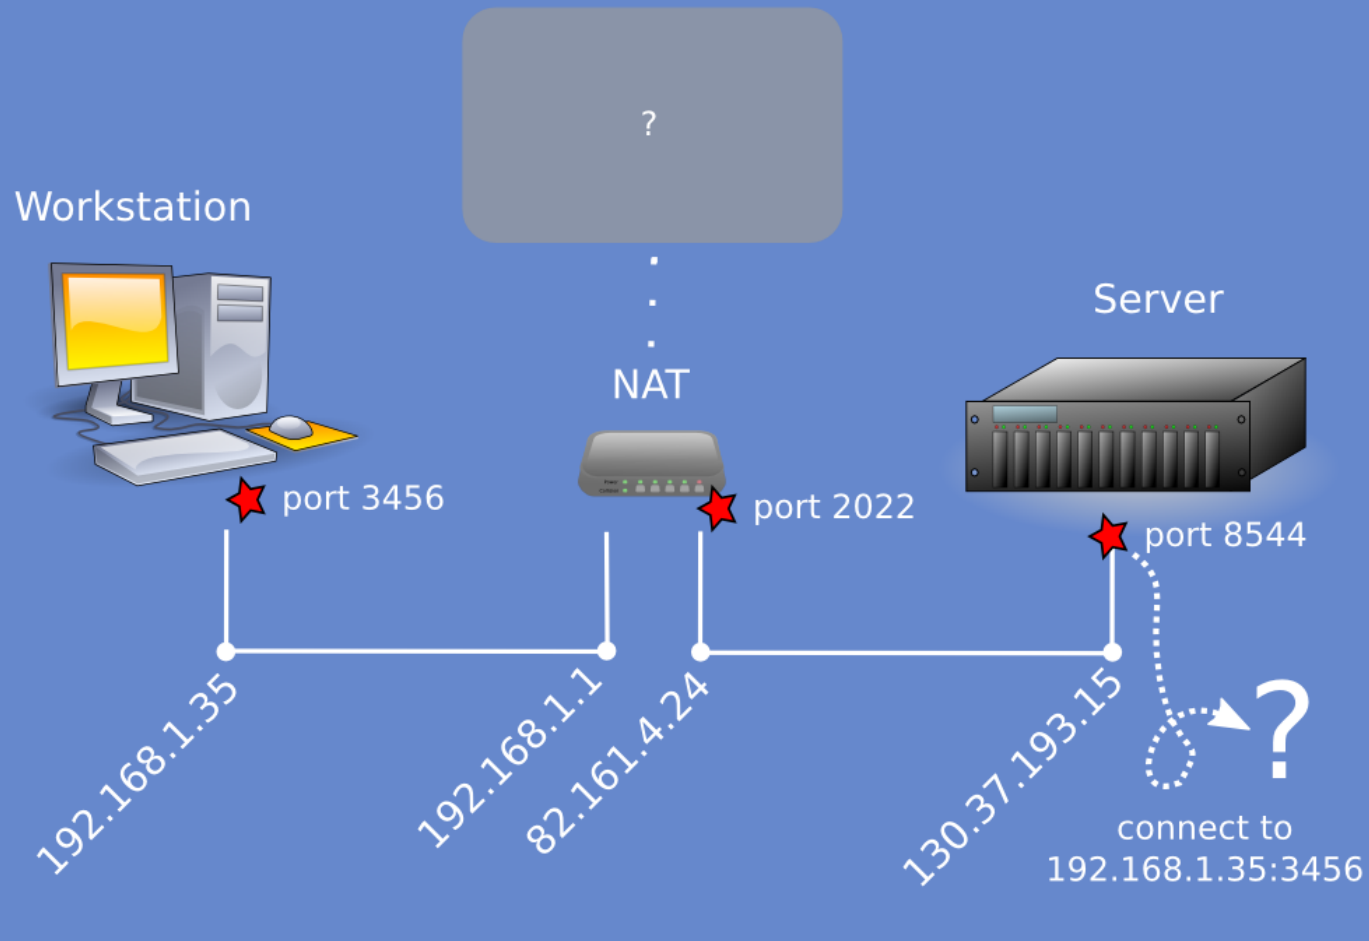
Problem 2:

Network Address Translation

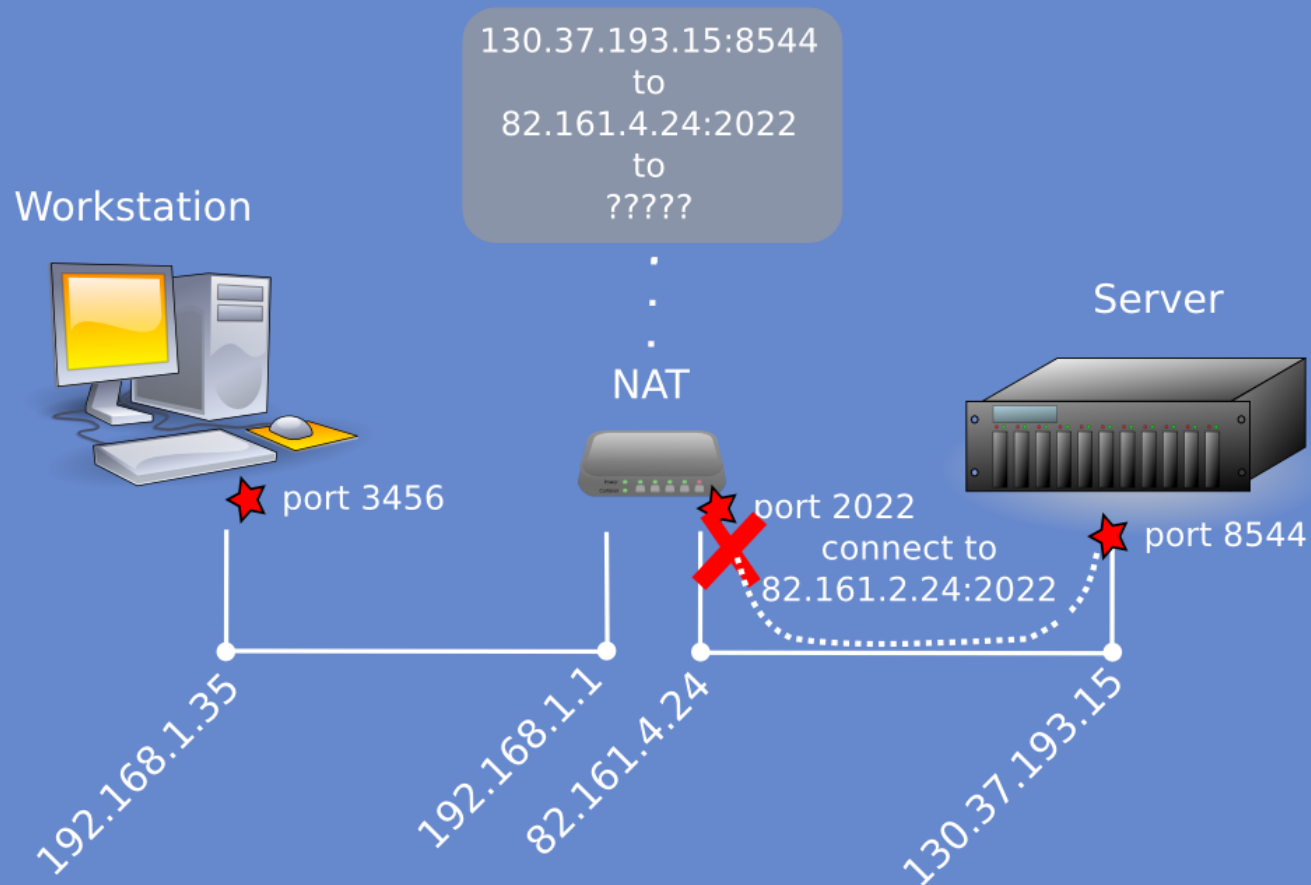


Problem 2:

Network Address Translation

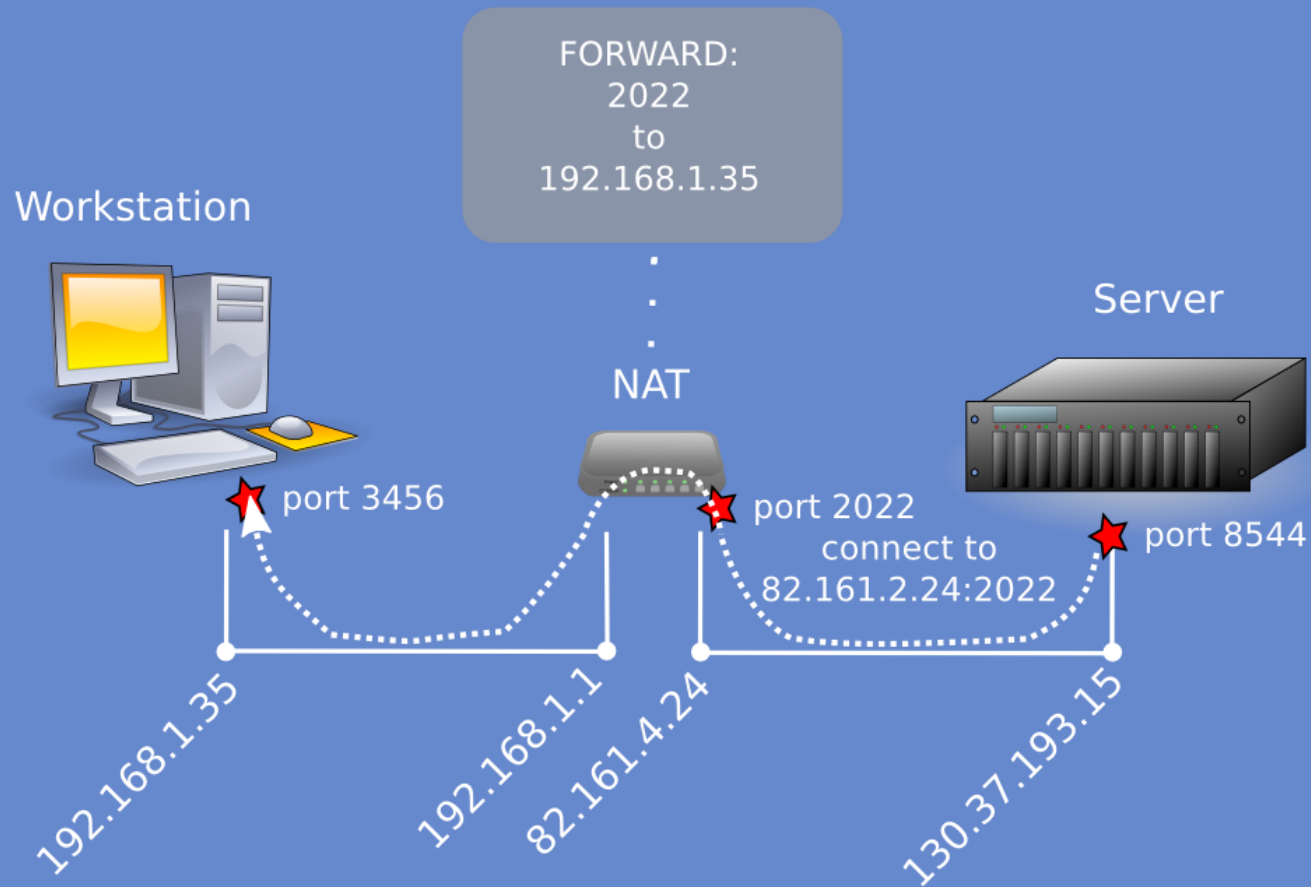


Network Address Translation




Problem 2:

Network Address Translation




Problem 3: Multi Homing

- Some sites have multiple networks
 - The target address depends on the source of the connection



Problem 4: Non-routed Networks


- No route between local network and internet
 - Only the frontend is reachable



Problem 5:

Machine Identification

- Private IPs (NAT/non-routed) lead to machine identification problems



SmartSockets Solutions

- Detect and solve connectivity problems using:
 - Smart Addressing
 - Side channel
 - SSH Tunneling (pass through firewalls)
 - STUN (detect external IP of NAT)
 - UPnP (automatic port forwarding)
 - ...
- Solutions integrated into single socket-like library
 - Mostly transparent to user



Smart Addressing

- Instead of using a single IP:port combination for each process we use:
 - All machine addresses
 - Add extra information
 - External address + port for NAT (STUN, UPnP)
 - SSH contact information
 - UUID (if entire address is private)
 - ...



Addressing Examples

130.37.193.15-42611

public address

130.37.193.15-42611~jason

public address

SSH user

130.37.197.201-42611/10.153.0.201/10.141.0.73-42611

public address

private addresses

{82.161.4.24-20122}/192.168.1.36-42611#a6.e5.01.1a.ad.f1

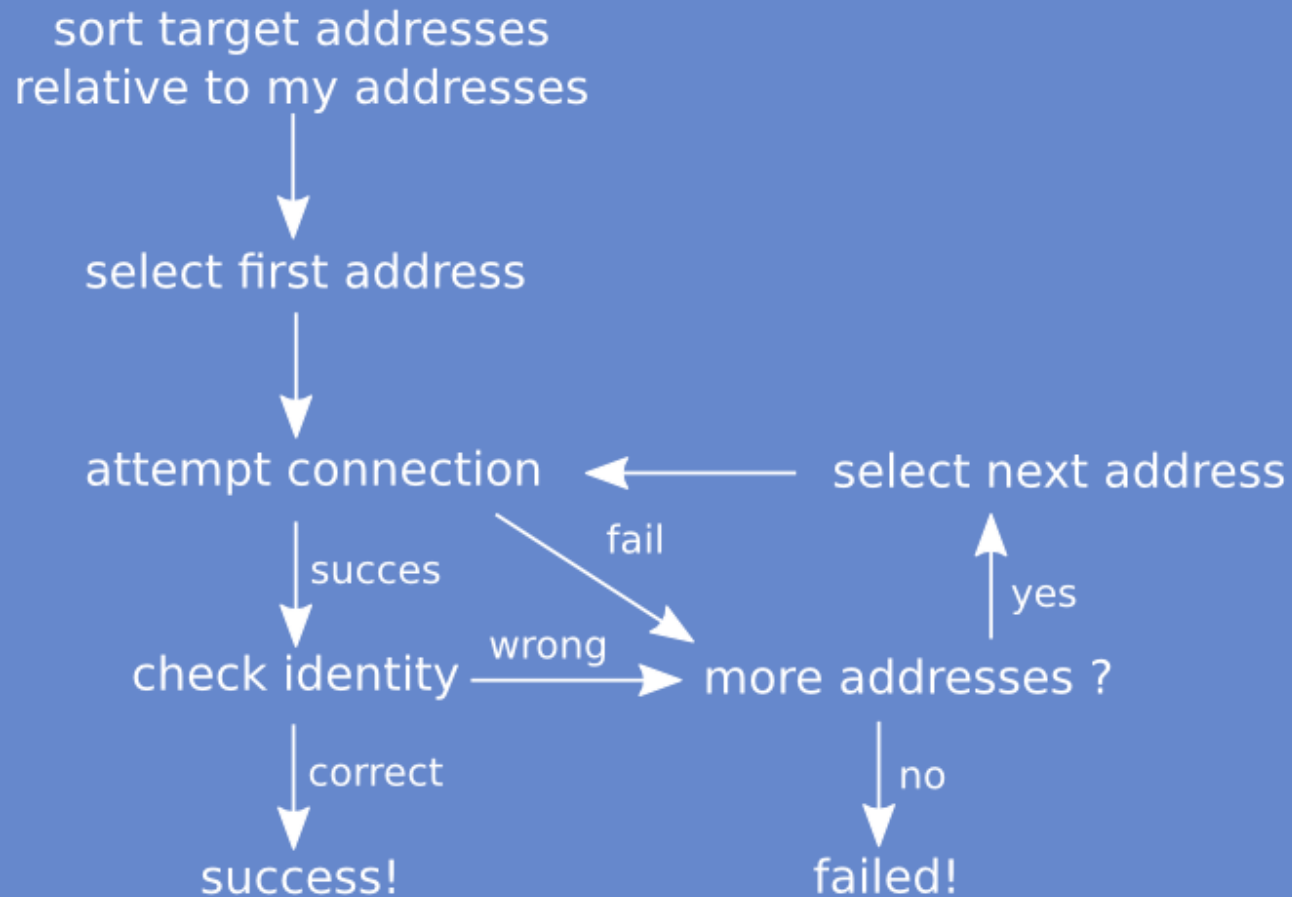
external address
with port forwarding

private
address

UUID
unique to host



Creating a Connection



Using Smart Addresses

- This solves **machine identification** problems
 - All addresses are known with multi-homing
 - Each identity is unique, even with private IPs
 - The identity is always checked on connection setup
- Still assumes anyone can create a connection
 - This will not help when target is behind NAT/Firewall
 - To solve this we need a **side channel**



Side channel

- Overlay network implemented using a set of **hubs**
 - **Support processes** for the application
 - Started in advance
- Hubs are run on machines with 'more connectivity'
 - Such as cluster frontends, 'open' machines, etc.
- How / where you start them is a separate problem
 - Solved by IbisDeploy (shown later)




Hubs

- Similar to a **peer-to-peer overlay** network
- Hubs connect to each other
 - Gossip information about other hubs
 - Automatically discover new hubs and routes
 - Need to set up **spanning tree** (or better)
 - Use direct connections and **SSH tunnels**
- Clients connect to a 'local' hub
 - Use as **side channel** for connection setup



Hub Overlay Network




Advanced Connection Setup

- **Reverse** direction of connection setup
 - Instruct target to set up connection (using hub)
 - Results in **direct connection**
- **Splicing** (switched off by default)
 - Connection setup from both sides (using hub)
 - Results in **direct connection**
- **Route** via overlay
 - Create **virtual connection** using hubs
 - Forward all data over side channel
 - Results in **indirect connection**



Advanced Connection Setup




SmartSockets

Overview of solutions


- Machine identification:
 - Smart addressing
 - Identity check at connection setup
- Unreachable machines:
 - SSH tunnels
 - Reverse connection setup
 - (Splicing)
 - Routing over hubs





Overlay Network

Created by hubs



Evaluation

Table 3: Connection setup time of SmartSockets (time in milliseconds).

<i>Target</i>	<i>Source</i>					
	DAS3-V	DAS3-D	Rockstar	Grid5000	Hiroshi	Desktop
DAS3-V		4.9 ^d (2.4)	332 ^d (166)	68 ^v	595 ^v	33 ^d (17)
DAS3-D	4.9 ^d (2.4)		335 ^d (167)	70 ^v	595 ^v	33 ^d (18)
Rockstar	500 ^r	503 ^r		206 ^v	718 ^v	182 ^v
Grid5000	35 ^v	38 ^v	206 ^v		593 ^v	54 ^v
Hiroshi	630 ^v	603 ^v	750 ^v	670 ^v		640 ^v
Desktop	49 ^r	52 ^r	183 ^v	84 ^v	606 ^v	

Annotations indicate connection style: *d* for direct, *r* for reverse, *s* for splicing, and *v* for routed.

When applicable, the connection setup time of regular sockets is shown between brackets.

- Regular TCP/IP only worked in 6 out of 30
- SmartSockets worked in 30 out of 30



Evaluation

Table 4: Roundtrip latency of SmartSockets (time in milliseconds).

<i>Target</i>	<i>Source</i>					
	DAS3-V	DAS3-D	Rockstar	Grid5000	Hiroshi	Desktop
DAS3-V		2.3 (2.3)	166 (166)	56	528	14 (14)
DAS3-D	2.3 (2.3)		167 (167)	57	533	15 (15)
Rockstar	166	167		205	590	195
Grid5000	56	57	205		524	50
Hiroshi	528	529	590	522		539
Desktop	14	15	190	43	522	


When applicable, the roundtrip latency of regular sockets is shown between brackets.

Table 5: Throughput of SmartSockets (in Mbit/second).

<i>Target</i>	<i>Source</i>					
	DAS3-V	DAS3-D	Rockstar	Grid5000	Hiroshi	Desktop
DAS3-V		182 (183)	2.6 (2.5)	2.5	0.25	0.65 (0.65)
DAS3-D	185 (186)		2.6 (2.5)	2.6	0.26	0.65 (0.65)
Rockstar	2.8	2.7		6.9	0.23	0.65
Grid5000	7.6	8.2	2.4		0.20	0.65
Hiroshi	0.73	0.73	0.70	0.73		0.61
Desktop	3.3	3.3	2.2	2.2	0.25	

When applicable, the throughput of regular sockets is shown between brackets.

However....



- SmartSockets is great, but too low level
- For Jungle computing we need support for
 - Malleability
 - Fault Tolerance
 - Robust and globally unique naming
 - Flexible communication primitives
- Provided by the Ibis Portability Layer (IPL)




Ibis Portability Layer (IPL)

- Simple API for Jungle Communication
 - Flexible communication model
 - Connection oriented messaging
 - Abstract addressing scheme
 - Resource tracking (JEL model)
 - Notifications when machines join/leave/crash
 - Efficient serialization
 - Send bytes, doubles, objects, etc.
 - Portable:
 - SmartSockets, TCP, UDP, MPI, MX, BlueTooth, ...



Communication Model





- Simple communication model
 - Unidirectional pipes
 - Two end points (send and receive ports)



- Connection oriented
 - Allows streaming (good with high latency)




Flexible Communication Model



Port Types

- All send and receive ports have a type
 - Defined at runtime
 - Specify set of capabilities
- Types must match when connecting!



Port Types

- Consists of a set of capabilities:
 - Connection patterns:
 - Unicast, many-to-one, one-to-many, many-to-many
 - Communication properties:
 - Fifo ordering, numbering, reliability
 - Serialization properties:
 - Bytes, primitive types, objects
 - Message delivery:
 - Explicit receipt, automatic upcalls, polling



Port Types

- Forces programmer to specify how each communication channel is used
 - Prevents bugs
 - Exception when contract is breached
- Allows efficient implementation to be selected
 - Unicast only ?
 - Transfer bytes only ?
 - Can save a lot complexity!




Messages

- Ports communicate using 'messages'
- Contain read or write methods for
 - Primitive types (byte, int, ...)
 - Object
 - Arrays slices (partial write / read in place)
- Unlimited message size
 - Streaming



Messages Example



Abstract addressing

- **IbisIdentifier:**
 - Abstract 'proces address' object
 - Hides network specific details
 - Examples: SmartSockets addresses, hostnames, IP addresses, MPI ranks, etc.
- Results in more **portable** applications
 - Independent of network infrastructure
- Why don't we use ranks ?
 - Hard to support malleability and fault-tolerance!

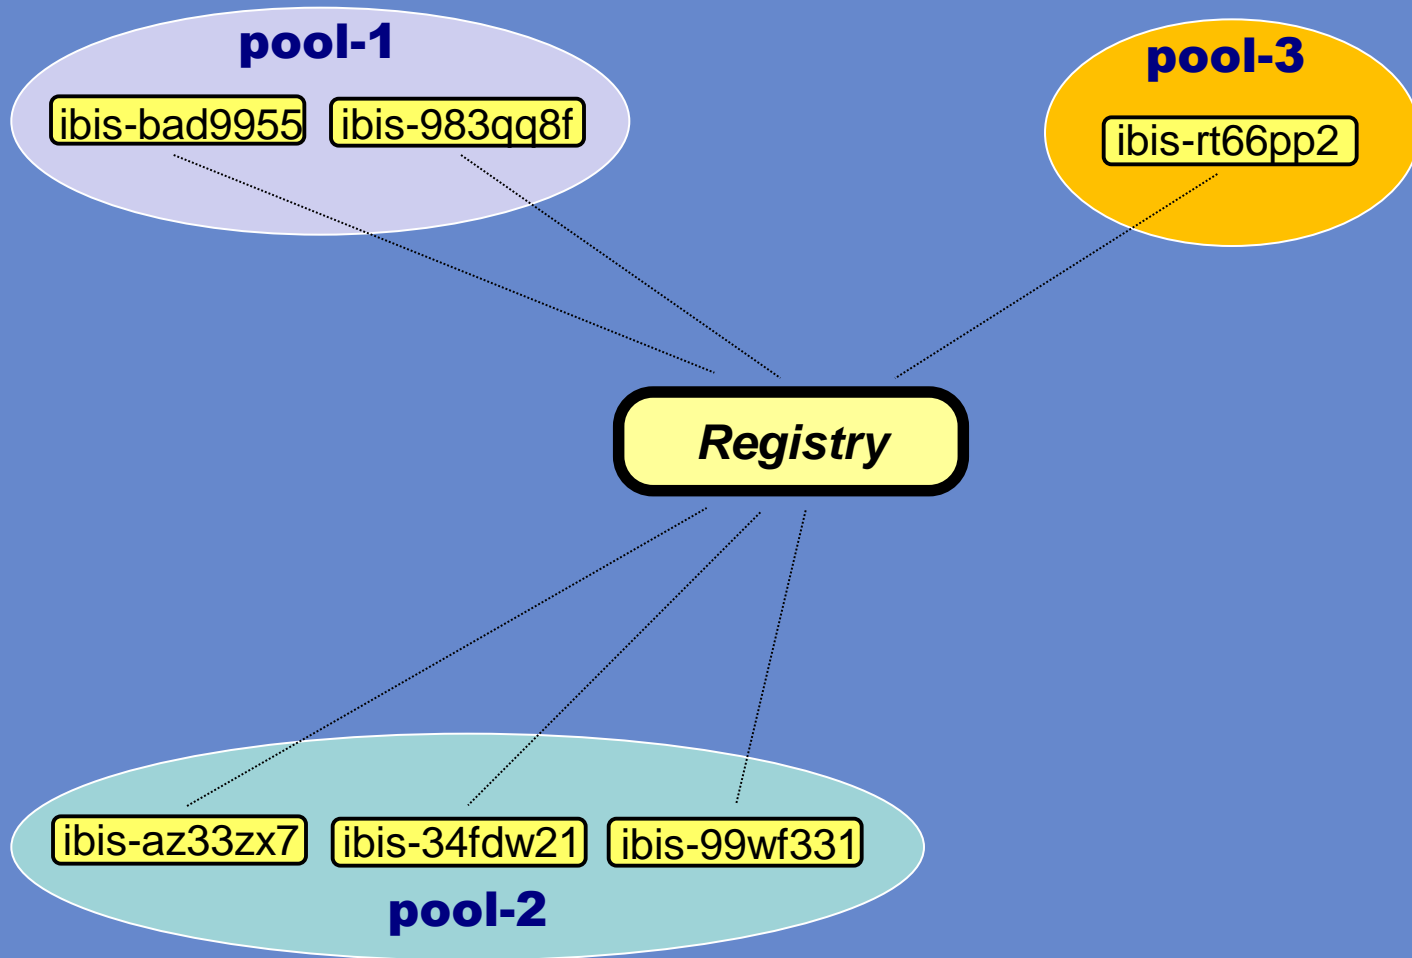


Resource Tracking

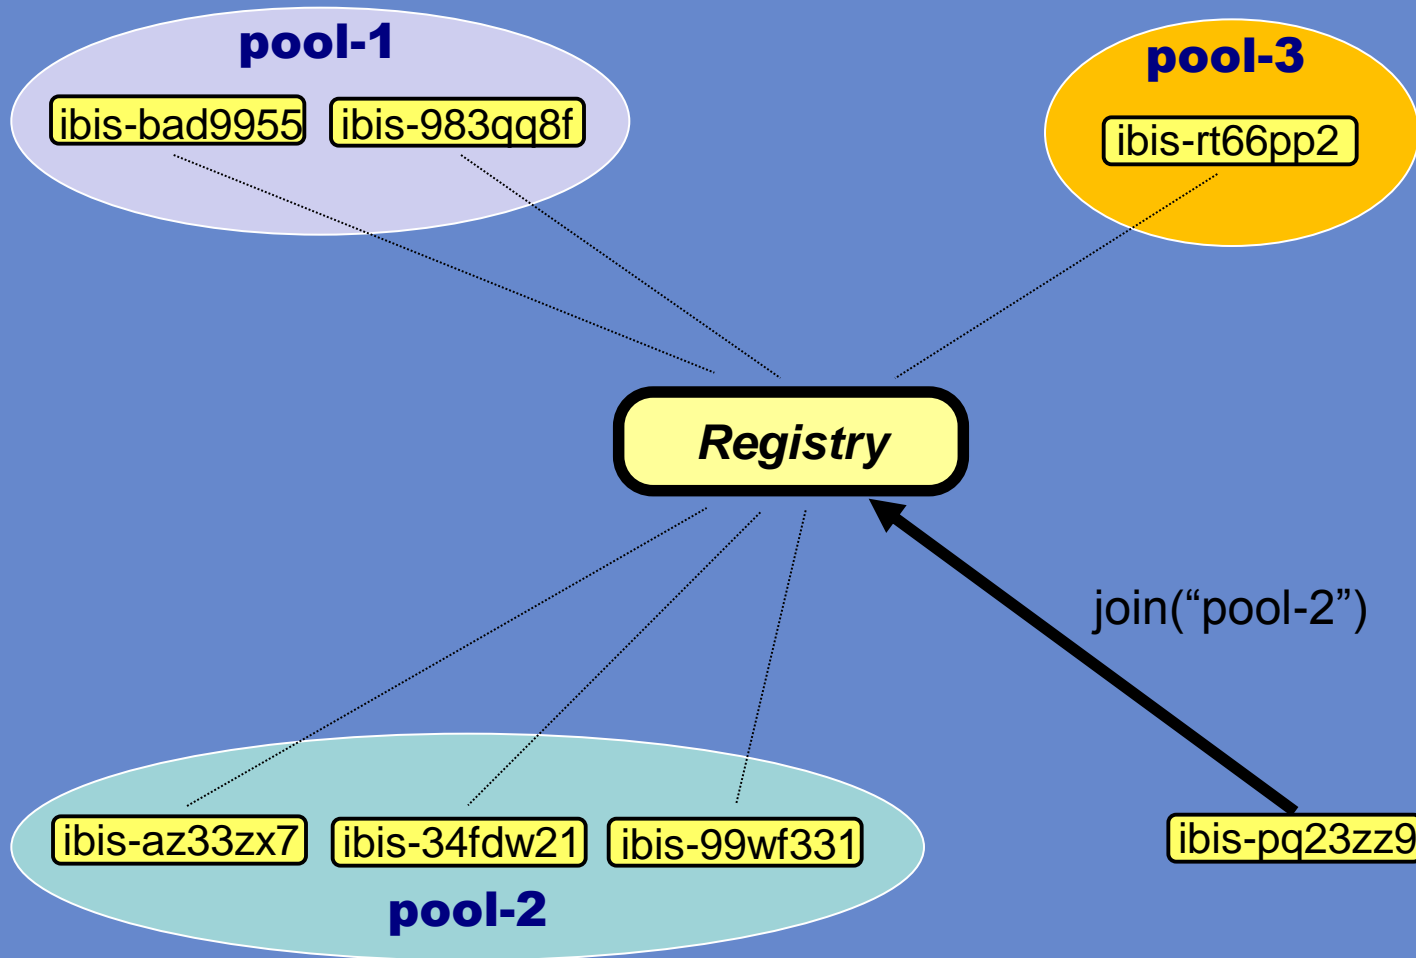
- IPL offers **JEL** (join, elect, leave) model
 - Application gets signal when someone **joins** or **leaves**
 - Supports **elections** for distributed decision making
 - Allows machines to be elected as “master”
 - Can **ensure totally ordered** notifications
- Implemented using separate **registry** component
 - Server that tracks application participants
 - Can track multiple applications simultaneously, each in its own **pool**




Pools & Malleability



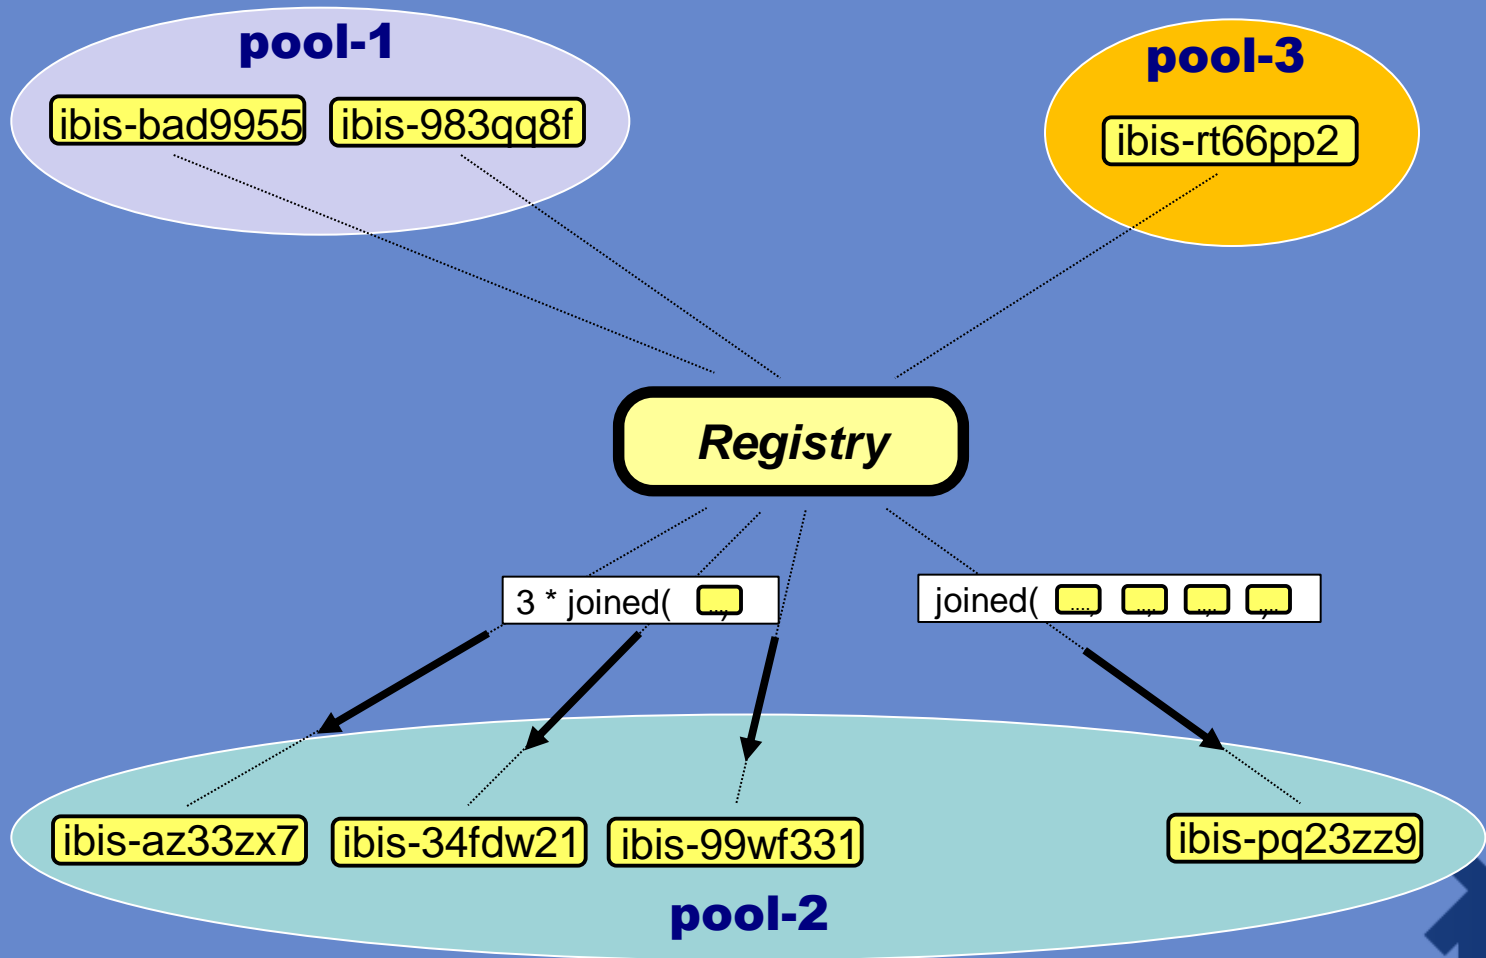
Pools & Malleability



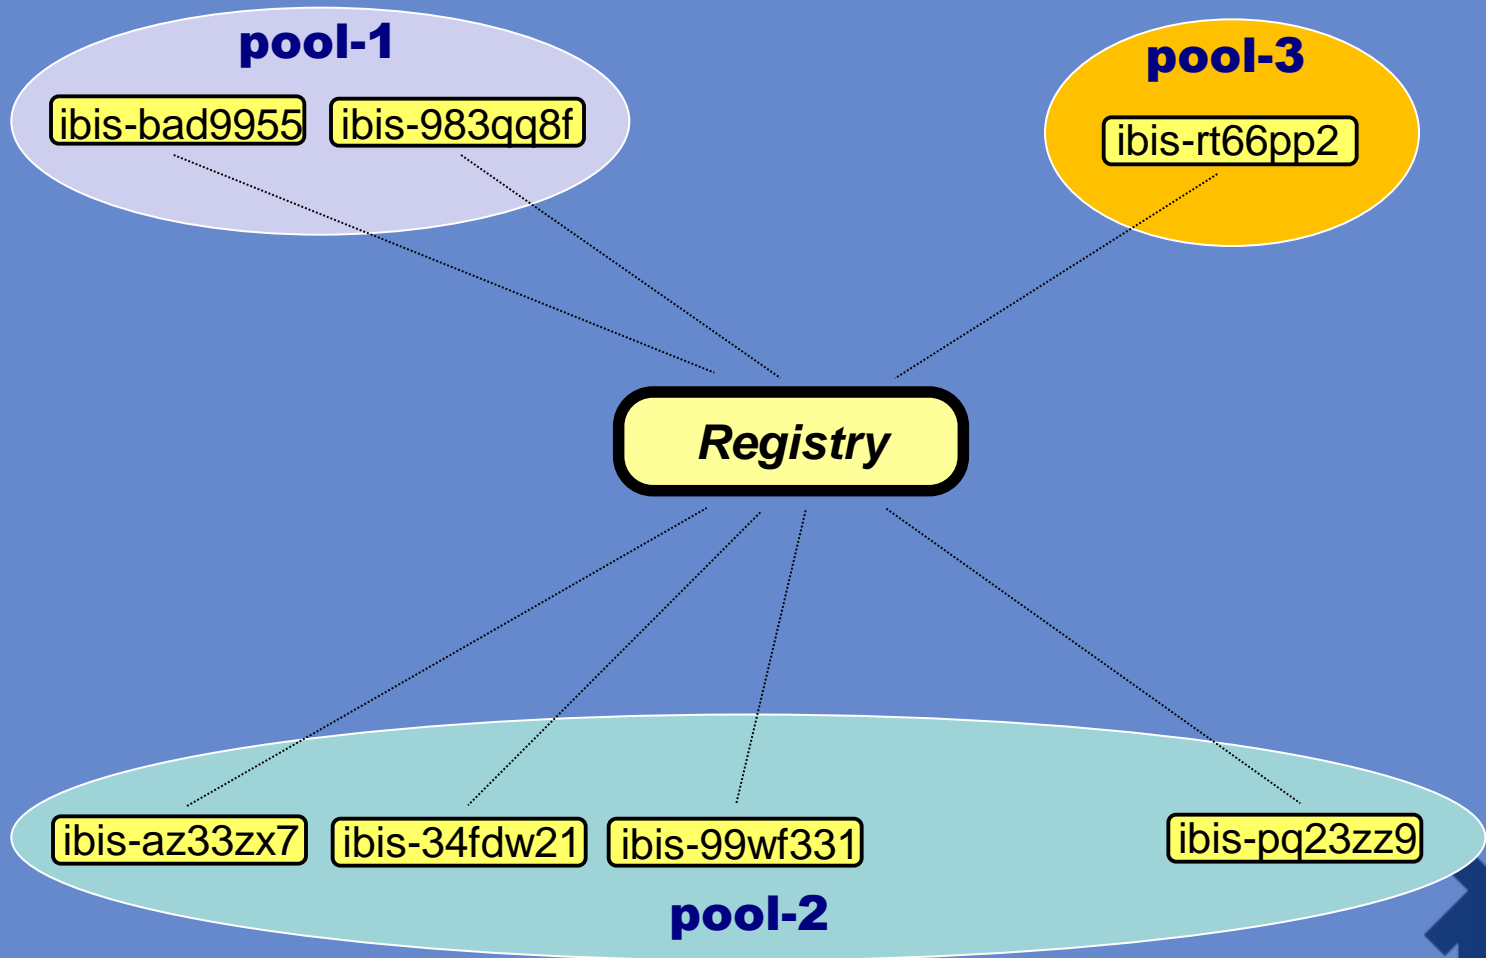
Pools & Malleability



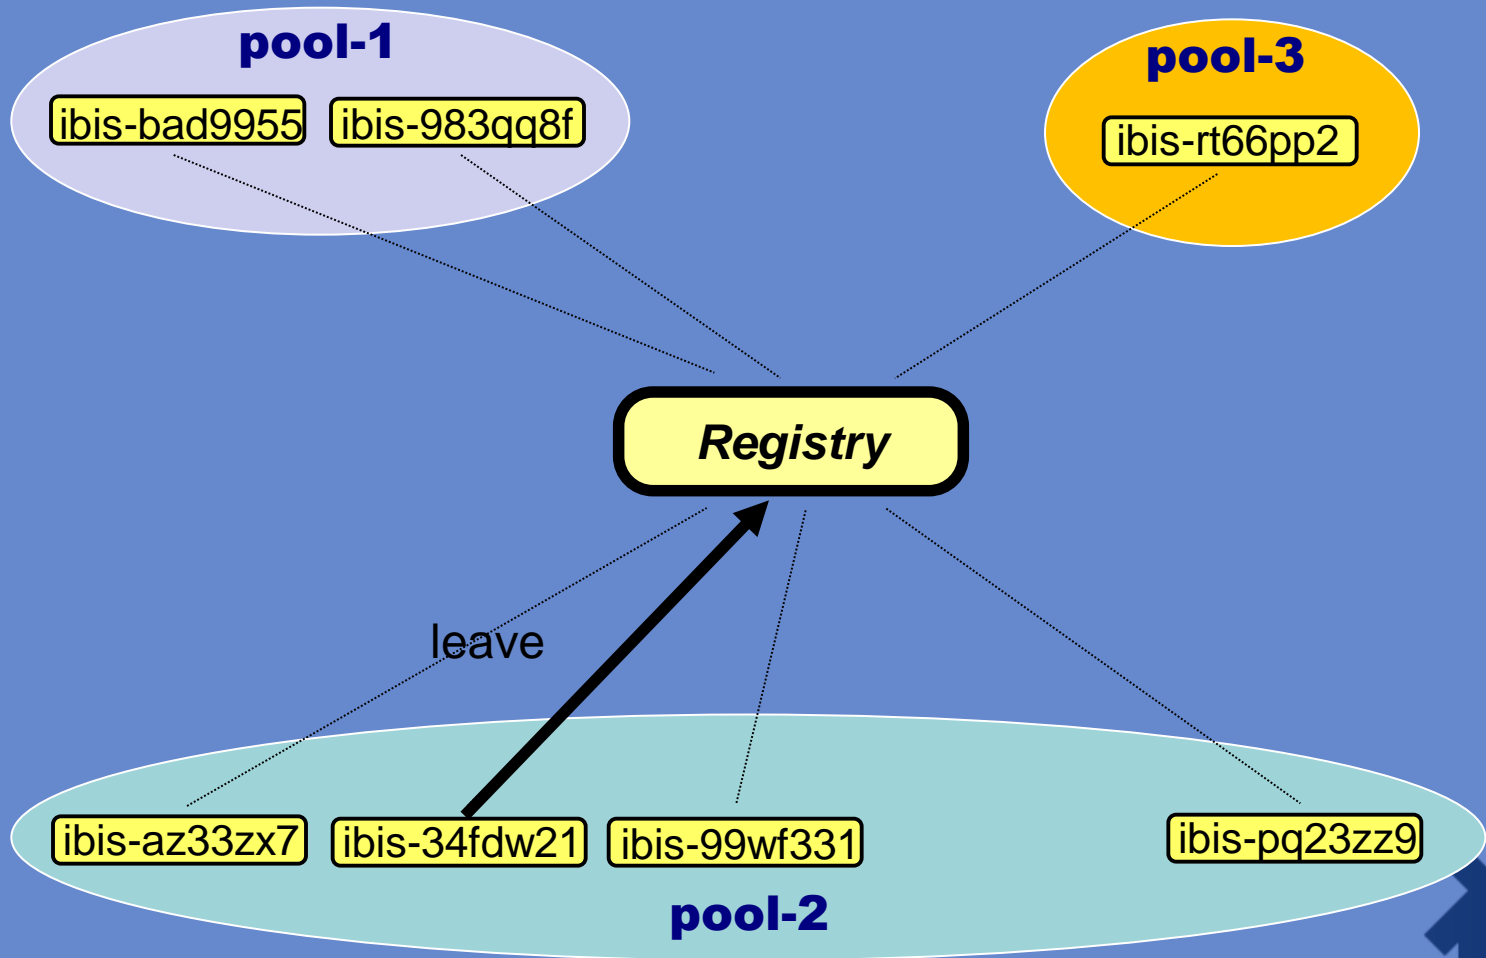
Pools & Malleability



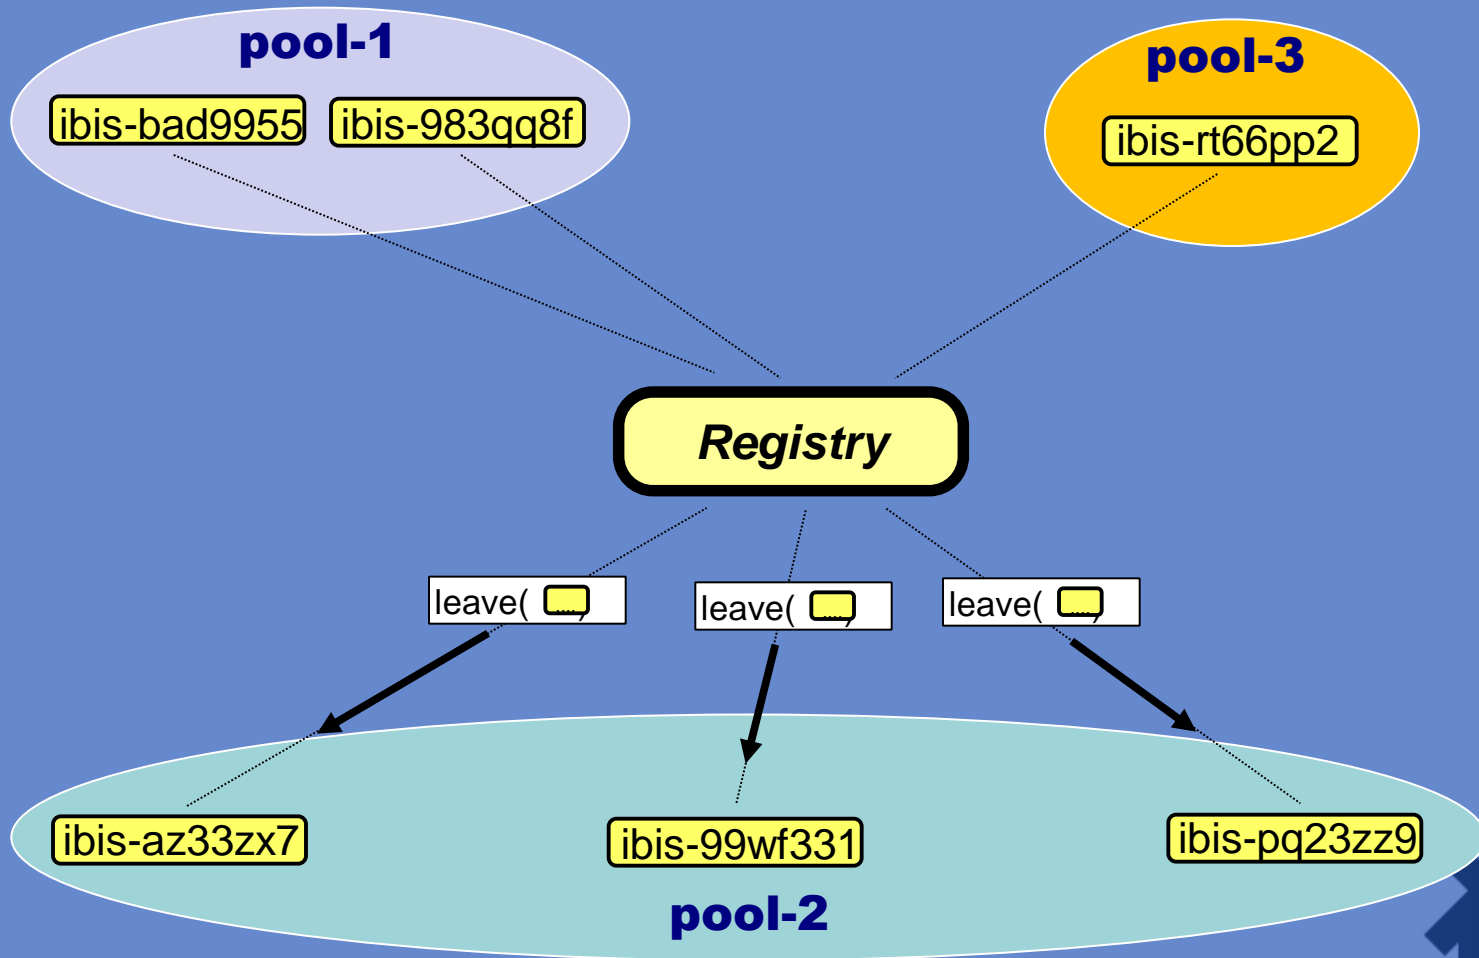
Pools & Malleability



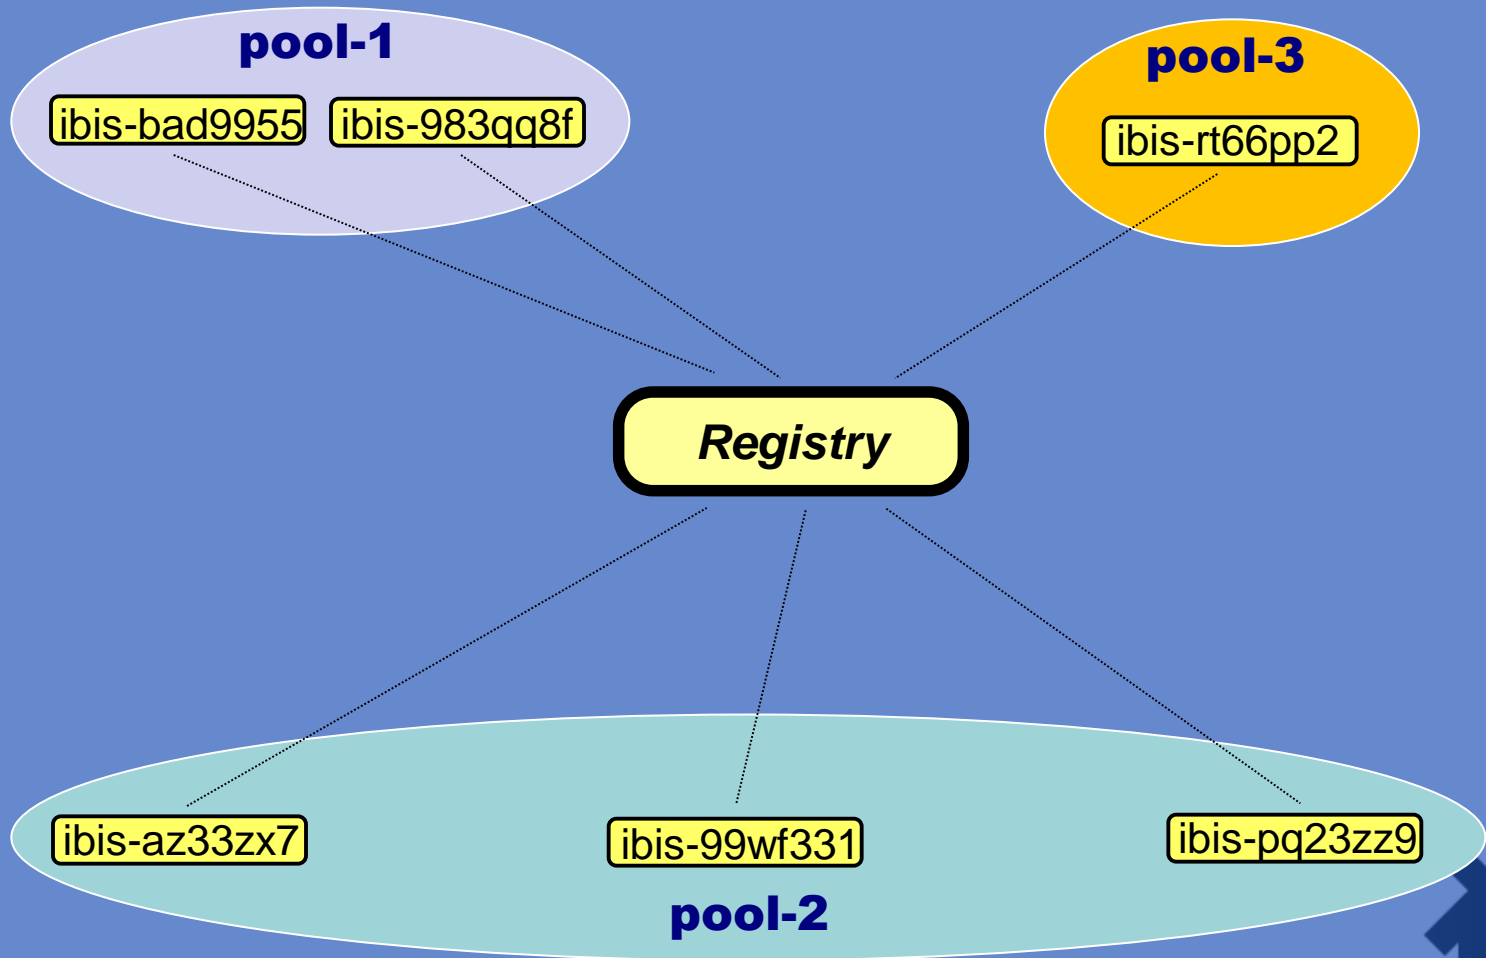
Pools & Malleability



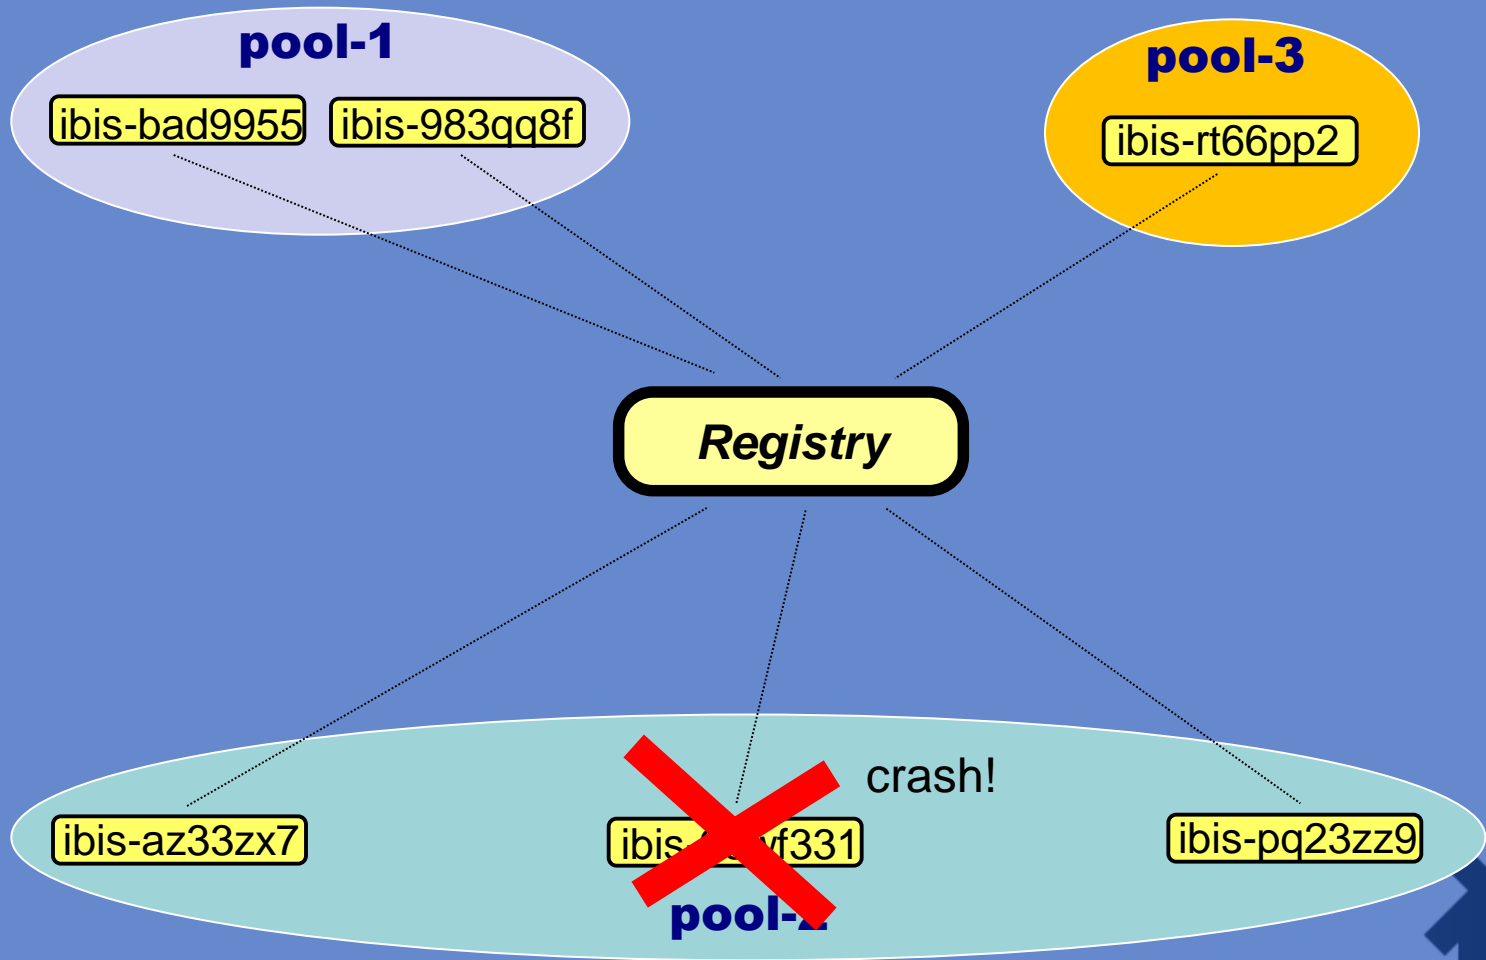
Pools & Malleability



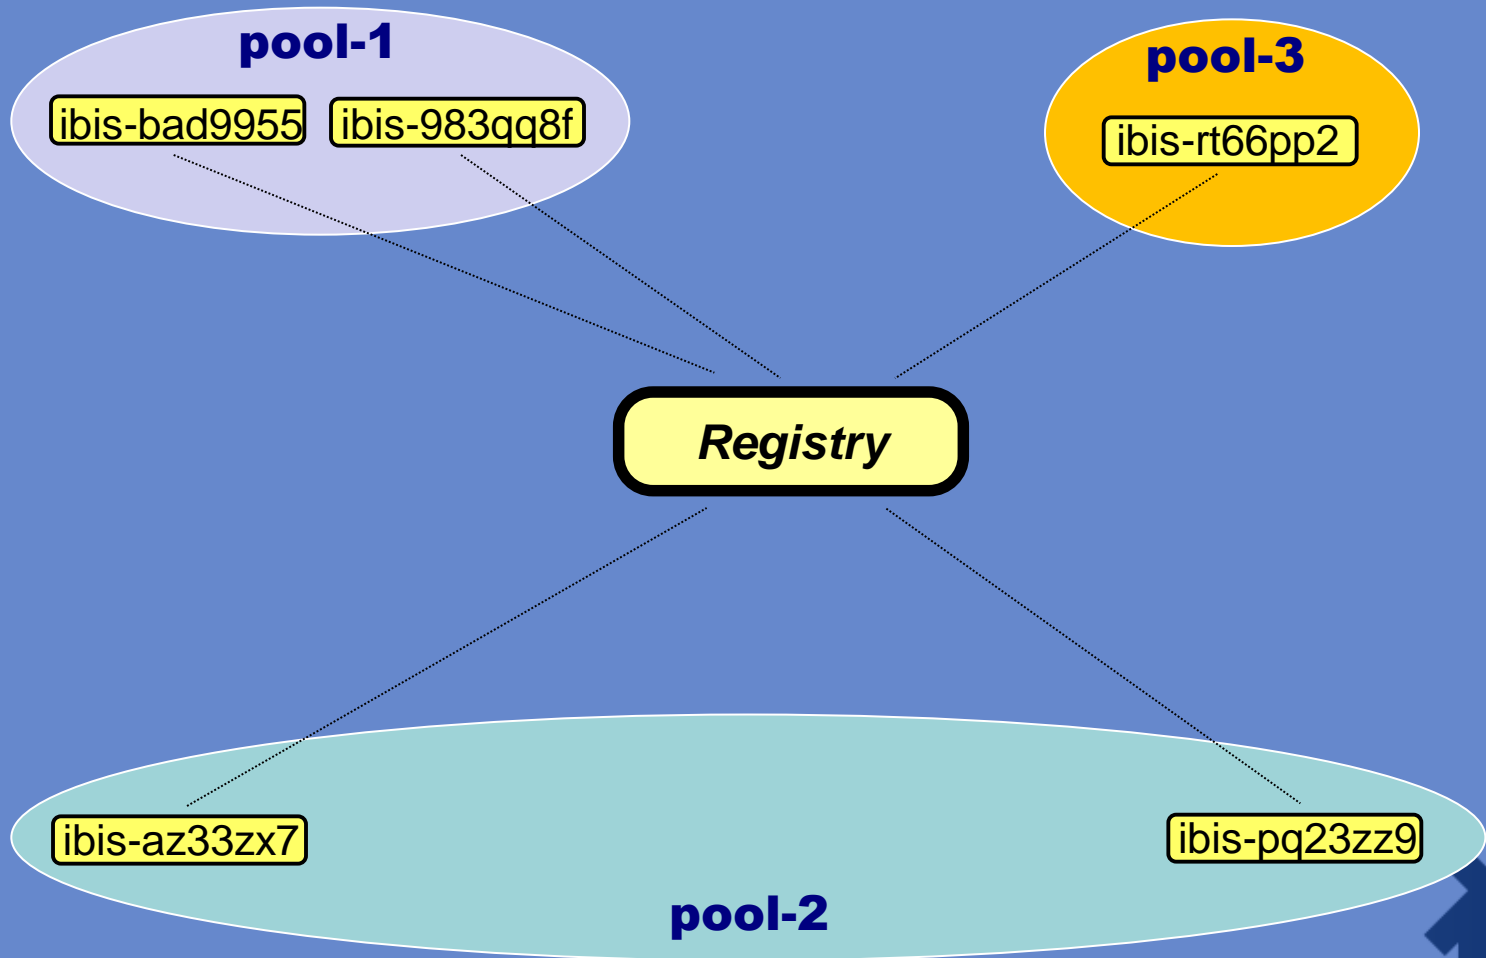
Pools & Malleability



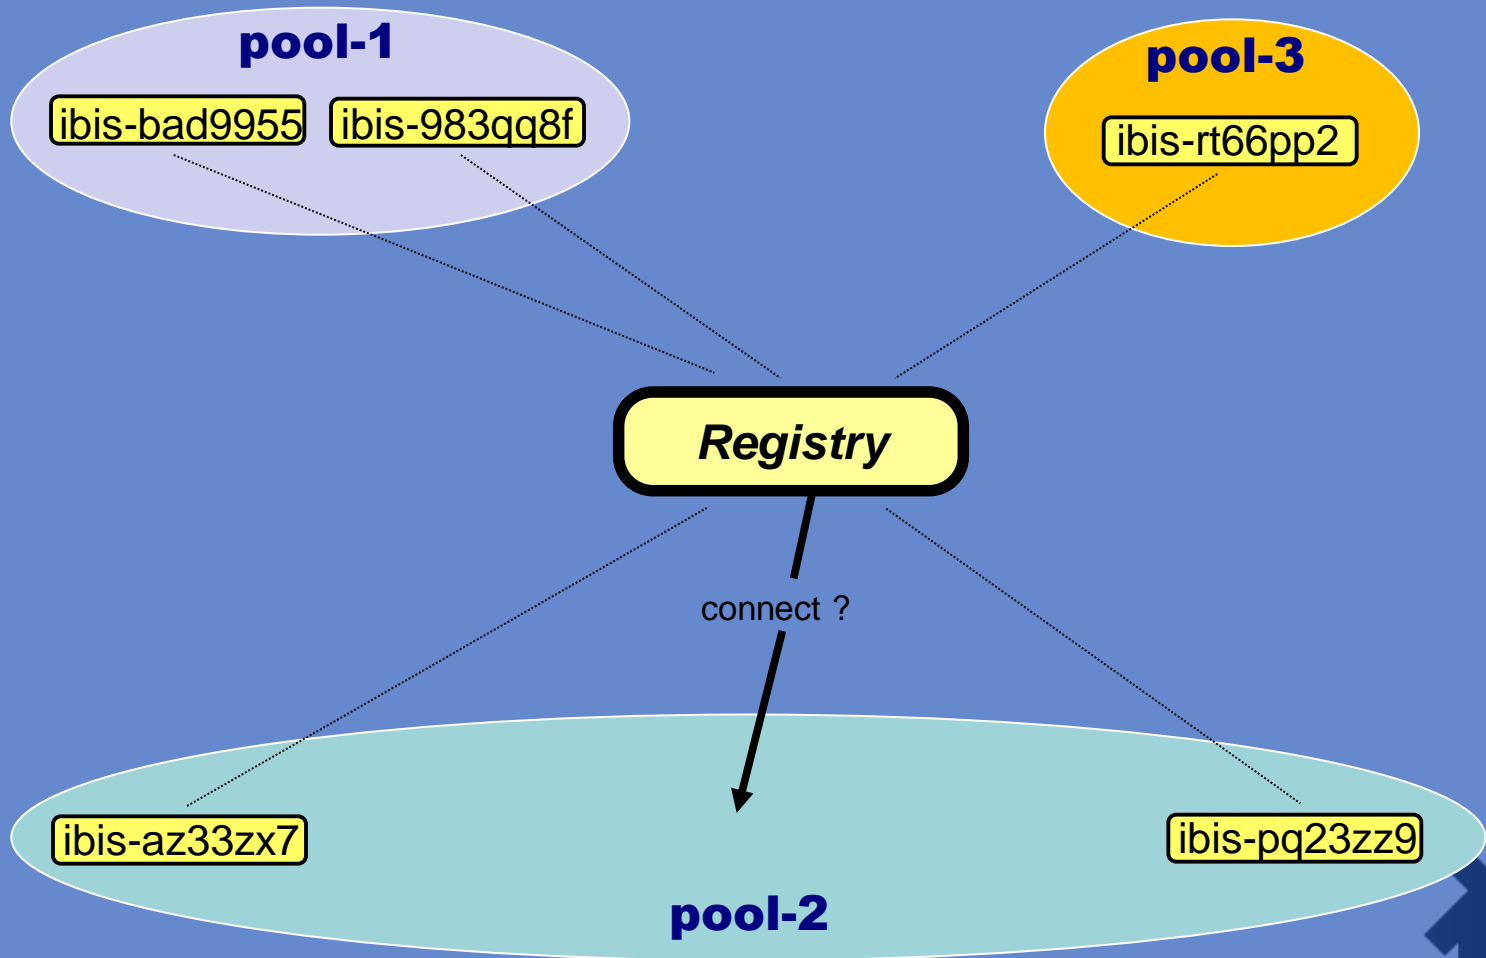
Pools & Malleability



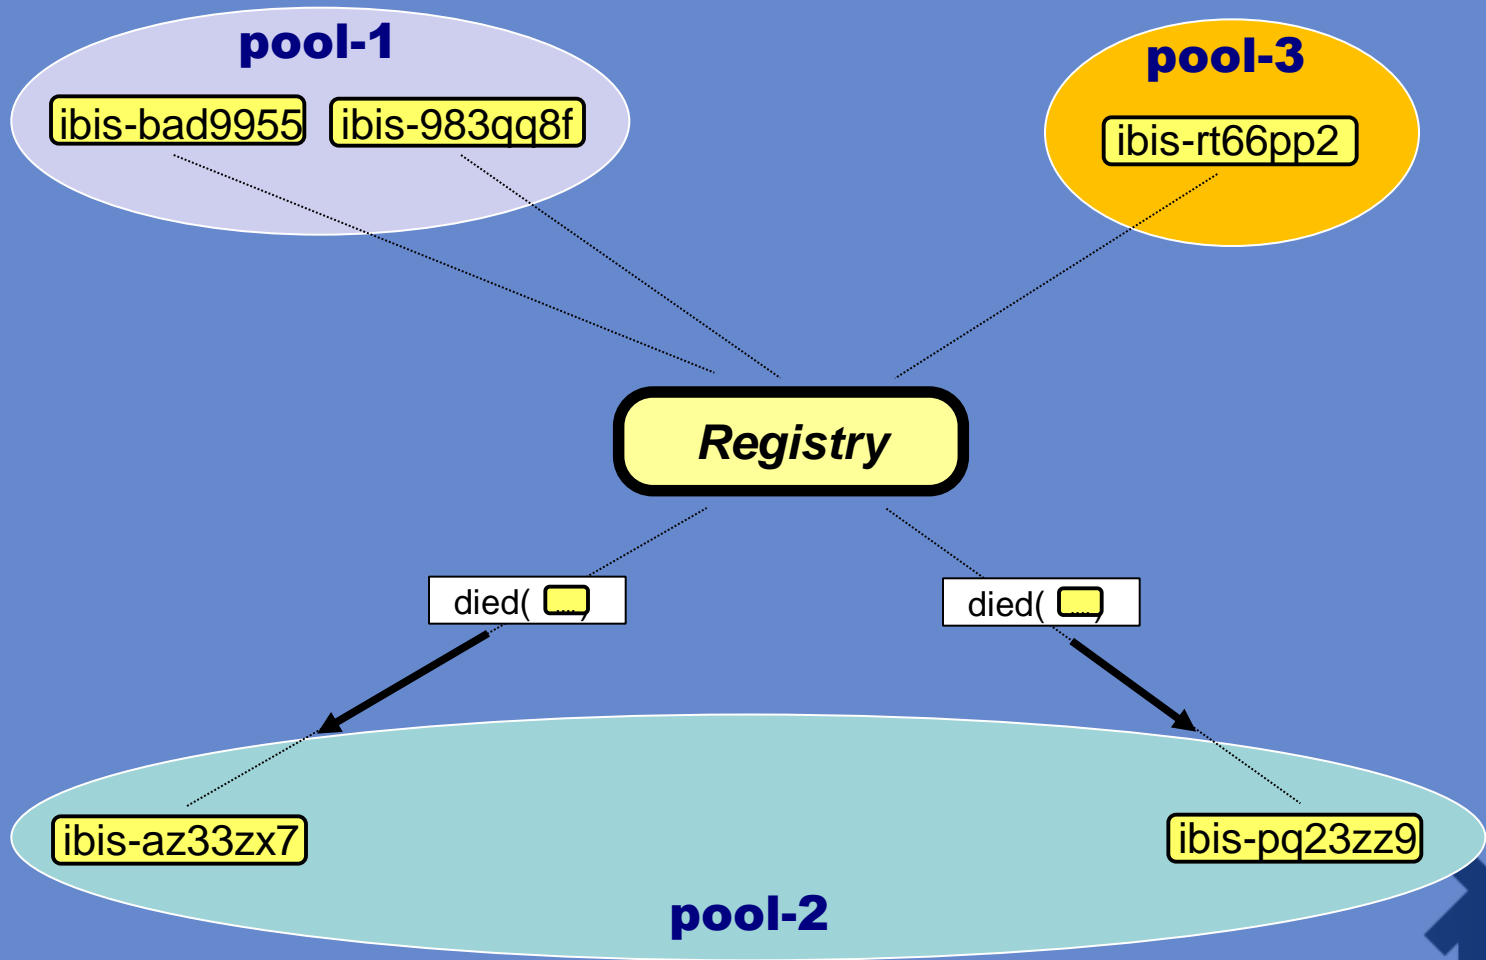
Pools & Malleability




Pools & Malleability



Pools & Malleability



Pools & Malleability



Registry

- Many implementations
 - Centralized, broadcast, gossiping, etc.
 - Different tradeoffs in **functionality**, **complexity**, **robustness**, **scalability** and **consistency**
- Application can select the functionality and consistency that is needed
 - Reducing functionality or consistency further improves scalability




IPL Example:

Back to the Pilot Jobs

- Now we have Jungle proof communication it should be easy to create a pilot job framework
- Recipy:
 - Create a Pilot Job framework using the IPL
 - Submit this to the resources using JavaGAT
 - Distribute work using the Pilot Job framework
- We'll use this as an IPL code example!



Pilot Job Framework Design



Example Code

Deployment

```
package tutorial20.glue;

import ibis.ipl.server.Server;

public class Main {

    static class Resource {
        String brokerURI;
        String javaLocation;

        Resource(String brokerURI, String javaLocation) {
            this.brokerURI = brokerURI;
            this.javaLocation = javaLocation;
        }
    }

    static int number = 0;

    public static JobDescription prepareJob(String serverAddress,
    public static void cleanup(LinkedList<Job> jobs)
    public static void main(String[] args) throws Exception {
    }
```



Example Code

```
package tutorial2;  
  
import ibis.ipl.  
  
public class Main  
  
    static class  
        String b  
        String j  
  
        Resource  
            this  
            this  
        }  
    }  
  
    static int n  
  
    public static  
  
    public static  
  
    public static  
  
}
```

```
public static void main(String[] args) throws Exception {  
  
    LinkedList<Resource> resources = new LinkedList<Resource>();  
  
    String[] arguments = null;  
    String executable = null;  
    String inputdir = null;  
    String outputdir = null;  
  
    // ...parse parameters here....  
  
    Server reg = new Server(new Properties());  
  
    Properties p = new Properties();  
    p.put("ibis.server.address", reg.getAddress());  
    p.put("ibis.pool.name", "PILOT");  
  
    JobServer jobServer = new JobServer(executable, arguments, inputdir,  
        outputdir, p);  
  
    LinkedList<Job> gatJobs = new LinkedList<Job>();  
  
    for (Resource resource : resources) {  
        ResourceBroker broker = GAT.createResourceBroker(new URI(  
            resource.brokerURI));  
        Job gatJob = broker.submitJob(prepareJob(reg.getAddress(),  
            resource.javaLocation));  
        gatJobs.add(gatJob);  
    }  
  
    jobServer.run();  
  
    cleanup(gatJobs);  
    GAT.end();  
}
```

Example Code

Deployment

```
package tutorial20.glue;

import ibis.ipl.server.Server;

public class PilotJob {

    static JobDescription prepareJob(String serverAddress,
                                     String javaLocation) throws GATObjectCreationException {

        JavaSoftwareDescription sd = new JavaSoftwareDescription();

        HashMap<String, String> properties = new HashMap<String, String>();
        properties.put("ibis.server.address", serverAddress);
        properties.put("ibis.pool.name", "PILOT");

        sd.setExecutable(javaLocation);
        sd.setJavaClassPath("ipl/*:GAT-examples.jar:.");
        sd.setJavaSystemProperties(properties);
        sd.setJavaMain("tutorial20.glue.PilotJob");

        sd.setStdout(GAT.createFile("stdout-" + number + ".txt"));
        sd.setStderr(GAT.createFile("stderr-" + number + ".txt"));

        sd.addPreStagedFile(GAT.createFile("lib/GAT-examples.jar"));
        sd.addPreStagedFile(GAT.createFile("lib/ipl"), GAT.createFile("ipl"));
        sd.addPreStagedFile(GAT.createFile("log4j.properties"));

        number++;

        return new JobDescription(sd);
    }
}
```

Example Code

Pilot Job

```
package tutorial20.glue;

import ibis.ipl.Ibis;

public class PilotJob {
    Ibis ibis;
    ReceivePort rp;
    SendPort sp;

    PilotJob() throws Exception {}

    Job getWork(Result previousResult) throws Exception {}

    void run() throws Exception {}

    public static void main(String[] args) {}
}
```



Example Code

Pilot Job

```
package tutorial20.glue;

import ibis.ipl.Ibis;

public class PilotJob {
    Ibis ibis;
    Receiver receiver;
    SendPort sendPort;

    public static void main(String[] args) {
        try {
            new PilotJob().run();
        } catch (Exception e) {
            System.err.println("PilotJob failed: " + e);
            e.printStackTrace(System.err);
        }
    }

    void run() throws Exception {
        // ...
    }

    public static void main(String[] args) {
        // ...
    }
}
```



Example Code

Pilot Job

```
package ...  
import ...  
public class PilotJob {  
    IbisFactory ibis;  
    ReceivePort rp;  
    SendPort sp;  
    PilotJob() throws Exception {  
        ibis = IbisFactory.createIbis(Shared.ibisCapabilities, null,  
            Shared.portTypeServer, Shared.portTypeSlave);  
        IbisIdentifier server = ibis.registry().getElectionResult("JobServer");  
        rp = ibis.createReceivePort(Shared.portTypeSlave, "receiver");  
        rp.enableConnections();  
        sp = ibis.createSendPort(Shared.portTypeServer);  
        sp.connect(server, "receiver");  
    }  
    Job getWork(Result previousResult) throws Exception {  
        ...  
    }  
    void run() throws Exception {  
        ...  
    }  
    public static void main(String[] args) {  
        ...  
    }  
}
```



Example Code

Pilot Job

```
package ...
import ...
public class PilotJob {
    IbisIdentifier server;
    ReceivePort rp;
    SendPort sp;

    PilotJob() throws Exception {
        ibis = IbisFactory.createIbis(Shared.ibisCapabilities, null,
            Shared.portTypeServer, Shared.portTypeSlave);

        IbisIdentifier server = ibis.registry().getElectionResult("JobServer");

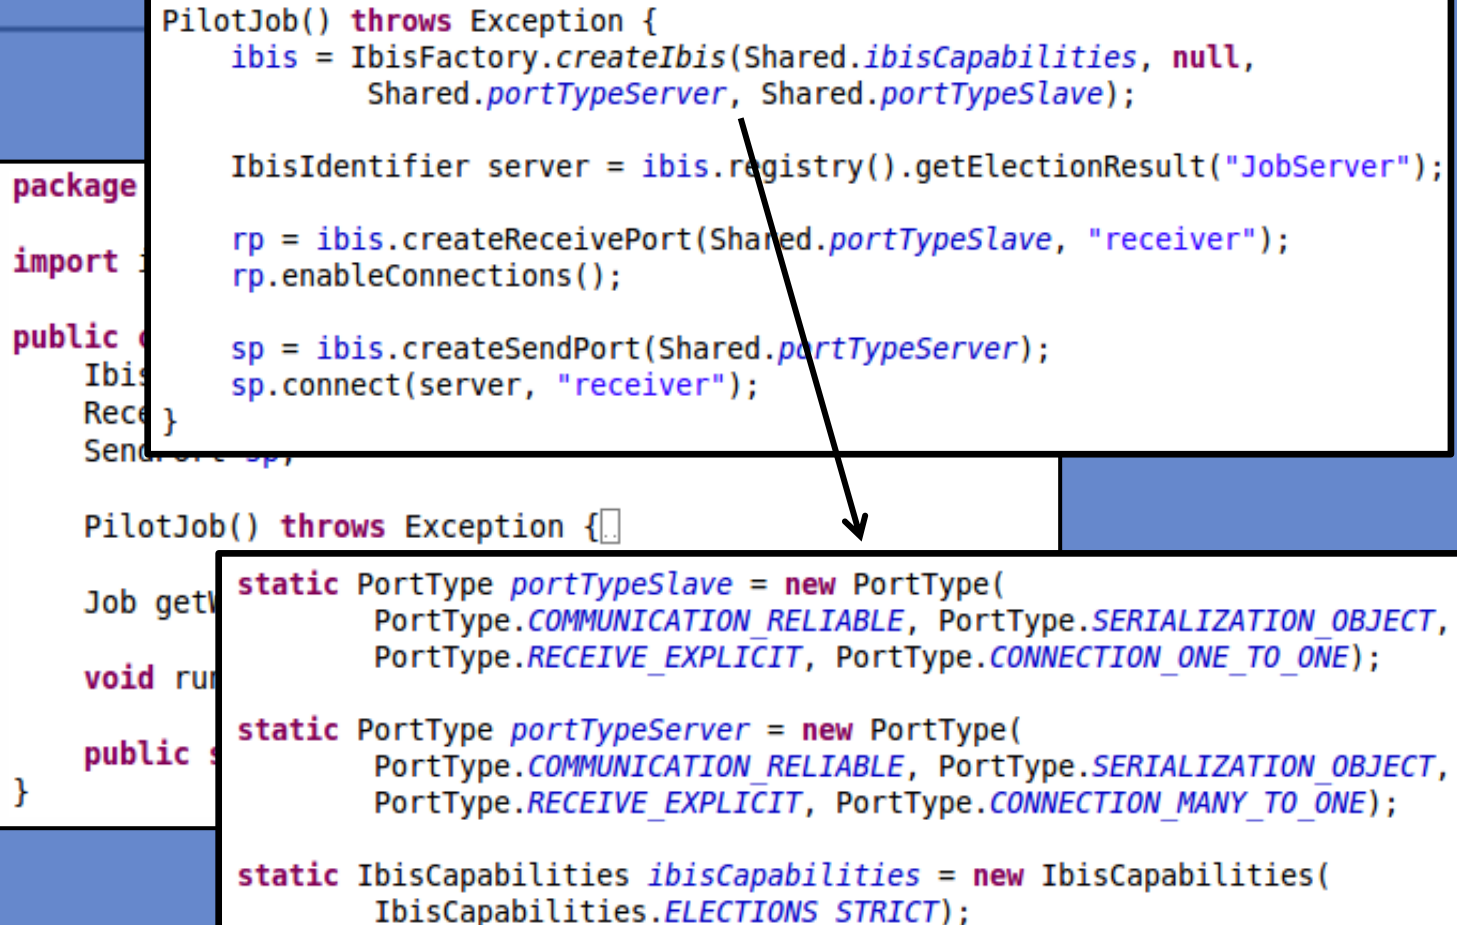
        rp = ibis.createReceivePort(Shared.portTypeSlave, "receiver");
        rp.enableConnections();

        sp = ibis.createSendPort(Shared.portTypeServer);
        sp.connect(server, "receiver");
    }

    static PortType portTypeSlave = new PortType(
        PortType.COMMUNICATION_RELIABLE, PortType.SERIALIZATION_OBJECT,
        PortType.RECEIVE_EXPLICIT, PortType.CONNECTION_ONE_TO_ONE);

    static PortType portTypeServer = new PortType(
        PortType.COMMUNICATION_RELIABLE, PortType.SERIALIZATION_OBJECT,
        PortType.RECEIVE_EXPLICIT, PortType.CONNECTION_MANY_TO_ONE);

    static IbisCapabilities ibisCapabilities = new IbisCapabilities(
        IbisCapabilities.ELECTIONS_STRICT);
}
```



Example Code

Pilot Job

```
package tutorial20.glue;

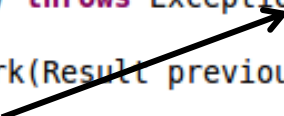
import ibis.ipl.Ibis;

public class PilotJob {
    Ibis ibis;
    ReceivePort rp;
    SendPort sp;

    PilotJob() throws Exception {
        Job getWork(Result previous) {
            void run() throws Exception {
                while (!job.empty) {
                    Result result = job.execute();
                    job = getWork(result);
                }
                ibis.end();
            }
        }

        void run() throws Exception {}

        public static void main(String[] args) {}
    }
}
```



Example Code

Pilot Job

```
package tutorial20.glu
```

```
import ibis.ipl.Ibis;
```

```
public class PilotJob {  
    Ibis ibis;  
    ReceivePort rp;  
    SendPort sp;
```

```
PilotJob() throws Exception {
```

```
    Job getWork(Result previousResult) throws Exception {
```

```
        void run() throws Exception {
```

```
            public static void main(String[] args) {
```

```
        }
```

```
        Job getWork(Result previousResult) throws Exception {  
            WriteMessage wm = sp.newMessage();  
            wm.writeObject(previousResult);  
            wm.finish();  
  
            ReadMessage rm = rp.receive();  
            Job job = (Job) rm.readObject();  
            rm.finish();  
            return job;  
        }
```



Example Code

Job Server Code

```
package tutorial20.glue;

import java.io.File;

public class JobServer {
    String inputDir, outputDir;
    LinkedList<Job> jobs = new LinkedList<Job>();
    Ibis ibis;
    ReceivePort rp;
    HashMap<IbisIdentifier, SendPort> workers =
        new HashMap<IbisIdentifier, SendPort>();

    JobServer(String executable, String[] arguments, String inputDir,
        SendPort getSendPort(IbisIdentifier target) throws IOException {
    void removeSendPort(IbisIdentifier target) throws IOException {
    void sendReply(IbisIdentifier target, Job job) throws IOException {
    void processResult(Result result) throws IOException {
    void handleRequest() throws IOException, ClassNotFoundException {
    void run() throws Exception {
}
```



Example Code

```
package tutorial20
```

```
import java.io.Fil
```

```
public class JobSe
```

```
String inputDir
```

```
LinkedList<Job
```

```
Ibis ibis;
```

```
ReceivePort rp
```

```
HashMap<IbisId
```

```
new Ha
```

```
JobServer(String executable, String[] arguments, String inputDir,  
String outputDir, Properties p) throws Exception {  
    this.inputDir = inputDir;  
    this.outputDir = outputDir;  
  
    for (String file : Shared.listFiles(inputDir, ".jpg"))  
        jobs.add(new Job(executable, arguments, file, "out-" + file));  
  
    ibis = IbisFactory.createIbis(Shared.ibisCapabilities, p, true, null,  
        Shared.portTypeServer, Shared.portTypeSlave);  
    ibis.registry().elect("JobServer");  
  
    rp = ibis.createReceivePort(Shared.portTypeServer, "receiver");  
    rp.enableConnections();  
}
```

```
JobServer(String executable, String[] arguments, String inputDir,
```

```
SendPort getSendPort(IbisIdentifier target) throws IOException {
```

```
void removeSendPort(IbisIdentifier target) throws IOException {
```

```
void sendReply(IbisIdentifier target, Job job) throws IOException {
```

```
void processResult(Result result) throws IOException {
```

```
void handleRequest() throws IOException, ClassNotFoundException {
```

```
void run() throws Exception {
```

```
}
```



Example Code

Job Server Code

```
package tutorial20.glue;

import java.io.File;

public class JobServer {
    String inputDir, outputDir;
    LinkedList<Job> jobs = new LinkedList<Job>();
    Ibis ibis;
    ReceivePort rp;
    HashMap<IbisIdentifier, SendPort> workers =
        new HashMap<IbisIdentifier, SendPort>();

    JobServer(String executable, String[] arguments, String inputDir,
        SendPort getSendPort(IbisIdentifier target) throws IOException {

    void removeSendPort(IbisIdentifier target) throws IOException {

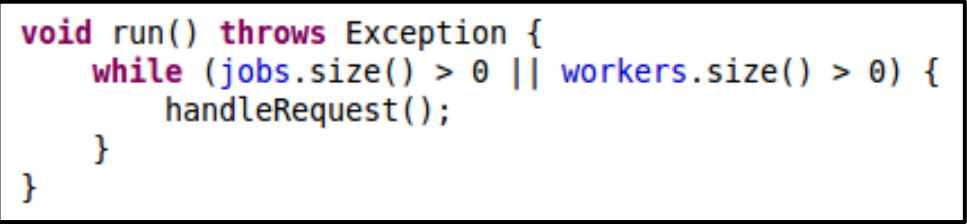
    void run() throws Exception {
        while (jobs.size() > 0 || workers.size() > 0) {
            handleRequest();
        }
    }

    void sendRep

    void process

    void handleR

    void run() throws Exception {
}
```



Example Code

Job Server Code

```
package tutorial20.glue;

import java.io.File;

public class JobServer {
    String inputDir, outputDir;
    LinkedList<Job> jobs = new LinkedList<Job>();
    Ibis ibis;
    ReceivePort rp;
    HashMap<IbisIdentifier, SendPort> workers =
        new HashMap<IbisIdentifier, SendPort>();

    JobServer(String inputDir, String outputDir) {
        this.inputDir = inputDir;
        this.outputDir = outputDir;
    }

    SendPort getSendPort(IbisIdentifier id) {
        return workers.get(id);
    }

    void removeSendPort(IbisIdentifier id) {
        workers.remove(id);
    }

    void sendReply(IbisIdentifier target, Job job) {
        SendPort sp = getSendPort(target);
        if (sp != null) {
            sp.send(job);
        }
    }

    void processResult(Result result) {
        Job job = (jobs.size() > 0 ? jobs.removeFirst() : new Job());
        job.setResult(result);
        sendReply(job.getIbisIdentifier(), job);
        processResult(result);
    }

    void handleRequest() throws IOException, ClassNotFoundException {
        ReadMessage rm = rp.receive();
        IbisIdentifier target = rm.origin().ibisIdentifier();
        Result result = (Result) rm.readObject();
        rm.finish();

        processResult(result);
    }

    void run() throws Exception {
        while (true) {
            handleRequest();
        }
    }
}
```



Example Code

Job Server Code

```
package tutorial20.glue;

import java.io.*;
import java.util.*;

public class JobServer {
    String inputDir;
    LinkedHashSet<IbisIdentifier> sendPorts;
    ReceivePort receivePort;
    HashMap<IbisIdentifier, SendPort> sendPortMap;

    JobServer(String inputDir) {
        this.inputDir = inputDir;
        sendPorts = new LinkedHashSet<IbisIdentifier>();
        receivePort = new ReceivePort();
        sendPortMap = new HashMap<IbisIdentifier, SendPort>();
    }

    void sendReply(IbisIdentifier target, Job job) throws IOException {
        SendPort sp = getSendPort(target);
        if (!job.empty()) {
            job.setInput(Shared.read(inputDir + File.separator + job.inputFile));
        }
        WriteMessage wm = sp.newMessage();
        wm.writeObject(job);
        wm.finish();
        if (job.empty()) {
            removeSendPort(target);
        }
    }

    void removeSendPort(IbisIdentifier target) {
        sendPorts.remove(target);
        sendPortMap.remove(target);
    }

    void sendReply(IbisIdentifier target, Job job) throws IOException {}

    void processResult(Result result) throws IOException {}

    void handleRequest() throws IOException, ClassNotFoundException {}

    void run() throws Exception {}
}
```



Example Code

```
package tutorial20.
```

```
import java.io.File
```

```
public class JobSer
```

```
String inputDir
```

```
LinkedList<Job> }
```

```
Ibis ibis;
```

```
ReceivePort rp;
```

```
HashMap<IbisIde
```

```
new Has
```

```
JobServer(Swin }
```

```
SendPort getSendPort(IbisIdentifier target) throws IOException {  
    SendPort sp = workers.get(target);  
    if (sp == null) {  
        sp = ibis.createSendPort(Shared.portTypeSlave);  
        sp.connect(target, "receiver");  
        workers.put(target, sp);  
    }  
    return sp;  
}
```

```
void removeSendPort(IbisIdentifier target) throws IOException {  
    SendPort sp = workers.remove(target);  
  
    if (sp != null) sp.close();  
}
```

```
SendPort getSendPort(IbisIdentifier target) throws IOException {  
      
    void removeSendPort(IbisIdentifier target) throws IOException {  
      
    void sendReply(IbisIdentifier target, Job job) throws IOException {  
      
    void processResult(Result result) throws IOException {  
      
    void handleRequest() throws IOException, ClassNotFoundException {  
      
    void run() throws Exception {  
      
    }  
}
```



Live DEMO



Conclusions

- IPL+SmartSockets = **Jungle proof** communication
 - **SmartSockets** solves **low-level** connectivity problems
 - **IPL** offers **high-level** communication primitives and resource tracking
- Adding JavaGAT makes it very easy to create applications that can on a jungle of resources
 - Pilot Job framework is a very simple example



Further improvements ?

- Improve hubs placement for IPL applications
 - Need automatic hub deployments
- Most **information** needed by JavaGAT is **static**
 - Individual resources don't change much over time
 - Neither do applications
 - Can't we store this info in **configuration files**?
- Can't we do all this in a “**point and click**” style ?
- Solution: Ibis Deploy (after the break)





Elections

- JEL also offers an 'election'
 - Allows a group to determine who's special
 - Ranks don't work in a malleable set of resources!
- Each election
 - Has a name (String)
 - Produces IbisIdentifier of the winner
 - Is not democratic
 - You can also be 'an observer'



Ibis Capabilities

- When initializing the application must specify:
 - The **PortTypes** it is going to use
 - Defines what kind of communication you need
 - The **Resource tracking behaviour** it needs
 - Defines what level of malleability you need
 - Optional: the preferred IPL implementation
 - SmartSockets, MX, MPI, etc.
- This allows the runtime to check if the requested combination is feasible

