



# Ibis as Master Key

## Hands-on – advanced examples

Jason Maassen

Computer Systems Group  
Department of Computer Science  
VU University, Amsterdam, The Netherlands



# Request

## Workflows


---

- Workflows (NBIC)
  - Jobs consisting of multiple stages (with dependencies)
  - Requirements vary per stage
- Ibis does not contain workflow system (yet), but:
  - JavaGAT is a perfect backend for a workflow system
    - Easy access to resources, files, etc.
- Question: how hard is it to implement a workflow ?



# Example

## Workflows



# Implementation

## Workflows

---

- Implementing the frontend is a lot of work
  - User friendly, flexible, GUI ?
  - Lots of systems exists!
- A JavaGAT backend is quite easy
  - We have already shown most of what you need (job submission and monitoring, file transfer, etc).
  - Still need the logic that determines what to run where and when.



# The code

## Workflows

---

- See masterkey.workflow package
- Main differences with previous examples:
  - Runs scripts instead of executables (more flexible)
  - Assumes data is remote
  - Reads work from a simple description file.

**disclaimer:** this is only a very simple example!  
(it took about 2 hours to implement and test)



# The code

## Workflows

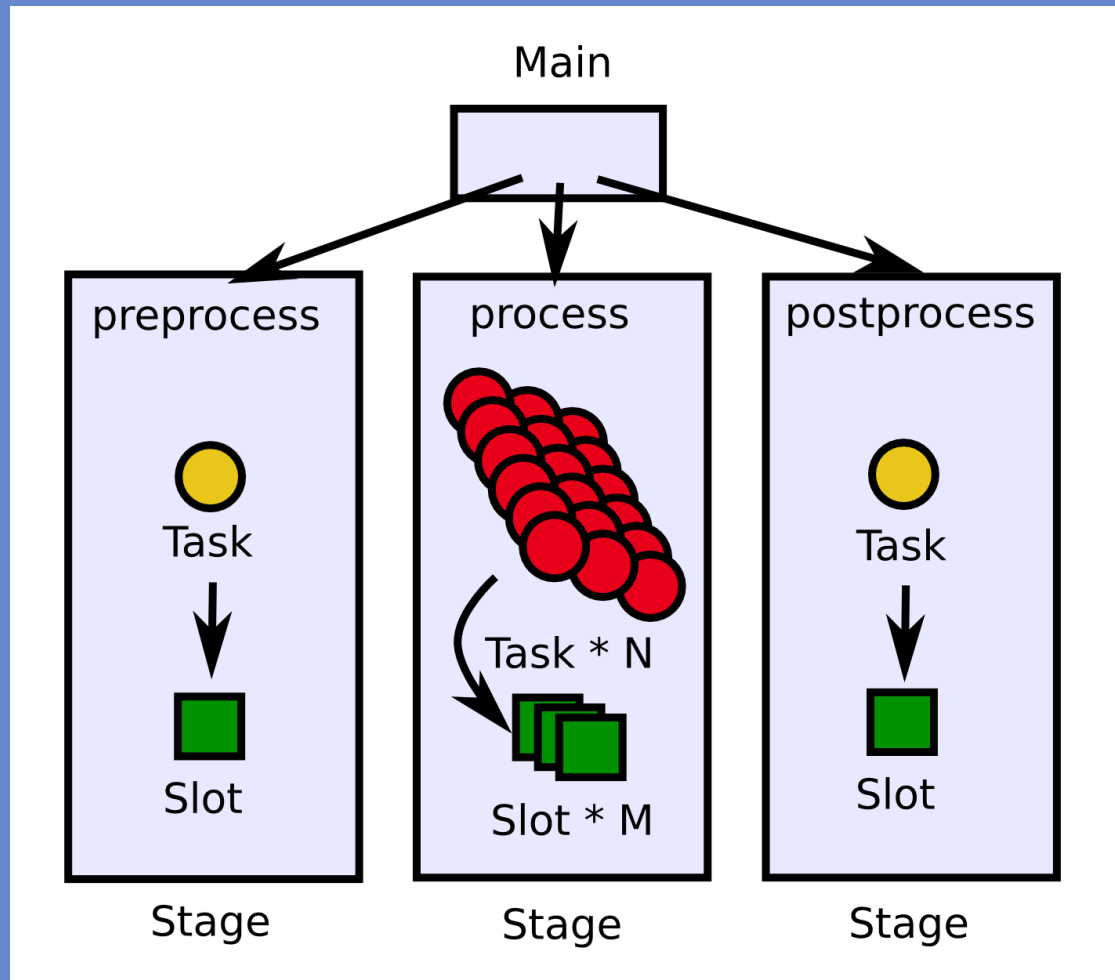
---

- Code consists of 4 classes:
  - **Main** is in control and creates a ...
  - **Stage** for each workflow stage that creates ...
  - **Slots** for each 'execution slot' and a list of ...
  - **Tasks** for each task to run.



# The code

## Workflows



# The code

## Workflows

- The workflow description looks like this:

#	STAGE	RESOURCE	SLOTS	SCRIPT	INPUT DATA	FILTER	OUTPUT DATA
#							
	preprocess,	ssh://fs0.das4.cs.vu.nl,	1,	/home/jason/pre.sh,	/home/jason/input.tgz,	-,	/home/jason/temp/stage1
	crunch,	sshpbs://fs0.das4.cs.vu.nl,	4,	/home/jason/crunch.sh,	/home/jason/temp/stage1,	.jpg,	/home/jason/temp/stage2
	postprocess,	ssh://fs0.das4.cs.vu.nl,	1,	/home/jason/post.sh,	/home/jason/temp/stage2,	-,	/home/jason/done.tgz

- Please change the input/output directories!
  - You can reuse the scripts and input file!

`./scripts/run masterkey.workflow.Main example.workflow`





# Requests

Bag of tasks

---

- Alternative to PilotJobs
  - PilotJob framework **bypasses** regular scheduler (sort of)
  - PilotJob framework **requires** IPL/SmartSockets
  - Both make sysadmins a bit nervous! 😊
- There is an alternative way of doing this
  - JavaGAT only
  - Use only regular schedulers
  - No additional communication
  - Add multicore support



# The Problem

## Bag of tasks

---


- Large number of data files to be processed
  - Pick a number – any number – but make it large!
- Processing each file is sequential
  - Oops. Modern machines are multi-core and schedulers often hand out whole machines.
- Varying processing time per data file
  - Varying data sizes, processing speeds (heuristics), core speeds, etc.



# The Problem

## Bag of tasks

---




Modern machines are multi-core and schedulers often hand out whole machines.



# The Problem

## Bag of tasks




Running one job per core does not help  
if the job sizes vary significantly  
(and you must now the number of cores)



# The Problem

## Bag of tasks




You would like to run **enough** jobs to keep all cores busy for an **equal** amount of time  
(difficult to do determine schedule in advance)



# The Solution

## Bag of tasks




Submit **application** that can run a **bag of tasks**  
Schedules jobs **dynamically** on the available cores  
Similar to a pilot job, but **no communication**



# The Solution

## Bag of tasks



Limit the time a bag-of-tasks may run (nice for queue)


Provide more than enough jobs to keep it busy

Kills any running jobs when it runs out of time



# The Solution

## Bag of tasks



When finished we collect all results and  
**resubmit the unprocessed ones**  
in the **next** bag of tasks





# The code

## Bag of tasks

---

- See masterkey.bagoftasks package
- Main differences with previous examples:
  - Runs scripts instead of executables (more flexible)

**disclaimer:** this is only a very simple example!  
(it only took a day to design, implement and test)



# The code

## Bag of tasks

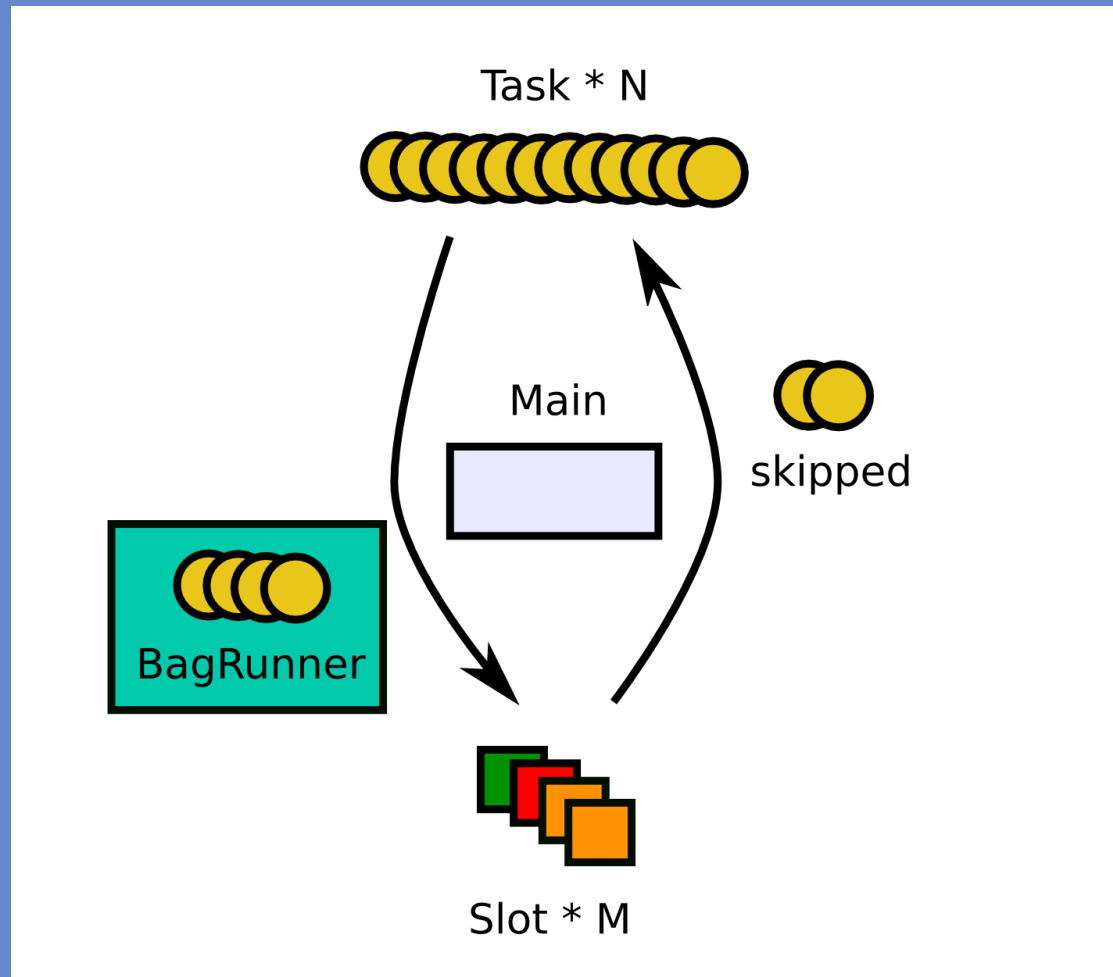
---

- Code consists of 5 classes:
  - **Main** is in control and creates ...
  - **Resources** containing info about resources and a ...
  - **Slot** for each 'execution slot' on a resource and a ...
  - **Tasks** for each task to run and deploys a ...
  - **BagOfTasksRunner** to a resource to dynamically schedule a bag of tasks on the available cores.



# The code

## Bag of tasks



# The code

## Bag of tasks

---

```
./scripts/run masterkey.bagoftasks.Main \  
  --resource sshsge://fs0.das4.cs.vu.nl /usr/bin/java 1 1 40 \  
  --script ./scripts/blur.sh \  
  --input ./images --output ./output
```

- The numbers after the resource are:
  - [slots] [minutes] [bagsize]
- Make sure the output directory exists!



# The code

## Bag of tasks

---

- After a run have a look in one of the tasks.in.\* and tasks.out.\* files
  - These are the in- and output of BagOfTasksRunner
- Also have a look at the stdout-\*.txt files
  - These show what jobs the BagOfTasksRunner has processed and how efficient it was.

